

BoltzMM: an R package for maximum pseudolikelihood estimation of fully-visible Boltzmann machines

Andrew T. Jones¹, Jessica J. Bagnall², and Hien D. Nguyen²

¹ School of Mathematics and Physics, University of Queensland, St. Lucia 4072, Queensland Australia ² Department of Mathematics and Statistics, La Trobe University, Bundoora 3086, Victoria Australia

DOI: [10.21105/joss.01193](https://doi.org/10.21105/joss.01193)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Submitted: 14 January 2019

Published: 14 February 2019

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC-BY](#)).

Summary

The *Boltzmann machine* (BM), first considered by Ackley, Hinton, & Sejnowski (1985), is a rich parametric probabilistic artificial neural network that is able to densely represent any *probability mass function* (PMF) over the support of *spin-binary strings* $\mathbb{X} = \{-1, +1\}^d$ (see, e.g., Le Roux & Bengio, 2008). Unfortunately, due to its complex form, and latent variable construction, it is often difficult to efficiently and meaningfully conduct parameter estimation when d is large and sample size n . Furthermore, it is often not necessary to consider such elaborate forms for practical modeling.

In Hyvarinen (2006), the *fully-visible BM* (FVBM) was first considered as a tractable simplification of the BM. Let $\mathbf{X} \in \mathbb{X}$ be such that the PMF can be written as

$$f(\mathbf{x}; \boldsymbol{\theta}) = \mathbb{P}(\mathbf{X} = \mathbf{x}) = \exp\left(\frac{1}{2}\mathbf{x}^\top \mathbf{M}\mathbf{x} + \mathbf{b}^\top \mathbf{x}\right) / z(\boldsymbol{\theta}),$$

where

$$z(\boldsymbol{\theta}) = \sum_{\boldsymbol{\xi} \in \mathbb{X}} \exp\left(\frac{1}{2}\boldsymbol{\xi}^\top \mathbf{M}\boldsymbol{\xi} + \mathbf{b}^\top \boldsymbol{\xi}\right),$$

is a normalizing constant, $\mathbf{b} \in \mathbb{R}^d$, and $\mathbf{M} \in \mathbb{R}^{d \times d}$ is a symmetric matrix with zero diagonal. We put the unique elements of the *bias vector* \mathbf{b} and the *interaction matrix* \mathbf{M} into the *parameter vector* $\boldsymbol{\theta}$. Interpretations of the parameter vector elements can be found in Section 2 of Bagnall, Jones, Karavarsamis, & Nguyen (2018).

The FVBM model can be viewed as a multivariate extension of the classical Bernoulli distribution, on spin-binary random variable strings (i.e. random variables $\mathbf{X} \in \mathbb{X}$, rather than $\mathbf{X} \in \{0, 1\}^d$), and can be demonstrated to be equivalent to the logistic multivariate binary model that was proposed by Cox (1972), and can be considered as a fully-connected binary graphical model (see Bagnall et al., 2018 for more details).

Let $\mathbf{X}_1, \dots, \mathbf{X}_n \in \mathbb{X}$ be a sample of n independent and identically distributed (IID) observations from an FVBM with unknown parameter vector $\boldsymbol{\theta}_0$. In Hyvarinen (2006), it was proved that $\boldsymbol{\theta}_0$ could be consistently estimated from $\mathbf{X}_1, \dots, \mathbf{X}_n \in \mathbb{X}$ via the so-called *maximum pseudolikelihood estimator* (MPLE) of the type described in Lindsay (1988) and Arnold & Strauss (1991). Furthermore, in H. D. Nguyen & Wood (2016a), an efficient and globally convergent *minorization-maximization* (MM; see Hunter & Lange, 2004; H. D. Nguyen, 2017) was proposed for the computation of the MPLE. The asymptotic normality of the MPLE was proved in H. D. Nguyen & Wood (2016b), which allows for the use of the MPLE of the FVBM as both a point estimator and a method for constructing useful hypothesis tests regarding relationships between binary random variables.

The **BoltzMM** package (Version 0.1.3; <https://CRAN.R-project.org/package=BoltzMM>) for the R statistical programming environment (R Core Team, 2018) provides a complete suite of functions for application and estimation of FVBM models. The most fundamental of the functions are **pvfbm** and **rfvbm**, which allow for the computation of the probability of any binary string, given valid input bias vector and interaction matrix, and the random generation of binary strings, given input parameter elements, respectively. Since the support \mathbb{X} is finite, it is often desirable to compute the probabilities of all possible outcomes, which can be achieved using the function **allpvfbm**. It is also often desirable to compute the marginal probability of any particular string index

$$f_i(x_i; \boldsymbol{\theta}) = \mathbb{P}(X_i = x_i),$$

where $i = 1, \dots, d$, $X_i \in \{-1, 1\}$, and $\mathbf{X}^\top = (X_1, \dots, X_d)$, for any parameter vector. This can be achieved using the function **marginpvfbm**.

The function **fitfvbm** implements the MM algorithm of H. D. Nguyen & Wood (2016a) in order to compute the MPLE of an unknown FVBM parameter vector $\boldsymbol{\theta}_0$, from an n observation realization of the sample $\mathbf{X}_1, \dots, \mathbf{X}_n$ (i.e., $\mathbf{x}_1, \dots, \mathbf{x}_n$). That is, **fitfvbm** takes in the data as a $n \times d$ dimensional matrix, where each row is a spin-binary string of length d . Using the **fitfvbm** function output, and the asymptotic normality results of H. D. Nguyen & Wood (2016b), standard errors can be computed using the function **fvbmstderr**, with the functions **fvbmccov**, **fvbmHess**, and **fvbmpartiald** providing intermediate calculations required by **fvbmstderr**.

As described in H. D. Nguyen & Wood (2016b) and demonstrated in Bagnall et al. (2018), hypothesis testing can be conducted using Wald-type hypothesis tests. These Wald-type tests can be constructed for testing of hypotheses of the form

$$H_0: \theta_{0k} = \theta_k, \text{ versus } H_1: \theta_{0k} \neq \theta_k,$$

where θ_{0k} is the k th element of the unknown parameter vector $\boldsymbol{\theta}_0$, and θ_k is some hypothesised value that it may take. The p -values and z -scores for these hypothesis tests can be computed using the outputs from **fitfvbm** and **fvbmstderr**, via the function **fvbmtests**.

All of the functions that have been comprehensively documented, and the documentation and example applications of each function can be accessed via the **help** function in R. Furthermore, the analysis of data arising from the voting patterns of the Senate of the 45th Australian Parliament in 2016, using the **BoltzMM** package, can be found in Bagnall et al. (2018). The report clearly describes the manner in which the FVBM can be used to conduct meaningful inference regarding multivariate binary data. The data that were analysed in the report are included in the package and can be accessed via the command **data(senate)**.

The **BoltzMM** package is programmed in native R, with particularly computationally intensive subroutines programmed in C and integrated via the **Rcpp** and **RcppArmadillo** packages of Eddelbuettel (2013). The package is therefore sufficiently suitable for analysis of moderately large dimensional data sets. On <https://cran.r-project.org>, the only other package that provides facilities for estimation of BMs, of any kind, is the **deepnet** package of Rong (2014). Here, facilities for estimation of the so-called *restricted BM* (RBM) are included, via the functions **rbm.down**, **rbm.train**, and **rbm.up**. The RBM is an alternative simplification to the BM, which preserves the denseness properties BM but can be parameterized by fewer parameter elements. However, the estimation of RBM models is as intractable as for BM models, and due to the latent variable construction preventing the establishment of distributional results for the estimators of the RBM parameter elements, its use for statistical inference remains limited. Thus we view the functions from **deepnet** as being complimentary but not overlapping with **BoltzMM**.

Users can obtain the latest build of **BoltzMM** on GitHub (<https://github.com/andrewthomasjones/BoltzMM>). The latest stable build can be obtained from CRAN

(<https://CRAN.R-project.org/package=BoltzMM>), and an archival build can be obtained from Zenodo (<http://doi.org/10.5281/zenodo.2563411>). Algorithm derivations and theoretical results regarding the implemented functions can be found in H. D. Nguyen & Wood (2016b) and H. D. Nguyen & Wood (2016a). An example application of the BoltzMM package for the analysis of real data is comprehensively described in Bagnall et al. (2018). Thorough descriptions of the package functions appear in the manual, which can be accessed at <https://cran.r-project.org/web/packages/BoltzMM/BoltzMM.pdf>. Bug reports and other feedback can be directed to the GitHub issues page (<https://github.com/andrewthomasjones/BoltzMM/issues>).

Acknowledgements

Jessica Bagnall is funded by a Research Training Program Stipend scholarship from La Trobe University. Hien Nguyen is funded under Australian Research Council (ARC) grant numbers DE170101134 and DP180101192.

References

- Ackley, D. H., Hinton, G. E., & Sejnowski, T. J. (1985). A learning algorithm for Boltzmann machines. *Cognitive Science*, 9, 147–169. doi:[https://doi.org/10.1016/S0364-0213\(85\)80012-4](https://doi.org/10.1016/S0364-0213(85)80012-4)
- Arnold, B. C., & Strauss, D. (1991). Pseudolikelihood estimation: some examples. *Sankhya B*, 53, 233–243. Retrieved from <https://www.jstor.org/stable/25052695>
- Bagnall, J., Jones, A., Karavarsamis, N., & Nguyen, H. (2018). *The fully-visible Boltzmann machine and the Senate of the 45th Australian Parliament in 2016*. SSRN. doi:<http://dx.doi.org/10.2139/ssrn.3287903>
- Cox, D. R. (1972). The analysis of multivariate binary data. *Journal of the Royal Statistical Society C*, 21(113-120). doi:<https://doi.org/10.2307/2346482>
- Eddelbuettel, D. (2013). *Seamless R and C++ Integration with Rcpp*. New York: Springer. doi:<https://doi.org/10.1007/978-1-4614-6868-4>
- Hunter, D. R., & Lange, K. (2004). A tutorial on MM algorithms. *The American Statistician*, 58, 30–37. doi:<https://doi.org/10.1198/0003130042836>
- Hyvarinen, A. (2006). Consistency of pseudolikelihood estimation of fully visible Boltzmann machines. *Neural Computation*, 18, 2283–2292. doi:<https://doi.org/10.1162/neco.2006.18.10.2283>
- Le Roux, N., & Bengio, Y. (2008). Representation power of restricted Boltzmann machines and deep belief networks. *Neural Computation*, 20, 1631–1649. doi:<https://doi.org/10.1162/neco.2008.04-07-510>
- Lindsay, B. (1988). Composite likelihood methods. *Contemporary Mathematics*, 8, 221–239. doi:<http://dx.doi.org/10.1090/conm/080/999014>
- Nguyen, H. D. (2017). An introduction to MM algorithms for machine learning and statistical estimation. *WIREs Data Mining and Knowledge Discovery*, 7, e1198. doi:<https://doi.org/10.1002/widm.1198>
- Nguyen, H. D., & Wood, I. A. (2016a). A block successive lower-bound maximization algorithm for the maximum pseudolikelihood estimation of fully visible Boltzmann machines. *Neural Computation*, 28, 485–492. doi:https://doi.org/10.1162/NECO_a_00813

Nguyen, H. D., & Wood, I. A. (2016b). Asymptotic normality of the maximum pseudo-likelihood estimator for fully visible Boltzmann machines. *IEEE Transactions on Neural Networks and Learning Systems*, 27, 897–902. doi:<https://doi.org/10.1109/TNNLS.2015.2425898>

R Core Team. (2018). *R: a language and environment for statistical computing*. R Foundation for Statistical Computing.

Rong, X. (2014). *deepnet: deep learning toolkit in R*. Retrieved from <https://CRAN.R-project.org/package=deepnet>