Surprise: A Python library for recommender systems

Nicolas Hug[1]

1 Columbia University

Summary

Recommender systems aim at providing users with a list of recommendations of items that a service offers. For example, a video streaming service will typically rely on a recommender system to propose a personalized list of movies or series to each of its users. A typical problem in recommendation is that of *rating prediction*: given an incomplete dataset of user-item interations which take the form of numerical ratings (e.g. on a scale from 1 to 5), the goal is to predict the missing ratings for all remaining user-item pairs.

`Surprise` is a Python library for building and analyzing rating prediction algorithms. It was designed to closely follow the `scikit-learn` API (Buitinck et al., 2013; Pedregosa et al., 2011) , which should be familiar to users acquainted with the Python machine learning ecosystem.

`Surprise` provides a collection of estimators (or prediction algorithms) for rating prediction. Among others, classical algorithms are implemented such as the main similarity-based algorithms (Aggarwal & others, 2016), as well as algorithms based on matrix factorization like SVD (Koren, Bell, & Volinsky, 2009) or NMF (Lee & Seung, 2001). It also supports tools for model evaluation like cross-validation iterators and built-in metrics *à la* `scikit-learn`, as well as tools for model selection and automatic hyper-parameter search, namely grid search and randomized search. Thanks to simple primitives and a light API, users can also implement their own recommendation technique with a minimal amount of code.

Classical datasets such as the MovieLens datasets (Harper & Konstan, 2015) are directly available in the package, but user-defined datasets are also supported either by loading `csv` files, or by using `pandas` dataframes (McKinney, 2010).

`Surprise` is mainly written in Python, while the computationally intensive parts are optimized with `Cython` (Behnel et al., 2011). Internally, `Surprise` relies on built-in Python data structures (mainly dictionaries) as well as `numpy` arrays (Walt, Colbert, & Varoquaux, 2011).

`Surprise` was designed to be useful to researchers who want to quickly explore new recommendation ideas by supporting the creation of custom prediction algorithms, but can also serve as a learning resource for students and less experienced users thanks to its detailed documentation.

Other popular recommendation libraries with similar functionalities include `LibRec` (Guo, Zhang, Sun, & Yorke-Smith, n.d.) (Java) or `MyMediaLite` (Gantner, Rendle, Freudenthaler, & Schmidt-Thieme, 2011) (C#). In Python, `OpenRec` (Yang, Bagdasaryan, Gruenstein, Hsieh, & Estrin, 2018) and `Spotlight` (Kula, 2017) support neural-network inspired algorithms; `implicit`[1] is specialized in implicit feedback recommendation, and `LightFM` (Kula, 2015) implements a hybrid algorithm based on matrix factorization. To the best of our knowledge, `Surprise` is the only library to provide a `scikit-learn` like API with model selection tools, and with a focus on explicit rating prediction.

---

[1] https://github.com/benfred/implicit

Example

Here is a simple example showing how to (down)load a dataset, split it into five folds for cross-validation, and compute the Mean Average Error (MAE) and the Root Mean Squared Error (RMSE) of the SVD algorithm.

```python
from surprise import SVD
from surprise import Dataset
from surprise.model_selection import cross_validate

# Load the movielens-100k dataset (download it if needed).
data = Dataset.load_builtin('ml-100k')

# Use the famous SVD algorithm, with default parameters.
algo = SVD()

# Run 5-fold cross-validation and print results. They can also be returned.
cross_validate(algo, data, measures=['RMSE', 'MAE'], cv=5, verbose=True)

# printed output:
# Evaluating RMSE, MAE of algorithm SVD on 5 split(s).
#               Fold 1  Fold 2  Fold 3  Fold 4  Fold 5  Mean    Std
# RMSE          0.9311  0.9370  0.9320  0.9317  0.9391  0.9342  0.0032
# MAE           0.7350  0.7375  0.7341  0.7342  0.7375  0.7357  0.0015
# Fit time      6.53    7.11    7.23    7.15    3.99    6.40    1.23
# Test time     0.26    0.26    0.25    0.15    0.13    0.21    0.06
```

Acknowledgements

References

Aggarwal, C. C., & others. (2016). *Recommender systems* (Vol. 1). Springer. doi:10.1007/978-3-319-29659-3

Behnel, S., Bradshaw, R., Citro, C., Dalcin, L., Seljebotn, D. S., & Smith, K. (2011). Cython: The best of both worlds. *Computing in Science & Engineering, 13*(2), 31–39. doi:10.1109/MCSE.2010.118

Buitinck, L., Louppe, G., Blondel, M., Pedregosa, F., Mueller, A., Grisel, O., Niculae, V., et al. (2013). API design for machine learning software: Experiences from the scikit-learn project. In *ECML pkdd workshop: Languages for data mining and machine learning* (pp. 108–122).

Gantner, Z., Rendle, S., Freudenthaler, C., & Schmidt-Thieme, L. (2011). MyMediaLite: A free recommender system library. In *Proceedings of the 5th acm conference on recommender systems (recsys 2011)*.

Guo, G., Zhang, J., Sun, Z., & Yorke-Smith, N. (n.d.). LibRec: A java library for recommender systems. In.

Harper, F. M., & Konstan, J. A. (2015). The movielens datasets: History and context. *Acm transactions on interactive intelligent systems (tiis)*, *5*(4), 1–19.

Koren, Y., Bell, R., & Volinsky, C. (2009). Matrix factorization techniques for recommender systems. *Computer*, *42*(8), 30–37.

Kula, M. (2015). Metadata embeddings for user and item cold-start recommendations. In T. Bogers & M. Koolen (Eds.), *Proceedings of the 2nd workshop on new trends on content-based recommender systems co-located with 9th ACM conference on recommender systems (recsys 2015), vienna, austria, september 16-20, 2015.*, CEUR workshop proceedings (Vol. 1448, pp. 14–21). CEUR-WS.org. Retrieved from http://ceur-ws.org/Vol-1448/paper4.pdf

Kula, M. (2017). Spotlight. https://github.com/maciejkula/spotlight; GitHub.

Lee, D. D., & Seung, H. S. (2001). Algorithms for non-negative matrix factorization. In *Advances in neural information processing systems* (pp. 556–562).

McKinney, W. (2010). Data structures for statistical computing in python. In S. van der Walt & J. Millman (Eds.), *Proceedings of the 9th python in science conference* (pp. 51–56). doi:10.25080/Majora-92bf1922-00a

Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., et al. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, *12*, 2825–2830.

Walt, S. van der, Colbert, S. C., & Varoquaux, G. (2011). The numpy array: A structure for efficient numerical computation. *Computing in Science & Engineering*, *13*(2), 22–30. doi:10.1109/MCSE.2011.37

Yang, L., Bagdasaryan, E., Gruenstein, J., Hsieh, C.-K., & Estrin, D. (2018). OpenRec: A modular framework for extensible and adaptable recommendation algorithms. In *Proceedings of the eleventh acm international conference on web search and data mining*, WSDM '18 (pp. 664–672). New York, NY, USA: Association for Computing Machinery. doi:10.1145/3159652.3159681