

c-lasso - a Python package for constrained sparse and robust regression and classification

Léo Simpson¹, Patrick L. Combettes², and Christian L. Müller^{3,4,5}

¹ Technische Universität München ² Department of Mathematics, North Carolina State University, Raleigh ³ Center for Computational Mathematics, Flatiron Institute, New York ⁴ Institute of Computational Biology, Helmholtz Zentrum München ⁵ Department of Statistics, Ludwig-Maximilians-Universität München

DOI: [10.21105/joss.02844](https://doi.org/10.21105/joss.02844)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Matthew Sottile](#) ↗

Reviewers:

- [@jbytecode](#)
- [@glemaitre](#)

Submitted: 02 November 2020

Published: 17 January 2021

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

We introduce c-lasso, a Python package that enables sparse and robust linear regression and classification with linear equality constraints. The underlying statistical forward model is assumed to be of the following form:

$$y = X\beta + \sigma\epsilon \quad \text{subject to} \quad C\beta = 0$$

Here, $X \in \mathbb{R}^{n \times d}$ is a given design matrix and the vector $y \in \mathbb{R}^n$ is a continuous or binary response vector. The matrix C is a general constraint matrix. The vector $\beta \in \mathbb{R}^d$ contains the unknown coefficients and σ an unknown scale. Prominent use cases are (sparse) log-contrast regression with compositional data X , requiring the constraint $1_d^T \beta = 0$ ([Aitchison & Bacon-Shone, 1984](#)) and the Generalized Lasso which is a *special case* of the described problem (see, e.g. ([James et al., 2020](#)), Example 3). The c-lasso package provides estimators for inferring unknown coefficients and scale (i.e., perspective M-estimators ([Combettes & Müller, 2020a](#))) of the form

$$\min_{\beta \in \mathbb{R}^d, \sigma \in \mathbb{R}_0} f(X\beta - y, \sigma) + \lambda \|\beta\|_1 \quad \text{subject to} \quad C\beta = 0$$

for several convex loss functions $f(\cdot, \cdot)$. This includes the constrained Lasso, the constrained scaled Lasso, sparse Huber M-estimators with linear equality constraints, and constrained (Huberized) Square Hinge Support Vector Machines (SVMs) for classification.

Statement of need

Currently, there is no Python package available that can solve these ubiquitous statistical estimation problems in a fast and efficient manner. c-lasso provides algorithmic strategies, including path and proximal splitting algorithms, to solve the underlying convex optimization problems with provable convergence guarantees. The c-lasso package is intended to fill the gap between popular Python tools such as [scikit-learn](#) which cannot solve these constrained problems and general-purpose optimization solvers such as [cvxpy](#) that do not scale well for these problems and/or are inaccurate. c-lasso can solve the estimation problems at a single regularization level, across an entire regularization path, and includes three model selection strategies for determining the regularization parameter: a theoretically-derived fixed regularization, k-fold cross-validation, and stability selection. We show several use cases of the package, including an application of sparse log-contrast regression tasks for compositional microbiome data, and highlight the seamless integration into R via [reticulate](#).

Functionalities

Installation and problem instantiation

c-lasso is available on pip and can be installed in the shell using

```
pip install c-lasso
```

c-lasso is a stand-alone package and not yet compatible with the scikit-learn API. The central object in the c-lasso package is the instantiation of a c-lasso problem.

```
# Import the main class of the package
from classo import classo_problem

# Define a c-lasso problem instance with default setting,
# given data X, y, and constraints C.
problem = classo_problem(X, y, C)
```

We next describe what type of problem instances are available and how to solve them.

Statistical problem formulations

Depending on the type of and the prior assumptions on the data, the noise ϵ , and the model parameters, c-lasso allows for different estimation problem formulations. More specifically, the package can solve the following four regression-type and two classification-type formulations:

R1 Standard constrained Lasso regression:

$$\min_{\beta \in \mathbb{R}^d} \|X\beta - y\|^2 + \lambda \|\beta\|_1 \quad \text{subject to} \quad C\beta = 0$$

This is the standard Lasso problem with linear equality constraints on the β vector. The objective function combines Least-Squares (LS) for model fitting with the L_1 -norm penalty for sparsity.

```
# Formulation R1
problem.formulation.huber = False
problem.formulation.concomitant = False
problem.formulation.classification = False
```

R2 Constrained sparse Huber regression:

$$\min_{\beta \in \mathbb{R}^d} h_\rho(X\beta - y) + \lambda \|\beta\|_1 \quad \text{subject to} \quad C\beta = 0$$

This regression problem uses the [Huber loss](#) h_ρ as objective function for robust model fitting with an L_1 penalty and linear equality constraints on the β vector. The default parameter ρ is set to 1.345 ([Huber, 1981](#)).

```
# Formulation R2
problem.formulation.huber = True
problem.formulation.concomitant = False
problem.formulation.classification = False
```

R3 Constrained scaled Lasso regression:

$$\min_{\beta \in \mathbb{R}^d, \sigma \in \mathbb{R}_0} \frac{\|X\beta - y\|^2}{\sigma} + \frac{n}{2}\sigma + \lambda \|\beta\|_1 \quad \text{subject to} \quad C\beta = 0$$

This formulation is the default problem formulation in c-lasso. It is similar to [R1](#) but allows for joint estimation of the (constrained) β vector and the standard deviation σ in a concomitant fashion ([Combettes & Müller, 2020a, 2020b](#)).

Formulation R3

```
problem.formulation.huber = False
problem.formulation.concomitant = True
problem.formulation.classification = False
```

R4 Constrained sparse Huber regression with concomitant scale estimation:

$$\min_{\beta \in \mathbb{R}^d, \sigma \in \mathbb{R}_0} \left(h_\rho \left(\frac{X\beta - y}{\sigma} \right) + n \right) \sigma + \lambda \|\beta\|_1 \quad \text{subject to} \quad C\beta = 0$$

This formulation combines [R2](#) and [R3](#) allowing robust joint estimation of the (constrained) β vector and the scale σ in a concomitant fashion ([Combettes & Müller, 2020a, 2020b](#)).

Formulation R4

```
problem.formulation.huber = True
problem.formulation.concomitant = True
problem.formulation.classification = False
```

C1 Constrained sparse classification with Square Hinge loss:

$$\min_{\beta \in \mathbb{R}^d} \sum_{i=1}^n l(y_i x_i^\top \beta) + \lambda \|\beta\|_1 \quad \text{subject to} \quad C\beta = 0$$

where x_i denotes the i^{th} row of X , $y_i \in \{-1, 1\}$, and $l(\cdot)$ is defined for $r \in \mathbb{R}$ as:

$$l(r) = \begin{cases} (1-r)^2 & \text{if } r \leq 1 \\ 0 & \text{if } r \geq 1 \end{cases}$$

This formulation is similar to [R1](#) but adapted for classification tasks using the Square Hinge loss with (constrained) sparse β vector estimation ([Lee & Lin, 2013](#)).

Formulation C1

```
problem.formulation.huber = False
problem.formulation.concomitant = False
problem.formulation.classification = True
```

C2 Constrained sparse classification with Huberized Square Hinge loss:

$$\min_{\beta \in \mathbb{R}^d} \sum_{i=1}^n l_{\rho}(y_i x_i^{\top} \beta) + \lambda \|\beta\|_1 \quad \text{subject to} \quad C\beta = 0.$$

This formulation is similar to [C1](#) but uses the Huberized Square Hinge loss l_{ρ} for robust classification with (constrained) sparse β vector estimation ([Rosset & Zhu, 2007](#)):

$$l_{\rho}(r) = \begin{cases} (1-r)^2 & \text{if } \rho \leq r \leq 1 \\ (1-\rho)(1+\rho-2r) & \text{if } r \leq \rho \\ 0 & \text{if } r \geq 1 \end{cases}$$

This formulation can be selected in c-lasso as follows:

```
# Formulation C2
problem.formulation.huber = True
problem.formulation.concomitant = False
problem.formulation.classification = True
```

Optimization schemes

The problem formulations *R1-C2* require different algorithmic strategies for efficiently solving the underlying optimization problems. The c-lasso package implements four published algorithms with provable convergence guarantees. The package also includes novel algorithmic extensions to solve Huber-type problems using the mean-shift formulation ([Mishra & Müller, 2019](#)). The following algorithmic schemes are implemented:

- Path algorithms (*Path-Alg*): This algorithm follows the proposal in ([Gaines et al., 2018](#); [Jeon et al., 2020](#)) and uses the fact that the solution path along λ is piecewise-affine ([Rosset & Zhu, 2007](#)). We also provide a novel efficient procedure that allows to derive the solution for the concomitant problem *R3* along the path with little computational overhead.
- Douglas-Rachford-type splitting method (*DR*): This algorithm can solve all regression problems *R1-R4*. It is based on Douglas-Rachford splitting in a higher-dimensional product space and makes use of the proximity operators of the perspective of the LS objective ([Combettes & Müller, 2020a, 2020b](#)). The Huber problem with concomitant scale *R4* is reformulated as scaled Lasso problem with mean shift vector ([Mishra & Müller, 2019](#)) and thus solved in $(n + d)$ dimensions.
- Projected primal-dual splitting method (*P-PDS*): This algorithm is derived from ([Briceño-Arias & López Rivera, 2019](#)) and belongs to the class of proximal splitting algorithms, extending the classical Forward-Backward (FB) (aka proximal gradient descent) algorithm to handle an additional linear equality constraint via projection. In the absence of a linear constraint, the method reduces to FB.
- Projection-free primal-dual splitting method (*PF-PDS*): This algorithm is a special case of an algorithm proposed in ([Combettes & Pesquet, 2012](#)) (Eq. 4.5) and also belongs to the class of proximal splitting algorithms. The algorithm does not require projection operators which may be beneficial when C has a more complex structure. In the absence of a linear constraint, the method reduces to the Forward-Backward-Forward scheme.

The following table summarizes the available algorithms and their recommended use for each problem:

	<i>Path-Alg</i>	<i>DR</i>	<i>P-PDS</i>	<i>PF-PDS</i>
<i>R1</i>	use for large λ and path computation	use for small λ	possible	use for complex constraints
<i>R2</i>	use for large λ and path computation	use for small λ	possible	use for complex constraints
<i>R3</i>	use for large λ and path computation	use for small λ	-	-
<i>R4</i>	-	only option	-	-
<i>C1</i>	only option	-	-	-
<i>C2</i>	only option	-	-	-

The following Python snippet shows how to select a specific algorithm:

```
problem.numerical_method = "Path-Alg"
# Alternative options: "DR", "P-PDS", and "PF-PDS"
```

Computation modes and model selection

The c-lasso package provides several computation modes and model selection schemes for tuning the regularization parameter.

- *Fixed Lambda*: This setting lets the user choose a fixed parameter λ or a proportion $l \in [0, 1]$ such that $\lambda = l \times \lambda_{\max}$. The default value is a scale-dependent tuning parameter that has been derived in (Shi et al., 2016) and applied in (Combettes & Müller, 2020b).
- *Path Computation*: This setting allows the computation of a solution path for λ parameters in an interval $[\lambda_{\min}, \lambda_{\max}]$. The solution path is computed via the *Path-Alg* scheme or via warm-starts for other optimization schemes.
- *Cross Validation*: This setting allows the selection of the regularization parameter λ via k-fold cross validation for $\lambda \in [\lambda_{\min}, \lambda_{\max}]$. Both the Minimum Mean Squared Error (or Deviance) (MSE) and the “One-Standard-Error rule” (1SE) are available (Hastie et al., 2009).
- *Stability Selection*: This setting allows the selection of the λ via stability selection (Combettes & Müller, 2020b; Lin et al., 2014; Meinshausen & Bühlmann, 2010). Three modes are available: selection at a fixed λ (Combettes & Müller, 2020b), selection of the q first variables entering the path (default setting), and of the q largest coefficients (in absolute value) across the path (Meinshausen & Bühlmann, 2010).

The Python syntax to use a specific computation mode and model selection is exemplified below:

```
# Example how to perform path computation and cross-validation:
problem.model_selection.LAMfixed = False
problem.model_selection.PATH = True
problem.model_selection.CV = True
problem.model_selection.StabSel = False

# Example how to add stability selection to the problem instance
problem.model_selection.StabSel = True
```

Each model selection procedure has additional meta-parameters that are described in the [Documentation](#).

Numerical benchmarks

To evaluate optimization accuracy and running time of the different algorithms available in `c-lasso`, we provide [micro-benchmark](#) experiments which also include `cvxpy`, an open source convex optimization software, for baseline comparison. All experiments have been computed using Python 3.9.1 on a MacBook Air with a 1.8 GHz Intel Core i5 processor and 8 Gb 1600 MHz DDR3 memory, operating on macOS High Sierra.

Figure 1 summarizes the results for the *Path-Alg*, *DR*, and *P-PDS* algorithms solving the regression formulation [R1](#) for different samples sizes n and problem dimensions p on synthetic data (using `c-lasso`'s data generator). We observe that `c-lasso`'s algorithms are faster and more accurate than the `cvx` baseline. For instance, for $d = 500$ features and $n = 500$ samples, the *Path-Alg* algorithm is about 70 times faster than `cvx`.




Figure 1: Average running times (left panel) of *Path-Alg* (blue), *P-PDS* (yellow), *DR* (green), and `cvx` (red) at fixed $\lambda = 0.1$ and corresponding average objective function value differences (with respect to the function value obtained by the *Path-Alg* solution as baseline) (right panel). Mean (and standard deviation) running time is calculated over 20 data replications for each sample size/dimension scenario (n, d) . On a single data set, the reported running time of an algorithm is the average time of five algorithm runs (to guard against system background process fluctuations).

The complete reproducible micro-benchmark is available [here](#).

Computational examples

Toy example using synthetic data

We illustrate the workflow of the `c-lasso` package on synthetic data using the built-in routine `random_data` which enables the generation of test problem instances with normally distributed data X , sparse coefficient vectors β , and constraints $C \in \mathbb{R}^{k \times d}$.

Here, we use a problem instance with $n = 100$, $d = 100$, a β with five non-zero components, $\sigma = 0.5$, and a zero-sum constraint.

```
from classo import classo_problem, random_data

n, d, d_nonzero, k, sigma = 100, 100, 5, 1, 0.5
(X, C, y), sol = random_data(
    n, d, d_nonzero, k, sigma,
    zerosum = True, seed = 123
)
print("Relevant variables : {}".format(numpy.nonzero(sol)[0]))

problem = classo_problem(X, y, C)

problem.formulation.huber = True
problem.formulation.concomitant = False
problem.formulation.rho = 1.5

problem.model_selection.LAMfixed = True
problem.model_selection.PATH = True
problem.model_selection.LAMfixedparameters.rescaled_lam = True
problem.model_selection.LAMfixedparameters.lam = 0.1

problem.solve()

print(problem.solution)
```

We use [formulation \$R2\$](#) with $\rho = 1.5$, [computation mode and model selections](#) *Fixed Lambda* with $\lambda = 0.1\lambda_{\max}$, *Path computation*, and *Stability Selection* (as per default).

The corresponding output reads:

```
Relevant variables : [43 47 74 79 84]

LAMBDA FIXED :
  Selected variables : 43    47    74    79    84
  Running time : 0.294s

PATH COMPUTATION :
  Running time : 0.566s

STABILITY SELECTION :
  Selected variables : 43    47    74    79    84
  Running time : 5.3s
```

c-lasso allows standard visualization of the computed solutions, e.g., coefficient plots at fixed λ , the solution path, the stability selection profile at the selected λ , and the stability selection profile across the entire path.

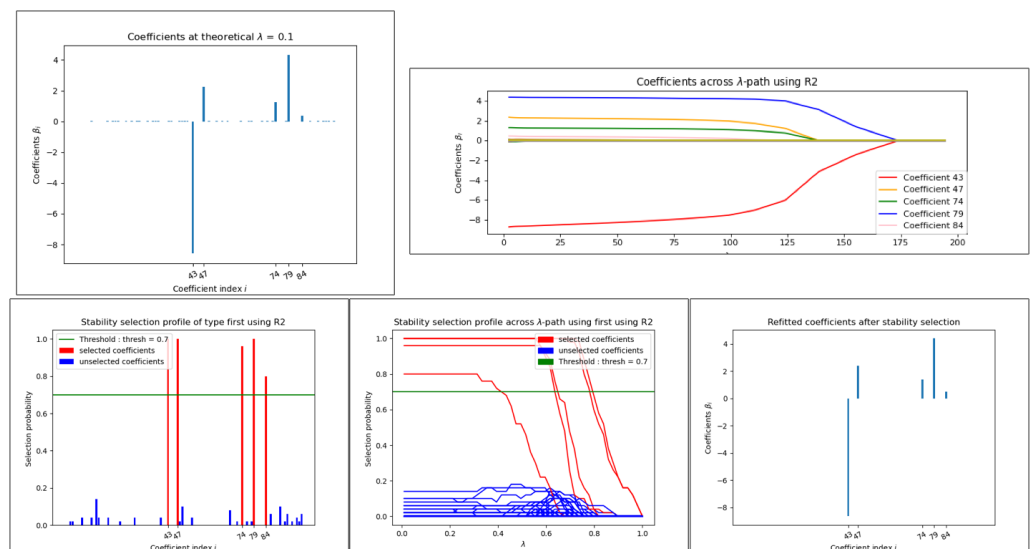


Figure 2: Visualizations after calling `problem.solution`

For this tuned example, the solutions at the fixed lambda and with stability selection recover the oracle solution. The solution vectors are stored in `problem.solution` and can be directly accessed for each mode/model selection.

Access to the estimated coefficient vector at a fixed lambda
`problem.solution.LAMfixed.beta`

Note that the run time for this $d = 100$ -dimensional example for a single path computation is about 0.5 seconds on a standard laptop.

Log-contrast regression on gut microbiome data

We next illustrate the application of `c-lasso` on the [COMBO microbiome dataset](#) (Combettes & Müller, 2020b; Lin et al., 2014; Shi et al., 2016). Here, the task is to predict the Body Mass Index (BMI) of $n = 96$ participants from $d = 45$ relative abundances of bacterial genera, and absolute calorie and fat intake measurements. The code snippet for this example is available in the [README.md](#) and the [example notebook](#).




Figure 3: Stability selection profiles of problems R3/R4 on the COMBO data

Stability selection profiles using [formulation R3](#) (left) and [R4](#) (right) on the COMBO dataset, reproducing Figure 5a in ([Combettes & Müller, 2020b](#)).

Calling c-lasso in R

The c-lasso package also integrates with R via the R package [reticulate](#). We refer to [reticulate](#)'s manual for technical details about connecting python environments and R. A successful use case of c-lasso is available in the R package [trac](#) ([Bien et al., 2020](#)), enabling tree-structured aggregation of predictors when features are rare.

The code snippet below shows how c-lasso is called in R to perform regression at a fixed $\lambda = 0.1\lambda_{\max}$. In R, X and C need to be of matrix type, and y of array type.

```
problem <- classo$classo_problem(X = X, C = C, y = y)
problem$model_selection$LAMfixed <- TRUE
problem$model_selection$StabSel <- FALSE
problem$model_selection$LAMfixedparameters$rescaled_lam <- TRUE
problem$model_selection$LAMfixedparameters$lam <- 0.1
problem$solve()

# Extract coefficient vector with tidy-verse
beta <- as.matrix(map_dfc(problem$solution$LAMfixed$beta, as.numeric))
```

Acknowledgements

The work of LS was conducted at and financially supported by the Center for Computational Mathematics (CCM), Flatiron Institute, New York, and the Institute of Computational Biology, Helmholtz Zentrum München. We thank Dr. Leslie Greengard (CCM and Courant Institute, NYU) for facilitating the initial contact between LS and CLM. The work of PLC was supported by the National Science Foundation under grant DMS-1818946.

References

- Aitchison, J., & Bacon-Shone, J. (1984). Log contrast models for experiments with mixtures. *Biometrika*, 71(2), 323–330. <https://doi.org/10.1093/biomet/71.2.323>
- Bien, J., Yan, X., Simpson, L., & Müller, C. L. (2020). Tree-Aggregated Predictive Modeling of Microbiome Data. *bioRxiv*. <https://doi.org/10.1101/2020.09.01.277632>
- Briceño-Arias, L., & López Rivera, S. (2019). A projected primal–dual method for solving constrained monotone inclusions. *Journal of Optimization Theory and Applications*, 180(3), 907–924. <https://doi.org/10.1007/s10957-018-1430-2>
- Combettes, P. L., & Müller, C. L. (2020a). Perspective maximum likelihood-type estimation via proximal decomposition. *Electron. J. Statist.*, 14(1), 207–238. <https://doi.org/10.1214/19-EJS1662>
- Combettes, P. L., & Müller, C. L. (2020b). Regression models for compositional data: General log-contrast formulations, proximal optimization, and microbiome data applications. *Statistics in Biosciences*. <https://doi.org/10.1007/s12561-020-09283-2>
- Combettes, P. L., & Pesquet, J.-C. (2012). Primal-Dual Splitting Algorithm for Solving Inclusions with Mixtures of Composite, Lipschitzian, and Parallel-Sum Type Monotone

- Operators. *Set-Valued and Variational Analysis*, 20, 307–320. <https://doi.org/10.1007/s11228-011-0191-y>
- Gaines, B. R., Kim, J., & Zhou, H. (2018). Algorithms for fitting the constrained lasso. *Journal of Computational and Graphical Statistics*, 27(4), 861–871. <https://doi.org/10.1080/10618600.2018.1473777>
- Hastie, T., Tibshirani, R., & Friedman, J. H. (2009). *The elements of statistical learning: Data mining, inference, and prediction*. Springer. <https://doi.org/10.1007/bf02985802>
- Huber, P. (1981). *Robust statistics*. John Wiley & Sons Inc. ISBN: 0-471-41805-6
- James, G. M., Paulson, C., & Rusmevichientong, P. (2020). Penalized and constrained optimization: An application to high-dimensional website advertising. *Journal of the American Statistical Association*, 115(529), 107–122. <https://doi.org/10.1080/01621459.2019.1609970>
- Jeon, J.-J., Kim, Y., Won, S., & Choi, H. (2020). Primal path algorithm for compositional data analysis. *Computational Statistics & Data Analysis*, 148, 106958. <https://doi.org/10.1016/j.csda.2020.106958>
- Lee, C.-P., & Lin, C.-J. (2013). A study on L2-loss (squared hinge-loss) multiclass SVM. *Neural Computation*, 25. https://doi.org/10.1162/NECO_a_00434
- Lin, W., Shi, P., Feng, R., & Li, H. (2014). Variable selection in regression with compositional covariates. *Biometrika*, 101(4), 785–797. <https://doi.org/10.1093/biomet/asu031>
- Meinshausen, N., & Bühlmann, P. (2010). Stability selection. *Journal of the Royal Statistical Society: Series B (Statistical Methodology)*, 72(4), 417–473. <https://doi.org/10.1111/j.1467-9868.2010.00740.x>
- Mishra, A., & Müller, C. L. (2019). *Robust regression with compositional covariates*. <http://arxiv.org/abs/1909.04990>
- Rosset, S., & Zhu, J. (2007). Piecewise linear regularized solution paths. *Annals of Statistics*, 35(3), 1012–1030. <https://doi.org/10.1214/009053606000001370>
- Shi, P., Zhang, A., & Li, H. (2016). Regression analysis for microbiome compositional data. *Annals of Applied Statistics*, 10(2), 1019–1040. <https://doi.org/10.1214/16-AOAS928>