The targets R package: a dynamic Make-like function-oriented pipeline toolkit for reproducibility and high-performance computing

William Michael Landau[1]

1 Eli Lilly and Company

Summary

The `targets` R package (Landau, 2020) is a pipeline toolkit for computationally intense reproducible research. It reduces the time and effort required to develop a data analysis project and maintain a trustworthy set of results. `targets` uses static code analysis to detect dependency relationships among interconnected computational tasks and construct a directed acyclic graph (DAG), which researchers can visualize in order to understand and communicate the structure of a complicated workflow. To run the pipeline at scale, `targets` leverages implicit parallel computing and optional cloud storage. In subsequent runs, `targets`, skips tasks that are already synchronized with their upstream dependencies, which not only reduces the runtime of rapidly developing workflows, but also provides tangible evidence of reproducibility.

In high-performance computing scenarios, `targets` uses its DAG to discern which targets can run concurrently and which targets are still waiting for other upstream targets to finish processing. As soon as a target's dependency requirements are met, the target is deployed to the next available parallel worker. Internally, `targets` leverages the `clustermq` package (Schubert, 2019) for persistent workers and the `future` package (Bengtsson, 2020) for transient workers. Both `clustermq` and `future` are powerful and versatile frameworks capable of submitting R workloads not only to multiple cores on a single machine, but also to popular resource managers on shared computing clusters.

`targets` is the successor to `drake` (Landau, 2018), which in turn originated from `remake` (FitzJohn, 2017), an R package modeled after GNU Make (Stallman, 1998). Unlike Make, `targets` and `drake` and `remake` focus on the R language, encourage an idiomatic function-oriented style of programming, and abstract each target as an R object. Relative to `remake` and `drake`, `targets` is friendlier and more efficient, surpassing the permanent architectural limitations of both predecessors. The data storage system of `targets` is lighter and more transparent, which helps users diagnose issues, move projects to different file systems, work with multiple contributors, and leverage seamless Metaflow-like cloud storage integration (Ge & Goyal, n.d.). In addition, `targets` supports stronger user-side guardrails, more introspective dependency graph visualizations, parallel efficient dynamic branching, and an interface more amenable to metaprogramming and third-party extensions.

References

Bengtsson, H. (2020). *A unifying framework for parallel and distributed processing in r using futures*. https://arxiv.org/abs/2008.00553

FitzJohn, R. (2017). *remake: Make-like build management*. https://github.com/richfitz/remake.

Ge, J., & Goyal, S. (n.d.). *Metaflow: Metaflow for r-lang*.

Landau, W. M. (2018). The drake r package: A pipeline toolkit for reproducibility and high-performance computing. *Journal of Open Source Software*, *3*(21), 550. https://doi.org/10.21105/joss.00550

Landau, W. M. (2020). *targets: dynamic function-oriented Make-like declarative pipelines for R*. https://github.com/ropensci/targets.

Schubert, M. (2019). clustermq enables efficient parallelization of genomic analyses. *Bioinformatics*, *35*(21), 4493–4495. https://doi.org/10.1093/bioinformatics/btz284

Stallman, R. (1998). *GNU Make, Version 3.77*. Free Software Foundation. ISBN: 1882114809