

gym-saturation: an OpenAI Gym environment for saturation provers

Boris Shminke ¹

¹ Laboratoire J.A. Dieudonné, CNRS and Université Côte d'Azur, France

DOI: [10.21105/joss.03849](https://doi.org/10.21105/joss.03849)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: ↗

Reviewers:

- [@lutzhamel](#)
- [@quickbeam123](#)

Submitted: 01 October 2021

Published: 03 March 2022

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

gym-saturation is an OpenAI Gym ([Brockman et al., 2016](#)) environment for reinforcement learning (RL) agents capable of proving theorems. Currently, only theorems written in a formal language of the Thousands of Problems for Theorem Provers (TPTP) library ([Sutcliffe, 2017](#)) in clausal normal form (CNF) are supported. gym-saturation implements the 'given clause' algorithm (similar to the one used in Vampire ([Kovács & Voronkov, 2013](#)) and E Prover ([Schulz et al., 2019](#))). Being written in Python, gym-saturation was inspired by PyRes ([Schulz & Pease, 2020](#)). In contrast to the monolithic architecture of a typical Automated Theorem Prover (ATP), gym-saturation gives different agents opportunities to select clauses themselves and train from their experience. Combined with a particular agent, gym-saturation can work as an ATP. Even with a non trained agent based on heuristics, gym-saturation can find refutations for 688 (of 8257) CNF problems from TPTP v7.5.0.

Statement of need

Current applications of RL to saturation-based ATPs like Enigma ([Jakubuv et al., 2020](#)) or Deepire ([Suda, 2021](#)) are similar in that the environment and the agent are not separate pieces of software but parts of larger systems that are hard to disentangle. The same is true for non saturation-based RL-friendly provers too (e.g. lazyCoP, Rawson & Reger ([2021](#))). This monolithic approach hinders free experimentation with novel machine learning (ML) models and RL algorithms and creates unnecessary complications for ML and RL experts willing to contribute to the field. In contrast, for interactive theorem provers, projects like HOList ([Bansal, Loos, Rabe, Szegedy, & Wilcox, 2019](#)) or GamePad ([Huang et al., 2019](#)) separate the concepts of environment and agent. Such modular architecture may lead to the development of easily comparable agents based on diverse approaches (see, e.g. Paliwal et al. ([2020](#)) or Bansal, Loos, Rabe, & Szegedy ([2019](#))). gym-saturation is an attempt to implement a modular environment-agent architecture of an RL-based ATP. In addition, some RL empowered saturation ATPs are not accompanied with their source code ([Abdelaziz et al., 2022](#)), while gym-saturation is open-source software.

Usage example

Suppose we want to prove an extremely simple theorem with a very basic agent. We can do that in the following way:

```
# first we create and reset a OpenAI Gym environment
from importlib.resources import files
import gym

env = gym.make(
    "gym_saturation:saturation-v0",
    # we will try to find a proof shorter than 10 steps
    step_limit=10,
```

```
# for a classical syllogism about Socrates
problem_list=[
    files("gym_saturation").joinpath(
        "resources/TPTP-mock/Problems/TST/TST003-1.p"
    )
],
)
env.reset()
# we can render the environment (that will become the beginning of the proof)
print("starting hypotheses:")
print(env.render("human"))
# our 'age' agent will always select clauses for inference
# in the order they appeared in current proof attempt
action = 0
done = False
while not done:
    observation, reward, done, info = env.step(action)
    action += 1
# SaturationEnv has an additional method
# for extracting only clauses which became parts of the proof
# (some steps were unnecessary to find the proof)
print("refutation proof:")
print(env.tstp_proof)
print(f"number of attempted steps: {action}")
```

The output of this script includes a refutation proof found:

```
starting hypotheses:
cnf(p_imp_q, hypothesis, ~man(X0) | mortal(X0)).
cnf(p, hypothesis, man(socrates)).
cnf(q, hypothesis, ~mortal(socrates)).
refutation proof:
cnf(_0, hypothesis, mortal(socrates), inference(resolution, [], [p_imp_q, p])).
cnf(_2, hypothesis, $false, inference(resolution, [], [q, _0])).
number of attempted steps: 6
```

Architecture

gym-saturation includes several sub-packages:

- parsing (happens during env.reset() in example code snippet)
- logic operations (happen during env.step(action) in the example)
- AI Gym environment implementation
- agent testing (a bit more elaborated version of the while loop from the example)

gym-saturation relies on a deduction system of four rules which is known to be refutationally complete (Brand, 1975):

$$\frac{C_1 \vee A_1, C_2 \vee \neg A_2}{\sigma(C_1 \vee C_2)}, \sigma = mgu(A_1, A_2) \quad (\text{resolution})$$

$$\frac{C_1 \vee s \approx t, C_2 \vee L[r]}{\sigma(L[t] \vee C_1 \vee C_2)}, \sigma = mgu(s, r) \quad (\text{paramodulation})$$

$$\frac{C \vee A_1 \vee A_2}{\sigma(C \vee L_1)}, \sigma = mgu(A_1, A_2) \quad (\text{factoring})$$

$$\frac{C \vee s \not\approx t}{\sigma(C)}, \sigma = mgu(s, t) \quad (\text{reflexivity resolution})$$

where C, C_1, C_2 are clauses, A_1, A_2 are atomic formulae, L is a literal, r, s, t are terms, and σ is a substitution (most general unifier). $L[t]$ is a result of substituting the term t in $L[r]$ for the term r at only one chosen position.

For parsing, we use the LARK parser (Shinan, 2021). We represent the clauses as Python classes forming tree-like structures. gym-saturation also includes a JSON serializer/deserializer for those trees. For example, a TPTP clause

```
cnf(a2,hypothesis,
    ( ~ q(a) | f(X) = X )).
```

becomes

```
Clause(
    literals=[
        Literal(
            negated=True,
            atom=Predicate(
                name="q", arguments=[Function(name="a", arguments=[])]
            ),
        ),
        Literal(
            negated=False,
            atom=Predicate(
                name "=",
                arguments=[
                    Function(name="f", arguments=[Variable(name="X")]),
                    Variable(name="X"),
                ],
            ),
        ),
    ],
    label="a2",
)
```

This grammar serves as the glue for gym-saturation sub-packages, which are, in principle, independent of each other. After switching to another parser or another deduction system, the agent testing script won't break, and RL developers won't need to modify their agents for compatibility (for them, the environment will have the same standard OpenAI Gym API).

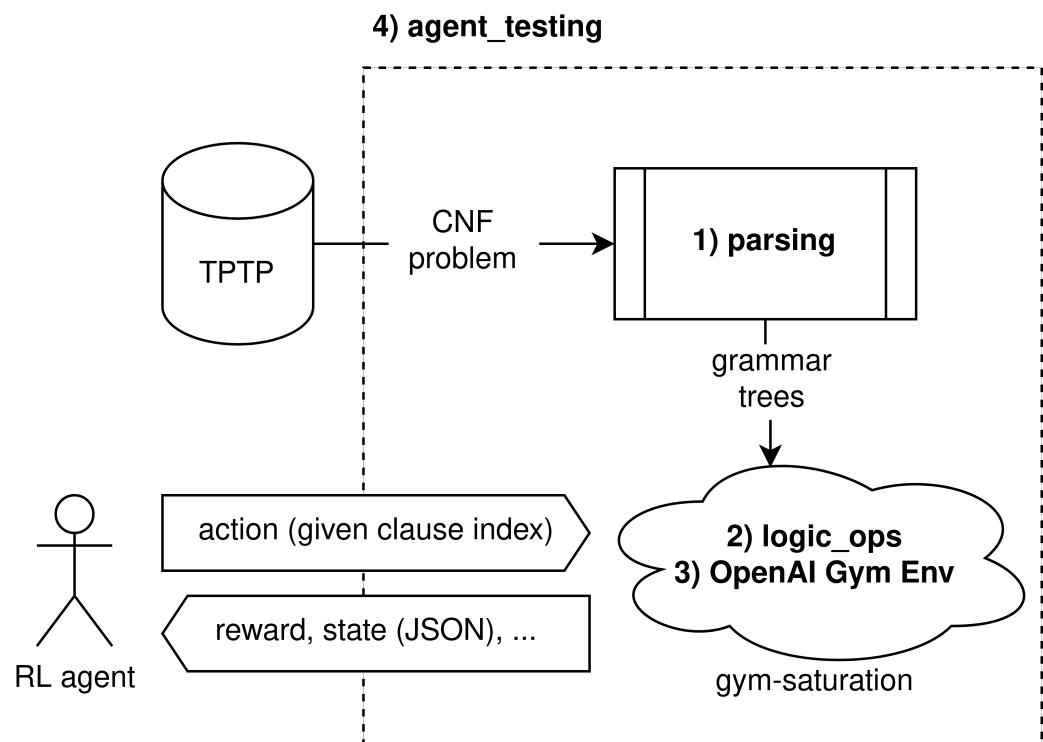


Figure 1: A diagram showing interactions between four main subpackages of gym-saturation: 1) parsing; 2) logic operations (including the given clause algorithm); 3) OpenAI Gym Env implementation; 4) the agent testing script.

Agent testing is a simple episode pipeline (see [Figure 1](#)). It is supposed to be run in parallel (e.g. using GNU Parallel, Tange (2021)) for a testing subset of problems. See the following table for the testing results of two popular heuristic-based agents on TPTP v7.5.0 (trained RL agents should strive to be more successful than those primitive baselines):

	size agent	age agent	size&age agent
proof found	509	206	688
step limit	1385	35	223
out of memory	148	149	148
5 min time out	6215	7867	7198
total	8257	8257	8257

size agent is an agent which always selects the shortest clause.

age agent is an agent which always selects the clause which arrived first to the set of unprocessed clauses ('the oldest one').

size&age agent is an agent which selects the shortest clause five times in a row and then one time — the oldest one.

'Step limit' means an agent didn't find proof after 1000 steps (the longest proof found consists of 287 steps). This can work as a 'soft timeout'.

Mentions

At the moment of writing this paper, gym-saturation was used by its author during their PhD studies for creating experimental RL-based ATPs.

Acknowledgements

This work has been supported by the French government, through the 3IA Côte d’Azur Investments in the Future project managed by the National Research Agency (ANR) with the reference number ANR-19-P3IA-0002. This work was performed using HPC resources from GENCI-IDRIS (Grant 2021-AD011013125).

References

- Abdelaziz, I., Crouse, M., Makni, B., Austel, V., Cornelio, C., Ikbal, S., Kapanipathi, P., Makondo, N., Srinivas, K., Witbrock, M., & Fokoue, A. (2022). Learning to guide a saturation-based theorem prover. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 1–1. <https://doi.org/10.1109/TPAMI.2022.3140382>
- Bansal, K., Loos, S. M., Rabe, M. N., & Szegedy, C. (2019). Learning to reason in large theories without imitation. *CoRR*, abs/1905.10501. <http://arxiv.org/abs/1905.10501>
- Bansal, K., Loos, S. M., Rabe, M. N., Szegedy, C., & Wilcox, S. (2019). HOList: An environment for machine learning of higher order logic theorem proving. In K. Chaudhuri & R. Salakhutdinov (Eds.), *Proceedings of the 36th international conference on machine learning, ICML 2019, 9-15 june 2019, long beach, california, USA* (Vol. 97, pp. 454–463). PMLR. <http://proceedings.mlr.press/v97/bansal19a.html>
- Brand, D. (1975). Proving theorems with the modification method. *SIAM Journal on Computing*, 4(4), 412–430. <https://doi.org/10.1137/0204036>
- Brockman, G., Cheung, V., Pettersson, L., Schneider, J., Schulman, J., Tang, J., & Zaremba, W. (2016). OpenAI gym. *CoRR*, abs/1606.01540. <http://arxiv.org/abs/1606.01540>
- Huang, D., Dhariwal, P., Song, D., & Sutskever, I. (2019). GamePad: A learning environment for theorem proving. *7th International Conference on Learning Representations, ICLR 2019, New Orleans, LA, USA, May 6-9, 2019*. <https://openreview.net/forum?id=r1xwKoR9Y7>
- Jakubuv, J., Chvalovský, K., Olsák, M., Piotrowski, B., Suda, M., & Urban, J. (2020). ENIGMA anonymous: Symbol-independent inference guiding machine (system description). In N. Peltier & V. Sofronie-Stokkermans (Eds.), *Automated reasoning - 10th international joint conference, IJCAR 2020, paris, france, july 1-4, 2020, proceedings, part II* (Vol. 12167, pp. 448–463). Springer. https://doi.org/10.1007/978-3-030-51054-1/_29
- Kovács, L., & Voronkov, A. (2013). First-order theorem proving and vampire. In N. Sharygina & H. Veith (Eds.), *Computer aided verification - 25th international conference, CAV 2013, saint petersburg, russia, july 13-19, 2013. proceedings* (Vol. 8044, pp. 1–35). Springer. https://doi.org/10.1007/978-3-642-39799-8/_1
- Paliwal, A., Loos, S. M., Rabe, M. N., Bansal, K., & Szegedy, C. (2020). Graph representations for higher-order logic and theorem proving. *The Thirty-Fourth AAAI Conference on Artificial Intelligence, AAAI 2020, the Thirty-Second Innovative Applications of Artificial Intelligence Conference, IAAI 2020, the Tenth AAAI Symposium on Educational Advances in Artificial Intelligence, EAAI 2020, New York, NY, USA, February 7-12, 2020*, 2967–2974. <https://aaai.org/ojs/index.php/AAAI/article/view/5689>
- Rawson, M., & Reger, G. (2021). lazyCoP: Lazy paramodulation meets neurally guided search. In A. Das & S. Negri (Eds.), *Automated reasoning with analytic tableaux and related*

- methods - 30th international conference, TABLEAUX 2021, birmingham, UK, september 6-9, 2021, proceedings* (Vol. 12842, pp. 187–199). Springer. https://doi.org/10.1007/978-3-030-86059-2/_11
- Schulz, S., Cruanes, S., & Vukmirovic, P. (2019). Faster, higher, stronger: E 2.3. In P. Fontaine (Ed.), *Automated deduction - CADE 27 - 27th international conference on automated deduction, natal, brazil, august 27-30, 2019, proceedings* (Vol. 11716, pp. 495–507). Springer. https://doi.org/10.1007/978-3-030-29436-6/_29
- Schulz, S., & Pease, A. (2020). Teaching automated theorem proving by example: PyRes 1.2 - (system description). In N. Peltier & V. Sofronie-Stokkermans (Eds.), *Automated reasoning - 10th international joint conference, IJCAR 2020, paris, france, july 1-4, 2020, proceedings, part II* (Vol. 12167, pp. 158–166). Springer. https://doi.org/10.1007/978-3-030-51054-1/_9
- Shinan, E. (2021). *Lark-parser* (Version 0.12.0). <https://pypi.org/project/lark-parser/>
- Suda, M. (2021). Improving ENIGMA-style clause selection while learning from history. In A. Platzer & G. Sutcliffe (Eds.), *Automated deduction - CADE 28 - 28th international conference on automated deduction, virtual event, july 12-15, 2021, proceedings* (Vol. 12699, pp. 543–561). Springer. https://doi.org/10.1007/978-3-030-79876-5/_31
- Sutcliffe, G. (2017). The TPTP Problem Library and Associated Infrastructure. From CNF to TH0, TPTP v6.4.0. *Journal of Automated Reasoning*, 59(4), 483–502.
- Tange, O. (2021). *GNU parallel 20210822 ('kabul')*. Zenodo. <https://doi.org/10.5281/zenodo.5233953>