

MetaWards: A flexible metapopulation framework for modelling disease spread

Christopher Woods^{*1}, Lester Hedges¹, Christopher Edsall², Ellen Brooks-Pollock³, Christopher Parton-Fenton⁴, Trevelyan J McKinley⁵, Matt J Keeling⁶, and Leon Danon³

¹ Research Software Engineering, Advanced Computing Research Centre, University of Bristol, UK ² Research Software Engineering, Research Computing Services, University of Cambridge, UK ³ Department of Engineering Mathematics, University of Bristol, UK ⁴ College of Engineering, Mathematics and Physical Sciences, University of Exeter, UK ⁵ College of Medicine and Health, University of Exeter, UK ⁶ Zeeman Institute (SBIDER), University of Warwick, Coventry, UK

DOI: [10.21105/joss.03914](https://doi.org/10.21105/joss.03914)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Mark A. Jensen](#) ↗

Reviewers:

- [@klmedeiros](#)
- [@acolum](#)

Submitted: 20 October 2021

Published: 11 February 2022

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Understanding how disease spreads through populations is important when designing and implementing control measures. [MetaWards](#) implements a stochastic metapopulation model of disease transmission that enables geographical modelling of disease spread that can scale all the way from modelling local transmission up to full national- or international-scale outbreaks. It is built in Python and has a flexible plugin architecture to support complex scenario modelling. This enables the code to be adapted to model new situations and new control measures as they arise, e.g. emergence of new variants of disease, enactment of different types of movement restrictions, availability of different types of vaccines etc. It implements a user-definable compartmental transmission model, such as an SIR model, that can be extended multi-dimensionally via multiple demographics or sub-populations, and multiple geographical regions. Models can be constructed from the various sources of movement and demographic data that are available, and are accelerated via Cython ([Behnel et al., 2020](#)), OpenMP, Scoop ([Hold & Gagnon, 2019](#)) and MPI4Py ([Dalcin & Fang, 2021](#)) to scale efficiently from running on personal laptops to large supercomputers. Python, R and command line interfaces and a complete set of tutorials empower researchers to adapt their models to a variety of scenarios.

Statement of need

MetaWards was originally developed as a C program to support modelling of disease transmission in Great Britain ([Danon et al., 2009](#); [Keeling et al., 2010](#)). The original C code has been published on GitHub separately ([Danon et al., 2020](#)) and is included in the main MetaWards repository for reference. In a metapopulation model, individuals are divided into different geographical regions, e.g. into electoral wards as in the original code. Disease transmission between regions occurs via the daily movement of individuals.

This model was recently adapted to model COVID-19 transmission in England and Wales ([Danon et al., 2021](#)). Upon adaption it became clear that the original C code needed to be made more robust, more trustable, faster, and easier to extend to model new scenarios. This Python version was created to meet all of those needs. This Python version of MetaWards can identically reproduce the outputs from the original C code. It has been optimised and parallelised, with extensive unit testing added to build trust and robustness. Use of Cython

^{*}corresponding author

(Behnel et al., 2020) enabled the ported code to run as quickly as the original C code, and then parallelisation via OpenMP provided a significant speed-up on multi-core machines. Re-architecting of the code, including creating a new multi-network (demographic) framework, enabled extension of the traditional compartmental transmission model across multiple dimensions; horizontally by adding disease stages, vertically by adding multiple networks to represent different demographics or sub-populations, and geographically via movement of modelled individuals to different custom locations, wards or regions. Finally, a flexible plugin architecture together with a tutorial-driven development style has created a flexible modelling framework with a rich set of tutorials and documentation. These demonstrate how MetaWards can be adapted easily via a Python, R or command-line interface to model a variety of complex scenarios. MetaWards is being used for on-going modelling of COVID-19 transmission, e.g. via the model (McKinley, 2021) developed by the UQ4Covid project (Williamson et al., 2020).

Tutorial-Driven Development

A tutorial-driven development process was used for MetaWards. This process was inspired by test-driven development, where the tests are written first. For tutorial-driven development, we aimed to write the tutorials first.

Development started with a discussion of the new scenarios that needed modelling. This was followed with tutorials that described how MetaWards would be used to model those scenarios. Finally, the code (and unit tests) were written so that MetaWards behaved as described in the tutorials. The project team took the view that if functionality was not described in a tutorial, then it didn't exist, as users would not know how to make use of it. The tutorials (Woods, 2021) evolved in lock-step with the code. Development was rapid, with new versions released regularly during 2020. To manage this change, MetaWards rigidly used the semantic versioning (SemVer) scheme (Preston-Werner, 2021) and provided strong guarantees of backwards compatibility. The website was versioned, so that tutorials for older versions of the software were easily accessible via a drop-down selector. Today, this versioning can be used to retroactively view how the tutorials developed alongside the code.

Features

MetaWards is designed to enable researchers to model how an infection may spread geographically (Figure 1), and what impact different control measures or individual behaviours may make.

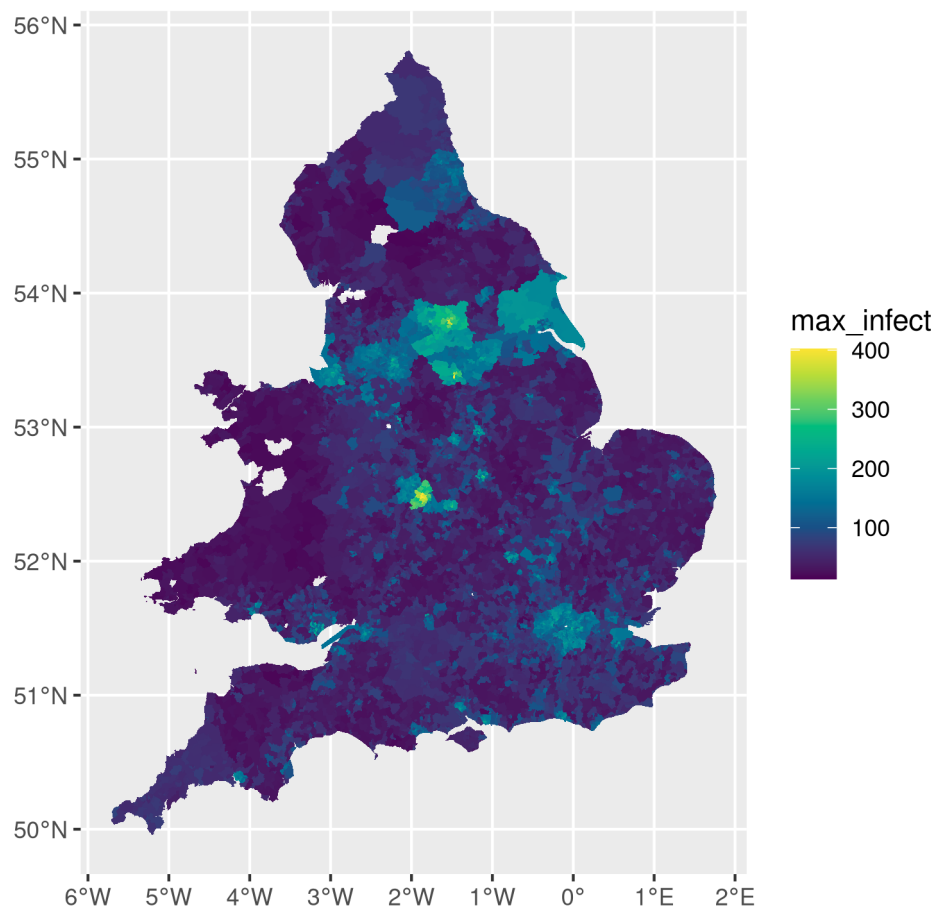


Figure 1: Analysis of a simulation run (Siegert, 2020) that used MetaWards to chart disease spread across England and Wales.

To this end, MetaWards features;

- a generic compartmental transmission model, with user-definable compartments and transmission algorithms. Supports both traditional (e.g. SIR, SEIR, SI, SIS) and custom models,
- a flexible plugin architecture that makes it easy to implement new control measures. For example, the tutorial shows how this can be used to model shielding, different lockdown scenarios and to investigate necessary durations of quarantine or self-isolation,
- multi-network demographic support. Multiple networks can be run as a single combined group, with custom plugins used to merge data between networks, and conditionally move individuals between different demographics. Tutorials demonstrate how this can be used to model, for example, shielding and self-isolation, hospital admissions, impact of individuals returning from holidays,
- per-ward custom parameter support. Different wards can have different parameters, meaning that local behaviour can be easily modelled (e.g. local lockdowns, changes in local control measures),
- complete-detail and full control over horizontal and vertical movements through disease stages or across demographics. Tutorials show how this can be used this to model vaccination and also to model waning immunity, with individuals returned from the R or V stages back to S,

- flexible data output support - again handled using an array of in-built or user-supplied data extraction plugin functions. Data can be output in whatever format is needed for analysis, either to text files or to databases,
- full reproducibility support baked in throughout. The code records enough data to make reproduction easy, with results designed to be the same given the same inputs, random number seed and number of threads,
- flexible input files that would enable modelling of any region or country to be undertaken, based on the various sources of movement and demographic data that are available. Models of England and Wales, and the UK have been created, and a Python and R API are provided to make it easy to create custom networks. These support everything from individual wards or local geographies, up to full national- or international-scale metapopulation models,
- support for scanning design files for optimisation or sensitivity analysis of nearly all input parameters, plus any user custom parameters used in the main code or any plugins. These scans can use as much compute as available, parallelising individual runs over multiple cores, and scaling multiple runs up to full supercomputers (Figure 2).

```

Preparing to run
Running 4 jobs using 4 process(es)
Running models in parallel using multiprocessing
Computing model run ✓

Completed job 1 of 4
(NO_CHANGE)[repeat 1]
2022-10-07: DAY: 352 S: 12190461 E: 0 I: 0 R: 47809539 IW: 0 TOTAL POPULATION 60000000

Computing model run ✓

Completed job 2 of 4
(NO_CHANGE)[repeat 2]
2022-10-29: DAY: 374 S: 12176223 E: 0 I: 0 R: 47823777 IW: 0 TOTAL POPULATION 60000000

Computing model run ✓

Completed job 3 of 4
(NO_CHANGE)[repeat 3]
2022-10-03: DAY: 348 S: 12199094 E: 0 I: 0 R: 47800906 IW: 0 TOTAL POPULATION 60000000

Computing model run ✓

Completed job 4 of 4
(NO_CHANGE)[repeat 4]
2022-10-11: DAY: 356 S: 12193262 E: 0 I: 0 R: 47806738 IW: 0 TOTAL POPULATION 60000000

```

Figure 2: Multiple scans performed in parallel, with live summary updates as jobs finish.

MetaWards is optimised and can run happily on small laptops. Individual national-scale networks fit in approximately 80 MB of memory, and model runs can take 15-90 seconds to perform. Compute and memory costs scale with the number of demographics that are added, but high performance and low memory consumption are design goals. Models using only a few wards are kilobytes, and take less than a second to run. Finally, a built-in `metawards-plot` tool is included in the package for simple visualisation of key results (Figure 3).

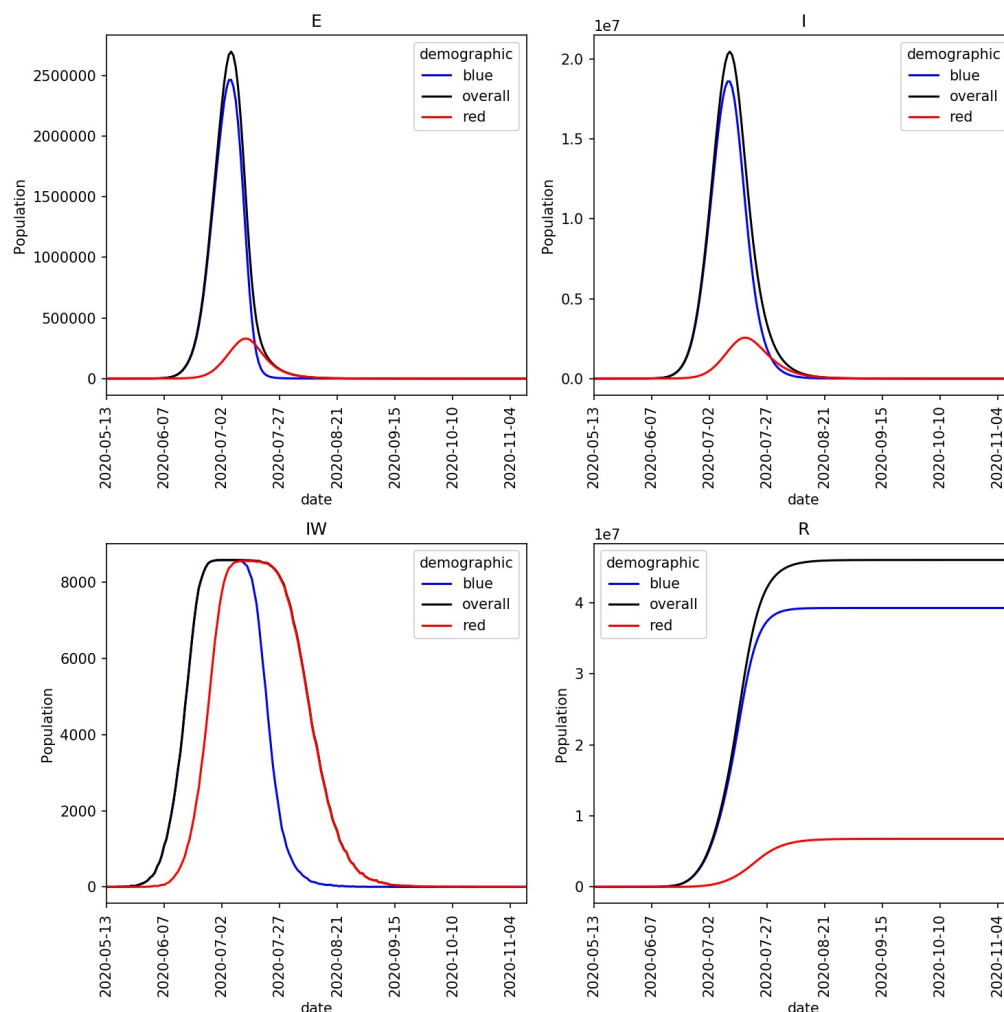


Figure 3: metawards-plot visualisation of results, including across multiple networks.

Acknowledgements

CW acknowledges funding via an EPSRC Research Software Engineering Fellowship (EP/N018591/1). TJM is supported by an “Expanding Excellence in England” award from Research England and UKRI grants: EP/V051555/1 (UQ4Covid) and MR/V038613/ (JUNIPER Consortium). LD, EBP and MJK are supported by MRC (grant number MC/PC/19067) and UKRI through the JUNIPER consortium (grant number MR/V038613/1). LD is further supported by EPSRC (EP/V051555/1 and The Alan Turing Institute, grant EP/N510129/1), and Pfizer through investigator-led grants on respiratory tract infections.

References

- Behnel, Stefan., Bradshaw, Robert., Dalcín, Lisandro., Florisson, Mark., Makarov, Vitja., & Seljebotn, D. Sverre. (2020). *Cython: C-extensions for python*. <https://cython.org>
- Dalcin, L., & Fang, Y.-L. L. (2021). mpi4py: Status Update After 12 Years of Development. *Computing in Science & Engineering*, 23(4), 47–54. <https://doi.org/10.1109/mcse.2021.3914>

3083216

- Danon, Leon., Brooks-Pollock, Ellen., Bailey, Mick., & Keeling, Matt. (2021). A spatial model of COVID-19 transmission in england and wales: Early spread, peak timing and the impact of seasonality. *Philisophical Transactions of the Royal Society B*, 376. <https://doi.org/10.1098/rstb.2020.0272>
- Danon, Leon., Brooks-Pollock, Ellen., & Keeling, M. J. (2020). MetaWards. In *GitHub repository*. GitHub. <https://github.com/ldanon/MetaWards>
- Danon, Leon., House, Thomas., & Keeling, Matt. J. (2009). The role of routine versus random movements on the spread of disease in great britain. *Epidemics*, 1(4), 250–258. <https://doi.org/10.1016/j.epidem.2009.11.002>
- Hold, Yannick., & Gagnon, Oliver. (2019). SCOOP (scalable concurrent operations in python). In *GitHub repository*. GitHub. <https://scoop.readthedocs.io>
- Keeling, M. J., Danon, Leon., Vernon, M. C., & House, T. A. (2010). Individual identity and movement networks for disease metapopulations. *Proceedings of the National Academy of Sciences*, 107(19), 8866–8870. <https://doi.org/10.1073/pnas.1000416107>
- McKinley, T. J. (2021). *UQ4Covid model*. GitHub. <https://uq4covid.github.io/vignettes/data/MetaWards/vignette/ModelDescription>
- Preston-Werner, Tom. (2021). *Semantic versioning 2.0.0*. GitHub. <https://semver.org>
- Siebert, Stefan. (2020). *MetaWards plotting tutorial*. GitHub. https://uq4covid.github.io/vignettes/metawards_plot
- Williamson, Danny., Danon, Leon., McKinley, T. J., Mcneall, Doug., Youngman, Ben., Siebert, Stefan., Challenor, Peter., Astfalck, Lachlan., Baker, Evan., Parton-Fenton, Christopher., Salter, James., & Xiang, Xiaoyu. (2020). *Uncertainty quantification for expensive COVID-19 simulation models (UQ4Covid)*. GitHub. <https://uq4covid.github.io>
- Woods, Christopher. (2021). *MetaWards tutorials*. GitHub. <https://metawards.org/tutorial>