

adaptr: an R package for simulating and comparing adaptive clinical trials

Anders Granholm¹, Aksel Karl Georg Jensen², Theis Lange², and Benjamin Skov Kaas-Hansen^{1,2}

¹ Department of Intensive Care, Copenhagen University Hospital - Rigshospitalet, Copenhagen, Denmark ² Section of Biostatistics, Department of Public Health, University of Copenhagen, Copenhagen, Denmark

DOI: [10.21105/joss.04284](https://doi.org/10.21105/joss.04284)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: Fabian Scheipl ↗

Reviewers:

- [@elimillera](#)
- [@stmcg](#)

Submitted: 16 March 2022

Published: 27 April 2022

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

The adaptr R package facilitates simulation and comparison of adaptive clinical trial designs using Bayesian statistical methods. The package supports a flexible number of arms, use of a common control arm, pre-specified and user-defined outcome- and posterior probability distribution-generating functions, fixed- and response-adaptive randomisation (RAR), various adaptation rules for arm dropping and stopping, calculation of trial design performance metrics, and visualisation of results. Simulations are relatively fast, can run in parallel, and can be reloaded from previous sessions.

Statement of need

Adaptive clinical trials are more flexible and can be more efficient than conventional randomised clinical trials (RCTs) with fixed sample sizes and allocation ratios ([Pallmann et al., 2018](#)). Because more efficient trials may be preferable for economic, logistic and ethical reasons, the interest in adaptive clinical trials is growing ([Granholm et al., 2021](#)). Planning adaptive trials and comparing adaptive trials designs are complex, however, and require statistical simulation because simple, closed-form sample size calculations are infeasible. Existing software is either closed source/commercial, has limited features or flexibility, or is based on graphical user interfaces that may hamper reproducibility (unlike scripting-based tools) ([Meyer et al., 2021](#)).

Overview and features

The adaptr R package ([R Core Team, 2021](#)) focuses on multi-arm adaptive clinical trials using adaptive stopping, arm dropping and/or response-adaptive randomisation ([Pallmann et al., 2018](#); [Viele, Broglio, et al., 2020](#)). Many realistic trial designs are supported by the built-in functions and more advanced trial designs can be simulated with user-defined functions.

Table 1 gives an overview of the functionality of adaptr:

Table 1: Overview of functionality in the adaptr package.

Category	Description
General	Allows for two or more intervention arms and comparison of all arms or pairwise comparisons against a common control arms. Simulations use Bayesian statistics.

Category	Description
Outcomes and posteriors	Supports multiple outcome types specified by user-defined functions, with built-in convenience functions for easy use with binary, binomially distributed and continuous, normally distributed outcomes. For more flexibility, user-defined functions for generating outcomes and sampling from posterior distributions are supported; these may use more complex functions provided by other packages (e.g., Markov chain Monte Carlo or variational inference).
Random allocation	Supports fixed and response-adaptive randomisation, and combinations thereof, including special handling of control group allocation in trials with a common control (e.g., matching highest non-control group allocation ratio and square-root-ratio-based allocation (Park et al., 2020; Viele, Broglio, et al., 2020)). Supports minimum/maximum allocation limits and <i>softening</i> non-fixed allocation probabilities by raising them to a power (Ryan et al., 2020) (which may vary throughout the trial) and normalising probabilities to sum to 100%.
Adaptation rules	Supports probabilistic trial-stopping/arm-dropping rules for superiority, inferiority, and practical equivalence between all arms in trial designs without a common control. Supports probabilistic trial-stopping/arm-dropping rules for superiority, inferiority, practical equivalence, and futility compared to a common control in such designs. In trial designs with a common control, a superior arm becomes the new control. In these cases equivalence/futility testing may be continued (using the new control) or disabled.
Performance metrics	Supports the calculation of various trial design performance metrics (Viele, Broglio, et al., 2020; Viele, Saville, et al., 2020) according to different arm selection strategies in trials not ending with superiority.
Visualisation	Offers graphical summaries of the overall or arm-specific statuses over time across multiple simulations (such as still <i>recruiting</i> or <i>trial stopped/arm dropped</i> for various reasons), and relevant arm-specific metrics (e.g., allocation probabilities) over time across single/multiple trial simulations.
Technical	Supports parallel execution on multiple cores using R's built-in <code>parallel</code> package. Allows reproducible results with seeding. Supports saving/loading large trial objects and <i>growing</i> previous simulation objects. Uses base R with no required external dependencies; optional plotting functionality does, however, require the <code>ggplot2</code> -package (Wickham et al., 2019).

Adaptive analyses are conducted at specified recruitment numbers, after random allocation of simulated patients to currently active arms. Both fixed and response-adaptive randomisation (RAR) as well as combinations (including several restrictions of RAR) are supported (Pallmann et al., 2018; Park et al., 2020; Ryan et al., 2020; The Adaptive Platform Trials Coalition, 2019). Outcomes are then generated, followed by drawing samples from the posterior distributions; the packages comes with built-in functions for generating outcomes and posteriors using fast models with conjugate, flat priors (Ryan et al., 2019), but more advanced estimation are allowed using user-defined functions.

Posterior draws are used to calculate probabilities for (i) enforcing adaptive trial-stopping/arm-

dropping rules with respect to superiority, inferiority, equivalence and/or futility, (ii) assessing the trial's final status (superiority, inferiority, equivalence, futility, or stopped at a maximum pre-specified sample size), and (iii) adjusting allocation probabilities if RAR is used. Single or multiple trial simulations can be run, saved and grown, and summarised using multiple relevant performance metrics (Viele, Broglio, et al., 2020; Viele, Saville, et al., 2020). Performance metrics can be calculated under different assumptions about the preferable intervention in clinical practice should the trials not end with superiority.

Functions and workflow

An overview of the principal functions in `adaptr` is given in **Table 2**:

Table 2: Overview of user-facing functions in `adaptr`.

Function name	Action
setup_trial , setup_trial_binom , setup_trial_norm	Defines the adaptive trial specification including arms, true outcome rates, stopping rules, outcome/posterior generating functions, and summary settings. Full flexibility is provided by setup_trial ; the convenience functions setup_trial_binom and setup_trial_norm provide sensible defaults for binary, binomially distributed and continuous, normally distributed outcomes, respectively.
run_trial	Runs a single simulation using a trial specification defined with setup_trial .
run_trials	Runs multiple trial simulations using a trial specification defined with setup_trial . Can save simulation objects to external files and load/grow previous simulations.
extract_results	Extracts key results and (partial) performance metrics from multiple simulations conducted by run_trials according to a specified arm selection strategy for simulations not ending in superiority. Returns a tidy data.frame with one simulation per row (Wickham et al., 2019).
summary	Summarises, in a human-friendly way, key results and calculates overall performance metrics for multiple simulations according to a specified arm selection strategy for simulations not ending in superiority.
plot_status	Plots the overall statuses or arm-specific statuses for multiple simulations over the course of the simulated trials.
plot_history	Plots the history of relevant metrics in each arm for one or multiple simulations over the course of the simulated trial(s).
print	Print methods for printing the outputs of most of the included functions in a human-friendly way.

`adaptr` is available on [CRAN](#) and [GitHub](#), where a [stand-alone website](#) with the package documentation is also available. `adaptr` can be installed using one of the following commands:

```
install.packages("adaptr") # CRAN

# install.packages("remotes")
remotes::install_github("INCEPTdk/adaptr") # GitHub
```

Once installed, as any other R package, it is loaded with:

```
library(adaptr)
#> Loading adaptr package (version 1.0.0).
```

```
#> See 'help("adaptr")' or 'vignette("Overview", "adaptr")' for help.
#> Further information available on https://github.com/INCEPTdk/adaptr/.
```

An adaptive trial design is specified with the generic `setup_trial` function (requires user-defined functions to generate/compare outcomes), or one of the helper functions `setup_trial_binom` or `setup_trial_norm` (wrappers around `setup_trial` with appropriate functions for generating outcomes and sampling from posterior distributions). For example, a simple four-arm trial with a binary, binomially distributed outcome can be specified with:

```
binom_trial <- setup_trial_binom(
  # Treatment arms
  arms = c("Control", "Experimental A", "Experimental B", "Experimental C"),

  # True event rates
  true_ys = c(0.25, 0.30, 0.22, 0.19),

  # Time of adaptive analyses
  # - first analysis when 400 patients are included, then after every 100
  data_looks = seq(from = 400, to = 2000, by = 100),

  # Name of the common control (leave undefined if no common control is desired)
  control = "Control",

  # Use square-root-ratio-based allocation (sqrt(number of non-control arms):1
  # for each non-control arm), with fixed control group allocation and RAR in
  # the non-control arms
  control_prob_fixed = "sqrt-based",

  # Define stopping rules
  superiority = 0.99, # Superiority probability threshold
  inferiority = 0.01, # Inferiority probability threshold
  equivalence_prob = 0.85, # Equivalence probability threshold
  equivalence_diff = 0.05, # Equivalence difference
  equivalence_only_first = TRUE, # Only assess equivalence for first control

  # Restrict non-fixed allocation ratios (limit extreme RAR)
  soften_power = 0.5
)

print(binom_trial)
#> Trial specification: generic binomially distributed outcome trial
#> * Undesirable outcome
#> * Common control arm: Control
#> * Control arm probability fixed at 0.366 (for 4 arms), 0.414 (for 3 arms),
#> 0.5 (for 2 arms)
#> * Best arm: Experimental C
#>
#> Arms, true outcomes, starting allocation probabilities
#> and allocation probability limits:
#>
#>      arms true_ys start_probs fixed_probs min_probs max_probs
#>      Control   0.25     0.366     0.366      NA      NA
#> Experimental A   0.30     0.211      NA      NA      NA
#> Experimental B   0.22     0.211      NA      NA      NA
#> Experimental C   0.19     0.211      NA      NA      NA
#>
#> Maximum sample size: 2000
```

```
#> Maximum number of data looks: 17
#> Planned data looks after: 400, 500, 600, 700, 800, 900, 1000, 1100,
#> 1200, 1300, 1400, 1500, 1600, 1700, 1800, 1900, 2000 patients
#>
#> Superiority threshold: 0.99
#> Inferiority threshold: 0.01
#> Equivalence threshold: 0.85 (only checked for first control)
#> Absolute equivalence difference: 0.05
#> No futility threshold
#> Soften power for all analyses: 0.5
```

Once specified, simulating a single trial is easy:

```
run_trial(binom_trial, seed = 202203) # Fixed random seed
#> Single simulation result: generic binomially distributed outcome trial
#> * Undesirable outcome
#> * Initial/final common control arms: Control/Experimental C
#>
#> Final status: inconclusive, stopped at maximum sample size
#> Final/maximum allowed sample sizes: 2000/2000 (100.0%)
#>
#> Final trial results:
#>      arms true_ys sum_ys ns raw_ests post_ests post_errs lo_cri hi_cri
#>      Control    0.25   42 145   0.290    0.293   0.0375  0.223  0.366
#> Experimental A    0.30   27  87   0.310    0.314   0.0492  0.223  0.415
#> Experimental B    0.22  192 888   0.216    0.217   0.0138  0.190  0.245
#> Experimental C    0.19  165 880   0.188    0.188   0.0133  0.163  0.215
#> final_status status_look status_probs final_alloc
#>      inferior      400      0.0032      0.366
#>      inferior      400      0.0036      0.211
#>      active        NA        NA      0.500
#>      control        NA        NA      0.500
#>
#> Simulation details:
#> * Random seed: 202203
#> * Credible interval width: 95%
#> * Number of posterior draws: 5000
#> * Posterior estimation method: medians with MAD-SDs
```

—so is running multiple trials, for example 1000. If `cores > 1`, simulations are runs in parallel on the number of cores specified. Setting `sparse = FALSE` keeps all intermediate simulation results; this is required for `plot_history` (see below) but makes the object a lot larger (in this case, by a factor 15).

```
sims <- run_trials(trial_spec = binom_trial, n_rep = 1000, cores = 4,
                  base_seed = 202204, sparse = FALSE)
```

Numerical results

We summarise the results, assuming the control arm would be used in clinical practice in case of inconclusive trials (unless dropped early):

```
summary(sims, select_strategy = "control")
#> Multiple simulation results: generic binomially distributed outcome trial
#> * Undesirable outcome
#> * Number of simulations: 1000
#> * Number of simulations summarised: 1000 (all trials)
```

```
#> * Common control arm: Control
#> * Selection strategy: first control if available (otherwise no selection)
#> * Treatment effect compared to: no comparison
#>
#> Performance metrics (using posterior estimates):
#> * Sample sizes: mean 1708.3 (SD: 438.0) | median 2000.0 (IQR: 1500.0 to
#> 2000.0)
#> * Total summarised outcomes: mean 382.5 (SD: 96.4) | median 422.0 (IQR:
#> 338.8 to 449.0)
#> * Total summarised outcome rates: mean 0.225 (SD: 0.014) | median 0.225
#> (IQR: 0.216 to 0.234)
#> * Conclusive: 45.2%
#> * Superiority: 39.0%
#> * Equivalence: 6.2%
#> * Futility: 0.0% [not assessed]
#> * Inconclusive at max sample size: 54.8%
#> * Selection probabilities: Control: 24.0% | Experimental A: 0.0% |
#> Experimental B: 1.3% | Experimental C: 33.1% | None: 41.6%
#> * RMSE: 0.02742
#> * RMSE treatment effect: 0.04741
#> * Ideal design percentage: 77.0%
#>
#> Simulation details:
#> * Simulation time: 19.9 secs
#> * Base random seed: 202204
#> * Credible interval width: 95%
#> * Number of posterior draws: 5000
#> * Estimation method: posterior medians with MAD-SDs
```

Alternatively, we can get *tidy* results (Wickham et al., 2019) for each simulation using the same selection strategy as above:

```
extr_res <- extract_results(sims, select_strategy = "control")

head(extr_res) # Print results for first six simulations
#>   sim final_n sum_ys ratio_ys final_status superior_arm selected_arm
#> 1     1    2000   436 0.2180000          max          <NA>          <NA>
#> 2     2    2000   452 0.2260000          max          <NA>          <NA>
#> 3     3    1800   389 0.2161111 superiority Experimental C Experimental C
#> 4     4    2000   440 0.2200000          max          <NA>          <NA>
#> 5     5     500   118 0.2360000 superiority Experimental B Experimental B
#> 6     6     900   182 0.2022222 superiority Experimental C Experimental C
#>
#>   sq_err sq_err_te
#> 1      NA      NA
#> 2      NA      NA
#> 3 0.0004297484 0.004184831
#> 4      NA      NA
#> 5 0.0078199784 0.015558136
#> 6 0.0019097953 0.001558565
```

Visual summaries

We can plot the overall trial statuses at each adaptive analysis using:

```
plot_status(sims)
```




Figure 1: Status plot

The overall allocation probabilities in each arm (y axis) against the total number of patients included in the trial (x axis) are plotted using:

```
# Requires sparse = FALSE in run_trials
plot_history(sims, x_value = "total n", y_value = "prob")
```

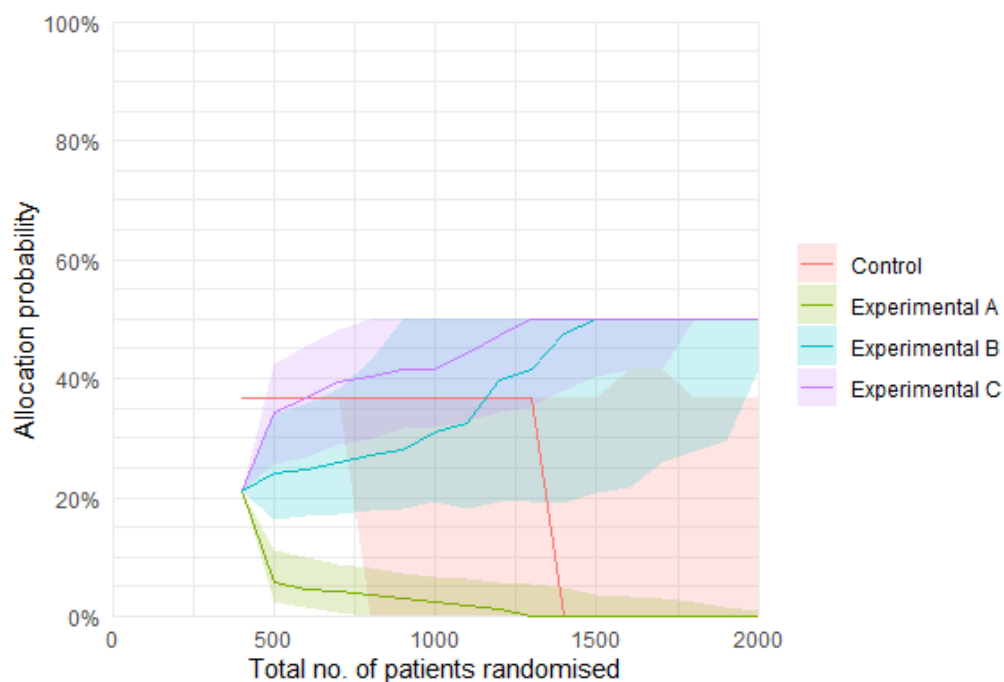


Figure 2: Allocation probabilities history plot

Arm-specific statuses and additional metrics over time may be plotted as well.

Discussion

We have developed a flexible, extensible and comprehensive **R** package that allows relatively easy and efficient simulation of adaptive clinical trials using various features with few, well-documented user-facing functions.

The package facilitates planning of adaptive clinical trials and comparison of trial designs using combinations of the most important and common adaptive features. In addition to the efficient convenience functions provided for common outcome types, more complex models are supported but may increase simulation time substantially and require (highly) parallel execution on multiple cores locally and/or remotely using cloud computing.

Compared to other open-source solutions ([Meyer et al., 2021](#)), `adaptr` comprises a completely functional and well-documented **R** package with support for multiple complex randomisation strategies, different stopping rules with detailed control, full support for user-defined outcome-/posterior-generating functions, parallel computation, easy calculation of various trial design performance metrics according to different arm selection strategies for trials not ending in superiority, and visualisation of relevant metrics across single or multiple simulations. While relatively feature-rich, the package currently has some limitations. These primarily include the lack of a ‘flooring’ option (stopping allocation to arms when their probabilities of being the best fall below a certain threshold, possibly with subsequent resumption of allocation ([Viele, Broglio, et al., 2020](#))), inability to add arms during the conduct (as done in adaptive platform trials ([The Adaptive Platform Trials Coalition, 2019](#)); this is complex and not supported in most other software ([Meyer et al., 2021](#))) and the lack of separate stopping rules/allocation ratios in subgroups or enrichment (which is similarly complex ([The Adaptive Platform Trials Coalition, 2019](#))). Finally, the package uses Bayesian statistical methods as these are well-suited for adaptive trials and easily extended; frequentist statistical methods are not supported.

In conclusion, the `adaptr` **R** package provides a feature-rich and extensible open-source, scripting-based solution for planning, simulating and comparing adaptive clinical trials.

Acknowledgement

`adaptr` has been developed as part of the [Intensive Care Platform Trial](#) programme, which has primarily been funded by a grant from [Sygeforsikringen “danmark”](#).

References

- Granholm, A., Alhazzani, W., Derde, L., Angus, D., Zampieri, F., Hammond, N., Sweeney, R., Myatra, S., Azoulay, E., Rowan, K., Young, P., Perner, A., & Møller, M. (2021). Randomised clinical trials in critical care: Past, present and future. *Intensive Care Medicine*. <https://doi.org/10.1007/s00134-021-06587-9>
- Meyer, E., Mesenbrink, P., Mielke, T., Parke, T., Evans, D., König, F., & on behalf of EU-PEARL (EU Patient-cEntric clinicAl tRial pLatforms) Consortium. (2021). Systematic review of available software for multi-arm multi-stage and platform clinical trial design. *Trials*, 22, 183. <https://doi.org/10.1186/s13063-021-05130-x>
- Pallmann, P., Bedding, A., Choodari-Oskooei, B., Dimairo, M., Flight, L., Hampson, L., Holmes, J., Mander, A., Odoni, L., Sydes, M., Villar, S., Wason, J., Weir, C., Wheeler, G., Yap, C., & Jaki, T. (2018). Adaptive designs in clinical trials: Why use them, and how to run and report them. *BMC Medicine*, 16, 29. <https://doi.org/10.1186/s12916-018-1017-7>

- Park, J., Harari, O., Dron, L., Lester, R., Thorlund, K., & Mills, E. (2020). An overview of platform trials with a checklist for clinical readers. *Journal of Clinical Epidemiology*, 125, 1–8. <https://doi.org/10.1016/j.jclinepi.2020.04.025>
- R Core Team. (2021). *R: A language and environment for statistical computing*. R Foundation for Statistical Computing. <https://www.R-project.org/>
- Ryan, E., Harrison, E., Pearse, R., & Gates, S. (2019). Perioperative haemodynamic therapy for major gastrointestinal surgery: The effect of a bayesian approach to interpreting the findings of a randomised controlled trial. *BMJ Open*, 9, e024256. <https://doi.org/10.1136/bmjopen-2018-024256>
- Ryan, E., Lamb, S., Williamson, E., & Gates, S. (2020). Bayesian adaptive designs for multi-arm trials: An orthopaedic case study. *Trials*, 21, 83. <https://doi.org/10.1186/s13063-019-4021-0>
- The Adaptive Platform Trials Coalition. (2019). Adaptive platform trials: Definition, design, conduct and reporting considerations. *Nature Reviews Drug Discovery*, 18, 797–807. <https://doi.org/10.1038/s41573-019-0034-3>
- Viele, K., Broglio, K., McGlothlin, A., & Saville, B. (2020). Comparison of methods for control allocation in multiple arm studies using response adaptive randomization. *Clinical Trials*, 17, 52–60. <https://doi.org/10.1177/1740774519877836>
- Viele, K., Saville, B., McGlothlin, A., & Broglio, K. (2020). Comparison of response adaptive randomization features in multiarm clinical trials with control. *Pharmaceutical Statistics*, 19, 602–612. <https://doi.org/10.1002/pst.2015>
- Wickham, H., Averick, M., Bryan, J., Chang, W., D'Agostino, L., Francois, R., Grolemond, G., Hayes, A., Henry, L., Hester, J., Kuhn, M., Pedersen, T., Miller, E., Bache, S., Müller, K., Ooms, J., Robinson, D., Seidel, D., Spinu, V., ... Yutani, H. (2019). Welcome to the tidyverse. *Journal of Open Source Software*, 4, 1686. <https://doi.org/10.21105/joss.01686>