












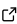

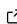
1 GLUE Code: A framework handling communication 2 and interfaces between scales

3 Aleksandra Pachalieva ^{1,2}, Robert S. Pavel ³✉, Javier E. Santos ^{1,2},
4 Abdourahmane Diaw ⁴, Nicholas Lubbers ³, Mohamed Mehana ²,
5 Jeffrey R. Haack ³, Hari S. Viswanathan ², Daniel Livescu ³, Timothy
6 C. Germann ⁵, and Christoph Junghans ³

7 1 Center for Non-Linear Studies, Los Alamos National Laboratory, Los Alamos, 87545 NM, USA 2 Earth
8 and Environmental Sciences (EES) Division, Los Alamos National Laboratory, Los Alamos, 87545 NM,
9 USA 3 Computer, Computational and Statistical Sciences (CCS) Division, Los Alamos National
10 Laboratory, Los Alamos, 87545 NM, USA 4 Fusion Energy Division, Oak Ridge National Laboratory, 1
11 Bethel Valley Road, Oak Ridge, TN 37831, USA 5 Theoretical Division, Los Alamos National Laboratory,
12 Los Alamos, 87545 NM, USA ✉ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Daniel S. Katz](#) 

Reviewers:

- [@govarguz](#)
- [@keipertk](#)

Submitted: 20 September 2022

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#))

13 Summary

14 Many scientific applications are inherently multiscale in nature. Such complex physical phenom-
15 ena often require simultaneous execution and coordination of simulations spanning multiple
16 time and length scales. This is possible by combining expensive small-scale simulations (such
17 as molecular dynamics simulations) with larger scale simulations (such continuum limit/hydro
18 solvers) to allow for considerably larger systems using task and data parallelism. However, the
19 granularity of the tasks can be very large and often leads to load imbalance. Traditionally, we
20 use approximations to streamline the computation of the more costly interactions and this
21 introduces trade-offs between simulation cost and accuracy. In recent years, the available
22 computational power and the advances in machine learning have made computing these
23 scale-bridging interactions and multiscale simulations more feasible.

24 One driving application has been plasma modeling in inertial confinement fusion (ICF) which
25 is fundamentally multiscale in nature. This requires deep understanding of how to extrapolate
26 microscopic information into macroscopically relevant scales. For example, in ICF one needs an
27 accurate understanding of the connection between experimental observables and the underlying
28 microphysics. The properties of the larger scales are often affected by the microscale behavior
29 incorporated usually into the equations of state and ionic and electronic transport coefficients
30 (Liboff, 1959; Rinderknecht et al., 2014; Rosenberg et al., 2015; Ross et al., 2017). Instead of
31 incorporating this information using reliable molecular dynamics (MD) simulations, one often
32 needs to use theoretical models, due to the inability of MD to reach engineering scales (Glosli
33 et al., 2007; Marinak et al., 1998). One approach to resolve this issue is by coupling two MD
34 simulations of different scales via force interpolation, e.g. the AdResS method (Krekeler et al.,
35 2018; Nagarajan et al., 2013). Another approach, which we will pursue in the scope of this
36 work, is by enabling scale bridging between MD simulations and meso/macro-scale models.

37 State of the art

38 Traditionally, multiscale simulations combine multiple simulation methods that need to be
39 simultaneously executed and coordinated. To achieve this, we use asynchronous task-based
40 runtime systems that include load balancers. Such load balancers can schedule and migrate
41 tasks to maintain throughput; however, the issue of fault tolerance remains (Cappello, 2009).

42 To avoid this, checkpointing is often used (Koo & Toueg, 1987), but it could be prohibitively
43 expensive when its output frequency is high. Other works such as the mystic framework
44 (Michael McKerns & Aivazis, 2019) is more geared towards large-scale machine learning
45 techniques (McKerns et al., 2011). Such codes are not designed toward coupling multiscale
46 simulations but have been demonstrated to be highly effective at solving hard optimization
47 problems.

48 In this work, we propose the Generic Learning User Enablement (GLUE) Code to facilitate the
49 coupling between scales. The GLUE Code builds upon previous work on multiscale coupling
50 (Haack et al., 2021; Karra et al., 2022; Lubbers et al., 2020; R. Pavel et al., 2017; R. S.
51 Pavel et al., 2015). At its simplest, we determine what physical properties must be exchanged
52 between the various scales and derive application programming interfaces (APIs) from these.

53 To take advantage of modern supercomputing and exascale platforms and we couple our
54 framework directly with high-performance-computing-friendly job schedulers. This allows us
55 to use one framework to generate training data with a minimal footprint (task scheduler
56 can fire off all jobs) and perform “hero-runs” of multiscale simulations. At every point, we
57 perform a fine grain simulation or utilize an upscaled result from the active learning model.
58 The GLUE Code uses active learning to evaluate the level of uncertainty and executes fine grain
59 simulations when more training data is required.

60 Statement of need

61 The GLUE Code is a modular framework designed to couple different scientific applications to
62 support use cases, including multiscale methods and various forms of machine learning. The
63 workflow was initially designed around supporting active learning as an alternative to coupled
64 applications to support multiscale methods but has been written in a sufficiently modular way
65 that different machine learning methods can be utilized. Application programming interfaces
66 (API) are available for C, C++, Fortran, and Python with support for various storage formats.
67 The code also provides direct coupling with high performance computing job schedulers such
68 as SLURM (Yoo et al., 2003) and Flux (Ahn et al., 2018). An overview of the GLUE Code
69 implementation is shown in Fig. 1, where the meso/macro-scale solvers BGK, CFDNS, and
70 LBM are also defined.

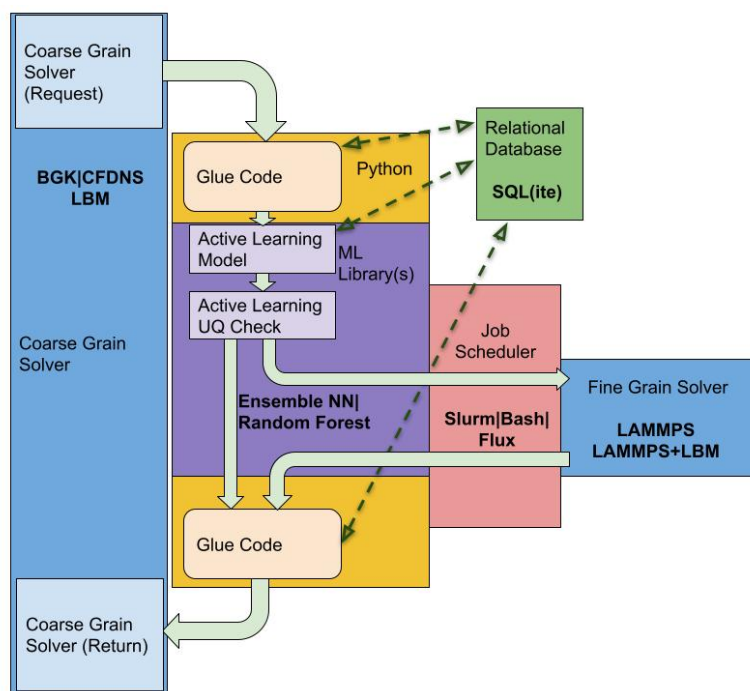


Figure 1

71 Figure 1: Sample of the GLUE Code implementation of our microscale-macroscopic coupling.
 72 On the macroscale simulation (left) sending a request, the GLUE Code (center) uses the active
 73 learning algorithms (center purple) to determine if the model's uncertainty quantification is
 74 such that a new fine grain simulation (right) needs to be called. Then either the result of the
 75 fine grain simulation or the model's prediction is returned to the macroscale simulation. This
 76 figure is adapted from Fig. 2 in Diaw et al. (Diaw et al., 2020).

77 The aim of the GLUE Code is to efficiently handle the communication and interfaces between
 78 the computing platform, surrogate model, coarse-scale code, and fine-scale code (Diaw et al.,
 79 2020). At its core, the GLUE Code determines what physical properties need to be exchanged
 80 between the scales of interest and spawns APIs using this information. The GLUE Code ensures
 81 that these physical properties are communicated between the various scientific codes needed
 82 for the given multiscale problem. The GLUE Code uses active learning algorithms to evaluate
 83 if the model has a high level of uncertainty, and thus, a fine grain simulation needs to be
 84 executed. At the end of each cycle, the framework returns either the result of the fine grain
 85 simulation or the model's prediction of the macroscopic simulation.

86 Characteristics of the GLUE code:

- 87 ■ The GLUE Code has a modular lightweight structure, where each component can be
 88 swapped out for a different implementation. This is possible due to the APIs structure
 89 of the framework.
- 90 ■ The modularity of the GLUE code allows us to replace a scientific solver with a machine
 91 learning solution. The scale bridging structures are explicitly defined, which allows us to
 92 switch between a fine grain MD simulation and an ML solution. If a neural network is
 93 capable of reproducing the same outputs as an MD simulation from the same inputs,
 94 then they would be functionally identical.
- 95 ■ As shown in Fig. 1, the main components of the GLUE Code are a Coarse Grain Solver,
 96 an active learning backend, a relational database, a job scheduler, and the coupling and
 97 coordination logic.

- 98
- 99
- 100
- 101
- 102
- 103
- 104
- 105
- 106
- 107
- 108
- 109
- 110
- 111
- 112
- 113
- 114
- 115
- 116
- 117
- 118
- 119
- 120
- The GLUE Code is written with a collection of commodity software as the backend to most of these components. Currently, the relational database is supported via SQL(ite), while the job scheduler has support for SLURM (Yoo et al., 2003), Lawrence Livermore National Laboratory's Flux scheduler (Ahn et al., 2018), and a rudimentary serialization model for debugging. Active learning is provided through PyTorch (Paszke et al., 2019) and Scikit-Learn (Pedregosa et al., 2011).
 - This approach relies on SQL(ite) for the communication as it is a guaranteed atomic read and write that simplifies a lot of our efforts at the cost of performance.
 - For coarse grain simulations, we currently couple with the Multi-BGK (Los Alamos National Laboratory, n.d.) as well as additional ICF codes. Preliminary studies have been made of coupling with more general Lattice Boltzmann simulations for other fields of study.
 - For fine scale simulations, we use a mixture of LAMMPS (Plimpton, 1995) and proven analytic solutions; however, one could easily switch to another MD solution as long as they have similar capabilities.
 - Since the effectiveness of machine learning methods varies depending on the available training data and problem, an additional check for any machine learning solution is added to the framework. In the GLUE Code, we query the models provided by the machine learning solution for the output associated with our inputs. The surrogate model then provides both the expected outputs and the uncertainty quantification to indicate its confidence that this specific model gave a valid answer. If the confidence is too low, the code falls back to calling the actual MD simulation and providing the data for the machine learning solution to retrain and generate a new surrogate model for future use.

121 Code structure

122 The GLUE code is organized as follows:

- 123
- 124
- 125
- 126
- 127
- 128
- 129
- 130
- 131
- 132
- 133
- 134
- 135
- 136
- 137
- 138
- GLUECode_Library contains the C++ library that is meant to be linked to the coarse grain solver. This allows existing applications to couple to the GLUECode_Service with minimal code alterations.
 - GLUECode_Service contains the python scripts that the library communicates with and uses a combination of active learning and spawning of fine grain simulation jobs to enable and accelerate multiscale scientific applications.
 - docs contains documentation for the Flux scheduler.
 - examples contains sniff tests for the serial and the MPI versions of the code, and utility scripts, such as pullICFData.py that demonstrate how the GLUE Code framework can be used to access training data to perform deeper analysis.
 - jsonFiles contains example input decks for the GLUE Code service.
 - lammepsFiles contains LAMMPS example scripts (in.Argon_Deuterium_masses and in.Argon_Deuterium_pLasma) that compute the mutual diffusion for a plasma composed of argon and deuterium at different concentrations and temperatures. More information about the test cases can be found within the files.
 - training contains training data required for active learning.

139 Acknowledgements

140 The research presented in this article was supported by the Laboratory Directed Research
141 and Development program of Los Alamos National Laboratory (LANL) under project number

- 142 20190005DR and used resources provided by the LANL Institutional Computing Program.
143 We acknowledge the support of the U.S. Department of Energy through the LANL/LDRD
144 Program for this work. AP and JS also acknowledge the Center for Non-Linear Studies at
145 LANL. LANL is operated by Triad National Security, LLC, for the National Nuclear Security
146 Administration of U.S. Department of Energy (Contract No. 89233218CNA000001). Assigned
147 LA-UR-22-29746.
- 148 Ahn, D. H., Bass, N., Chu, A., Garlick, J., Grondona, M., Herbein, S., Koning, J., Patki, T.,
149 Scogland, T. R., & Springmeyer, B. (2018). Workflows in support of large-scale science
150 (WORKS'18). *Proceedings of the 2018 ACM/IEEE Workflows in Support of Large-Scale
151 Science (WORKS'18)*, 10–19.
- 152 Cappello, F. (2009). Fault tolerance in petascale/exascale systems: Current knowledge,
153 challenges and research opportunities. *The International Journal of High Performance
154 Computing Applications*, 23(3), 212–226. <https://doi.org/10.1177/1094342009106189>
- 155 Diaw, A., Barros, K., Haack, J., Junghans, C., Keenan, B., Li, Y., Livescu, D., Lubbers, N.,
156 McKerns, M., Pavel, R., & others. (2020). Multiscale simulation of plasma flows using
157 active learning. *Physical Review E*, 102(2), 023310. [https://doi.org/10.1103/physreve.
158 102.023310](https://doi.org/10.1103/physreve.102.023310)
- 159 Glosli, J. N., Richards, D. F., Caspersen, K. J., Rudd, R. E., Gunnels, J. A., & Streitz, F. H.
160 (2007). Extending stability beyond CPU millennium: A micron-scale atomistic simulation of
161 Kelvin-Helmholtz instability. *SC'07: Proceedings of the 2007 ACM/IEEE Conference on
162 Supercomputing*, 1–11. <https://doi.org/10.1145/1362622.1362700>
- 163 Haack, J., Diaw, A., Pavel, R., Sagert, I., Keenan, B., Livescu, D., Lubbers, N., McKerns, M.,
164 Junghans, C., & Germann, T. (2021). Enabling predictive scale-bridging simulations through
165 active learning. *APS Division of Plasma Physics Meeting Abstracts, 2021*, ZO03–012.
- 166 Karra, S., Mehana, M., Lubbers, N., Chen, Y., Diaw, A., Santos, J. E., Pachaliev, A., Pavel,
167 R. S., Haack, J. R., McKerns, M., & others. (2022). Predictive scale-bridging simulations
168 through active learning. *arXiv Preprint arXiv:2209.09811*. [https://doi.org/10.48550/
169 ARXIV.2209.09811](https://doi.org/10.48550/ARXIV.2209.09811)
- 170 Koo, R., & Toueg, S. (1987). Checkpointing and rollback-recovery for distributed systems.
171 *IEEE Transactions on Software Engineering*, 1, 23–31.
- 172 Krekeler, C., Agarwal, A., Junghans, C., Praprotnik, M., & Delle Site, L. (2018). Adaptive
173 resolution molecular dynamics technique: Down to the essential. *The Journal of Chemical
174 Physics*, 149(2), 024104. <https://doi.org/10.1063/1.5031206>
- 175 Liboff, R. L. (1959). Transport coefficients determined using the shielded Coulomb potential.
176 *The Physics of Fluids*, 2(1), 40–46. <https://doi.org/10.1063/1.1724389>
- 177 Los Alamos National Laboratory. (n.d.). LANL/Multi-BGK: Conservative multispecies kinetic
178 equation solver. In *GitHub*. <https://github.com/lanl/Multi-BGK>
- 179 Lubbers, N., Agarwal, A., Chen, Y., Son, S., Mehana, M., Kang, Q., Karra, S., Junghans,
180 C., Germann, T. C., & Viswanathan, H. S. (2020). Modeling and scale-bridging using
181 machine learning: Nanoconfinement effects in porous media. *Scientific Reports*, 10(1),
182 1–13. <https://doi.org/10.1038/s41598-020-69661-0>
- 183 Marinak, M., Haan, S., Dittrich, T., Tipton, R., & Zimmerman, G. (1998). A comparison of
184 three-dimensional multimode hydrodynamic instability growth on various National Ignition
185 Facility capsule designs with HYDRA simulations. *Physics of Plasmas*, 5(4), 1125–1132.
186 <https://doi.org/10.1063/1.872643>
- 187 McKerns, M. M., Strand, L., Sullivan, T., Fang, A., & Aivazis, M. A. (2011). Building a
188 framework for predictive science. *Proceedings of the 10th Python in Science Conference*,
189 *arXiv Preprint arXiv:1202.1056*. <https://doi.org/10.48550/arXiv.1202.1056>

- 190 Michael McKerns, P. H., & Aivazis, M. (2019). *Mystic: Highly-constrained non-convex*
191 *optimization and UQ*. <https://uqfoundation.github.io/project/mystic>
- 192 Nagarajan, A., Junghans, C., & Matysiak, S. (2013). Multiscale simulation of liquid water using
193 a four-to-one mapping for coarse-graining. *Journal of Chemical Theory and Computation*,
194 *9*(11), 5168–5175. <https://doi.org/10.1021/ct400566j>
- 195 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z.,
196 Gimelshein, N., Antiga, L., Desmaison, A., Kopf, A., Yang, E., DeVito, Z., Raison, M.,
197 Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., ... Chintala, S. (2019). *PyTorch: An*
198 *imperative style, high-performance deep learning library*. 8024–8035. [http://papers.neurips.](http://papers.neurips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf)
199 [cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf](http://papers.neurips.cc/paper/9015-pytorch-an-imperative-style-high-performance-deep-learning-library.pdf)
- 200 Pavel, R. S., McPherson, A. L., Germann, T. C., & Junghans, C. (2015). Database assisted
201 distribution to improve fault tolerance for multiphysics applications. *Proceedings of the 2nd*
202 *International Workshop on Hardware-Software Co-Design for High Performance Computing*,
203 1–8. <https://doi.org/10.1145/2834899.2834908>
- 204 Pavel, R., Junghans, C., Mniszewski, S. M., & Germann, T. C. (2017). *Using charm++ to*
205 *support multiscale multiphysics*.
- 206 Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M.,
207 Prettenhofer, P., Weiss, R., Dubourg, V., & others. (2011). Scikit-learn: Machine learning
208 in Python. *The Journal of Machine Learning Research*, *12*, 2825–2830.
- 209 Plimpton, S. (1995). Fast parallel algorithms for short-range molecular dynamics. *Journal of*
210 *Computational Physics*, *117*(1), 1–19. <https://doi.org/10.2172/10176421>
- 211 Rinderknecht, H., Sio, H., Li, C., Zylstra, A., Rosenberg, M., Amendt, P., Delettrez, J., Bellei,
212 C., Frenje, J., Johnson, M. G., & others. (2014). First observations of nonhydrodynamic
213 mix at the fuel-shell interface in shock-driven inertial confinement implosions. *Physical*
214 *Review Letters*, *112*(13), 135001. <https://doi.org/10.1103/physrevlett.112.135001>
- 215 Rosenberg, M., Séguin, F., Amendt, P., Atzeni, S., Rinderknecht, H., Hoffman, N., Zylstra, A.,
216 Li, C., Sio, H., Gatu Johnson, M., & others. (2015). Assessment of ion kinetic effects in
217 shock-driven inertial confinement fusion implosions using fusion burn imaging. *Physics of*
218 *Plasmas*, *22*(6), 062702. <https://doi.org/10.1063/1.4921935>
- 219 Ross, J., Higginson, D., Ryutov, D., Fiuza, F., Hatarik, R., Huntington, C., Kalantar, D., Link,
220 A., Pollock, B., Remington, B., & others. (2017). Transition from collisional to collisionless
221 regimes in interpenetrating plasma flows on the national ignition facility. *Physical Review*
222 *Letters*, *118*(18), 185003. <https://doi.org/10.1103/PhysRevLett.118.185003>
- 223 Yoo, A. B., Jette, M. A., & Grondona, M. (2003). Slurm: Simple Linux utility for resource
224 management. *Workshop on Job Scheduling Strategies for Parallel Processing*, 44–60.
225 https://doi.org/10.1007/10968987_3