

FuseMedML: a framework for accelerated discovery in machine learning based biomedicine

Alex Golts^{1*}¶, Moshe Raboh^{1*}, Yoel Shoshan^{1*}, Sagi Polaczek¹, Simona Rabinovici-Cohen¹, and Efrat Hexter¹

¹ IBM Research - Haifa, Israel ¶ Corresponding author * These authors contributed equally.

DOI: [10.21105/joss.04943](https://doi.org/10.21105/joss.04943)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: [Jacob Schreiber](#) ↗ 

Reviewers:

- [@anupamajha1](#)
- [@suragnair](#)

Submitted: 16 November 2022

Published: 19 January 2023

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Machine Learning is at the forefront of scientific progress in Healthcare and Medicine. To accelerate scientific discovery, it is important to have tools that allow progress iterations to be collaborative, reproducible, reusable and easily built upon without “reinventing the wheel” for each task.

FuseMedML, or *fuse*, is a Python framework designed for accelerated Machine Learning (ML) based discovery in the medical domain. It is highly flexible and designed for easy collaboration, encouraging code reuse. Flexibility is enabled by a generic data object design where data is kept in a nested (hierarchical) Python dictionary (NDict), allowing to efficiently process and fuse information from multiple modalities. Functional components allow to specify input and output keys, to be read from and written to the nested dictionary.

Easy code reuse is enabled through key components implemented as standalone packages under the main *fuse* repo using the same design principles. These include *fuse.data* - a flexible data processing pipeline, *fuse.dl* - reusable Deep Learning (DL) model architecture components and loss functions, and *fuse.eval* - a library for evaluating ML models.

Statement of need

Medical research often involves multiple modalities (e.g., imaging, clinical data, biochemical representations) and tasks (e.g., classification, segmentation, clinical condition prediction). In our experience working on numerous such projects, we have identified three key challenges: 1. Setting up or implementing a new baseline model can be time-consuming, even when similar projects have already been completed by the same lab. 2. Transferring individual components across projects can be difficult, leading to researchers frequently “reinventing the wheel.” 3. Collaborating between projects across modalities and domains, such as imaging and molecules, is often challenging.

To address these challenges, FuseMedML was developed with the goal of simplifying and streamlining medical research projects.

Before open sourcing it, we used *fuse* internally in multiple research projects ([Raboh, Levanony, et al., 2022](#)), ([Rabinovici-Cohen, Tlusty, et al., 2022](#)), ([Rabinovici-Cohen, Fernández, et al., 2022](#)), ([Jubran et al., 2021](#)), ([Tlusty et al., 2021](#)), ([Golts et al., 2022](#)), ([Barros et al.](#)) and experienced significant improvement in development time, reusability and collaboration. We were also able to meaningfully measure our progress and statistical significance of our results with off-the-shelf *fuse.eval* components that facilitate metrics' confidence interval calculation and model comparison. These tools have enabled us to organize two challenges as part of the 2022 International Symposium on Biomedical Imaging (ISBI) ([Raboh, Golts, et al., 2022](#)), ([Pati et al., 2022](#)).

State of the field

FuseMedML is a comprehensive machine learning library that focuses on the biomedical domain. It offers a range of tools covering the entire development process, including data preparation, model training, and evaluation. Built on top of popular machine learning frameworks such as PyTorch (Paszke et al., 2019) and PyTorch Lightning (Falcon & The PyTorch Lightning team, 2019), FuseMedML also includes flexible domain-specific capabilities to complement these frameworks. Overall, FuseMedML aims to facilitate machine learning discoveries within the healthcare and life science sectors. One way in which *fuse* can complement PyTorch is through its generic design concept (See Figure 1) of storing arbitrary types of data in a specialized nested dictionary. This is a key driver of flexibility, allowing minimal code modifications when moving building blocks between different projects. Concretely, *fuse* has a dataset class that extends the PyTorch dataset, and a model wrapper class that enables PyTorch models to operate on batch_dicts rather than tensors.

In the case of PyTorch Lightning, *fuse* integrates with it directly as it builds upon its comprehensive trainer class, also allowing users to define their models and data modules in PyTorch Lightning style, with flexible levels of customizability.




Figure 1: This figure illustrates FuseMedML's design concept. A *fuse* component is instantiated with input and output keys. These keys refer to the *sample_dict*, the basic data sample structure of *fuse* represented by a special nested Python dictionary called "NDict".

There are existing PyTorch-based ML libraries that similarly to *fuse* cater to researchers in the biomedical domain. Two examples of such prominent libraries are MONAI (Cardoso et al., 2022) and PyHealth (Zhao et al., 2021). MONAI is primarily focused on medical imaging applications. PyHealth on the other hand mainly focuses on health records data. *fuse* is designed to support different types of medical data and multimodal use cases involving imaging, clinical and biochemical data.

As with generic ML frameworks like PyTorch and PyTorch Lightning, *fuse* can also coexist with the more specific libraries like MONAI, PyHealth or others. A user may opt to borrow parts from different libraries and complement them with components from *fuse*. As another example, a user may want to use the data *ops* of *fuse* which are generic and flexible, or its data caching mechanism, which allows to separate processing into a static and dynamic pipelines, controlling the desired stages to be cached.

Packages

fuse.data

FuseMedML's data package is designed for building a flexible and powerful data pipeline with reusable building blocks called *ops*. See Figure 2 for a simple example for how such a building block can be used across different projects.

Each *op* class's `__call__` function gets as an input a *sample_dict*, a dictionary that stores all the necessary information about a sample processed so far. Typically, an *op*'s constructor gets keys that specify what it should consider in *sample_dict* and where to store the output. Similarly, a minibatch is represented by a *batch_dict*.

A special kind of *ops* are "Meta *ops*". They can be thought of as a form of wrapper *op* around

a regular, lower level *op* or function, to help achieve a special behavior such as repeating that low level *op*, applying it with random values and more. “Meta *ops*” also help avoid writing boilerplate code.

A data pipeline may consist of a static_pipeline and a dynamic_pipeline. The output of the static_pipeline can be cached to optimize running time and GPU utilization. The dynamic_pipeline is responsible for “online” processing that we don’t want to cache, such as random augmentations. An instance of a *fuse* dataset class, which inherits from the PyTorch dataset class is then created from defined static and dynamic pipelines.

The data package also includes generic utilities such as a PyTorch based sampler enabling batch class balancing and a tool for splitting data into folds according to predefined criteria.

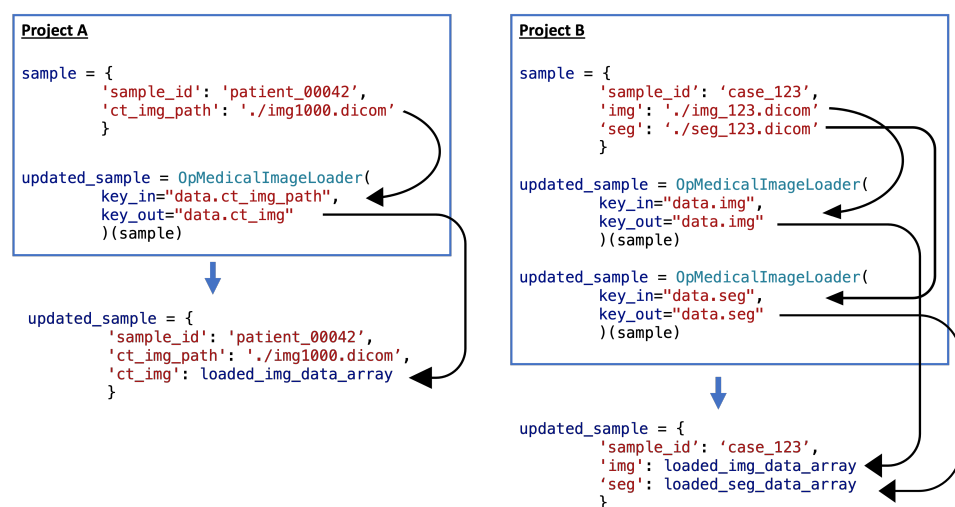


Figure 2: In this example a medical image loader is the *fuse* component reused in projects A and B. Different projects can have different formats for their data samples, but they can all use *OpMedicalImageLoader* by providing the appropriate key names when calling it. In Project B the same key name is used for the input and output, resulting in the loaded image data overriding the image paths in the updated sample.

fuse.dl

FuseMedML’s DL package works with PyTorch models, only modifying them to interact with a batch_dict. For training, *fuse.dl* utilizes PyTorch-Lightning, either through an already made *LightningModuleDefault* class that inherits from Pytorch-Lightning’s *LightningModule* class, or by allowing users who seek maximal customizability to implement their own custom *LightningModule* and operate in close resemblance to the standard PyTorch-Lightning workflow or use alternative training loop implementations.

fuse.dl also offers generic core DL components such as model architectures and losses, implemented in *fuse* style. See an example model architecture definition in [Figure 3](#).

```
ModelMultiHead(
    conv_inputs=('data.input.img',),
    backbone=BackboneResnet3D(in_channels=1),
    heads=[
        Head3D(head_name='classification',
                mode="classification",
                conv_inputs=[("model.backbone_features", 512)]
                ),
    ]
)
```

Figure 3: In this example a model architecture is defined using the `ModelMultiHead` class. It contains of a 3D ResNet backbone represented by the `BackboneResnet3D` class and a 3D classification head represented by the `Head3D` class. Note the user can define a list of heads, to support a multi task use case. The inputs to the backbone and classification heads are defined in the *fuse* style described earlier, using the `batch_dict` key names with the relevant data. This enables easy reuse of similar model architectures between projects.

fuse.eval

FuseMedML's evaluation package is a standalone library for evaluating machine learning models using various performance metrics and comparing the results between models. It offers advanced capabilities such as a generic confidence interval wrapper for any metric, a generic one-versus-all extension for converting any binary metric to a multi-class scenario, and metrics for comparing models while considering statistical significance. The package also includes model calibration tools and a pipeline for combining a sequence of metrics with possible dependencies. In addition, the evaluation package supports automatic per-fold evaluation and subgroup analysis, and can handle large data sets through batching and multiprocessing. See [Figure 4](#) for an example of an evaluation metric pipeline that can be reused across projects.

```
data = {"pred": [0.1, 0.2, 0.6, 0.7, 0.8, ...],
        "target": [0, 0, 1, 1, 1, ...] }

metrics = OrderedDict(
    [("auc", CI(MetricAUROC(pred="pred", target="target", stratum="target"))),
     ("auc_pr", CI(MetricAUCPR(pred="pred", target="target", stratum="target")))]
)

evaluator = EvaluatorDefault()
results = evaluator.eval(ids=None, data=data, metrics=metrics)
```



```
results = {
    'metrics': {
        'auc': {'org': 0.84, 'mean': 0.85, 'std': 0.06, 'conf_interval': 95,
                'conf_lower': 0.71, 'conf_upper': 0.95},
        'auc_pr': {'org': 0.80, 'mean': 0.81, 'str': 0.05, 'conf_interval': 95,
                   'conf_lower': 0.68, 'conf_upper': 0.91}
    }
}
```

Figure 4: In this example a pipeline of evaluation metric components is shown. It consists of two metrics: the Area Under the receiver operating characteristic Curve and the Area Under the Precision-Recall Curve. Both metrics are wrapped with a Confidence Interval (CI) metric, resulting in a lower and upper bound for each metric. The metrics are executed by an instance of the `EvaluatorDefault` class, the basic *fuse.eval* class that combines input sources, evaluates using the specified metrics, generates a report and returns a dictionary with all the metrics results.

Extensions

The core technology of FuseMedML and its component packages is general, while domain-specific functionality is contained within extensions. One such extension, *fuse-imaging*, is currently available and extends the FuseMedML data package with operations useful for medical imaging, as well as implementations of public medical datasets.

References

- Barros, V., Tlustý, T., Barkan, E., Hexter, E., Gruen, D., Guindy, M., & Rosen-Zvi, M. Virtual biopsy by using artificial intelligence-based multimodal modeling of binational mammography data. *Radiology*, 0(0), 220027. <https://doi.org/10.1148/radiol.220027>
- Cardoso, M. J., Li, W., Brown, R., Ma, N., Kerfoot, E., Wang, Y., Murray, B., Myronenko, A., Zhao, C., Yang, D., Nath, V., He, Y., Xu, Z., Hatamizadeh, A., Myronenko, A., Zhu, W., Liu, Y., Zheng, M., Tang, Y., ... Feng, A. (2022). *MONAI: An open-source framework for deep learning in healthcare*. <https://doi.org/10.48550/arXiv.2211.02701>
- Falcon, W., & The PyTorch Lightning team. (2019). *PyTorch Lightning* (Version 1.4) [Computer software]. <https://doi.org/10.5281/zenodo.3828935>
- Golts, A., Khapun, D., Shats, D., Shoshan, Y., & Gilboa-Solomon, F. (2022). An ensemble of 3D u-net based models for segmentation of kidney and masses in CT scans. In N. Heller, F. Isensee, D. Trofimova, R. Tejpaul, N. Papanikolopoulos, & C. Weight (Eds.), *Kidney and kidney tumor segmentation* (pp. 103–115). Springer International Publishing. https://doi.org/10.1007/978-3-030-98385-7_14
- Jubran, I., Raboh, M., Perek, S., Gruen, D., & Hexter, E. (2021). A glimpse into the future: Disease progression simulation for breast cancer in mammograms. In D. Svoboda, N. Burgos, J. M. Wolterink, & C. Zhao (Eds.), *Simulation and synthesis in medical imaging* (pp. 34–43). Springer International Publishing. https://doi.org/10.1007/978-3-030-87592-3_4
- Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., Desmaison, A., Kopf, A., Yang, E., DeVito, Z., Raison, M., Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., ... Chintala, S. (2019). PyTorch: An imperative style, high-performance deep learning library. In H. Wallach, H. Larochelle, A. Beygelzimer, F. dAlché-Buc, E. Fox, & R. Garnett (Eds.), *Advances in neural information processing systems* (Vol. 32). Curran Associates, Inc. <https://proceedings.neurips.cc/paper/2019/file/bdbca288fee7f92f2bfa9f7012727740-Paper.pdf>
- Pati, P., Jaume, G., Brancati, N., Müller, H., Riccio, D., De Pietro, G., Foncubierta, A., Raboh, M., Anniciello, A. M., Scognamiglio, G., Feroce, F., Frucci, M., & Gabrani, M. (2022). *BRIGHT Challenge - BReast tumor Image classification on Gigapixel Histopathological images*. <https://research.ibm.com/haifa/Workshops/BRIGHT/>.
- Rabinovici-Cohen, S., Fernández, X. M., Grandal Rejo, B., Hexter, E., Hijano Cubelos, O., Pajula, J., Pölönen, H., Rey, F., & Rosen-Zvi, M. (2022). Multimodal prediction of five-year breast cancer recurrence in women who receive neoadjuvant chemotherapy. *Cancers*, 14(16). <https://doi.org/10.3390/cancers14163848>
- Rabinovici-Cohen, S., Tlustý, T., Fernández, X. M., & Rejo, B. G. (2022). Early prediction of metastasis in women with locally advanced breast cancer. In K. Drukker, K. M. Iftekharuddin, H. Lu, M. A. Mazurowski, C. Muramatsu, & R. K. Samala (Eds.), *Medical imaging 2022: Computer-aided diagnosis* (Vol. 12033, p. 120330F). International Society for Optics; Photonics; SPIE. <https://doi.org/10.1117/12.2613169>
- Raboh, M., Golts, A., Heller, N., Tejpaul, R., Abdallah, N., Benidir, T., Campbell, S. C., Remer, E., Foncubierta, A., Gabrani, M., Müller, H., Hexter, E., Rabinovici-Cohen, S.,

- Shoshan, Y., Weight, C., & Rosen-Zvi, M. (2022). *KNIGHT Challenge - Kidney clinical Notes and Imaging to Guide and Help personalize Treatment and biomarkers discovery*. <https://research.ibm.com/haifa/Workshops/KNIGHT/>.
- Raboh, M., Levanony, D., Dufort, P., & Sitek, A. (2022). Context in medical imaging: the case of focal liver lesion classification. In O. Colliot & I. Išgum (Eds.), *Medical imaging 2022: Image processing* (Vol. 12032, p. 1203200). International Society for Optics; Photonics; SPIE. <https://doi.org/10.1117/12.2609385>
- TLusty, T., Ozery-Flato, M., Barros, V., Barkan, E., Amit, M., Gruen, D., Guindy, M., Arazi, T., Rozin, M., Rosen-Zvi, M., & Hexter, E. (2021). Pre-biopsy multi-class classification of breast lesion pathology in mammograms. In C. Lian, X. Cao, I. Rekik, X. Xu, & P. Yan (Eds.), *Machine learning in medical imaging* (pp. 277–286). Springer International Publishing. https://doi.org/10.1007/978-3-030-87589-3_29
- Zhao, Y., Qiao, Z., Xiao, C., Glass, L., & Sun, J. (2021). PyHealth: A python library for health predictive models. *CoRR*, *abs/2101.04209*. <https://arxiv.org/abs/2101.04209>