








Mantik: A Workflow Platform for the Development of Artificial Intelligence on High-Performance Computing Infrastructures

Thomas Seidler^{1,2*}, Fabian Emmerich^{2,3*}, Kristian Ehlert³, Rico Berner¹, Oliver Nagel-Kanzler¹, Norbert Schultz⁵, Markus Quade¹, Martin G. Schultz⁴, and Markus Abel¹

¹ Ambrosys GmbH, Potsdam, Germany ² Institute of Physics and Astronomy, University of Potsdam, Germany ³ 4Cast GmbH & Co. KG, Potsdam, Germany ⁴ Jülich Supercomputing Centre (JSC), Forschungszentrum Jülich GmbH, Germany ⁵ Reactive Core GmbH, Berlin, Germany ¶ Corresponding author * These authors contributed equally.

DOI: [10.21105/joss.06136](https://doi.org/10.21105/joss.06136)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Arfon Smith](#)  

Reviewers:

- [@zhaozhang](#)
- [@glofst](#)

Submitted: 18 September 2023

Published: 02 June 2024

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

The use of machine learning (ML) approaches is exponentially increasing, and for many scientific applications, high-performance computing (HPC) infrastructure is used to train large models. However, the tooling for an easy deployment of models for training or inference on HPC infrastructures is not satisfactory, e.g. reproducibility, collaboration and monitoring of ML models are not addressed in existing toolsets. With Mantik, we provide an open-source cloud platform, which simplifies the development of and collaboration on ML models on HPC facilities, and enhances reproducibility by supporting data and code versioning as well as experiment tracking. The users are able to develop their applications in the environment they are most comfortable with – their local machine. Usage of the best-choice IDE and most recent software versions allow to leverage the full potential of the software stack for their research. Using Mantik's remote file service allows for simple management of data in remote storages and keeping track of it. As soon as an application is ready for training or inference, users can immediately submit it to an HPC cluster. During application development, users can train and/or evaluate their models on HPC clusters via CLI on their local machine or our browser-based Mantik cloud platform. The latter only requires an internet browser such that e.g., ML training from your phone becomes feasible. Once training or inference has begun, a user is able to monitor the application in real time on the Mantik cloud platform.

Statement of need

With Mantik, we cover a variety of features that are at the heart of data science and ML. Mantik streamlines fundamental features that form the core of a data scientist's workflow including easy access to external computing resources, data management and collaborative model development. The interest in platforms for ML operations (MLOps) is on the rise, resulting in a rapidly growing market for solutions. Due to this, we will here not provide references to other existing platforms but refer the interested reader to a recent survey by Kreuzberger et.al. ([Kreuzberger et al., 2023](#)).

The Mantik platform provides a secure interface for an easy and unified access to remote computing resources. The interface is extensible to a variety of infrastructures. Currently, the [UNICORE](#) connector is implemented for resources of the Jülich Supercomputing Centre. Further, Mantik supports the connection to commercial clouds by using their respective

interfaces. The unified interface for HPC infrastructures simplifies training and deployment of ML models on different sites. The flexible deployment in Mantik is realized by containerization of the developed ML models, which is standard for commercial clouds to avoid dependencies on hardware. This allows to make use of the desired software packages and versions to fully exploit the potential of state-of-the-art ML frameworks.

For data handling Mantik provides a remote file service that allows for storing data remotely, e.g. in a cloud, and tracking the data by keeping its location as a reference and supporting versioning. The data is therefore always available to the users and changes can be traced. Mantik assures that the actual data will only be accessed if necessary to optimize network load and wait times. Analogously, code and ML models can be versioned using git and their references can be stored. Whereas training often involves large data sets and big infrastructure, inference may run on small devices which may require optimized, or compiled code. In order to allow for these different use cases, the deployment and versioning of data, code, and models is fully supported by Mantik. For experiment tracking, our platform integrates the existing open-source framework MLflow ([Chen et al., 2020](#)) to provide researchers with the ability to track their modeling results independent of user hardware. Storage of training workflows, experiments, and models, including records of input parameters, metrics, models, artifacts, ensures reproducibility and enables users to share their ML solutions with the community.

A [browser-based GUI](#) is provided for hosting and maintaining your ML projects. Here, ML tasks can be executed codeless on remote computing resources, including scheduling and interacting with submitted jobs (status, properties, log messages, access to file system). The platform allows the user to securely save credentials required to quickly access remote services. All references to existing data, code, models and experiments that are stored in a particular project are displayed on the Mantik web-platform. Project owners can invite other users to their projects to share their results and ML solutions effortlessly, which fosters collaboration and knowledge transfer.

Mantik is designed to be used by researchers, data scientists as well as ML engineers. For example, project KI:STE ([Seidler et al., 2021](#)) aims for building a platform particularly based on Mantik for promoting the application of ML in earth system sciences. In the [project Maelstrom](#) Mantik is already used to develop large-scale ML applications for weather and climate science. Due to the simple structure and the variety of built-in features, Mantik provides a kick-start for all users interested in running ML projects on HPC. Consequently, Mantik saves time allowing the user to focus on research and improving their ML approaches rather than bothering about tooling. Mantik allows for easy bookkeeping of key entities (data, models, experiments) that constitute the ML life cycle, therefore providing reproducibility that forms the basis of scientific work and supports its credibility. Moreover, shareable and open ML projects enable students and junior data scientists to familiarize themselves with state-of-the-art approaches. The ease of use and inherent reproducibility therefore renders Mantik also a great choice for education purposes in the field of ML.

The Mantik platform architecture

In Mantik the organizational entity is the Mantik project. In [Figure 1](#), we provide our vision for the Mantik platform and how the different core units within the Mantik project relate to the typical development lifecycle of ML solutions.

- Data and AI modeling: In Mantik projects you can store references to all ingredients that you would need to develop your ML solution (data, code, infrastructure).
- Training: A training session of an ML algorithm (code) with some given data on a particular infrastructure, e.g. HPC, is organized in a run.
- AI engineering: Many runs form an experiment that can be used to find the best performing ML solution for your particular use case.

- **Deployment:** Your evaluation of the experiments results in models that are trained algorithms with fixed sets of parameters, trained on known architectures and meant be used for specific types of data, e.g. video data or texts.

Reproducibility

Collaboration

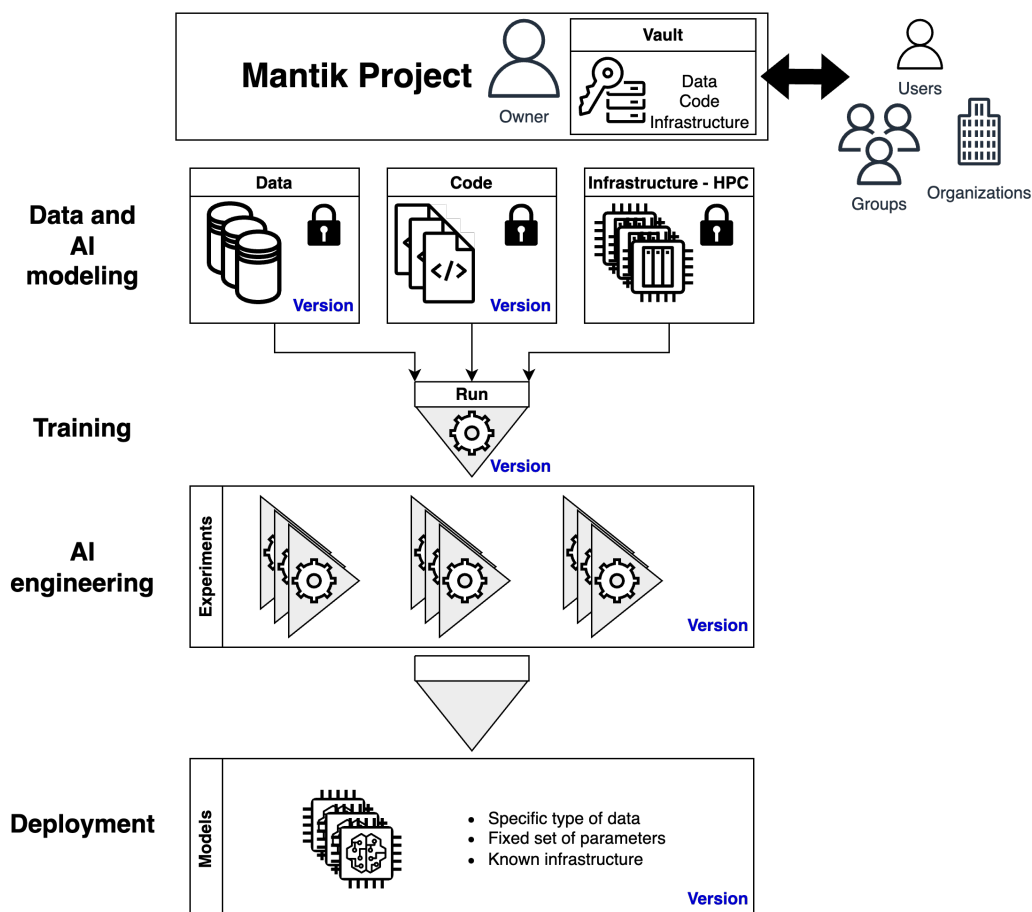


Figure 1: Overview of the organization of the Mantik platform.

Acknowledgements

T.S. and M.A. acknowledge financial support from the project KI:STE funded by Bundesministerium für Bildung und Forschung (BMBF). R.B., M.Q. and M.A. acknowledge financial support EMDRIVE and CECAS funded by BMBF. F.E., K.E., M.S. and M.A. gratefully acknowledge funding from the MAELSTROM EuroHPC-JU project (JU) under 955513. The JU receives support from the European Union's Horizon Research and Innovation Programme and United Kingdom, Germany, Italy, Luxembourg, Switzerland, and Norway. We are grateful for all contributions to the Mantik platform and in particular to the outstanding people involved in its development: the [Mantik contributors](#).

References

- Chen, A., Chow, A., Davidson, A., DCunha, A., Ghodsi, A., Hong, S. A., Konwinski, A., Mewald, C., Murching, S., Nykodym, T., Ogilvie, P., Parkhe, M., Singh, A., Xie, F., Zaharia, M., Zang, R., Zheng, J., & Zumar, C. (2020). Developments in MLflow: A system to accelerate the machine learning lifecycle. *Proceedings of the Fourth International Workshop on Data Management for End-to-End Machine Learning*. <https://doi.org/10.1145/3399579.3399867>
- Kreuzberger, D., Kühl, N., & Hirschl, S. (2023). Machine learning operations (mlops): Overview, definition, and architecture. *IEEE Access*, 11, 31866–31879. <https://doi.org/10.1109/ACCESS.2023.3262138>
- Seidler, T., Schultz, N., Quade, Dr. M., Autermann, C., Gräler, Dr. B., & Abel, P. Dr. M. (2021). *Easing and promoting the application of ML and AI in earth system sciences - introducing the KI:STE platform*. <https://doi.org/10.5194/egusphere-egu21-9632>