

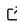


scores: A Python package for verifying and evaluating models and predictions with xarray

Tennessee Leeuwenburg ^{1¶}, Nicholas Loveday ¹, Elizabeth E. Ebert¹, Harrison Cook ¹, Mohammadreza Khanarmuei ¹, Robert J. Taggart ¹, Nikeeth Ramanathan ¹, Maree Carroll ¹, Stephanie Chong ², Aidan Griffiths³, and John Sharples¹

¹ Bureau of Meteorology, Australia ² Independent Contributor, Australia ³ Work undertaken while at the Bureau of Meteorology, Australia ¶ Corresponding author

DOI: [10.21105/joss.06889](https://doi.org/10.21105/joss.06889)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Chris Vernon](#) 

Reviewers:

- [@calebweinreb](#)
- [@savante93](#)

Submitted: 12 June 2024

Published: 09 July 2024

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

scores is a Python package containing mathematical functions for the verification, evaluation and optimisation of forecasts, predictions or models. It supports labelled n-dimensional (multidimensional) data, which is used in many scientific fields and in machine learning. At present, scores primarily supports the geoscience communities; in particular, the meteorological, climatological and oceanographic communities.

scores not only includes common scores (e.g., Mean Absolute Error), it also includes novel scores not commonly found elsewhere (e.g., Fixed Risk Multicategorical (FIRM) score, Flip-Flop Index), complex scores (e.g., threshold-weighted continuous ranked probability score), and statistical tests (such as the Diebold Mariano test). It also contains isotonic regression which is becoming an increasingly important tool in forecast verification and can be used to generate stable reliability diagrams. Additionally, it provides pre-processing tools for preparing data for scores in a variety of formats including cumulative distribution functions (CDF). At the time of writing, scores includes over 50 metrics, statistical techniques and data processing tools.

All of the scores and statistical techniques in this package have undergone a thorough scientific and software review. Every score has a companion Jupyter Notebook tutorial that demonstrates its use in practice.

scores supports xarray datatypes, allowing it to work with Earth system data in a range of formats including NetCDF4, HDF5, Zarr and GRIB among others. scores uses Dask for scaling and performance. Support for pandas is being introduced.

The scores software repository can be found at <https://github.com/nci/scores/>.

Statement of Need

Labelled, n-dimensional data is widely used in many scientific fields. The Earth system science community makes heavy use of physics-based and machine learning models, both to process observations (such as identifying land use from satellite data) and to make predictions about the future (such as forecasting the weather). These models, predictions and forecasts undergo verification and evaluation to assess their correctness.

The purpose of scores is (a) to mathematically verify and validate models and predictions and (b) to foster research into new scores and metrics.

scores handles dimensionality and weighting (e.g., latitude weighting) more effectively than commonly-used data science packages. While there are existing open source Python verification packages for labelled n-dimensional data (see “Related Software Packages” further below), none of these packages offer all of the key benefits of scores.

Key Benefits of scores

To meet the needs of researchers and other users, scores provides the following key benefits.

Data Handling

- Works with labelled, n-dimensional data (e.g., geospatial, vertical and temporal dimensions) for both point-based and gridded data. scores can effectively handle the dimensionality, data size and data structures commonly used for:
 - gridded Earth system data (e.g., numerical weather prediction models)
 - tabular, point, latitude/longitude or site-based data (e.g., forecasts for specific locations).
- Handles missing data, masking of data and weighting of results.
- Supports xarray ([Hoyer & Hamman, 2017](#)) datatypes, and works with NetCDF4 ([Unidata, 2024](#)), HDF5 ([The HDF Group & Koziol, 2020](#)), Zarr ([Miles et al., 2024](#)) and GRIB ([World Meteorological Organization, 2024](#)) data formats among others.

Usability

- A companion Jupyter Notebook ([Jupyter Team, 2024](#)) tutorial for each metric and statistical test that demonstrates its use in practice.
- Novel scores not commonly found elsewhere (e.g., FIRM ([Taggart et al., 2022](#)), Flip-Flop Index ([Griffiths et al., 2019, 2021](#))).
- Commonly-used scores are also included, meeting user requests to use scores as a standalone package.
- All scores and statistical techniques have undergone a thorough scientific and software review.
- An area specifically to hold emerging scores which are still undergoing research and development. This provides a clear mechanism for people to share, access and collaborate on new scores, and be able to easily re-use versioned implementations of those scores.

Compatibility

- Highly modular - provides its own implementations, avoids extensive dependencies and offers a consistent API.
- Easy to integrate and use in a wide variety of environments. It has been used on workstations, servers and in high performance computing (supercomputing) environments.
- Maintains 100% automated test coverage.
- Uses Dask ([Dask Development Team, 2016](#)) for scaling and performance.
- Support for pandas ([McKinney, 2010](#); [The pandas development team, 2024](#)) is being introduced.

Metrics, Statistical Techniques and Data Processing Tools Included in scores

At the time of writing, scores includes over 50 metrics, statistical techniques and data processing tools. For an up to date list, please see the scores documentation.

The ongoing development roadmap includes the addition of more metrics, tools, and statistical tests.

Table 1: A curated selection of the metrics, tools and statistical tests currently included in scores

	Description	A Selection of the Functions Included in scores
Continuous	Scores for evaluating single-valued continuous forecasts.	Mean Absolute Error (MAE), Mean Squared Error (MSE), Root Mean Squared Error (RMSE), Additive Bias, Multiplicative Bias, Pearson's Correlation Coefficient, Flip-Flop Index (Griffiths et al., 2019 , 2021), Quantile Loss, Murphy Score (Ehm et al., 2016).
Probability	Scores for evaluating forecasts that are expressed as predictive distributions, ensembles, and probabilities of binary events.	Brier Score (Brier, 1950), Continuous Ranked Probability Score (CRPS) for Cumulative Distribution Functions (CDFs) (including threshold-weighting, see Gneiting & Ranjan (2011)), CRPS for ensembles (Ferro, 2014 ; Gneiting & Raftery, 2007), Receiver Operating Characteristic (ROC), Isotonic Regression (reliability diagrams) (Dimitriadis et al., 2021).
Categorical	Scores (including contingency table metrics) for evaluating forecasts of categories.	Probability of Detection (POD), Probability of False Detection (POFD), False Alarm Ratio (FAR), Success Ratio, Accuracy, Peirce's Skill Score (Peirce, 1884), Critical Success Index (CSI), Gilbert Skill Score (Gilbert, 1884), Heidke Skill Score, Odds Ratio, Odds Ratio Skill Score, F1 Score, Symmetric Extremal Dependence Index (Ferro & Stephenson, 2011), Flxed Risk Multicategorical (FIRM) Score (Taggart et al., 2022).
Spatial	Scores that take into account spatial structure.	Fractions Skill Score (Roberts & Lean, 2008).
Statistical Tests	Tools to conduct statistical tests and generate confidence intervals.	Diebold-Mariano (Diebold & Mariano, 1995) with both the Harvey et al. (1997) and Hering & Genton (2011) modifications.
Processing Tools	Tools to pre-process data.	Data matching, discretisation, cumulative density function manipulation.

Use in Academic Work

In 2015, the Australian Bureau of Meteorology began developing a new verification system called Jive, which became operational in 2022. For a description of Jive see Loveday, Griffiths, et al. ([2024](#)). The Jive verification metrics have been used to support several publications ([Foley & Loveday, 2020](#); [Griffiths et al., 2017](#); [Taggart, 2022a, 2022b, 2022c](#)).

scores has arisen from the Jive verification system and provides Jive verification functions as a modular, open source package. scores also includes additional metrics that Jive does not contain.

scores has been used to explore user-focused approaches to evaluating probabilistic and categorical forecasts (Loveday, Taggart, et al., 2024).

Related Software Packages

There are multiple open source verification packages in a range of languages. Below is a comparison of scores to other open source Python verification packages. None of these include all of the metrics implemented in scores (and vice versa).

xskillscore (Bell et al., 2021) provides many but not all of the same functions as scores. The Jupyter Notebook tutorials in scores cover a wider array of metrics.

climpred (Brady & Spring, 2021) uses xskillscore combined with data handling functionality, and is focused on ensemble forecasts for climate and weather. climpred makes some design choices related to data structure (specifically associated with climate modelling) which may not generalise effectively to broader use cases. Releasing scores separately allows the differing design philosophies to be considered by the community.

METplus (Brown et al., 2021) is a substantial verification system used by weather and climate model developers. METplus includes a database and a visualisation system, with Python and shell script wrappers to use the MET package for the calculation of scores. MET is implemented in C++ rather than Python. METplus is used as a system rather than providing a modular Python API.

Verif (Nipen et al., 2023) is a command line tool for generating verification plots whereas scores provides a Python API for generating numerical scores.

Pysteps (Imhoff et al., 2023; Pulkkinen et al., 2019) is a package for producing short-term ensemble predictions, focusing on probabilistic nowcasting of radar precipitation fields. It includes a significant verification submodule with many useful verification scores. Pysteps does not provide a standalone verification API.

PyForecastTools (Morley & Burrell, 2020) is a Python package for model and forecast verification which supports darray rather than xarray data structures and does not include Jupyter Notebook tutorials.

Acknowledgements

We would like to thank Jason West and Robert Johnson from the Bureau of Meteorology for their feedback on an earlier version of this manuscript.

We would like to thank and acknowledge the National Computational Infrastructure (nci.org.au) for hosting the scores repository within their GitHub organisation.

References

- Bell, R., Spring, A., Brady, R., Huang, A., Squire, D., Blackwood, Z., Sitter, M. C., & Chegini, T. (2021). *xarray-contrib/xskillscore: Metrics for verifying forecasts* (Version v0.0.23). Zenodo. <https://doi.org/10.5281/zenodo.5173153>
- Brady, R. X., & Spring, A. (2021). climpred: Verification of weather and climate forecasts. *Journal of Open Source Software*, 6(59), 2781. <https://doi.org/10.21105/joss.02781>

- Brier, G. W. (1950). Verification of forecasts expressed in terms of probability. *Monthly Weather Review*, 78(1), 1–3. [https://doi.org/10.1175/1520-0493\(1950\)078%3C0001:vofeit%3E2.0.co;2](https://doi.org/10.1175/1520-0493(1950)078%3C0001:vofeit%3E2.0.co;2)
- Brown, B., Jensen, T., Gotway, J. H., Bullock, R., Gilleland, E., Fowler, T., Newman, K., Adriaansen, D., Blank, L., Burek, T., Harrold, M., Hertneky, T., Kalb, C., Kucera, P., Nance, L., Opatz, J., Vigh, J., & Wolff, J. (2021). The Model Evaluation Tools (MET): More than a decade of community-supported forecast verification. *Bulletin of the American Meteorological Society*, 102(4), E782–E807. <https://doi.org/10.1175/BAMS-D-19-0093.1>
- Dask Development Team. (2016). *Dask: Library for dynamic task scheduling*. <http://dask.pydata.org>
- Diebold, F. X., & Mariano, R. S. (1995). Comparing predictive accuracy. *Journal of Business & Economic Statistics*, 13(3), 253–263. <https://doi.org/10.1080/07350015.1995.10524599>
- Dimitriadis, T., Gneiting, T., & Jordan, A. I. (2021). Stable reliability diagrams for probabilistic classifiers. *Proceedings of the National Academy of Sciences*, 118(8), e2016191118. <https://doi.org/10.1073/pnas.2016191118>
- Ehm, W., Gneiting, T., Jordan, A., & Krüger, F. (2016). Of quantiles and expectiles: Consistent scoring functions, Choquet representations and forecast rankings. *Journal of the Royal Statistical Society. Series B (Statistical Methodology)*, 78(3), 505–562. <https://doi.org/10.1111/rssb.12154>
- Ferro, C. A. T. (2014). Fair scores for ensemble forecasts. *Quarterly Journal of the Royal Meteorological Society*, 140(683), 1917–1923. <https://doi.org/10.1002/qj.2270>
- Ferro, C. A. T., & Stephenson, D. B. (2011). Extremal dependence indices: Improved verification measures for deterministic forecasts of rare binary events. *Weather and Forecasting*, 26(5), 699–713. <https://doi.org/10.1175/WAF-D-10-05030.1>
- Foley, M., & Loveday, N. (2020). Comparison of single-valued forecasts in a user-oriented framework. *Weather and Forecasting*, 35(3), 1067–1080. <https://doi.org/10.1175/waf-d-19-0248.1>
- Gilbert, G. K. (1884). Finley's tornado predictions. *American Meteorological Journal*, 1(5), 166–172.
- Gneiting, T., & Raftery, A. E. (2007). Strictly proper scoring rules, prediction, and estimation. *Journal of the American Statistical Association*, 102(477), 359–378. <https://doi.org/10.1198/016214506000001437>
- Gneiting, T., & Ranjan, R. (2011). Comparing density forecasts using threshold-and quantile-weighted scoring rules. *Journal of Business & Economic Statistics*, 29(3), 411–422. <https://doi.org/10.1198/jbes.2010.08110>
- Griffiths, D., Foley, M., Ioannou, I., & Leeuwenburg, T. (2019). Flip-Flop Index: Quantifying revision stability for fixed-event forecasts. *Meteorological Applications*, 26(1), 30–35. <https://doi.org/10.1002/met.1732>
- Griffiths, D., Jack, H., Foley, M., Ioannou, I., & Liu, M. (2017). *Advice for automation of forecasts: A framework*. Bureau of Meteorology. <https://doi.org/10.22499/4.0021>
- Griffiths, D., Loveday, N., Price, B., Foley, M., & McKelvie, A. (2021). Circular Flip-Flop Index: Quantifying revision stability of forecasts of direction. *Journal of Southern Hemisphere Earth Systems Science*, 71(3), 266–271. <https://doi.org/10.1071/es21010>
- Harvey, D., Leybourne, S., & Newbold, P. (1997). Testing the equality of prediction mean squared errors. *International Journal of Forecasting*, 13(2), 281–291. [https://doi.org/10.1016/S0169-2070\(96\)00719-4](https://doi.org/10.1016/S0169-2070(96)00719-4)
- Hering, A. S., & Genton, M. G. (2011). Comparing spatial predictions. *Technometrics*, 53(4),

- 414–425. <https://doi.org/10.1198/tech.2011.10136>
- Hoyer, S., & Hamman, J. (2017). xarray: N-D labeled arrays and datasets in Python. *Journal of Open Research Software*, 5(1). <https://doi.org/10.5334/jors.148>
- Imhoff, R. O., De Cruz, L., Dewettinck, W., Brauer, C. C., Uijlenhoet, R., Heeringen, K.-J. van, Velasco-Forero, C., Nerini, D., Van Genderachter, M., & Weerts, A. H. (2023). Scale-dependent blending of ensemble rainfall nowcasts and numerical weather prediction in the open-source pysteps library. *Quarterly Journal of the Royal Meteorological Society*, 149(753), 1335–1364. <https://doi.org/10.1002/qj.4461>
- Jupyter Team. (2024). *Jupyter interactive notebook*. GitHub. <https://github.com/jupyter/notebook>
- Loveday, N., Griffiths, D., Leeuwenburg, T., Taggart, R., Pagano, T. C., Cheng, G., Plastow, K., Ebert, E., Templeton, C., Carroll, M., Khanarmuei, M., & Nagpal, I. (2024). The Jive verification system and its transformative impact on weather forecasting operations. *arXiv preprint arXiv:2404.18429*. <https://doi.org/10.48550/arXiv.2404.18429>
- Loveday, N., Taggart, R., & Khanarmuei, M. (2024). A user-focused approach to evaluating probabilistic and categorical forecasts. *Weather and Forecasting*. <https://doi.org/10.1175/waf-d-23-0201.1>
- McKinney, W. (2010). Data structures for statistical computing in Python. In S. van der Walt & J. Millman (Eds.), *Proceedings of the 9th Python in Science Conference* (pp. 56–61). <https://doi.org/10.25080/Majors-92bf1922-00a>
- Miles, A., jakirkham, Bussonnier, M., Moore, J., Orfanos, D. P., Bennett, D., Bourbeau, J., Hamman, J., Stansby, D., Fulton, A., Lee, G., Patel, Z., Abernathey, R., Kristensen, M. R. B., Rzepka, N., Chopra, S., Awa, A. B., Verma, S., Rocklin, M., ... shikharsg. (2024). *zarr-developers/zarr-python: v2.18.2* (Version v2.18.2). Zenodo. <https://doi.org/10.5281/zenodo.11320255>
- Morley, S., & Burrell, A. (2020). *drsteve/PyForecastTools: Version 1.1.1* (Version v1.1.1). Zenodo. <https://doi.org/10.5281/zenodo.3764117>
- Nipen, T. N., Stull, R. B., Lussana, C., & Seierstad, I. A. (2023). Verif: A weather-prediction verification tool for effective product development. *Bulletin of the American Meteorological Society*, 104(9), E1610–E1618. <https://doi.org/10.1175/bams-d-22-0253.1>
- Pearce, C. S. (1884). The numerical measure of the success of predictions. *Science*, ns-4(93), 453–454. <https://doi.org/10.1126/science.ns-4.93.453.b>
- Pulkkinen, S., Nerini, D., Pérez Hortal, A. A., Velasco-Forero, C., Seed, A., Germann, U., & Foresti, L. (2019). Pysteps: An open-source Python library for probabilistic precipitation nowcasting (v1.0). *Geoscientific Model Development*, 12(10), 4185–4219. <https://doi.org/10.5194/gmd-12-4185-2019>
- Roberts, N. M., & Lean, H. W. (2008). Scale-selective verification of rainfall accumulations from high-resolution forecasts of convective events. *Monthly Weather Review*, 136(1), 78–97. <https://doi.org/10.1175/2007MWR2123.1>
- Taggart, R. (2022a). *Assessing calibration when predictive distributions have discontinuities*. <http://www.bom.gov.au/research/publications/researchreports/BRR-064.pdf>
- Taggart, R. (2022b). Evaluation of point forecasts for extreme events using consistent scoring functions. *Quarterly Journal of the Royal Meteorological Society*, 148(742), 306–320. <https://doi.org/10.1002/qj.4206>
- Taggart, R. (2022c). Point forecasting and forecast evaluation with generalized Huber loss. *Electronic Journal of Statistics*, 16(1), 201–231. <https://doi.org/10.1214/21-ejs1957>
- Taggart, R., Loveday, N., & Griffiths, D. (2022). A scoring framework for tiered warnings and

- multicategorical forecasts based on fixed risk measures. *Quarterly Journal of the Royal Meteorological Society*, 148(744), 1389–1406. <https://doi.org/10.1002/qj.4266>
- The HDF Group, & Koziol, Q. (2020). *HDF5-version 1.12.0*. <https://doi.org/10.11578/dc.20180330.1>
- The pandas development team. (2024). *pandas-dev/pandas: Pandas (Version v2.2.2)*. Zenodo. <https://doi.org/10.5281/zenodo.10957263>
- Unidata. (2024). *Network common data form (NetCDF)*. UCAR/Unidata Program Center. <https://doi.org/10.5065/D6H70CW6>
- World Meteorological Organization. (2024). *WMO no. 306 FM 92 GRIB (edition 2)*. World Meteorological Organization. <https://codes.wmo.int/grib2>