The Journal of Open Source Software

sleev: An R Package for Semiparametric Likelihood Estimation with Errors in Variables

Jiangmei Xiong [1¶], **Sarah C. Lotspeich** [2], **Joey B. Sherrill** [3], **Gustavo Amorim** [1], **Bryan E. Shepherd** [1], **and Ran Tao** [1,4]

1 Department of Biostatistics, Vanderbilt University Medical Center, USA **2** Department of Statistical Sciences, Wake Forest University, USA **3** Brigham Young University, USA **4** Vanderbilt Genetics Institute, Vanderbilt University Medical Center, USA ¶ Corresponding author

DOI: 10.xxxxxx/draft

Software

- Review
- Repository
- Archive

Editor: Mehmet Hakan Satman

Reviewers:

- @alemermartinez
- @aalfons

Submitted: 22 September 2024
Published: unpublished

License
Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License (CC BY 4.0).

DRAFT

Summary

Data with measurement error in the outcome, covariates, or both are not uncommon, particularly with the increased use of routinely collected data for biomedical research. In settings with error-prone data, often only a subsample of study data is validated, also known as two-phase studies. The sieve maximum likelihood estimator (SMLE), which combines the error-prone data on all records with the validated data on a subsample, is a highly efficient and robust estimator to analyze such data. However, given their complexity, a computationally efficient and user-friendly tool is needed to obtain SMLEs. R package sleev fills this gap by making semiparametric likelihood-based inference using SMLEs for error-prone two-phase data in settings with binary and continuous outcomes. Functions from this package can be used to analyze data with error-prone responses and covariates.

Statement of Need

Routinely collected data are being used frequently in biomedical research, such as electronic health records. However, these data tend to be error-prone, and careless use of these data could lead to biased estimates and misleading research findings (Duan et al., 2016). To avoid invalid study results, trained experts carefully verify and extract data elements. However, it is usually only feasible to validate data for a subset of records or variables. After validation, researchers have error-prone pre-validation data for all records (phase one data) error-free validated data on a subset of records (phase two data). Analyses aims to combine the two types of data to obtain estimates that have low bias and are as efficient as possible.

SMLE, which combines the error-prone data on all records with the validated data on a subsample, is a highly efficient and robust estimator to analyze such two-phase data (Lotspeich et al., 2022; Tao et al., 2021). Still, in practice these estimators can be difficult to implement, as they involve approximating nuisance conditional densities using B-splines (Schumaker, 2007) and then maximizing the semiparametric likelihood via a sophisticated EM algorithm (Tao et al., 2017). R package sleev makes this method readily applicable for practitioners in a user-friendly way. sleev integrates and extends R packages logreg2ph and TwoPhaseReg, two primitive R packages developed with the original methods papers (Lotspeich et al., 2022; Tao et al., 2021). These two packages lacked proper documentation, were computationally slow, and were difficult to use. To promote the use of SMLE in two-phase data, extensive work has been done to create sleev, a computationally efficient and user-friendly R package to analyze two-phase, error-prone data. Specifically, in sleev we rewrote the core algorithms in C++ to speed up the computation, and we unified the syntax across functions.

The Journal of Open Source Software

SMLE for linear regression

In this section, we briefly introduce the SMLEs for linear regression. The SMLEs for logistic regression are similar to linear regression and described in the package vignette.Suppose that we want to fit a standard linear regression model for a continuous outcome Y and covariates $\mathbf{X}$: $Y = \alpha + {}^T\mathbf{X} + \epsilon$, where $\epsilon \sim N(0, \sigma^2)$. Our goal is to obtain estimates of $\boldsymbol{\theta} = (\alpha,\ {}^T, \sigma^2)^T$. When we have error-prone data, Y and $\mathbf{X}$ are unobserved except for a subset of validated records. For the majority of unvalidated records, only the error-prone outcome $Y^* = Y + W$ and covariates $\mathbf{X}^* = \mathbf{X} + \mathbf{U}$ are observed in place of Y and $\mathbf{X}$, where W and $\mathbf{U}$ are the errors for the outcome and covariates, respectively. We assume that $W,\ \mathbf{U} \perp \epsilon$. With potential errors in our data, a naive regression analysis using error-prone variables Y^* and $\mathbf{X}^*$ could render misleading results.

We assume that the joint density of the complete data $(Y^*, \mathbf{X}^*, W, \mathbf{U})$ takes the form

$$P(Y^*, \mathbf{X}^*,\ W,\ \mathbf{U}) = P(Y^*|\mathbf{X}^*,\ W,\ \mathbf{U})\, P(W,\ \mathbf{U}|\mathbf{X}^*\ )\, P(\mathbf{X}^*)$$

$$= P_{\boldsymbol{\theta}}\left(Y|X\right) P(W,\ \mathbf{U}|\mathbf{X}^*)\, P(\mathbf{X}^*)$$

where $P(\cdot)$ and $P(\cdot|\cdot)$ denote density and conditional density functions, respectively. $P_{\boldsymbol{\theta}}\left(Y|\mathbf{X}\right)$ then refers to the conditional density function of $Y = \alpha + \mathbf{X} + \epsilon$. Denote the validation indicator variable by V, with $V = 1$ indicating that a record was validated and $V = 0$ otherwise. For records with $V = 0$, their measurement errors $(W, \mathbf{U})$ are missing, and therefore their contributions to the log-likelihood can be obtained by integrating out W and $\mathbf{U}$. Let $(Y_i^*, \mathbf{X}_i^*, W_i, \mathbf{U}_i, V_i, Y_i, \mathbf{X}_i)$ for $i = 1, \ldots, n$ denote independent and identically distributed realizations of $(Y^*, \mathbf{X}^*, W, \mathbf{U}, V, Y, \mathbf{X})$ in a sample of n subjects. Then, the observed-data log-likelihood takes the form

$$\sum_{i=1}^{n} V_i \{\log P_{\boldsymbol{\theta}}\left(Y_i|\mathbf{X}_i\right) + \log P(W_i, \mathbf{U}_i|\mathbf{X}_i^*)\}$$

$$+ \sum_{i=1}^{n} (1 - V_i) \log \left\{ \int \int P_{\boldsymbol{\theta}}\left(Y_i^* - w|\mathbf{X}_i^* - \mathbf{u}\right) P(w, \mathbf{u}|\mathbf{X}_i^*)\, dw d\mathbf{u} \right\}. \qquad (1)$$

We estimate the unknown measurement error model, $P(W_i, \mathbf{U}_i|\mathbf{X}_i^*)$ using B-spline sieves. Specifically, we approximate $P(w_i, \mathbf{u}_i|\mathbf{X}_i^*)$ and $\log P(W_i, \mathbf{U}_i|\mathbf{X}_i^*)$ by $\sum_{k=1}^{m} I\left(w_i = w_k, \mathbf{u}_i = \mathbf{u}_k\right) \sum_{j=1}^{s_n}$ and $\sum_{k=1}^{m} I\left(W_i = w_k, \mathbf{U}_i = \mathbf{u}_k\right) \sum_{j=1}^{s_n} B_j^q\left(\mathbf{X}_i^*\right) \log p_{kj}$, respectively. $\{(w_1, \mathbf{u}_1), \ldots, (w_m, \mathbf{u}_m)\}$ are the m distinct observed $(W, \mathbf{U})$ values, $B_j^q\left(\mathbf{X}_i^*\right)$ is the jth B-spline basis function of order q evaluated at $\mathbf{X}_i^*$, s_n is the dimension of the B-spline basis, and p_{kj} is the coefficient associated with $B_j^q\left(\mathbf{X}_i^*\right)$ and $(w_k, \mathbf{u}_k)$. The log-likelihood in expression (1) is now approximated by

$$\sum_{i=1}^{n} V_i \{\log P_{\boldsymbol{\theta}}\left(Y_i|\mathbf{X}_i\right) + \sum_{k=1}^{m} I\left(W_i = w_i,\ \mathbf{U}_i = \mathbf{u}_k\right) \sum_{j=1}^{S_n} B_j^q(\mathbf{X}_i^*) \log p_{kj}\}$$

$$+ \sum_{i=1}^{n} (1 - V_i) \log \left\{ \sum_{k=1}^{m} P_{\boldsymbol{\theta}}\left(Y_i^* - w|\mathbf{X}_i^* - \mathbf{u}\right) \sum_{j=1}^{S_n} B_j^q(\mathbf{X}_i^*) \log p_{kj} \} \right\}. \qquad (2)$$

The maximization of expression (2) is carried out through an EM algorithm to find the SMLEs $\hat{\boldsymbol{\theta}}$ and $\hat{p}_{kj}$. The covariance matrix of the SMLE $\hat{\boldsymbol{\theta}}$ is obtained through the method of profile

DRAFT

likelihood (Murphy & Van der Vaart, 2000). Full details on the SMLE method for logistic regression with error-prone data, including its theoretical properties, can be found in (Lotspeich et al., 2022).

Functionalities of the sleev R package

The sleev package provides a user-friendly way to obtain the SMLEs and their standard errors. The package can be installed and loaded from the R CRAN library or Github. The sleev package includes two main functions: linear2ph() and logistic2ph(), to fit linear and logistic regressions, respectively, under two-phase sampling with an error-prone outcome and covariates. The input arguments are similar for the two functions and listed in Table 1. From Table 1, we see that in addition to the arguments for error-prone and error-free outcome and covariates, the user needs to specify the B-spline matrix $B_j^q\left(\mathbf{X}_i^*\right)$ to be used in the estimation of the error densities.

The sleev package included a dataset constructed to mimic data from the Vanderbilt Comprehensive Care Clinic (VCCC) patient records from (Giganti et al., 2020). Table 1 lists the variables in this dataset to be used in subsequent analyses.

Example: Case study with mock data

We now illustrate how to obtain SMLEs using the sleev package with dataset mock.vccc. Specifically, we show how to fit a linear regression model in the presence of errors in both the outcome and covariates using the linear2ph() function. Situations with more covariates and examples with logistic regression are included in the package vignette.

This example fits a linear regression model with CD4 count at antiretroviral therapy (ART) initiation regressed on viral load (VL) at ART initiation, adjusting for sex at birth. Both CD4 and VL are error-prone, partially validated variables, whereas sex is error-free.

Because of skewness, we often transform both CD4 and VL. In our analysis, CD4 was divided by 10 and square root transformed and VL was $\log_{10}$ transformed:

```
data(mock.vccc)
mock.vccc$CD4_val_sq10 <- sqrt(mock.vccc$CD4_val/10)
mock.vccc$CD4_unval_sq10 <- sqrt(mock.vccc$CD4_unval/10)
mock.vccc$VL_val_l10 <- log10(mock.vccc$VL_val)
mock.vccc$VL_unval_l10 <- log10(mock.vccc$VL_unval)
```

To obtain the SMLEs, we first need to set up the B-spline basis for the covariates `VL_unval_l10` (the transformed error-prone VL variable from phase one) and `Sex`. Here, we use a cubic B-spline basis with the `degree=3` argument in our call to the `bs()` function from the `splines` package. The size of the basis s_n is set to be 20; this was specified through the `df = 20` argument. The B-spline basis is set up separately for the two `Sex` groups, and the size of the B-spline basis is assigned in proportion to the relative size of the two `Sex` groups. This allows the errors in `VL_unval_l10` to be heterogeneous between males and females. The described B-spline basis is constructed as follows.

```
n <- nrow(mock.vccc) # get the number of rows of the dataset
sn <- 20 # set the size of the B-spline basis
# portion the size of the B-spline basis according to the size of two sex groups
sex_ratio <- sum(mock.vccc$Sex==0)/n
sn_0 <- round(sn*sex_ratio, digits = 0)
sn_1 <- sn-sn_0
# create B-spline basis for each sex group separately through bs()
Bspline_0 <- splines::bs(x=mock.vccc$VL_unval_l10[mock.vccc$Sex==0], df=sn_0, degree=3,
```

DRAFT

```
                              intercept=TRUE)
Bspline_1 <- splines::bs(x=mock.vccc$VL_unval_l10[mock.vccc$Sex==1], df=sn_1, degree=3,
                              intercept=TRUE)
# create B-spline basis for this analysis by combining the two bases created above
Bspline <- matrix(data=0, nrow=n, ncol=sn)
Bspline[mock.vccc$Sex==0,1:(sn_0)] <- Bspline_0
Bspline[mock.vccc$Sex==1,(sn_0+1):sn] <- Bspline_1
# name B-spline basis matrix columns properly
colnames(Bspline) <- paste0("bs", 1:sn)
```

Alternatively, if the investigator has prior knowledge that the errors in `VL_unval_l10` are likely to be homogeneous, one may fit a simpler model by not stratifying the B-spline basis by `Sex`. As a final step before running the analysis, we add the B-spline basis to the analysis dataset to create a single data frame:

```
data <- data.frame(cbind(mock.vccc, Bspline))
```

The SMLEs can be obtained by running function `linear2ph()`, as shown in the code below. The fitted SMLEs are stored in a list object. Here, we assign the fitted SMLEs to the variable name res_linear.This list contains five slots: coefficients, covariance, sigma, converge, and converge_cov. We should first check if the EM algorithms for estimating the regression coefficients and their covariance matrix converged by checking if `res_linear$converge` and `res_linear$converge_conv` are TRUE.

The `res_linear$coefficient` slot gives the regression coefficient estimates and their corresponding standard error estimates, z-statistics, and p-values.

```
res_linear <- linear2ph(Y_unval="CD4_unval_sq10", Y="CD4_val_sq10", X_unval="VL_unval_l1
                        X="VL_val_l10", Z="Sex", Bspline=colnames(Bspline), data=data,
                        hn_scale = 1, noSE = FALSE, TOL = 1e-04, MAX_ITER = 1000,
                        verbose = FALSE)
```

Check convergence:

```
c(res_linear$converge, res_linear$converge_cov)
```

Result coefficients:

```
res_linear$coefficients

# Estimate SE Statistic p-value
# Intercept 4.821 0.1587 30.39 0.000000
# VL_val_l10 -0.141 0.0398 -3.55 0.000389
# Sex 0.273 0.1089 2.51 0.012229
```

Acknowledgement

This research was supported by the National Institute of Health grants R01AI131771, R01HL094786, and P30AI110527 and the 2022 Biostatistics Faculty Development Award from the Department of Biostatistics at Vanderbilt University Medical Center.

Reference

Table 1: Main arguments for linear2ph() and logistic2ph()

Argument	Description
Y_unval	Column name of unvalidated outcome in the input dataset.

DRAFT

Argument	Description
Y	Column name of validated outcome in the input dataset. NAs in the input will be counted as individuals not selected in phase two.
X_unval	Column names of unvalidated covariates in the input dataset.
X	Column names of validated covariates in the input dataset. NAs in the input will be counted as individuals not selected in phase two.
Z	Column names of error-free covariates in the input dataset.
Bspline	Column names of the B-spline basis in the input dataset.
data	Dataset
hn_scale	Scale of the perturbation constant in the variance estimation via the method of profile likelihood. The default is 1.
noSE	Standard errors of the parameter estimators will not be estimated when set to TRUE. The default is FALSE.
TOL Convergence criterion.	The default is 0.0001.
MAX_ITER	Maximum number of iterations in the EM algorithm. The default is 1000.
verbose	Print analysis details when set to TRUE. The default is FALSE.

Table 2: Data dictionary for mock.vccc

Name	Status	Type	Description
ID	error-free		Patient ID
VL_unval	error-prone	continuous	Viral load (VL) at antiretroviral therapy (ART) initiation
VL_val	validated		
ADE_unval	error-prone	binary	Had an AIDS-defining event (ADE) within one year of ART initiation: 1 - yes, 0 – no
ADE_val	validated		
CD4_unval	error-prone	continuous	CD4 count at ART initiation
CD4_val	validated		
prior_ART	error-free	binary	Experienced ART before enrollment: 1 - yes, 0 - no
Sex	error-free	binary	Sex at birth of patient: 1 - male, 0 - female
Age	error-free	continuous	Age of patient

Duan, R., Cao, M., Wu, Y., Huang, J., Denny, J. C., Xu, H., & Chen, Y. (2016). An empirical study for impacts of measurement errors on EHR based association studies. *AMIA Annual Symposium Proceedings*, *2016*, 1764.

Giganti, M. J., Shaw, P. A., Chen, G., Bebawy, S. S., Turner, M. M., Sterling, T. R., & Shepherd, B. E. (2020). Accounting for dependent errors in predictors and time-to-event outcomes using electronic health records, validation samples, and multiple imputation. *The Annals of Applied Statistics*, *14*(2), 1045. https://doi.org/10.1214/20-aoas1343

Lotspeich, S. C., Shepherd, B. E., Amorim, G. G., Shaw, P. A., & Tao, R. (2022). Efficient odds ratio estimation under two-phase sampling using error-prone data from a multi-national HIV research cohort. *Biometrics*, *78*(4), 1674–1685. https://doi.org/10.1111/biom.13512

Murphy, S. A., & Van der Vaart, A. W. (2000). On profile likelihood. *Journal of the American Statistical Association*, *95*(450), 449–465. https://doi.org/10.2307/2669386

DRAFT

Schumaker, L. (2007). *Spline functions: Basic theory*. Cambridge university press. https://doi.org/10.1017/CBO9780511618994

Tao, R., Lotspeich, S. C., Amorim, G., Shaw, P. A., & Shepherd, B. E. (2021). Efficient semiparametric inference for two-phase studies with outcome and covariate measurement errors. *Statistics in Medicine*, *40*(3), 725–738. https://doi.org/10.1002/sim.8799

Tao, R., Zeng, D., & Lin, D.-Y. (2017). Efficient semiparametric inference under two-phase sampling, with applications to genetic association studies. *Journal of the American Statistical Association*, *112*(520), 1468–1476. https://doi.org/10.1080/01621459.2017.1295864

DRAFT