

hetGPy: Heteroskedastic Gaussian Process Modeling in Python

David O’Gara¹✉, Mickaël Binois², Roman Garnett³, and Ross A Hammond⁴

¹ Division of Computational and Data Sciences, Washington University in St. Louis, USA ² Acumes team, Université Côte d’Azur, Inria, Sophia Antipolis, France ³ Department of Computer Science, McKelvey School of Engineering, Washington University in St. Louis, USA ⁴ School of Public Health, Washington University in St. Louis, USA ⁵ Center on Social Dynamics and Policy, Brookings Institution, Washington DC, USA ⁶ Santa Fe Institute, Santa Fe, USA ✉ Corresponding author

DOI: [10.xxxxx/draft](https://doi.org/10.xxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: ↗

Submitted: 18 October 2024

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/))

Summary

Computer experiments are ubiquitous in the physical and social sciences. When experiments are time-consuming to run, emulator (or surrogate) models are often used to map the input–output response surface of the simulator, treating it as a black box. The workhorse function of emulator models is Gaussian process regression (GPR). GPs provide flexible, non-linear regression targets with good interpolation properties and uncertainty quantification. However, it is well-known that naïve GPR scales cubically with input size, and specifically involves intensive computation of matrix determinants and solving linear systems (Garnett, 2023; Gramacy, 2020) when fitting hyperparameters. Further, naïve GPR with noisy observations typically assumes an independent, identically-distributed noise process, but many data-generating mechanisms, especially those found in stochastic computer simulation, exhibit input-dependent noise (also known as heteroskedasticity) (Baker et al., 2020). The software package hetGP (Binois & Gramacy, 2021) alleviates these both of these concerns: when the dataset of interest contains replicates, as it is possible to perform inference and prediction with cost growing cubically in the number of unique design locations n rather than the full dataset of size N and can jointly model the mean and noise process as two coupled GPs, allowing smooth noise dynamics over parameter space (Binois et al., 2018). The package has been used in a variety of contexts, such as statistics (Binois et al., 2018, 2019), biology (Lazaridis et al., 2022) and computational epidemiology (Shattock et al., 2022). We present a Python reimplementation hetGPy, developed in part due to Python’s widespread use in computer simulation (Downey, 2023; Kinser, 2022).

Statement of Need

Python is a popular language for software development, data science, and computer experimentation. Its object orientated framework, high-level functionality, and third-party libraries such as numpy (Harris et al., 2020), scipy (Virtanen et al., 2020), pytorch (Paszke et al., 2019) and scikit-learn (Pedregosa et al., 2011) make it a powerful tool for academic and industry professionals alike. Python is especially popular for computer simulation, with one particular example being the widespread use of Python models of COVID-19 spread (Aylett-Bullock et al., 2021; Kerr et al., 2021; O’Gara et al., 2023). hetGPy is well-posed for sequential design of Python models, mirroring the functionality of hetGP without having to rely on intermediate libraries such as reticulate (Ushey et al., 2023) or rpy2, or in a more laborious case, converting a Python simulation to R.

The state of the art for GPR in Python is GPyTorch (Gardner et al., 2018), facilitated by

43 black-box matrix–matrix multiplication (BBMM) which is extremely computationally efficient
 44 on GPUs. Other GPR routines for Python exist as well and can be found in libraries such as
 45 PyMC (Abril-Pla et al., 2023), GPflow (Matthews et al., 2017), and GPJax (Pinder & Dodd,
 46 2022). However, to our knowledge, these libraries, under their default behavior, do not jointly
 47 model the mean and variance as coupled GPs or take advantage of replication in datasets,
 48 meaning that under large degrees of replication and input-dependent noise, as in common
 49 in stochastic computer experiments (Baker et al., 2020), hetGPpy will be more computationally
 50 efficient. Figure 1 (a) shows a heteroskedastic GP fit to a simulated motorcycle accident
 51 dataset (Silverman, 1985). While it is possible to model input-dependent noise in GPyTorch or
 52 BoTorch (Balandat et al., 2020), specifying a smooth noise process in the method of (Binois
 53 et al., 2018) would require a custom implementation. Figure 1 (b) illustrates the results
 54 of a simple one-dimensional example where we compare the model fits using both hetGPpy
 55 and GPyTorch. While both models result in similar predictions, hetGPpy is several orders of
 56 magnitude faster, performing exact inference in less than 1 second, while GPyTorch takes over
 57 10 seconds.

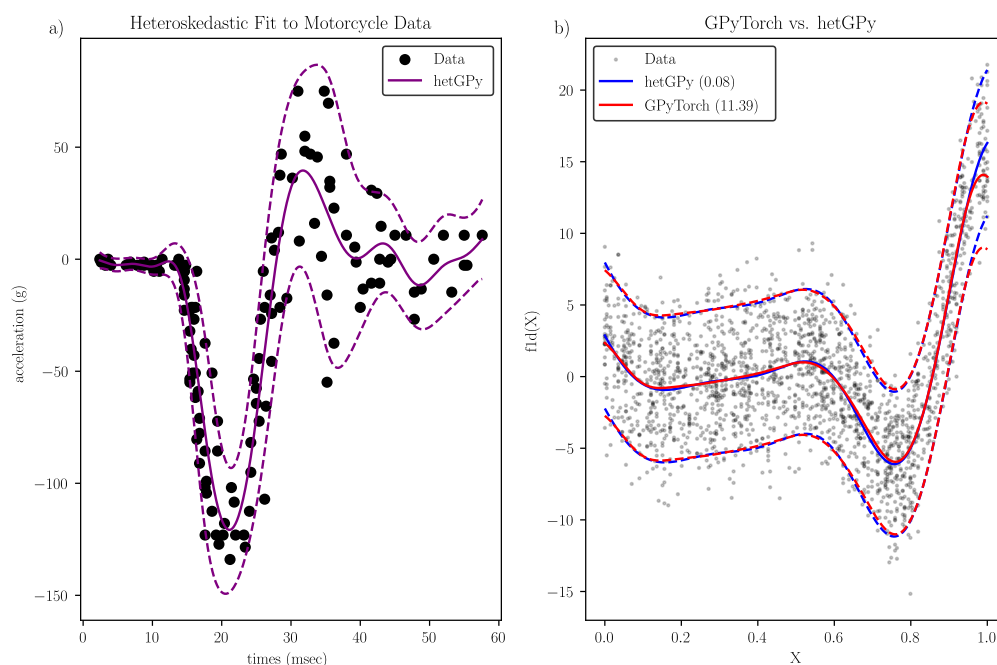


Figure 1: The main features of hetGPpy. Panel (a) shows a heteroskedastic fit to the motorcycle data (Silverman, 1985). Panel (b) shows that hetGPpy yields faster training than (naïve) GPyTorch with similar performance under high replication. Bolded and dashed lines indicate the predictive mean and 90% predictive intervals for homoscedastic GPR, respectively. Data were sampled from the f1d function $f(x) = (6x - 2)^2 \sin(12x - 4)$ (Forrester et al., 2008) in hetGP with between 1 and 20 replicates at each design location. Model training times in seconds are next to legend labels.

58 hetGPpy also has two intermediate goals: (1) efficient computations, accomplished via imple-
 59 mentation on numpy arrays and (2) minimal dependencies, the core of which are numpy for
 60 efficient array-based computation and scipy which contains the definitive implementation of
 61 the L-BFGS-B algorithm in Fortran (Byrd et al., 1995; Morales & Nocedal, 2011) used for
 62 maximum likelihood estimation of hyperparameters. Our experiments indicate hetGPpy is able
 63 to learn response surfaces efficiently, and in the case of high replication, do so on CPUs more
 64 efficiently than the default implementation in GPyTorch, as shown for a suite of test problems
 65 in Table 1. As a comparator, we also conduct a set of experiments using stochastic kriging
 66 (SK) (Ankenman et al., 2010), a precursor method to (Binois et al., 2018) that also allows for
 67 maximum likelihood estimation with replication, and under the case of homoskedasticity, is

68 nearly equivalent to homoscedastic GPR in hetGP and hetGPy (Gramacy, 2020). We implement
 69 SK with a custom GPyTorch likelihood that accounts for replication under homoskedasticity.
 70 Specifically, given a dataset X with unique designs (X_1, \dots, X_k) each replicated (n_1, \dots, n_k)
 71 times, we pre-average the outputs Y_i at each unique input location, and then estimate the
 72 diagonal of the noise matrix as (σ^2 / n_i) . We see that the for the SK case, hetGPy and
 73 GPyTorch have training times on a similar order of magnitude. The package hetGPy is under
 74 active development and is well-posed to engage with the wider Python community for future
 75 extension such as arbitrary kernel functions with auto-differentiation methods facilitated by
 76 PyTorch.

Covariance	Experiment	Time (seconds)				Number of Evaluations	
		GPyTorch (naïve)	GPyTorch (SK)	hetGPy	GPyTorch (naïve)	GPyTorch (SK)	hetGPy
Gaussian	Branin	98.449	5.622	2.693	13	18	13
	Goldstein-Price	119.019	3.688	1.895	21	12	10
	Hartmann-4D	98.566	4.578	4.936	18	15	16
	Hartmann-6D	466.381	19.845	21.472	95	69	40
	Sphere-6D	206.05	13.989	7.702	42	51	20
Matern $\nu = 5/2$	Branin	90.501	4.765	6.563	13	12	20
	Goldstein-Price	153.033	4.406	4.005	20	11	15
	Hartmann-4D	201.208	6.895	6.755	21	18	16
	Hartmann-6D	339.387	28.427	45.23	53	67	57
	Sphere-6D	147.118	16.034	16.345	24	38	14

77
 78 **Table 1:** Comparing training times across a suite of test problems and libraries. All experiments
 79 reflect exact GPR with homoscedastic noise. Optimization problems were selected from
 80 (Picheny et al., 2013) with implementations from (Surjanovic & Bingham, n.d.). Experiments
 81 consisted of a Latin Hypercube design of 1,000 unique locations, with between 1 and 10
 82 replicates. iid Gaussian noise was added to each resulting dataset.

83 References

- 84 Abril-Pla, O., Andreani, V., Carroll, C., Dong, L., Fannesbeck, C. J., Kochurov, M., Kumar,
 85 R., Lao, J., Luhmann, C. C., Martin, O. A., Osthege, M., Vieira, R., Wiecki, T., & Zinkov,
 86 R. (2023). PyMC: A modern, and comprehensive probabilistic programming framework in
 87 Python. *PeerJ Computer Science*, 9, e1516. <https://doi.org/10.7717/peerj-cs.1516>
- 88 Ankenman, B., Nelson, B. L., & Staum, J. (2010). Stochastic Kriging for Simulation
 89 Metamodeling. *Operations Research*, 58(2), 371–382. <https://doi.org/10.1287/opre.1090.0754>
- 90
 91 Aylett-Bullock, J., Cuesta-Lazaro, C., Quera-Bofarull, A., Icaza-Lizaola, M., Sedgewick, A.,
 92 Truong, H., Curran, A., Elliott, E., Caulfield, T., Fong, K., Vernon, I., Williams, J.,
 93 Bower, R., & Krauss, F. (2021). J une :
 94 Open-source individual-based epidemiology simulation. *Royal Society Open Science*, 8(7),
 95 210506. <https://doi.org/10.1098/rsos.210506>
- 96 Baker, E., Barbillon, P., Fadikar, A., Gramacy, R. B., Herbei, R., Higdon, D., Huang, J.,
 97 Johnson, L. R., Ma, P., Mondal, A., Pires, B., Sacks, J., & Sokolov, V. (2020). *Analyzing*
 98 *Stochastic Computer Models: A Review with Opportunities*. arXiv. <https://doi.org/10.48550/arXiv.2002.01321>
- 99
 100 Balandat, M., Karrer, B., Jiang, D., Daulton, S., Letham, B., Wilson, A. G., & Bakshy, E.
 101 (2020). BoTorch: A Framework for Efficient Monte-Carlo Bayesian Optimization. *Advances*
 102 *in Neural Information Processing Systems*, 33, 21524–21538. [https://proceedings.neurips.
 103 cc/paper/2020/hash/f5b1b89d98b7286673128a5fb112cb9a-Abstract.html](https://proceedings.neurips.cc/paper/2020/hash/f5b1b89d98b7286673128a5fb112cb9a-Abstract.html)
- 104 Binois, M., & Gramacy, R. B. (2021). **hetGP** : Heteroskedastic Gaussian Process Modeling
 105 and Sequential Design in *r*. *Journal of Statistical Software*, 98(13). [https://doi.org/10.
 106 18637/jss.v098.i13](https://doi.org/10.18637/jss.v098.i13)
- 107 Binois, M., Gramacy, R. B., & Ludkovski, M. (2018). Practical Heteroscedastic Gaussian

- 108 Process Modeling for Large Simulation Experiments. *Journal of Computational and*
109 *Graphical Statistics*, 27(4), 808–821. <https://doi.org/10.1080/10618600.2018.1458625>
- 110 Binois, M., Huang, J., Gramacy, R. B., & Ludkovski, M. (2019). Replication or Exploration?
111 Sequential Design for Stochastic Simulation Experiments. *Technometrics*, 61(1), 7–23.
112 <https://doi.org/10.1080/00401706.2018.1469433>
- 113 Byrd, R. H., Lu, P., Nocedal, J., & Zhu, C. (1995). A Limited Memory Algorithm for
114 Bound Constrained Optimization. *SIAM J. Sci. Comput.*, 16(5), 1190–1208. <https://doi.org/10.1137/0916069>
- 116 Downey, A. B. (2023). *Modeling and Simulation in Python: An Introduction for Scientists*
117 *and Engineers*. No Starch Press.
- 118 Forrester, A. I. J., Sóbester, A., & Keane, A. J. (2008). *Engineering Design via Surrogate*
119 *Modelling: A Practical Guide* (1st ed.). Wiley. <https://doi.org/10.1002/9780470770801>
- 120 Gardner, J., Pleiss, G., Weinberger, K. Q., Bindel, D., & Wilson, A. G. (2018). GPyTorch:
121 Blackbox Matrix-Matrix Gaussian Process Inference with GPU Acceleration. *Advances in*
122 *Neural Information Processing Systems*, 31. [https://proceedings.neurips.cc/paper/2018/](https://proceedings.neurips.cc/paper/2018/hash/27e8e17134dd7083b050476733207ea1-Abstract.html)
123 [hash/27e8e17134dd7083b050476733207ea1-Abstract.html](https://proceedings.neurips.cc/paper/2018/hash/27e8e17134dd7083b050476733207ea1-Abstract.html)
- 124 Garnett, R. (2023). *Bayesian Optimization*. Cambridge University Press.
- 125 Gramacy, R. B. (2020). *Surrogates: Gaussian Process Modeling, Design and Optimization for*
126 *the Applied Sciences*. Chapman Hall/CRC.
- 127 Harris, C. R., Millman, K. J., Walt, S. J. van der, Gommers, R., Virtanen, P., Cournapeau, D.,
128 Wieser, E., Taylor, J., Berg, S., Smith, N. J., Kern, R., Picus, M., Hoyer, S., Kerkwijk,
129 M. H. van, Brett, M., Haldane, A., Río, J. F. del, Wiebe, M., Peterson, P., ... Oliphant,
130 T. E. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362. <https://doi.org/10.1038/s41586-020-2649-2>
- 132 Kerr, C. C., Stuart, R. M., Mistry, D., Abeyesuriya, R. G., Rosenfeld, K., Hart, G. R., Núñez,
133 R. C., Cohen, J. A., Selvaraj, P., Hagedorn, B., George, L., Jastrzębski, M., Izzo, A. S.,
134 Fowler, G., Palmer, A., Delport, D., Scott, N., Kelly, S. L., Bennette, C. S., ... Klein, D. J.
135 (2021). Covasim: An agent-based model of COVID-19 dynamics and interventions. *PLOS*
136 *Computational Biology*, 17(7), e1009149. <https://doi.org/10.1371/journal.pcbi.1009149>
- 137 Kinser, J. M. (2022). *Modeling and Simulation in Python*. CRC Press.
- 138 Lazaridis, I., Alpaslan-Roodenberg, S., Acar, A., Açikkol, A., Agelarakis, A., Aghikyan, L.,
139 Akyüz, U., Andreeva, D., Andrijašević, G., Antonović, D., Armit, I., Atmaca, A., Avetisyan,
140 P., Aytek, A. İ., Bacvarov, K., Badalyan, R., Bakardzhiev, S., Balen, J., Bejko, L., ... Reich,
141 D. (2022). The genetic history of the Southern Arc: A bridge between West Asia and
142 Europe. *Science*, 377(6609), eabm4247. <https://doi.org/10.1126/science.abm4247>
- 143 Matthews, A. G. de G., Wilk, M. van der, Nickson, T., Fujii, K., Boukouvalas, A., Le{\o}n-
144 Villagr{a}, P., Ghahramani, Z., & Hensman, J. (2017). GPflow: A Gaussian Process
145 Library using TensorFlow. *Journal of Machine Learning Research*, 18(40), 1–6. <http://jmlr.org/papers/v18/16-537.html>
- 147 Morales, J. L., & Nocedal, J. (2011). Remark on “algorithm 778: L-BFGS-B: Fortran subrou-
148 tines for large-scale bound constrained optimization.” *ACM Transactions on Mathematical*
149 *Software*, 38(1), 7:1–7:4. <https://doi.org/10.1145/2049662.2049669>
- 150 O’Gara, D., Rosenblatt, S. F., Hébert-Dufresne, L., Purcell, R., Kasman, M., & Hammond, R.
151 A. (2023). TRACE-Omicron: Policy Counterfactuals to Inform Mitigation of COVID-19
152 Spread in the United States. *Advanced Theory and Simulations*, 2300147. <https://doi.org/10.1002/adts.202300147>
- 154 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z.,

- 155 Gimelshein, N., Antiga, L., Desmaison, A., Köpf, A., Yang, E., DeVito, Z., Raison, M.,
156 Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., ... Chintala, S. (2019). *PyTorch: An*
157 *Imperative Style, High-Performance Deep Learning Library*. arXiv. [https://doi.org/10.](https://doi.org/10.48550/arXiv.1912.01703)
158 [48550/arXiv.1912.01703](https://doi.org/10.48550/arXiv.1912.01703)
- 159 Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel,
160 M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau,
161 D., Brucher, M., Perrot, M., & Duchesnay, E. (2011). Scikit-learn: Machine Learning in
162 Python. *Journal of Machine Learning Research*, 12, 2825–2830.
- 163 Picheny, V., Wagner, T., & Ginsbourger, D. (2013). A benchmark of kriging-based infill criteria
164 for noisy optimization. *Structural and Multidisciplinary Optimization*, 48(3), 607–626.
165 <https://doi.org/10.1007/s00158-013-0919-4>
- 166 Pinder, T., & Dodd, D. (2022). GPJax: A Gaussian Process Framework in JAX. *Journal of*
167 *Open Source Software*, 7(75), 4455. <https://doi.org/10.21105/joss.04455>
- 168 Shattock, A. J., Le Rutte, E. A., Dünner, R. P., Sen, S., Kelly, S. L., Chitnis, N., & Penny, M.
169 A. (2022). Impact of vaccination and non-pharmaceutical interventions on SARS-CoV-2
170 dynamics in Switzerland. *Epidemics*, 38, 100535. [https://doi.org/10.1016/j.epidem.2021.](https://doi.org/10.1016/j.epidem.2021.100535)
171 [100535](https://doi.org/10.1016/j.epidem.2021.100535)
- 172 Silverman, B. W. (1985). Some Aspects of the Spline Smoothing Approach to Non-Parametric
173 Regression Curve Fitting. *Journal of the Royal Statistical Society: Series B (Methodological)*,
174 47(1), 1–21. <https://doi.org/10.1111/j.2517-6161.1985.tb01327.x>
- 175 Surjanovic, S., & Bingham, D. (n.d.). *Virtual Library of Simulation Experiments: Test*
176 *Functions and Datasets*.
- 177 Ushey, K., Allaire, J. J., & Tang, Y. (2023). *Reticulate: Interface to 'Python'*. [https:](https://CRAN.R-project.org/package=reticulate)
178 [//CRAN.R-project.org/package=reticulate](https://CRAN.R-project.org/package=reticulate)
- 179 Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D.,
180 Burovski, E., Peterson, P., Weckesser, W., Bright, J., Walt, S. J. van der, Brett, M.,
181 Wilson, J., Millman, K. J., Mayorov, N., Nelson, A. R. J., Jones, E., Kern, R., Larson, E.,
182 ... Mulbregt, P. van. (2020). SciPy 1.0: Fundamental algorithms for scientific computing in
183 Python. *Nature Methods*, 17(3), 261–272. <https://doi.org/10.1038/s41592-019-0686-2>