

¹ Raster Tools: An Open Source Toolbox for Raster Processing

³ Fredrick Bunt  ¹, Jesse Johnson  ¹, and John Hogland  ²

⁴ 1 University of Montana, USA 2 Rocky Mountain Research Station, USA 

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: Kanishka B. Narayan 

Reviewers:

- [@nagellette](#)
- [@kanishkan91](#)

Submitted: 10 January 2025

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))

⁵ Summary

The Raster Tools Python package provides an application programmers interface (API) for building raster data processing pipelines. Pipelines built with Raster Tools can efficiently process extremely large raster datasets, including those larger than the system's available memory. Raster Tools can be deployed on systems ranging from small laptops to high performance computing environments. To aid in the construction of raster processing pipelines, Raster Tools provides a consistent API to a suite of scalable and lazy raster processing functions including focal, zonal, clipping, convolution, and distance analysis operations. Raster Tools makes it much easier and more accessible to perform research using today's massive raster datasets.

¹⁵ Statement of Need

Python has become very popular as a language for data processing. This is because it provides a rich, general purpose, data stack, and simple, easy to understand syntax. Because of this, it is a good choice for tackling the expanding data problem. Some key libraries in the stack are Numpy ([Harris et al., 2020](#)), Xarray ([Hoyer & Hamman, 2017](#)), SciPy ([Virtanen et al., 2020](#)), and Dask ([Dask Development Team, 2016](#)). Numpy forms the foundation of Python's stack by providing a fast and powerful N-dimensional (ND) array implementation and API for manipulating them. The Xarray package builds on this by supplying an API for working with labelled ND datasets. This is especially relevant to raster data with labeled dimensions and coordinates. Its data structures can be backed by any array object that implements Numpy's array interface, with Numpy being the default. The SciPy package provides optimized implementations of a wide variety of scientific and engineering algorithms.

The Python data stack also includes packages for working with geospatial data. The two main packages are both Python bindings for GDAL ([GDAL/OGR contributors, 2024](#)). The first is the GDAL project's package of the same name. This package is a thin wrapper around the C/C++ library and can be cumbersome and difficult to use. The second is Rasterio ([Gillies & others, 2013--](#)), which provides a more modern and "Pythonic" API for working with geospatial data, but does not implement many of GDAL's features. A third package, rioxarray ([rioxarray Development Team, 2019](#)), is a plugin for Xarray that allows raster formats to be easily read into Xarray's data structures and for those data structures to use Rasterio operations on the underlying data. Finally, odc-geo ([OpenDataCube, 2025](#)) provides several geospatial tools and algorithms such as grid definitions and reprojection. It also can function as a plugin for Xarray.

Together, these packages and the wider Python data stack can be leveraged to carry out simple and advanced raster processing and geographic information systems (GIS) work. There are some limitations, however. Several merit consideration. First, in the GIS world, null or missing data is typically handled with sentinel values that are outside of the data domain (e.g. -9999), while in much of Python's data stack, this is handled with NaN values. These two paradigms

42 can be difficult to reconcile, and it requires boilerplate code to properly mask data. Second,
43 Numpy does not support parallel processing without writing C extensions, which is difficult
44 at best. Newer packages such as Numba ([Lam et al., 2015](#)) and Numexpr ([McLeod, 2018](#))
45 can help with this, but their scope is still limited. They also require more boilerplate code
46 to be written. Third, if the size of the target dataset approaches or exceeds the available
47 memory, Numpy and everything built on it will exhaust the available memory and fail. This is
48 because Numpy and most of the Python data stack is “in-core.” All the data is in memory
49 when processing occurs. To tackle this problem, more boilerplate code must be written to
50 manually chunk the data and process it in batches.

51 The last two issues can be solved by using the Dask package. Dask provides a relatively simple
52 interface for parallel computing in Python. It allows large arrays to be chunked into smaller
53 pieces and processed in parallel using its schedulers. This allows for out-of-core (OOC) or
54 external memory processing of datasets, where only part of the dataset is loaded into memory
55 at any given time. Dask implements much of Numpy’s API and can therefore be used to
56 replace Numpy arrays in data pipelines, with some boilerplate applied. Xarray and odc-geo
57 have first class support for Dask, while rioxarray has partial support for it.

58 With the packages described above, it is possible to perform both simple and advanced data
59 and GIS processing on large datasets. It can be time consuming and difficult, however, to
60 implement. Boilerplate code is needed to cover compatibility gaps, such as making Rasterio
61 and SciPy functions compatible with Dask. Code is also needed to properly handle missing
62 data and to implement common and advanced GIS operations. There is no package that does
63 all of this for the user.

64 Current State of the Field

65 Several open-source Python packages try to address these needs in the Python ecosystem, but
66 each does so differently and has its own limitations. Some examples are WhiteboxTools ([J.
67 Lindsay, 2018; J. B. Lindsay, 2016](#)), GeoUtils ([G. developers, 2020–2026](#)), and Xarray-Spatial
68 ([X.-S. developers, 2020–2026](#)).

69 WhiteboxTools is a geospatial library providing hundreds of raster and GIS analysis functions
70 through Python bindings. It is compiled and highly performant and capable of processing raster
71 data in parallel. However, it presents several limitations: it lacks integration with the Python
72 data stack and requires data to be fully loaded into memory for processing. Additionally,
73 operations cannot be chained in memory; each tool must read from storage, process the data,
74 and write the output back to storage before the next step can begin. A newer adaptation,
75 Whitebox Workflows for Python, has been released, which allows chaining. It is proprietary
76 freeware with a paid upgrade package, however.

77 Built on top of Rasterio, GeoUtils is a package designed for efficient geospatial analysis with
78 Python operator overloading and NumPy integration. It leverages Python’s multiprocessing
79 module to process rasters in parallel. However, the package has notable limitations: most
80 operations are eager and load the full dataset into memory, and it only implements a few
81 geospatial operations such as reproject and proximity. Additionally, GeoUtils currently lacks
82 compatibility with Dask and offers limited integration with Xarray. Work is ongoing to add
83 Dask support, and an Xarray extension is planned. xDEM is a package from the same authors,
84 built on top of GeoUtils, which provides a suite of digital elevation model (DEM) focused
85 functions.

86 Finally, Xarray-Spatial is a package designed for spatial analysis of rasters, built directly on
87 Xarray data objects. It provides a wide variety of analysis functions and features acceleration via
88 Numba and scalability through Dask support. Additionally, the library offers GPU capabilities
89 by leveraging CuPy ([Okuta et al., 2017](#)) and Numba’s just-in-time (JIT) cuda compilation.
90 However, the package has some limitations, including an inflexible and, sometimes, inconsistent
91 interface where support for Dask, GPU acceleration, and parallel processing varies across

92 different functions. Data type support is also restricted, with most functions converting inputs
 93 to float 32 and using NaNs for missing data, and the library currently lacks some common
 94 raster operations found in other tools.

95 Software Description

96 We introduce the Raster Tools package as a solution to the above problems and as a tool
 97 to build pipelines for processing of both large and small-scale raster datasets. Raster Tools
 98 seeks to bridge the gaps in the Python data stack by taking care of the necessary boilerplate.
 99 It also provides a platform for carrying out GIS processing and analysis by implementing a
 100 suite of common and advanced GIS operations that are OOC compatible. These operations
 101 would normally be pulled from several other packages, implemented from scratch, or some
 102 combination of the two. By taking care of these tasks, Raster Tools enables researchers to
 103 tackle large problems and datasets whose size may have been prohibitive before.

104 At the core of Raster Tools is the Raster class. This class provides several useful features that
 105 help with raster processing such as operator overloading, Numpy universal function (Ufunc) and
 106 aggregation compatibility, and the ability to open or save to any GDAL supported format. With
 107 operator overloading, Raster objects behave much like Numpy arrays and can be used with any
 108 of Python's arithmetic, logical, bitwise, and comparison operators. By implementing Numpy's
 109 Ufunc interface, Raster objects provide compatibility with Numpy's element-wise functions such
 110 as sin, cos, sqrt, exp, etc. The result of calling a Numpy Ufunc on a Raster object will always
 111 be another Raster object. These two features, together, make implementing mathematical
 112 models as easy as writing the equations as expressions in code. All these operations are
 113 missing-data aware. The Raster class implements a lazy evaluation model inherited from Dask.
 114 Computational graphs are constructed during operation chaining, deferring data loading and
 115 processing until an explicit trigger, such as saving to disk or explicit loading, is executed. This
 116 enables the definition of complex pipelines that exceed local memory limits. The code listing
 117 below shows an example of an index being calculated from Landsat 8 bands and saved to disk.

```
import numpy as np
import raster_tools as rts

# Create two Raster objects pointing to the red and near-IR bands
red = rts.Raster("data/landsat_stack.tif").get_bands(4)
nir = rts.Raster("data/landsat_stack.tif").get_bands(5)
# Calculate the MSAVI2 index. This is lazy and no computation has occurred.
msavi2 = (2 * (nir + 1) - np.sqrt((2 * nir + 1) ** 2 - 8 * (nir - red))) / 2
# Save to storage. This starts computation.
msavi2.save("outputs/msavi2.tif", tiled=True, bigtiff="IF_NEEDED", compress="LZW")
```

118 In conjunction with the Raster class, Raster Tools provides an API for raster processing. This
 119 API consists of the following modules:

- 120 ■ clipping: functions for clipping rasters using vector geometries.
- 121 ■ creation: functions for creating rasters based off a reference grid with analogs
 to Numpy's ones_like, zeros_like, etc, as well as for generating data from random
 distributions.
- 122 ■ distance: functions for computing proximity and cost-distance analysis.
- 123 ■ focal: functions for applying focal operations, including convolution, to rasters.
- 124 ■ general: general purpose functions such as raster value remapping.
- 125 ■ line_stats: advanced function for calculating statistics on line-vectors overlayed on a
 raster.
- 126 ■ rasterize: functions for rasterizing vector geometries.
- 127 ■ surface: functions for transforming DEMs, such as aspect, slope, and 3D surface area.
- 128 ■ vector: functions for handling vector data.

132 ■ warp: functions for reprojecting rasters.
133 ■ zonal: functions for computing zonal statistics.

134 All parts of the API are lazy and parallel with one exception. The cost-distance analysis
135 functions are fully in memory and non-parallel due to the difficulty in parallelizing the algorithm.
136 The results returned by the cost-distance analysis functions are Raster objects and thus still
137 compatible with the rest of the API. With the suite of operations provided in the API, it is
138 possible to carry out a wide range of raster analysis and processing tasks.

139 Because the API is lazy, it is possible to chain many operations together before saving results
140 back to disk. This can prevent the need to write an intermediate product after each operation
141 and greatly reduce read/write overhead in processing pipelines as well as storage needs.

142 Compatibility

143 Raster Tools offers a high degree of compatibility with the wider Python data stack. The Raster
144 class has methods and attributes for converting raster objects into Numpy and Dask arrays,
145 Xarray DataArray and Dataset objects, and dask-geopandas GeoDataFrame vectors ([Bossche et al., 2025](#)). This compatibility works in the opposite direction as well. The Raster Tools
146 classes and functions can ingest Numpy and Dask arrays and Xarray objects. This extends to
147 the overloaded operators. It is possible to multiply a Numpy array against a Raster object with
148 matching spatial dimensions and get a new Raster with the expected result. Because of Raster
149 Tools' compatibility, it is possible to easily convert data to an appropriate format for another
150 library and then ingest results back into Raster Tools, or vice versa.

152 Raster Tools also provides the Model Predict API. This API allows a model to be applied to
153 any raster in a lazy fashion, compatible with the rest of the package. Using the Model Predict
154 API allows models not normally compatible with Dask or Raster Tools to be easily folded into
155 pipelines.

156 Scaling

157 To evaluate the performance and scaling of Raster Tools, we benchmarked it in a controlled
158 computing environment. The environment used was a virtual machine with 400 GiB of memory
159 and a variable number of CPU cores. We created several scenarios for benchmarking. To test
160 scaling, we set the CPU count at 2, 4, 8, 16, and 32 and ran all the scenarios for each count.
161 The scenarios are described in the table below. The input rasters for the scenarios, labeled A,
162 B, and C below, were copies of the same 30m digital elevation model (DEM) of the continental
163 US (CONUS). These were single precision (float 32) rasters and roughly 60 GB uncompressed
164 on disk. The timing for the scenarios was done using Python's time module and includes both
165 loading of the input rasters and writing of the result rasters. A chunk size of 256 MiB was
166 used. Testing was performed in Q3-Q4 2025. At the time, Raster Tools was only compatible
167 with Dask's default scheduler, so that was the Dask scheduler used. The results shown here
168 should be taken as a baseline, since the new scheduler is expected to improve performance.

Scenario	Basic Forumula	Description
----------	----------------	-------------

Scenario	Basic Forumula	Description
Raster Math	$A + B + C$	Adding three rasters together
Transcendental Functions	$\sin(A) * \cos(B) / \exp(C)$	Combining three transcendental functions
Focal Mean	<code>focal_stats(A, "mean", 5)</code>	Calculate the focal mean with a 5x5 window
Focal Entropy	<code>focal_stats(A, "entropy", 5)</code>	Calculate the focal entropy with a 5x5 window
Slope	<code>slope(A)</code>	Calculate the slope of a DEM
Curvature	<code>curvature(A)</code>	Calculate the curvature of the DEM
Reproject	<code>reproject(A, new_crs)</code>	Calculate the curvature of the DEM

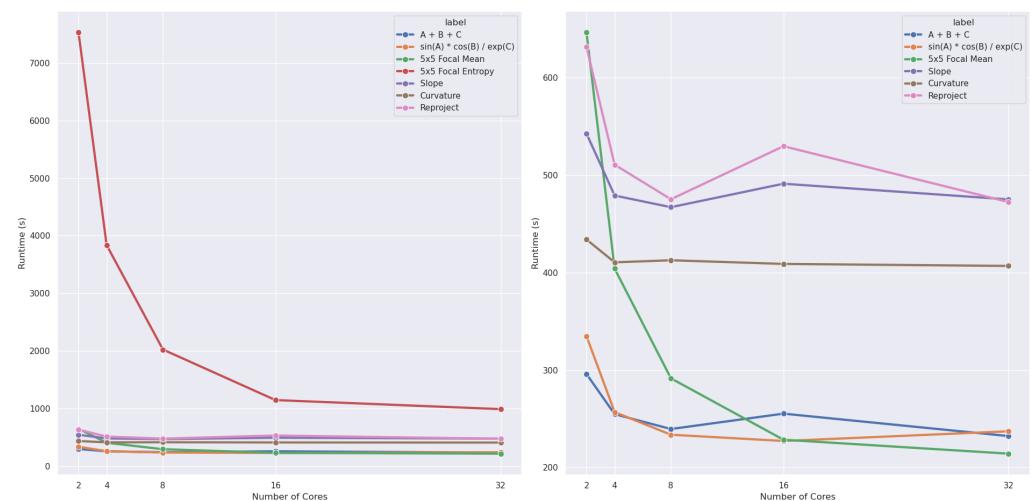


Figure 1: Left) Runtimes for all scenarios. Right) The runtimes for all scenarios, excluding the Focal Entropy scenario.

169 Figure 1 shows the runtimes of the scenarios as the number of cores being used increases.
170 Filtering out the Focal Entropy case for the right-hand plot allows the results for the other
171 scenarios to be viewed more clearly. The Focal Entropy scenario shows strong scaling with the
172 number of cores. Focal Mean shows moderate scaling. The rest show scaling that varies from
173 none to mild.

Use in Research

175 Raster Tools has been used in several research projects. It was used in Johnson et al. (2022)
176 to process 35 years of Monitoring Trends in Burn Severity (MTBS) fire severity raster data for
177 all of Montana into a tabular format for model training. It was used in Lahrichi et al. (2025)
178 to process burn mask rasters for use in training fire spread prediction models. Raster tools
179 was also used to process rasters, rasterize vectors, and complete a wide range of geospatial
180 analysis in Wiard (2025) and Wiard-Greene et al. (2026).

181 Raster Tools is also being used in production by the Forest Service for their Potential Control
182 Locations (PCL) product (Connor et al., 2017), (O'Connor et al., 2022). When initially created,

183 the PCL product was a static, yearly raster for the western half of the United States at 30m
184 resolution. The pipeline for producing PCL was large and time intensive. In 2023 and 2024,
185 the pipeline was completely rebuilt using Raster Tools for its raster processing. The current
186 PCL product is now dynamic and can be produced hourly for local fire support.

187 Limitations

188 Raster Tools has limitations which we acknowledge here. During the benchmarking period
189 (Q3–Q4 2025), Raster Tools was constrained to Dask's default threaded scheduler due to a
190 dependency being incompatible with the newer Distributed scheduler ([Dask developers, 2026](#)).
191 While this incompatibility has since been resolved, the performance data in [Figure 1](#) reflects
192 the overhead inherent to the threaded scheduler. As shown, scheduling overhead dominates
193 the runtime for lighter tasks. However, scenarios involving compute-intensive operations
194 (high computation-to-chunk-size ratio) essentially mask this overhead, demonstrating that the
195 package scales well when the scheduler is not the bottleneck.

196 Therefore, the presented benchmarks represent a performance baseline. Since the Distributed
197 scheduler is generally more performant and handles memory backpressure more effectively,
198 users can expect performance to match or exceed these figures.

199 A related limitation of the threaded scheduler is its prioritization of CPU utilization over
200 memory constraints. Memory consumption is proportional to the number of cores multiplied
201 by the data-chunk size. On systems with a high-core-count, this requires significant available
202 memory. While reducing Dask's target chunk size mitigates this, it increases the number of
203 chunks and adds scheduling overhead.

204 Importantly, Raster Tools is designed to be scheduler and configuration agnostic. It builds the
205 computational graph but does not enforce a specific execution engine or set of configurations.
206 This ensures that users retain full control over Dask configuration, including the ability to
207 switch to the Distributed scheduler. It also means that Raster Tools will automatically benefit
208 from future improvements in the Dask ecosystem.

209 Comparison

210 Except for cost-distance functions, Raster Tools provides a lazy, parallel, OOC, flexible,
211 consistent, and chainable API, along with a suite of raster processing functions, and fully
212 integrates with the wider Python data stack. This contrasts with other packages in the field
213 like those listed above, which force a trade-off in one way or another, between these traits.

214 Conclusion

215 Raster Tools fills the gaps in the Python ecosystem by simplifying raster and GIS workflows and
216 reducing the need for boilerplate code. Its compatibility with widely used libraries and support
217 for lazy, OOC operations make it easier for researchers to process large and complex datasets
218 without reinventing common geospatial functions. By providing a unified API for raster analysis,
219 Raster Tools accelerates the development of data pipelines and opens new opportunities for
220 research projects that previously faced technical or resource limitations. Raster Tools has
221 already proven valuable in real-world studies and operational contexts, and it continues to
222 expand the possibilities for geospatial science and applied research.

223 Bossche, J. V. den, Fleischmann, M., Statham, T., Augspurger, T., (dahn), D. J., Signell,
224 J., Lusk, D., Bunt, F., Gadomski, P., Hagen, R., Bell, R., Lumitz, S., bernardpazio,
225 RichardScottOZ, Morris, M., Miclat, J., Baker, J., Bourbeau, J., Truong, I., ... Zaidi,
226 A. A. (2025). *Geopandas/dask-geopandas: v0.5.0* (Version v0.5.0). Zenodo. <https://doi.org/10.5281/zenodo.15579702>

- 228 Connor, C. D. O., Calkin, D. E., & Thompson, M. P. (2017). An empirical machine learning
229 method for predicting potential fire control locations for pre-fire planning and operational
230 fire management. *International Journal of Wildland Fire*, 26(7), 587–597. <https://doi.org/10.1071/wf16135>
- 232 Dask developers. (2026). *Dask distributed documentation*. <https://distributed.dask.org/>
- 233 Dask Development Team. (2016). *Dask: Library for dynamic task scheduling*. <http://dask.pydata.org>
- 235 developers, G. (2020--2026). *GeoUtils: Consistent geospatial analysis in python*. <https://github.com/GlacioHack/geoutils>
- 237 developers, X.-S. (2020--2026). *Xarray-spatial*. <https://github.com/xarray-contrib/xarray-spatial>
- 239 GDAL/OGR contributors. (2024). *GDAL/OGR geospatial data abstraction software library*.
240 Open Source Geospatial Foundation. <https://doi.org/10.5281/zenodo.5884351>
- 241 Gillies, S., & others. (2013--). *Rasterio: Geospatial raster i/o for Python programmers*.
242 Mapbox. <https://github.com/rasterio/rasterio>
- 243 Harris, C. R., Millman, K. J., Walt, S. J. van der, Gommers, R., Virtanen, P., Cournapeau, D.,
244 Wieser, E., Taylor, J., Berg, S., Smith, N. J., Kern, R., Picus, M., Hoyer, S., Kerkwijk,
245 M. H. van, Brett, M., Haldane, A., Río, J. F. del, Wiebe, M., Peterson, P., ... Oliphant,
246 T. E. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362. <https://doi.org/10.1038/s41586-020-2649-2>
- 248 Hoyer, S., & Hamman, J. (2017). Xarray: N-D labeled arrays and datasets in Python. *Journal
249 of Open Research Software*, 5(1). <https://doi.org/10.5334/jors.148>
- 250 Johnson, J., Marcozzi, A., Bunt, F., Bova, J., & Hogland, J. (2022). Predicting fire severity in
251 montana using a random forest classification scheme. In D. X. Viegas & L. M. Ribeiro
252 (Eds.), *Advances in fire research 2022* (pp. 323–328). University of Coimbra Press.
253 https://doi.org/10.14195/978-989-26-2298-9_51
- 254 Lahrichi, S., Bova, J., Johnson, J., & Malof, J. (2025). *Improved wildfire spread prediction
255 with time-series data and the WSTS+ benchmark*. <https://arxiv.org/abs/2502.12003>
- 256 Lam, S. K., Pitrou, A., & Seibert, S. (2015). Numba: A llvm-based python jit compiler.
257 *Proceedings of the Second Workshop on the LLVM Compiler Infrastructure in HPC*, 1–6.
- 258 Lindsay, J. (2018). *WhiteboxTools user manual*. https://www.whiteboxgeo.com/manual/wbt_book/preface.html
- 260 Lindsay, J. B. (2016). Whitebox GAT: A case study in geomorphometric analysis. *Computers
261 & Geosciences*, 95, 75–84. <http://dx.doi.org/10.1016/j.cageo.2016.07.003>
- 262 McLeod, R. (2018). *NumExpr: Fast numerical expression evaluator for NumPy*. <https://doi.org/10.5281/zenodo.2483274>
- 264 O'Connor, C. D., Haas, J. R., Gannon, B. M., Dunn, C. J., Thompson, M. P., & Calkin, D. E.
265 (2022). Modelling potential control locations: Development and adoption of data-driven
266 analytics to support strategic and tactical wildfire containment decisions. *Environmental
267 Sciences Proceedings*, 17(1). <https://doi.org/10.3390/environsciproc2022017073>
- 268 Okuta, R., Unno, Y., Nishino, D., Hido, S., & Loomis, C. (2017). CuPy: A NumPy-
269 compatible library for NVIDIA GPU calculations. *Proceedings of Workshop on Machine
270 Learning Systems (LearningSys) in the Thirty-First Annual Conference on Neural Information
271 Processing Systems (NIPS)*. http://learningsys.org/nips17/assets/papers/paper_16.pdf
- 272 OpenDataCube. (2025). *Odc-geo*. <https://github.com/opendatacube/odc-geo>
- 273 rioxarray Development Team. (2019). *Rioxarray: Geospatial xarray extension powered by*

- 274 *rasterio*. Corteva, Inc. <https://github.com/corteva/rioxarray>
- 275 Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D.,
276 Burovski, E., Peterson, P., Weckesser, W., Bright, J., van der Walt, S. J., Brett, M., Wilson,
277 J., Millman, K. J., Mayorov, N., Nelson, A. R. J., Jones, E., Kern, R., Larson, E., ... SciPy
278 1.0 Contributors. (2020). SciPy 1.0: Fundamental Algorithms for Scientific Computing in
279 Python. *Nature Methods*, 17, 261–272. <https://doi.org/10.1038/s41592-019-0686-2>
- 280 Wiard, L. A. (2025). *Investigating the impact of aerial firefighting on rate of wildfire spread*
281 [Master's thesis, University of Montana; University of Montana]. <https://scholarworks.umt.edu/etd/12531>
- 283 Wiard-Greene, L., Johnson, J., Hogland, J., Bunt, F., & Bova, J. (2026). Investigating the
284 impact of aerial firefighting on rate of wildfire spread. *Fire*, 9(1). <https://doi.org/10.3390/fire9010002>

DRAFT