

# <sup>1</sup> DSWL package: a Python implementation of the <sup>2</sup> Debiased Spatial Whittle Likelihood

<sup>3</sup> Arthur P. Guillaumin  <sup>1</sup>, Thomas Goodwin  <sup>5</sup>, Olivia Walbert  <sup>4</sup>, Adam M.  
<sup>4</sup> Sykulski  <sup>2</sup>, Sofia C. Olhede  <sup>3</sup>, and Frederik J. Simons  <sup>4</sup>

<sup>5</sup> 1 Queen Mary University of London, United Kingdom 2 Imperial College London, United Kingdom 3

<sup>6</sup> Ecole Polytechnique Fédérale de Lausanne, Switzerland 4 Princeton University, United States of America

<sup>7</sup> 5 School of Economics, University of New South Wales, Australia

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

**Editor:** George K. Thiruvathukal 



## Reviewers:

- [@MarineChap](#)
- [@weiji14](#)

**Submitted:** 13 March 2025

**Published:** unpublished

## License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/))

## <sup>8</sup> Summary

<sup>9</sup> The Debiased Spatial Whittle Likelihood (DSWL) package is an open-source Python package  
<sup>10</sup> that implements the eponymous paper ([Guillaumin et al., 2022](#)). The methodology allows users  
<sup>11</sup> to efficiently infer the parameters of stationary / homogeneous spatial and spatio-temporal  
<sup>12</sup> covariance models for univariate or multivariate processes from gridded data with potential  
<sup>13</sup> missing observations, e.g. due to natural boundaries. It leverages the Fast Fourier Transform,  
<sup>14</sup> and therefore can benefit from further computational gains through GPU implementations  
<sup>15</sup> offered by PyTorch ([Paszke et al., 2019](#)) or Cupy ([Okuta et al., 2017](#)), both made available  
<sup>16</sup> within the package as alternative backends to Numpy ([Harris et al., 2020](#)). As such, DSWL  
<sup>17</sup> on GPU allows to fit covariance models to data observed on grids with tens of millions of  
<sup>18</sup> locations.

## <sup>19</sup> Statement of need

<sup>20</sup> Describing patterns of spatial and spatio-temporal covariance is of interest to practitioners in a  
<sup>21</sup> wide range of applied sciences such as geosciences, meteorology or climate science. Stationary  
<sup>22</sup> covariance modelling allows for a first-order approximation of the covariance structure, and  
<sup>23</sup> leads to many practical applications such as kriging ([Stein, 1999](#)) and forecasting via the  
<sup>24</sup> conditional Gaussian multivariate distribution. The inference of parameters for a physics-based  
<sup>25</sup> covariance model can also be of interest in its own right.

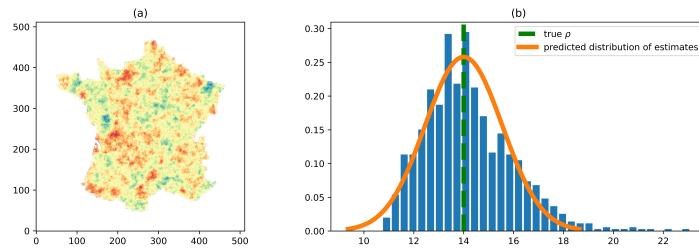
<sup>26</sup> A major hurdle in spatio-temporal modelling is the computational cost of the Gaussian  
<sup>27</sup> likelihood function. This is particularly relevant for modern spatio-temporal datasets, from  
<sup>28</sup> physics simulations to real-world data. The computational burden of parameter inference also  
<sup>29</sup> arises from complex spatio-temporal covariance models with a large number of parameters  
<sup>30</sup> which typically require a high number of likelihood evaluations during the optimization process  
<sup>31</sup> or when running an MCMC sampler.

<sup>32</sup> A common means to circumvent this computational burden is to use approximations to the  
<sup>33</sup> Gaussian likelihood. Among these methods, the Whittle likelihood is a standard spectral domain  
<sup>34</sup> method for gridded data. Along its computational benefits, the Whittle likelihood provides  
<sup>35</sup> robustness to departures from Gaussianity and allows to restrict the second-order model to a  
<sup>36</sup> specific range of spatio-temporal frequencies. However, for spatial and spatio-temporal data  
<sup>37</sup> where the dimension  $d$  is greater than 2, the standard Whittle likelihood suffers from a large  
<sup>38</sup> bias and typically does not allow for missing observations.

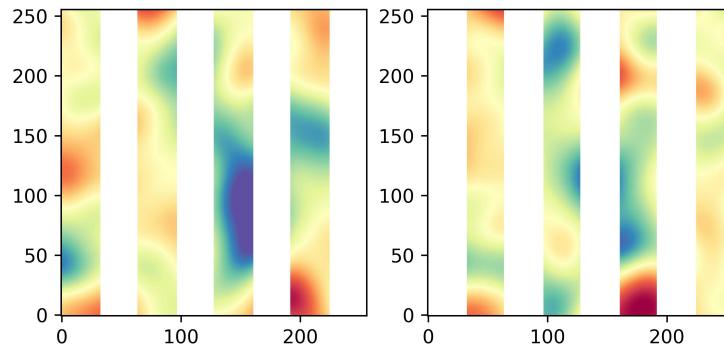
<sup>39</sup> DSWL is a Python implementation of the Debiased Spatial Whittle likelihood ([Guillaumin  
<sup>40</sup> et al., 2022](#)), a method that addresses the bias of the Whittle likelihood ([Sykulski et al.,](#)

41     2019). Although it requires gridded data as it relies on the Fast Fourier Transform, the  
 42     implemented method additionally allows for missing observations, making it amenable to  
 43     practical applications where a full hyperrectangle of data measurements might not be available.  
 44     Missing observations might occur due to natural boundaries or due to measurement constraints.  
 45     As an example, in Figure 1 we show a simulated sample from an exponential covariance model  
 46     observed on a domain with the shape of metropolitan France (sans Corsica), along with the  
 47     distribution of estimates obtained from 1000 independent samples generated from the same  
 48     model and the predicted distribution of estimates, which can be used to build confidence  
 49     intervals.

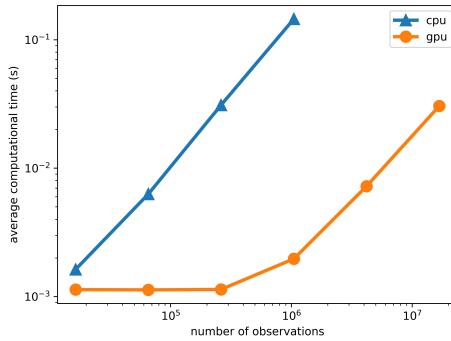
50     The package allows the user to treat multivariate data, including those cases where the  
 51     missingness patterns might differ between the variates. For instance, in Figure 2 we show a  
 52     realization of a bivariate random field with distinct patterns of missing observations between the  
 53     two variates, from which we can still infer the parameters of the model, such as the correlation  
 54     between the two fields. The code base also includes tapering, the use of which can further  
 55     alleviate boundary effects (Dahlhaus & Künsch, 1987). Finally, the user can switch between  
 56     several computing backends, Numpy, Cupy and PyTorch. This allows to further benefit from  
 57     computational gains via GPU implementations of the Fast Fourier Transform. In practice,  
 58     we observe computational GPU-versus-CPU speed-ups of order  $\times 10$  up to order  $\times 100$  as  
 59     reported in Figure 3 (CPU: Intel(R) Xeon(R) Platinum 8268 CPU @ 2.90GHz, GPU: NVIDIA  
 60     A100-PCIE 40GB, numpy 1.26.4, cupy 13.4.0).



**Figure 1:** A simulated sample from an exponential covariance kernel observed on a domain with the shape of metropolitan France (sans Corsica) (a), along with the distribution of estimates obtained from 1000 independent realizations from the same model with range parameter  $\rho = 14$  spatial units (b)



**Figure 2:** An example of a bivariate random field with distinct patterns of missing observations



**Figure 3:** Computational time of the Debiased Spatial Whittle Likelihood averaged over 1000 samples on square grids of increasing sizes, compared between CPU and GPU (Cupy backend)

Other approximation techniques are available for the inference of spatio-temporal covariance models. Among those, we can mention Vecchia-type likelihood approximations (Katzfuss & Guinness, 2021) implemented e.g. by Jurek (2023), Katzfuss (2023) and Joseph Guinness (2023), and covariance tapering (Kaufman et al., 2008), although for the latter we are not aware of open-source implementations.

## 66 Software structure

67 The software is organized around several modules that can be grouped into the following  
68 categories:

- 69     ▪ grids and sampling:
  - 70         – grids.py: this module is used to define the rectangular grids where the data sit  
71             via the class RectangularGrid. A mask of zeros (missing) and ones (not missing)  
72             can be set to specify potential missing observations, for instance to account for  
73             natural boundaries
  - 74         – simulation.py: this module allows to efficiently sample a realization from a model  
75             on a grid via circulant embedding (Dietrich & Newsam, 1997).
- 76     ▪ models:
  - 77         – models.py: this module allows to define a covariance model. Standard covariance  
78             models are pre-defined, such as the exponential covariance model, the squared  
79             exponential (Gaussian) covariance model and the Matérn covariance model. These  
80             standard covariance models can also be combined (e.g. via summation) to form  
81             more complex covariance models.
- 82     ▪ estimation:
  - 83         – periodogram.py: this module allows to compute the periodogram of the data, and  
84             to obtain the expected periodogram for a given model and grid combination.
  - 85         – multivariate\_periodogram.py: this module allows to compute the periodogram  
86             for multivariate data.
  - 87         – likelihood.py: this module allows to define the Debiased Whittle Likelihood and  
88             the corresponding estimator. The optimizer can be selected among those offered  
89             by the optimize package of the Scipy library (Virtanen et al., 2020).

90 A documentation including example notebooks is available, and issues can be raised on [Github](#).  
91 Example notebooks can also be run directly in the browser via [mybinder.org](#).

## 92 Acknowledgements

93 This research utilised Queen Mary University of London's Apocrita HPC facility, supported by  
94 QMUL Research-IT. doi:10.5281/zenodo.438045. In particular, this research made use of the  
95 OnDemand portal ([Hudak et al., 2018](#)).

## 96 References

- 97 Dahlhaus, R., & Künsch, H. (1987). Edge effects and efficient parameter estimation for  
98 stationary random fields. *Biometrika*, 74(4), 877–882. <https://doi.org/10.1093/biomet/74.4.877>
- 100 Dietrich, C. R., & Newsam, G. N. (1997). Fast and Exact Simulation of Stationary Gaussian  
101 Processes through Circulant Embedding of the Covariance Matrix. *SIAM Journal on  
102 Scientific Computing*, 18(4), 1088–1107. <https://doi.org/10.1137/S1064827592240555>
- 103 Guillaumin, A. P., Sykulski, A. M., Olhede, S. C., & Simons, F. J. (2022). The Debiased  
104 Spatial Whittle Likelihood. *Journal of the Royal Statistical Society Series B: Statistical  
105 Methodology*, 84(4), 1526–1557. <https://doi.org/10.1111/rssb.12539>
- 106 Harris, C. R., Millman, K. J., Walt, S. J. van der, Gommers, R., Virtanen, P., Cournapeau, D.,  
107 Wieser, E., Taylor, J., Berg, S., Smith, N. J., Kern, R., Picus, M., Hoyer, S., Kerkwijk,  
108 M. H. van, Brett, M., Haldane, A., Río, J. F. del, Wiebe, M., Peterson, P., ... Oliphant,  
109 T. E. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362. <https://doi.org/10.1038/s41586-020-2649-2>
- 111 Hudak, D., Johnson, D., Chalker, A., Nicklas, J., Franz, E., Dockendorf, T., & McMichael, B.  
112 L. (2018). Open OnDemand: A web-based client portal for HPC centers. *Journal of Open  
113 Source Software*, 3(25), 622. <https://doi.org/10.21105/joss.00622>
- 114 Joseph Guinness, Y. F., Matthias Katzfuss. (2023). *GpGp: Fast gaussian process computation  
115 using vecchia's approximation*. CRAN. <https://cran.r-project.org/web/packages/GpGp/index.html>
- 117 Jurek, M. (2023). pyMRA: Multi-resolution approximation for gaussian processes. In *GitHub  
118 repository*. GitHub. <https://github.com/marcinjurek/pyMRA>
- 119 Katzfuss, M. (2023). GPvecchia: Fast gaussian-process inference using vecchia approximations.  
120 In *GitHub repository*. GitHub. <https://github.com/katzfuss-group/GPvecchia>
- 121 Katzfuss, M., & Guinness, J. (2021). A General Framework for Vecchia Approximations of  
122 Gaussian Processes. *Statistical Science*, 36(1). <https://doi.org/10.1214/19-STS755>
- 123 Kaufman, C. G., Schervish, M. J., & Nychka, D. W. (2008). Covariance Tapering for  
124 Likelihood-Based Estimation in Large Spatial Data Sets. *Journal of the American Statistical  
125 Association*, 103(484), 1545–1555. <https://doi.org/10.1198/016214508000000959>
- 126 Okuta, R., Unno, Y., Nishino, D., Hido, S., & Loomis, C. (2017). CuPy: A NumPy-  
127 compatible library for NVIDIA GPU calculations. *Proceedings of Workshop on Machine  
128 Learning Systems (LearningSys) in the Thirty-First Annual Conference on Neural Information  
129 Processing Systems (NIPS)*. [http://learningsys.org/nips17/assets/papers/paper\\_16.pdf](http://learningsys.org/nips17/assets/papers/paper_16.pdf)
- 130 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T.,  
131 Lin, Z., Gimelshein, N., Antiga, L., & others. (2019). PyTorch: An imperative style,  
132 high-performance deep learning library. *Advances in Neural Information  
133 Processing Systems*, 8024–8035. <https://proceedings.neurips.cc/paper/2019/file/bdbca288fee7f92f2bfa9f7012727740-Paper.pdf>
- 135 Stein, M. L. (1999). *Interpolation of spatial data: Some theory for kriging*. Springer New  
136 York, NY. <https://doi.org/10.1007/978-1-4612-1494-6>

<sup>137</sup> Sykulski, A. M., Olhede, S. C., Guillaumin, A. P., Lilly, J. M., & Early, J. J. (2019). The debiased  
<sup>138</sup> Whittle likelihood. *Biometrika*, 106(2), 251–266. <https://doi.org/10.1093/biomet/asy071>

<sup>139</sup> Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D.,  
<sup>140</sup> Burovski, E., Peterson, P., Weckesser, W., Bright, J., & others. (2020). SciPy 1.0:  
<sup>141</sup> Fundamental algorithms for scientific computing in python. *Nature Methods*, 17(3),  
<sup>142</sup> 261–272. <https://doi.org/10.1038/s41592-019-0686-2>

DRAFT