

graph-pes: graph-based machine-learning models for potential-energy surfaces

John L. A. Gardner¹ and Volker L. Deringer¹

¹ Inorganic Chemistry Laboratory, Department of Chemistry, University of Oxford, Oxford OX1 3QR, United Kingdom

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [↗](#)

Submitted: 22 June 2025

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

We present graph-pes, an open-source toolkit for accelerating the development, training, and deployment of machine-learned interatomic potential (MLIP) models that act on graph representations of atomic structures. The graph-pes toolkit comprises three components:

- The graph-pes Python package:** a modular framework containing all functionality required to build, train, and evaluate graph-based MLIPs. The package includes a mature data pipeline for converting atomic structures into graph representations (AtomicGraphs), a fully featured base class for MLIP implementations (GraphPESModel), and a suite of common data manipulation routines and model building blocks. We provide independent (re-) implementations of common MLIP architectures out-of-the-box, as well as interfaces to several foundational MLIP models (see below).
- The graph-pes-train command-line interface (CLI):** a convenience tool for training graph-based MLIPs on datasets of labelled atomic structures directly from the command line. The tool is compatible with any GraphPESModel (i.e., those defined in graph-pes, user-designed ones, and foundation models) and is designed to be easily extensible via custom loss functions, optimisers, datasets, and more.
- Molecular-dynamics drivers** for popular MD engines that allow any GraphPESModel to be used in GPU-accelerated MD simulations. We currently provide a pair style for use in LAMMPS (Thompson et al., 2022), a GraphPESCalculator for use in ASE (Larsen et al., 2017), and an integration with the torch-sim package (Gangan et al., 2025).

Statement of need

In recent years, machine-learned PES models, commonly referred to as machine-learned interatomic potentials (MLIPs), have become central tools for computational chemistry and materials science (Deringer et al., 2019).

These models are trained on labels generated by quantum-mechanical methods, but scale much more favourably with system size, making it possible to simulate the dynamics of large systems (millions of atoms and more) over extended timescales. In this way, MLIPs are facilitating the study of complex physical and chemical phenomena at the atomic scale, in turn driving the generation of novel insight and understanding.

Many flavours of MLIPs exist, and with them have arisen a variety of software packages that are typically tailored to training specific architectures (see examples below). Given their unique specialisations, these individual software implementations do not normally conform to a common interface, making it difficult for practitioners to migrate their training and validation pipelines between different model architectures.

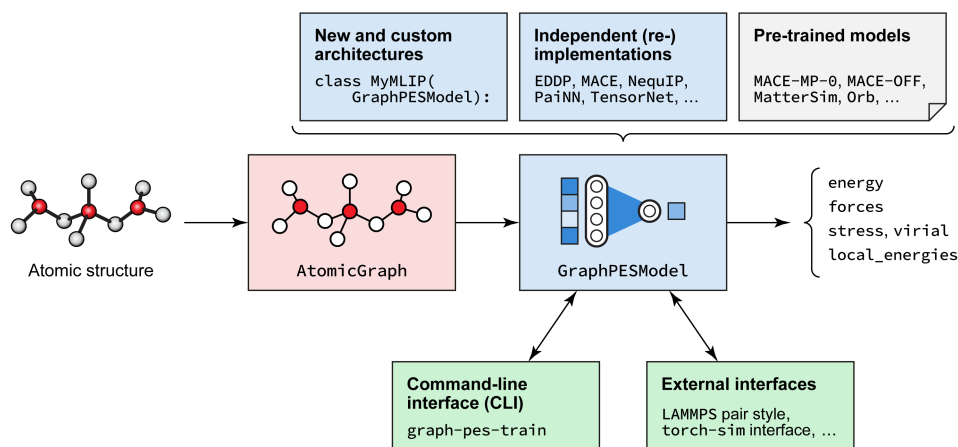


Figure 1: Schematic overview of the functionality of graph-pes. The core components are highlighted in colour. Red: The AtomicGraph class is used to represent atomic structures and incorporates the notion of locality via a neighbour list. Blue: The GraphPESModel class is the general base class for all models defined using graph-pes. We provide independent, stand-alone (re-) implementations of popular architectures, together with interfaces allowing access to several pre-trained and “foundational” MLIP models. Custom, user-defined MLIPs are easy to create, and are fully compatible with the rest of graph-pes’s functionality. Green: graph-pes includes a CLI for easy training, and interfaces to multiple external simulation tools for evaluating MLIP models.

graph-pes provides a **unified interface and software framework** for defining, training, and working with graph-based MLIP models. This reduces the barrier to entry for researchers wanting to implement new MLIP architectures, and allows practitioners to easily explore different MLIP architectures for their specific use cases: training scripts require as little as one line of code to swap between model architectures, while validation scripts can be written in an architecture-agnostic manner, with LAMMPS input scripts, ASE calculators, and torch-sim simulations requiring no changes other than pointing to a different model file.

Features and implementation

Below, we briefly summarise the key design choices, components, and features of graph-pes, and further emphasise the advantages of having a unified interface for all MLIP architectures. For an extended overview of the graph-pes framework, and comprehensive documentation, please visit [this URL](#).

Representing atomic structures with graphs

An atomic structure containing N atoms is completely defined by the positions of its atoms ($\mathbf{R} \in \mathbb{R}^{N \times 3}$) and their chemical identities ($Z \in \mathbb{Z}_+^N$).¹ A graph representation of the atomic structure incorporates this complete description, together with an edge list ($\mathbf{E} \in \mathbb{Z}^{E \times 2}$) indicating which atoms are within the local environment of others (defined, for instance, using a fixed cut-off radius). The resulting graph, $G = \{\mathbf{R}, Z, \mathbf{E}\}$, is thus an extremely general representation of chemical structure with a built-in definition of locality.

We therefore define the AtomicGraph class as the base data structure in graph-pes, and provide convenience methods to convert these to and from ase.Atoms objects (Hjorth Larsen et al., 2017).

¹This statement is assuming that the structure is isolated: defining a periodic structure requires the trivial addition of a unit cell and periodic boundary conditions.

Writing performant code to implement common graph-based operations can be challenging: we therefore provide optimised implementations to access many derived properties (such as `number_of_atoms`, `neighbour_distances`, and `number_of_neighbours`) as well as to perform common graph-based operations (such as `index_over_neighbours`, `sum_over_neighbours`, and `sum_per_structure`). All of these functions work for both single and batched graph instances, simplifying the implementation of new MLIP models, and making their forwards passes easily readable.

Model implementations

All MLIP models in `graph-pes` are implemented as subclasses of the `GraphPESModel` base class, which itself inherits from the `torch.nn.Module` class. These models take an (optionally batched) `AtomicGraph` as input, and are able to return a collection of PES property predictions, including the total energy, atomic forces, and cell stress tensors.

Implementations need only define a forward pass that returns a local energy for each atom in the graph, or a total energy for the entire structure; we use the functionality from `torch.autograd` to automatically calculate force, and stress tensors in a conservative manner (Paszke et al., 2019). For faster modelling, we also fully support models that return direct force and stress tensor predictions (e.g., `TensorNet` or `orb-v3-*` with their optional direct force readout heads).

Building on the `GraphPESModel` class, we provide independent (re-) implementations of popular MLIP architectures, including `PaiNN` (Schütt et al., 2021), `EDDP` (Pickard, 2022), `NequIP` (Batzner et al., 2022), `MACE` (Batatia et al., 2022), and `TensorNet` (Simeon & de Fabritiis, 2023). We use building blocks provided by the `e3nn` (Geiger & Smidt, 2022) package to implement models that act on spherical tensor decompositions.

Furthermore, we provide an `AdditionModel` implementation, which makes energy, force, and stress predictions as a sum over several independent models. This allows `graph-pes` to add the following features onto any other model architecture:

- **Offset energies.** A common feature of quantum-mechanical labelling methods is that the “reference energy” of an isolated atom is (i) non-zero, (ii) different for each element, and (iii) varying between different levels of theory/method. We provide the `EnergyOffset` model to account for this, removing the need to include these contributions in other model implementations within `graph-pes`.
- **Pair repulsions.** Adding smooth, short-ranged, repulsive pair repulsion contributions on top of many-bodied model predictions guarantees correct model behaviour in the short-range limit, and can act to stabilise MD simulations. We provide the `LennardJones`, `Morse`, and `ZBLCoreRepulsion`, all with smoothed cutoffs and optionally learnable parameters, to trivially add these repulsions to any other model implementation.

Training and validation

We provide the `graph-pes-train` CLI tool for training any `GraphPESModel` on datasets of labelled atomic structures. Configuration for this tool is specified via a hierarchically-structured YAML file, with separate sections for specifying the model, data, loss, and fitting parameters. As well as training from scratch, we also support the loading of pre-trained models, allowing for fine-tuning of existing models on new datasets. In this manner, a wide variety of pre-train/fine-tune strategies are supported, including synthetic pre-training (Gardner et al., 2024), foundation model fine-tuning (see below), and frozen transfer learning (Radova et al., 2025).

Under the hood, `graph-pes-train` builds upon the PyTorch Lightning (Falcon & The PyTorch Lightning team, 2019) training loop, allowing for a wide range of advanced features, including learning rate scheduling, stochastic weight averaging, gradient clipping, and more. By using the `data2objects` package (Gardner, 2024) to parse configuration files, we also support the

109 use of arbitrary, user-defined components, including custom loss functions, model architectures,
110 optimisers, and datasets.

111 Because all models conform to the same interface, all training features can be used with any
112 model architecture. Similarly, all downstream model uses can be written in an architecture-
113 agnostic manner, allowing for MD, relaxations, and other scripts to be written once, and then
114 used with any MLIP architecture, for example in defining extended validation beyond simple
115 error metrics (Morrow et al., 2023).

116 Easy access to foundation models

117 A topical and recent area of research is the development of universal or “foundational”
118 MLIPs that can describe the potential-energy surface of a wide range of systems. graph-pes
119 integrates directly with the mace-torch, mattersim, and orb-models packages to provide
120 access to, among others, the MACE-MP (Batatia et al., 2024), MatterSim (Yang et al., 2024),
121 orb-v2 (Neumann et al., 2024), MACE-OFF (Kovács et al., 2025), Egret-v1 (Mann et al., 2025),
122 and orb-v3 (Rhodes et al., 2025) families of models. Each of these integrations generates
123 GraphPESModels that are directly compatible with all graph-pes features, including fine-tuning,
124 validation pipelines, and MD simulations.

125 Related work

126 graph-pes is beginning to drive a substantial number of projects within our research group,
127 and we hope that it will be useful to many others. In recent preprints, we have described the
128 use of graph-pes for fitting NequIP models to datasets created using the autoplex software
129 (Liu et al., 2024), for assessing the zero-shot performance of different graph-network MLIP
130 models (Mahmoud et al., 2025), and for fine-tuning and distilling atomistic foundation models
131 (Gardner et al., 2025).

132 Relevant alternative packages that offer training and validation functionality for *specific* ML-
133 PES architectures include: schnetpack (Schütt et al., 2019, 2023), deepmd-kit (Wang et al.,
134 2018; Zeng et al., 2023), nequip (Batzner et al., 2022), mace-torch (Batatia et al., 2022),
135 torchmd-net (Pelaez et al., 2024), and fairchem (Shuaibi et al., 2025). The MatterTune
136 package (Kong et al., 2025) provides a unified interface for fine-tuning atomistic foundation
137 models, but does not create models with a common interface, or allow for training arbitrary
138 MLIP architectures from scratch.

139 Acknowledgements

140 We thank Zoé Faure Beaulieu, Krystian Gierczak, and Daniel Thomas du Toit for early testing
141 and feedback. J.L.A.G. acknowledges a UKRI Linacre - The EPA Cephalosporin Scholarship,
142 support from an EPSRC DTP award [grant number EP/T517811/1], and from the Department
143 of Chemistry, University of Oxford. This work was supported by UK Research and Innovation
144 [grant number EP/X016188/1].

145 References

- 146 Batatia, I., Benner, P., Chiang, Y., Elena, A. M., Kovács, D. P., Riebesell, J., Advincula,
147 X. R., Asta, M., Avaylon, M., Baldwin, W. J., Berger, F., Bernstein, N., Bhowmik, A.,
148 Blau, S. M., Cărare, V., Darby, J. P., De, S., Della Pia, F., Deringer, V. L., ... Csányi, G.
149 (2024). *A foundation model for atomistic materials chemistry* (No. arXiv:2401.00096).
150 <https://doi.org/10.48550/arXiv.2401.00096>

- 151 Batatia, I., Kovacs, D. P., Simm, G. N. C., Ortner, C., & Csanyi, G. (2022, October). MACE:
152 Higher Order Equivariant Message Passing Neural Networks for Fast and Accurate Force
153 Fields. *Advances in Neural Information Processing Systems*.
- 154 Batzner, S., Musaelian, A., Sun, L., Geiger, M., Mailoa, J. P., Kornbluth, M., Molinari,
155 N., Smidt, T. E., & Kozinsky, B. (2022). E(3)-equivariant graph neural networks for
156 data-efficient and accurate interatomic potentials. *Nature Communications*, 13(1), 2453.
157 <https://doi.org/10.1038/s41467-022-29939-5>
- 158 Deringer, V. L., Caro, M. A., & Csányi, G. (2019). Machine Learning Interatomic Potentials
159 as Emerging Tools for Materials Science. *Advanced Materials*, 31(46), 1902765. <https://doi.org/10.1002/adma.201902765>
- 161 Falcon, W., & The PyTorch Lightning team. (2019). *PyTorch Lightning* (Version 1.4).
162 <https://doi.org/10.5281/zenodo.3828935>
- 163 Gangan, A. S., Cohen, O. A., Riebesell, J., Goodall, R., Kolluru, A., & Falletta, S. (2025).
164 *Torch-Sim* (Version 0.1.0). <https://doi.org/10.5281/zenodo.7486816>
- 165 Gardner, J. L. A. (2024). data2objects: Self-documenting config files for use with python. In
166 *GitHub repository*. GitHub. <https://github.com/jla-gardner/data2objects>
- 167 Gardner, J. L. A., Baker, K. T., & Deringer, V. L. (2024). Synthetic pre-training for neural-
168 network interatomic potentials. *Machine Learning: Science and Technology*, 5(1), 015003.
169 <https://doi.org/10.1088/2632-2153/ad1626>
- 170 Gardner, J. L. A., Toit, D. F. T. du, Mahmoud, C. B., Beaulieu, Z. F., Juraskova, V., Paşca,
171 L.-B., Rosset, L. A. M., Duarte, F., Martelli, F., Pickard, C. J., & Deringer, V. L. (2025).
172 *Distillation of atomistic foundation models across architectures and chemical domains* (No.
173 arXiv:2506.10956). <https://doi.org/10.48550/arXiv.2506.10956>
- 174 Geiger, M., & Smidt, T. (2022). *E3nn: Euclidean Neural Networks* (No. arXiv:2207.09453).
175 <https://doi.org/10.48550/arXiv.2207.09453>
- 176 Hjorth Larsen, A., Jørgen Mortensen, J., Blomqvist, J., Castelli, I. E., Christensen, R., Duřak,
177 M., Friis, J., Groves, M. N., Hammer, B., Hargus, C., Hermes, E. D., Jennings, P. C., Bjerre
178 Jensen, P., Kermode, J., Kitchin, J. R., Leonhard Kolsbjerg, E., Kubal, J., Kaasbjerg, K.,
179 Lysgaard, S., ... Jacobsen, K. W. (2017). The atomic simulation environment—a Python
180 library for working with atoms. *Journal of Physics: Condensed Matter*, 29(27), 273002.
181 <https://doi.org/10.1088/1361-648X/aa680e>
- 182 Kong, L., Shoghi, N., Hu, G., Li, P., & Fung, V. (2025). *MatterTune: An Integrated, User-
183 Friendly Platform for Fine-Tuning Atomistic Foundation Models to Accelerate Materials
184 Simulation and Discovery* (No. arXiv:2504.10655). <https://doi.org/10.48550/arXiv.2504.10655>
- 186 Kovács, D. P., Moore, J. H., Browning, N. J., Batatia, I., Horton, J. T., Pu, Y., Kapil, V.,
187 Witt, W. C., Magdău, I.-B., Cole, D. J., & Csányi, G. (2025). *MACE-OFF: Transferable
188 Short Range Machine Learning Force Fields for Organic Molecules* (No. arXiv:2312.15211).
189 <https://doi.org/10.48550/arXiv.2312.15211>
- 190 Larsen, A. H., Mortensen, J. J., Blomqvist, J., Castelli, I. E., Christensen, R., Duřak, M.,
191 Friis, J., Groves, M. N., Hammer, B., Hargus, C., Hermes, E. D., Jennings, P. C.,
192 Jensen, P. B., Kermode, J., Kitchin, J. R., Kolsbjerg, E. L., Kubal, J., Kaasbjerg, K.,
193 Lysgaard, S., ... Jacobsen, K. W. (2017). The atomic simulation environment—a Python
194 library for working with atoms. *Journal of Physics: Condensed Matter*, 29(27), 273002.
195 <https://doi.org/10.1088/1361-648X/aa680e>
- 196 Liu, Y., Morrow, J. D., Ertural, C., Fragapane, N. L., Gardner, J. L. A., Naik, A. A., Zhou, Y.,
197 George, J., & Deringer, V. L. (2024). *An automated framework for exploring and learning
198 potential-energy surfaces* (No. arXiv:2412.16736). <https://doi.org/10.48550/arXiv.2412.16736>

199 16736

- 200 Mahmoud, C. B., El-Machachi, Z., Gierczak, K. A., Gardner, J. L. A., & Deringer, V. L. (2025).
201 *Assessing zero-shot generalisation behaviour in graph-neural-network interatomic potentials*
202 (No. arXiv:2502.21317). <https://doi.org/10.48550/arXiv.2502.21317>
- 203 Mann, E. L., Wagen, C. C., Vandezande, J. E., Wagen, A. M., & Schneider, S. C. (2025).
204 *Egret-1: Pretrained Neural Network Potentials For Efficient and Accurate Bioorganic*
205 *Simulation* (No. arXiv:2504.20955). <https://doi.org/10.48550/arXiv.2504.20955>
- 206 Morrow, J. D., Gardner, J. L. A., & Deringer, V. L. (2023). How to validate machine-
207 learned interatomic potentials. *The Journal of Chemical Physics*, 158(12), 121501. <https://doi.org/10.1063/5.0139611>
- 209 Neumann, M., Gin, J., Rhodes, B., Bennett, S., Li, Z., Choubisa, H., Hussey, A., & Godwin,
210 J. (2024). *Orb: A Fast, Scalable Neural Network Potential* (No. arXiv:2410.22570).
211 <https://doi.org/10.48550/arXiv.2410.22570>
- 212 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z.,
213 Gimelshein, N., Antiga, L., Desmaison, A., Kopf, A., Yang, E., DeVito, Z., Raison, M.,
214 Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., ... Chintala, S. (2019). PyTorch: An
215 Imperative Style, High-Performance Deep Learning Library. *Advances in Neural Information*
216 *Processing Systems*, 32.
- 217 Pelaez, R. P., Simeon, G., Galvelis, R., Mirarchi, A., Eastman, P., Doerr, S., Thölke, P.,
218 Markland, T. E., & Fabritiis, G. D. (2024). *TorchMD-net 2.0: Fast neural network*
219 *potentials for molecular simulations*. <https://arxiv.org/abs/2402.17660>
- 220 Pickard, C. J. (2022). Ephemeral data derived potentials for random structure search. *Physical*
221 *Review B*, 106(1), 014102. <https://doi.org/10.1103/PhysRevB.106.014102>
- 222 Radova, M., Stark, W. G., Allen, C. S., Maurer, R. J., & Bartók, A. P. (2025). *Fine-tuning*
223 *foundation models of materials interatomic potentials with frozen transfer learning* (No.
224 arXiv:2502.15582). <https://doi.org/10.48550/arXiv.2502.15582>
- 225 Rhodes, B., Vandenhaute, S., Šimkus, V., Gin, J., Godwin, J., Duignan, T., & Neumann,
226 M. (2025). *Orb-v3: Atomistic simulation at scale* (No. arXiv:2504.06231). <https://doi.org/10.48550/arXiv.2504.06231>
- 228 Schütt, K. T., Hessmann, S. S. P., Gebauer, N. W. A., Lederer, J., & Gastegger, M. (2023).
229 SchNetPack 2.0: A neural network toolbox for atomistic machine learning. *The Journal of*
230 *Chemical Physics*, 158(14), 144801. <https://doi.org/10.1063/5.0138367>
- 231 Schütt, K. T., Kessel, P., Gastegger, M., Nicoli, K. A., Tkatchenko, A., & Müller, K.-R. (2019).
232 SchNetPack: A Deep Learning Toolbox For Atomistic Systems. *Journal of Chemical Theory*
233 *and Computation*, 15(1), 448–455. <https://doi.org/10.1021/acs.jctc.8b00908>
- 234 Schütt, K. T., Unke, O. T., & Gastegger, M. (2021). *Equivariant message passing for*
235 *the prediction of tensorial properties and molecular spectra* (No. arXiv:2102.03150).
236 <https://doi.org/10.48550/arXiv.2102.03150>
- 237 Shuaibi, M., Das, A., Sriram, A., Misko, Barroso-Luque, L., Gao, R., Goyal, S., Ulissi, Z.,
238 Wood, B., Xie, T., Yoon, J., Wander, B., Kolluru, A., Barnes, R., Sunshine, E., Tran,
239 K., Xiang, Levine, D., Shoghi, N., ... Michel, K. (2025). *FAIRChem* (Version 2.2.0).
240 <https://doi.org/10.5281/zenodo.15587498>
- 241 Simeon, G., & de Fabritiis, G. (2023). *TensorNet: Cartesian Tensor Representations for*
242 *Efficient Learning of Molecular Potentials* (No. arXiv:2306.06482). <https://arxiv.org/abs/2306.06482>
- 244 Thompson, A. P., Aktulga, H. M., Berger, R., Bolintineanu, D. S., Brown, W. M., Crozier, P.
245 S., in 't Veld, P. J., Kohlmeyer, A., Moore, S. G., Nguyen, T. D., Shan, R., Stevens, M. J.,

- 246 Tranchida, J., Trott, C., & Plimpton, S. J. (2022). LAMMPS - a flexible simulation tool for
247 particle-based materials modeling at the atomic, meso, and continuum scales. *Computer*
248 *Physics Communications*, 271, 108171. <https://doi.org/10.1016/j.cpc.2021.108171>
- 249 Wang, H., Zhang, L., Han, J., & E, W. (2018). DeePMD-kit: A deep learning package for
250 many-body potential energy representation and molecular dynamics. *Computer Physics*
251 *Communications*, 228, 178–184. <https://doi.org/10.1016/j.cpc.2018.03.016>
- 252 Yang, H., Hu, C., Zhou, Y., Liu, X., Shi, Y., Li, J., Li, G., Chen, Z., Chen, S., Zeni, C., Horton,
253 M., Pinsler, R., Fowler, A., Zügner, D., Xie, T., Smith, J., Sun, L., Wang, Q., Kong, L., ... Lu,
254 Z. (2024). *MatterSim: A Deep Learning Atomistic Model Across Elements, Temperatures*
255 *and Pressures* (No. arXiv:2405.04967). <https://doi.org/10.48550/arXiv.2405.04967>
- 256 Zeng, J., Zhang, D., Lu, D., Mo, P., Li, Z., Chen, Y., Rynik, M., Huang, L., Li, Z., Shi, S.,
257 Wang, Y., Ye, H., Tuo, P., Yang, J., Ding, Y., Li, Y., Tisi, D., Zeng, Q., Bao, H., ... Wang,
258 H. (2023). DeePMD-kit v2: A software package for deep potential models. *The Journal of*
259 *Chemical Physics*, 159(5), 054801. <https://doi.org/10.1063/5.0155600>

DRAFT