

# HiGP: A high-performance Python package for Gaussian Processes

Hua Huang<sup>1\*</sup>, Tianshi Xu<sup>2\*</sup>, Yuanzhe Xi<sup>2</sup>, and Edmond Chow<sup>1¶</sup>

<sup>1</sup> School of Computational Science, Engineering, Georgia Institute of Technology, USA <sup>2</sup> Department of Mathematics, Emory University, USA ¶ Corresponding author \* These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- Review
- Repository
- Archive

Editor: James Gaboardi

## Reviewers:

- @Edenhofer
- @davidogara
- @Naeemkh

Submitted: 04 July 2025

Published: unpublished

## License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License (CC BY 4.0).

## Summary

Gaussian Processes (GPs) (Rasmussen & Williams, 2005) are flexible, nonparametric Bayesian models widely used for regression and classification because of their ability to capture complex data patterns and quantify predictive uncertainty. However, the  $\mathcal{O}(n^3)$  computational cost of kernel matrix operations poses a major obstacle to applying GPs at scale. HiGP is a high-performance Python package designed to overcome these scalability limitations through advanced numerical linear algebra and hierarchical kernel representations. It integrates  $\mathcal{H}^2$  matrices to achieve near-linear complexity in both storage and computation for spatial datasets, supports on-the-fly kernel evaluation to avoid explicit storage in large-scale problems, and incorporates a robust Adaptive Factorized Nyström (AFN) preconditioner (Zhao et al., 2024) that accelerates convergence of iterative solvers across a broad range of kernel spectra. These computational kernels are implemented in C++ for maximum performance and exposed through Python interfaces, enabling seamless integration with modern machine learning workflows. HiGP also includes analytically derived gradient computations for efficient hyperparameter optimization, avoiding the inefficiencies of automatic differentiation in iterative solvers. By serving as a reusable numerical engine, HiGP complements existing GP frameworks such as GPJax (Pinder & Dodd, 2022), KeOps (Charlier et al., 2021), and GaussianProcesses.jl (Fairbrother et al., 2022), providing a reliable and scalable computational backbone for large-scale Gaussian Process regression and classification.

## Gaussian Processes

For training points  $\mathbf{X} \in \mathbb{R}^{n \times d}$ , a noisy training observation set  $\mathbf{y} \in \mathbb{R}^n$ , and testing points  $\mathbf{X}_* \in \mathbb{R}^{m \times d}$ , a standard GP model assumes that the noise-free testing observations  $\mathbf{y}_* \in \mathbb{R}^m$  follow a joint Gaussian distribution that depends on a set of parameters, including scale  $f$ , noise level  $s$ , and kernel parameters  $l$ . The GP model finds the optimal parameters  $\Theta := (s, f, l)$  by minimizing the negative log marginal likelihood:

$$L(\Theta) = \frac{1}{2} \left( \mathbf{y}^\top \widehat{\mathbf{K}}^{-1} \mathbf{y} + \log |\widehat{\mathbf{K}}| + n \log 2\pi \right),$$

where  $\widehat{\mathbf{K}}$  denotes the regularized kernel matrix. An optimization process usually requires the gradient of  $L(\Theta)$ :

$$\frac{\partial L}{\partial \theta} = \frac{1}{2} \left( -\mathbf{y}^\top \widehat{\mathbf{K}}^{-1} \frac{\partial \widehat{\mathbf{K}}}{\partial \theta} \widehat{\mathbf{K}}^{-1} \mathbf{y} + \text{tr} \left( \widehat{\mathbf{K}}^{-1} \frac{\partial \widehat{\mathbf{K}}}{\partial \theta} \right) \right), \quad \theta \in \Theta.$$

Using preconditioned iterative methods with preconditioner  $\mathbf{M} \approx \widehat{\mathbf{K}}$  is a common option (Aune et al., 2014; Chen et al., 2023; Hensman et al., 2013; Pleiss et al., 2018; Wenger et al., 2022;

35 Wilson et al., 2015; Zhang et al., 2024). In this approach,  $\widehat{\mathbf{K}}^{-1}\mathbf{y}$  is approximated via the  
 36 preconditioned conjugate gradient (PCG) method (Saad, 2003). To handle the logarithmic  
 37 determinant and trace terms, they are first rewritten as

$$\log |\widehat{\mathbf{K}}| = \log |\mathbf{M}| + \log |\mathbf{M}^{-1/2} \widehat{\mathbf{K}} \mathbf{M}^{-1/2}|, \quad (1)$$

$$\text{tr}(\widehat{\mathbf{K}}^{-1} \frac{\partial \widehat{\mathbf{K}}}{\partial \theta}) = \text{tr} \left( \mathbf{M}^{-1} \frac{\partial \mathbf{M}}{\partial \theta} \right) + \text{tr} \left( \widehat{\mathbf{K}}^{-1} \frac{\partial \widehat{\mathbf{K}}}{\partial \theta} - \mathbf{M}^{-1} \frac{\partial \mathbf{M}}{\partial \theta} \right). \quad (2)$$

38 The second component of each new expression is then estimated using the stochastic Lanczos  
 39 quadrature (Ubaru et al., 2017) and the Hutchinson estimator [Hutchinson (1989); Meyer:2021],  
 40 respectively.

## 41 Statement of Need

42 The Gaussian Process (GP) community has advanced rapidly in recent years, developing  
 43 scalable inference frameworks and more efficient kernel representations. Modern libraries such  
 44 as GPyTorch (Gardner et al., 2018), GPflow (Matthews et al., 2017; van der Wilk et al.,  
 45 2020), GPJax (Pinder & Dodd, 2022), KeOps (Charlier et al., 2021), and GaussianProcesses.jl  
 46 (Fairbrother et al., 2022) leverage GPUs and automatic differentiation to perform GP inference  
 47 efficiently on moderately large datasets. Concurrently, new algorithms, including preconditioned  
 48 optimization methods (Wenger et al., 2022), alternating-projection solvers (Wu et al., 2024),  
 49 GPU-accelerated Vecchia approximations for spatial data (Pan et al., 2024), robust relevance-  
 50 pursuit inference (Ament et al., 2024), and latent Kronecker formulations for structured  
 51 covariance matrices (Lin et al., 2025), have further improved the scalability and robustness  
 52 of GP models. Yet, most existing frameworks emphasize modeling flexibility and seamless  
 53 integration with autodiff ecosystems, rather than optimizing the low-level numerical routines  
 54 that dominate runtime for very large or ill-conditioned kernel systems. HiGP is designed to  
 55 address this computational gap by focusing on the numerical core of GP inference. It provides  
 56 robust, scalable, and hardware-efficient implementations of kernel algebra, preconditioned  
 57 iterative solvers, and gradient computations, offering three primary contributions.

58 Firstly, HiGP addresses the efficiency of MatVec, the most performance-critical operation in  
 59 iterative methods. For large 2D or 3D datasets, the dense kernel matrix is compressed into a  
 60  $\mathcal{H}^2$  matrix (Hackbusch et al., 2000; Hackbusch & Börm, 2002) in HiGP, resulting in  $\mathcal{O}(n)$   
 61 storage and computation costs. For large high-dimensional datasets, HiGP computes small  
 62 kernel matrix blocks on-the-fly and immediately uses them in MatVec and discards them, which  
 63 allows HiGP to handle extremely large datasets with a moderate memory size.

64 Secondly, HiGP uses iterative solvers with the newly proposed AFN preconditioner (Zhao  
 65 et al., 2024), which is designed for robust preconditioning of kernel matrices. Experiments  
 66 demonstrate that AFN can significantly improve the accuracy and robustness of iterative solvers  
 67 for kernel matrix systems. Furthermore, AFN and  $\mathcal{H}^2$  matrix computation rely on evaluating  
 68 many small kernel matrices in parallel, which is easily handled in C++ but would incur large  
 69 overhead in Python, making implementation in other libraries such as GPyTorch or GPFlow  
 70 more challenging.

71 Lastly, HiGP uses accurate and efficient hand-coded gradient calculations. GPyTorch relies  
 72 on the automatic differentiation (autodiff) provided in PyTorch to calculate gradients (Equa-  
 73 tion 2). However, autodiff can be inefficient and inaccurate for computing the gradient of  
 74 the preconditioner, so we use hand-coded gradient calculations for better performance and  
 75 accuracy.

76 The HiGP documentation<sup>1</sup> provides a comparison of the accuracy and performance of HiGP  
 77 and GPyTorch.

<sup>1</sup><https://github.com/huanghua1994/HiGP/blob/main/docs/5-Performance-tests.md>

## Design and Implementation

We implemented HiGP in Python 3 and C++ with the goal of providing both a set of ready-to-use out-of-the-box Python interfaces for regular users and a set of reusable high-performance shared-memory multithreading computational primitives for advanced users. The HiGP C++ code implements all performance-critical operations. The HiGP Python code wraps the C++ units into four basic Python modules: `krnlmatmodule` for computing kernel matrices and its derivatives, `precondmodule` for PCG solver with AFN preconditioner, `gprproblemmodule` and `gpcproblemmodule` for computing the the loss and gradient for GP regression and classification. The two modules `gprproblemmodule` and `gpcproblemmodule` allow a user to train a GP model with any gradient-based optimizer.

We further implemented two high-level modules `GPRModel` and `GPCModel` using PyTorch parameter registration and optimizer to simplify the training and use of GP models. Listing 1 shows an example of defining and training a GP regression and using the trained model for prediction.

```
# Listing 1: HiGP example code of training and using a GPR model
gprproblem = hign.gprproblem.setup(data=train_x, label=train_y,
                                   kernel_type=hign.GaussianKernel)

model = hign.GPRModel(gprproblem)
optimizer = torch.optim.Adam(model.parameters(), lr=0.1)
for i in ranges(max_steps):
    loss = model.calc_loss_grad()
    optimizer.step()
params = model.get_params()
pred = hign.gpr_prediction(train_x, train_y, test_x,
                           hign.GaussianKernel, params)
```

We note that the HiGP Python interfaces (except for `GPRModel` and `GPCModel` models) are *stateless*. This design aims to simplify the interface and decouple different operations. A user can train and use different GP models with the same or different data and configurations in the same file.

## References

- Ament, S., Santorella, E., Eriksson, D., Letham, B., Balandat, M., & Bakshy, E. (2024). Robust gaussian processes via relevance pursuit. *The Thirty-Eighth Annual Conference on Neural Information Processing Systems*. <https://openreview.net/forum?id=5FATPIIWUJ>
- Aune, E., Simpson, D. P., & Eidsvik, J. (2014). Parameter estimation in high dimensional Gaussian distributions. *Statistics and Computing*, 24(2), 247–263. <https://doi.org/10.1007/s11222-012-9368-y>
- Charlier, B., Feydy, J., Glaunès, J. A., Collin, F.-D., & Durif, G. (2021). Kernel operations on the GPU, with autodiff, without memory overflows. *Journal of Machine Learning Research*, 22(74), 1–6.
- Chen, L., Chen, T., Detha, U., & Andersen, M. S. (2023). Towards scalable kernel-based regularized system identification. *2023 62nd IEEE Conference on Decision and Control (CDC)*, 1498–1504. <https://doi.org/10.1109/CDC49753.2023.10384051>
- Fairbrother, J., Nemeth, C., Rischard, M., Brea, J., & Pinder, T. (2022). GaussianProcesses.jl: A nonparametric bayes package for the julia language. *Journal of Statistical Software*, 102(1). <https://doi.org/10.18637/jss.v102.i01>
- Gardner, J., Pleiss, G., Weinberger, K. Q., Bindel, D., & Wilson, A. G. (2018). GPyTorch: Blackbox matrix-matrix Gaussian process inference with GPU acceleration. *Advances in*

- 114 *Neural Information Processing Systems*, 31.
- 115 Hackbusch, W., & Börm, S. (2002). Data-sparse approximation by adaptive  $\mathcal{H}^2$ -matrices.  
116 *Computing*, 69(1), 1–35.
- 117 Hackbusch, W., Khoromskij, B. N., & Sauter, S. A. (2000). On  $\mathcal{H}^2$ -matrices. In *Lectures on*  
118 *applied mathematics: Proceedings of the symposium organized by the sonderforschungsbereich 438 on the occasion of karl-heinz hoffmann's 60th birthday, munich, june 30 - july*  
119 *1, 1999* (pp. 9–29). Max Planck Institute for Mathematics in the Sciences; Springer.
- 120
- 121 Hensman, J., Fusi, N., & Lawrence, N. D. (2013). *Gaussian processes for big data*. 282–290.  
122 <https://dl.acm.org/doi/10.5555/3023638.3023667>
- 123 Hutchinson, M. F. (1989). A stochastic estimator of the trace of the influence matrix for  
124 laplacian smoothing splines. *Communications in Statistics - Simulation and Computation*,  
125 18(3), 1059–1076. <https://doi.org/10.1080/03610918908812806>
- 126 Lin, J. A., Ament, S., Balandat, M., Eriksson, D., Hernández-Lobato, J. M., & Bakshy, E.  
127 (2025). Scalable gaussian processes with latent kronecker structure. *Forty-Second Interna-*  
128 *tional Conference on Machine Learning*. <https://openreview.net/forum?id=Nv70EgUAA7>
- 129 Matthews, A. G. de G., van der Wilk, M., Nickson, T., Fujii, Keisuke., Boukouvalas, A.,  
130 León-Villagrà, P., Ghahramani, Z., & Hensman, J. (2017). GPflow: A Gaussian process  
131 library using TensorFlow. *Journal of Machine Learning Research*, 18(40), 1–6.
- 132 Pan, Q., Abdulah, S., Genton, M. G., Keyes, D. E., Ltaief, H., & Sun, Y. (2024). GPU-  
133 accelerated vecchia approximations of gaussian processes for geospatial data using batched  
134 matrix computations. *arXiv Preprint arXiv:2403.07412*.
- 135 Pinder, T., & Dodd, D. (2022). GPJax: A gaussian process framework in JAX. *Journal of*  
136 *Open Source Software*, 7(75), 4455. <https://doi.org/10.21105/joss.04455>
- 137 Pleiss, G., Gardner, J., Weinberger, K., & Wilson, A. G. (2018). Constant-time predictive  
138 distributions for Gaussian processes. *Proceedings of the 35th International Conference on*  
139 *Machine Learning*, 80, 4114–4123.
- 140 Rasmussen, C. E., & Williams, C. K. I. (2005). *Gaussian processes for machine learning*. The  
141 MIT Press. <https://doi.org/10.7551/mitpress/3206.001.0001>
- 142 Saad, Y. (2003). *Iterative methods for sparse linear systems*. Society for Industrial; Applied  
143 Mathematics. <https://doi.org/10.1137/1.9780898718003>
- 144 Ubaru, S., Chen, J., & Saad, Y. (2017). Fast estimation of  $\text{tr}(f(A))$  via stochastic lanczos  
145 quadrature. *SIAM Journal on Matrix Analysis and Applications*, 38(4), 1075–1099. <https://doi.org/10.1137/16M1104974>  
146
- 147 van der Wilk, M., Dutordoir, V., John, S., Artemev, A., Adam, V., & Hensman, J. (2020). A  
148 *framework for interdomain and multioutput Gaussian processes*. [https://arxiv.org/abs/](https://arxiv.org/abs/2003.01115)  
149 [2003.01115](https://arxiv.org/abs/2003.01115)
- 150 Wenger, J., Pleiss, G., Hennig, P., Cunningham, J. P., & Gardner, J. R. (2022). Preconditioning  
151 for scalable gaussian process hyperparameter optimization. *Proceedings of the 39th*  
152 *International Conference on Machine Learning*, 23751–23780.
- 153 Wilson, A. G., Dann, C., & Nickisch, H. (2015). *Thoughts on massively scalable Gaussian*  
154 *processes*.
- 155 Wu, K., Meinshausen, N., Cunningham, J. P., & Pozzi, M. (2024). Large-scale gaussian  
156 processes via alternating projection. *Proceedings of the 27th International Conference on*  
157 *Artificial Intelligence and Statistics*, 2810–2818.
- 158 Zhang, M., Chen, T., & Mu, B. (2024). Asymptotic properties of generalized maximum  
159 likelihood hyper-parameter estimator for regularized system identification. *2024 IEEE*

160 63rd Conference on Decision and Control (CDC), 4961–4966. [https://doi.org/10.1109/](https://doi.org/10.1109/CDC56724.2024.10885816)  
161 [CDC56724.2024.10885816](https://doi.org/10.1109/CDC56724.2024.10885816)

162 Zhao, S., Xu, T., Huang, H., Chow, E., & Xi, Y. (2024). An adaptive factorized Nyström  
163 preconditioner for regularized kernel matrices. *SIAM Journal on Scientific Computing*,  
164 46(4), A2351–A2376. <https://doi.org/10.1137/23M1565139>

DRAFT