

¹ Kleinkram: Open Robotic Data Management

² Cyril Püntener  ^{1*}, Johann Schwabe  ^{1*}, Dominique Garmier^{1†}, Jonas
³ Frey  ^{1,2†}, and Marco Hutter 

⁴ 1 Robotic Systems Lab, ETH Zurich, Switzerland 2 Max Planck Institute for Intelligent Systems,
⁵ Tübingen, Germany  Corresponding author * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review ↗](#)
- [Repository ↗](#)
- [Archive ↗](#)

Editor: Sébastien Boisgérault  ¹⁰

Reviewers:

- [@boisgera](#)
- [@c-joly](#)

Submitted: 16 June 2025

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#))¹⁹

⁶ Summary

⁷ Data is key to advancing robotic perception, navigation, locomotion, and reasoning. However,
⁸ managing diverse robotic datasets presents unique challenges, including scalability, quality
⁹ assurance, and seamless integration into diverse workflows. To address this gap, we introduce
¹⁰ Kleinkram, a free and open-source data management system tailored for robotics research.
¹¹ Kleinkram enables efficient storage, indexing, and sharing of datasets, from small experiments
¹² to large-scale collections. It supports essential workflows like validation, curation, and
¹³ benchmarking via customizable compute actions. Designed with a user-friendly web interface,
¹⁴ CLI integration, and scalable deployment options, Kleinkram empowers researchers to streamline
¹⁵ data management and accelerate robotics innovation.

¹⁶ Statement of need

¹⁷ To render robotic data useful for research, it is essential to store, organize, and index the
data in a way that makes it easily searchable and shareable. While large corporations have
developed internal tools, the broader robotics community often relies on fragmented solutions.
¹⁸ Current solutions typically fall into three categories, none of which fully satisfy academic
research needs. (1) Commercial SaaS platforms (e.g., Foxglove, Roboto AI) offer excellent
indexing and visualisations but are often closed-source and cost-prohibitive for academia. (2)
Cloud-native reference architectures (e.g., AWS ADDF, Azure DataOps) provide scalability but
require complex, vendor-locked infrastructure setups that are overkill for single labs. Finally, (3)
legacy open-source tools (e.g., Marv Robotics) or raw storage (S3, Google Drive) either lack
modern CI/CD integration or offer no content-based indexing, forcing researchers to manually
download large bags for inspection. Thus, a gap remains for an openly available, ready-to-use,
and easy-to-deploy solution exists for the robotics research community. Additionally, features
such as data verification and the ability to perform tailored compute jobs on newly generated
datasets are highly desirable, as they facilitate benchmarking, reproducibility and algorithmic
development.

¹⁹ To address these challenges, we introduce **Kleinkram**, a self-hosted web service designed for
scalable and efficient data management. Unlike traditional cloud storage, Kleinkram natively
integrates compute capabilities, automates data transfer, and eliminates the tedious manual
effort typically associated with data management workflows. By categorizing and structuring
data around common robotics use cases, Kleinkram facilitates the creation of large, diverse
datasets that can be easily shared and reused across multiple projects. Its intuitive web interface
ensures accessibility, while a command-line interface (CLI) enables seamless integration into
automated pipelines and headless systems. Kleinkram supports widely adopted standards,
building on ROS1 and ROS2 message definitions, and offers native compatibility with ROSbag
and MCAP data formats.

42 Data Structure

43 Kleinkram is designed around the typical data generation process in mobile robotics, assuming
44 data is collected and stored primarily in ROS1/ROS2-compatible ROSbag or MCAP file formats.
45 Once data recording for an experiment is complete, these files can be efficiently uploaded
46 and ingested into the Kleinkram system for centralised storage, indexing, and subsequent
47 post-processing. It is important to note that the current version of Kleinkram focuses on post-
48 recording and data management. It does not support real-time data streaming or processing
49 on the fly.

50 To provide structure and facilitate organization and retrieval, Kleinkram requires data to be
51 organized according to a strict three-layer hierarchy: Projects, Missions, and Files. Each layer
52 maintains a one-to-many relationship with the layer below it. While users have flexibility in
53 how they map their specific activities to this structure, the intended model is that a Project
54 represents a distinct research project, which requires data storage. A Mission corresponds to a
55 single, self-contained experiment or data collection run conducted within that project, and it
56 contains all the individual data Files (like ROSbag or MCAP files) recorded during that specific
57 deployment. This structured approach aids in navigating, managing, and understanding large
58 volumes of experimental data.

59 System Architecture

60 Kleinkram's system architecture is modular, comprising several interacting microservices.

- 61 ■ **Python Client Library and CLI** Provides programmatic access to Kleinkram's
62 functionalities, enabling efficient data transfer operations (upload, download) directly
63 from Python scripts or the command line. This allows seamless integration into robotic
64 workflows and automated data pipelines running on robots or workstations, removing
65 the need for manual browser interaction for data transfer.

66 The CLI is built using the typer library, sharing a core Python codebase with the client
67 library.

- 68 ■ **Web interface** Serves as the primary graphical user interface for users to interact with
69 Kleinkram. It allows for the browsing, managing and structuring files and their metadata.
70 It is implemented as a single-page application using the Vue3 framework and the Quasar
71 component library.

- 72 ■ **Backend API** Acts as the central communication layer between the client applications
73 (web UI and Python client/CLI) and the data storage and processing components. It
74 handles authentication, data indexing, metadata management, and schedules background
75 tasks.

76 The backend is implemented using the NestJS framework and utilises a PostgreSQL
77 database for storing all metadata related to projects, missions, files, users, and actions.

- 78 ■ **Data Store** The raw robotic data files (ROSbags, MCAPs) are stored on an S3-compatible
79 object storage backend. This provides scalability and flexibility. Users can easily deploy
80 and manage their own storage using self-hosted solutions like S3-compatible storage
81 (e.g. SeaweedFS), or utilise public cloud S3 services. Kleinkram interacts with the data
82 store via the S3 API.

- 83 ■ **Action Runner** This component enables the execution of customisable data processing and
84 analysis tasks directly on the data stored within Kleinkram. Users can define "Actions"
85 (e.g., validate data integrity, extract sensor metadata, generate preview visualisations,
86 convert formats, run benchmarking scripts).

87 These actions are packaged as Docker containers. The action runner orchestrates the
88 execution of these containers, providing them access to the necessary data from the data
89 store using the client library or CLI.

- 90 ▪ **Observability (Optional)** Monitoring and logging system performance, resource usage,
91 and task execution status are crucial for managing a scalable data system. Integration
92 with observability tools, such as the Grafana Stack (Loki for logs, Prometheus for metrics,
93 Tempo for traces, Grafana for dashboards) can provide insights into the system's health
94 and the progress of the data processing task.

95 Usecase

96 Kleinkram has been used internally at the Robotic Systems Lab at ETH Zurich over the past
97 year. During this time, it has stored over 20 TB of data collected from various robotic systems,
98 effectively replacing the lab's previous reliance on Google Drive for data storage. The largest
99 project supported by Kleinkram was the **GrandTour dataset** (Frey et al., 2025), in which
100 our legged robot **ANYmal** (Hutter et al., 2016), equipped with **Boxi** (Frey et al., 2025), a
101 multi-sensor payload, was deployed across various locations in Switzerland.

102 Following each data collection mission, raw data — recorded in the form of ROSbags and
103 MCAP files — was uploaded directly to Kleinkram via its command-line interface (CLI).
104 The intuitive Docker-based action integration allowed us to easily define and execute data
105 verification tests. These include, for example, checks to ensure that all sensor streams were
106 recorded at the expected frequencies and correct time synchronization was established during
107 the distributed recordings, as well as common sense checking for validity of data (e.g. images
108 are not black or the IMUs measure the gravity vector).

109 Beyond data verification, Kleinkram enabled us to run full SLAM pipelines retrospectively,
110 automatically producing standard Absolute and Relative Trajectory Error (ATE/RTE) metrics.
111 This compute integration was critical for development, benchmarking, and evaluation.

112 Equally important was Kleinkram's user-friendly CLI, which provided quick access to summary
113 statistics such as dataset counts, durations, and other key metrics — many of which were
114 directly used in associated publications. Given that for our use case, data has to be mainly
115 accessed within the ETH network infrastructure; datasets can be pulled on demand and deleted
116 afterwards, fully utilising the fast on-premise interconnect infrastructure without relying on
117 external servers.

118 Throughout the project, Kleinkram also enforced metadata submission during upload. Users
119 were required to include a YAML file describing the mission, which captured essential information
120 such as the robot operator, specific hardware configuration, location, and links to related
121 resources (e.g., associated Google Drive folders for images). This structured metadata was
122 essential for organizing and retrieving data at scale.

123 Acknowledgements

124 This work was primarily supported by the Open Research Data Grant at ETH Zurich. Jonas
125 Frey is supported by the Max Planck ETH Center for Learning Systems.

126 This work was supported and partially funded by Leica Geosystems, which is part of Hexagon.
127 In addition, this work was supported by the National Centre of Competence in Research
128 Robotics (NCCR Robotics), the ETH RobotX research grant funded through the ETH Zurich
129 Foundation, the European Union's Horizon 2020 research and innovation program under grant
130 agreement No 101016970, No 101070405, and No 101070596, and an ETH Zurich Research
131 Grant No. 21-1 ETH-27.

¹³² We extend our sincere appreciation to Noel Kampus for the initial design of the web interface,
¹³³ to Lars Leuthold and Marvin Lichtsteiner for their valuable contributions, and to William
¹³⁴ Talbot and Turcan Tuna for their efforts in the internal testing of Kleinkram.

¹³⁵ **References**

- ¹³⁶ Frey, J., Tuna, T., Fu, L. F. T., Weibel, C., Patterson, K., Krummenacher, B., Müller, M.,
¹³⁷ Nubert, J., Fallon, M., Cadena, C., & Hutter, M. (2025). Boxi: Design decisions in the
¹³⁸ context of algorithmic performance for robotics. *Proceedings of Robotics: Science and
¹³⁹ Systems*. <https://doi.org/10.48550/arXiv.2504.18500>
- ¹⁴⁰ Hutter, M., Gehring, C., Jud, D., Lauber, A., Bellicoso, C. D., Tsounis, V., Hwangbo, J., Bodie,
¹⁴¹ K., Fankhauser, P., Bloesch, M., Diethelm, R., Bachmann, S., Melzer, A., & Hoepflinger,
¹⁴² M. (2016). ANYmal - a highly mobile and dynamic quadrupedal robot [Conference Paper].
¹⁴³ *2016 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 38–44.
¹⁴⁴ <https://doi.org/10.3929/ethz-a-010686165>