

# Kleinkram: Open Robotic Data Management

Cyrrill Püntener<sup>1\*</sup>, Johann Schwabe<sup>1\*</sup>, Dominique Garmier<sup>1</sup>, Jonas Frey<sup>1,2</sup>, and Marco Hutter<sup>1</sup>

<sup>1</sup> Robotic Systems Lab, ETH Zurich, Switzerland <sup>2</sup> Max Planck Institute for Intelligent Systems, Tübingen, Germany ¶ Corresponding author \* These authors contributed equally.

DOI: 10.xxxxxx/draft

## Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: Sébastien Boisgérault

## Reviewers:

- @boisgera
- @c-joly

Submitted: 16 June 2025

Published: unpublished

## License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

## Summary

Data is key to advancing robotic perception, navigation, locomotion, and reasoning. However, managing diverse robotic datasets presents unique challenges, including scalability, quality assurance, and seamless integration into diverse workflows. To address this gap, we introduce Kleinkram, a free and open-source data management system tailored for robotics research. Kleinkram enables efficient storage, indexing, and sharing of datasets, from small experiments to large-scale collections. It supports essential workflows like validation, curation, and benchmarking via customizable compute actions. Designed with a user-friendly web interface, CLI integration, and scalable deployment options, Kleinkram empowers researchers to streamline data management and accelerate robotics innovation.

## Statement of need

To render robotic data useful for research, it is essential to store, organize, and index the data in a way that makes it easily searchable and shareable. While large corporations have developed internal tools or rely on closed-source third-party providers, no openly available, ready-to-use, and easy-to-deploy solution exists for the robotics research community. Additionally, features such as data verification and the ability to perform tailored compute jobs on newly generated datasets are highly desirable, as they facilitate benchmarking, reproducibility and algorithmic development.

To address these challenges, we introduce **Kleinkram**, an on-premise cloud solution designed for scalable and efficient data management. Unlike traditional cloud storage, Kleinkram natively integrates compute capabilities, automates data transfer, and eliminates the tedious manual effort typically associated with data management workflows. By categorizing and structuring data around common robotics use cases, Kleinkram facilitates the creation of large, diverse datasets that can be easily shared and reused across multiple projects. Its intuitive web interface ensures accessibility, while a command-line interface (CLI) enables seamless integration into automated pipelines and headless systems. Kleinkram supports widely adopted standards, building on ROS1 and ROS2 message definitions, and offers native compatibility with ROSbag and MCAP data formats.

## Data Structure

Kleinkram is designed around the typical data generation process in mobile robotics, assuming data is collected and stored primarily in ROS1/ROS2-compatible ROSbag or MCAP file formats. Once data recording for an experiment is complete, these files can be efficiently uploaded and ingested into the Kleinkram system for centralized storage, indexing, and subsequent

39 post-processing. It is important to note that the current version of Kleinkram focuses on post-  
40 recording and data management. It does not support real-time data streaming or processing  
41 on the fly.

42 To provide structure and facilitate organization and retrieval, Kleinkram requires data to be  
43 organized according to a strict three-layer hierarchy: Projects, Missions, and Files. Each layer  
44 maintains a one-to-many relationship with the layer below it. While users have flexibility in  
45 how they map their specific activities to this structure, the intended model is that a Project  
46 represents a distinct research project, which requires data storage. A Mission corresponds to a  
47 single, self-contained experiment or data collection run conducted within that project, and it  
48 contains all the individual data Files (like ROSbag or MCAP files) recorded during that specific  
49 deployment. This structured approach aids in navigating, managing, and understanding large  
50 volumes of experimental data.

## 51 System Architecture

52 Kleinkram's system architecture is modular, comprising several interacting microservices.

- 53 ■ **Python Client Library and CLI** Provides programmatic access to Kleinkram's functional-  
54 ities, enabling efficient data transfer operations (upload, download) directly from Python  
55 scripts or the command line. This allows seamless integration into robotic workflows  
56 and automated data pipelines running on robots or workstations, removing the need for  
57 manual browser interaction for data transfer.

58 The CLI is built using the typer library, sharing a core Python codebase with the client  
59 library.

- 60 ■ **Web interface** Serves as the primary graphical user interface for users to interact with  
61 Kleinkram. It allows for the browsing, managing and structuring files and their metadata.  
62 It is implemented as a single-page application using the Vue3 framework and the Quasar  
63 component library.

- 64 ■ **Backend API** Acts as the central communication layer between the client applications  
65 (web UI and Python client/CLI) and the data storage and processing components. It  
66 handles authentication, data indexing, metadata management, and schedules background  
67 tasks.

68 The backend is implemented using the NestJS framework and utilises a PostgreSQL  
69 database for storing all metadata related to projects, missions, files, users, and actions.

- 70 ■ **Data Store** The raw robotic data files (ROSBags, MCAPs) are stored on an S3-compatible  
71 object storage backend. This provides scalability and flexibility. Users can easily deploy  
72 and manage their own storage using self-hosted solutions like MinIO, or utilise public  
73 cloud S3 services. Kleinkram interacts with the data store via the S3 API.

- 74 ■ **Action Runner** This component enables the execution of customisable data processing and  
75 analysis tasks directly on the data stored within Kleinkram. Users can define "Actions"  
76 (e.g., validate data integrity, extract sensor metadata, generate preview visualisations,  
77 convert formats, run benchmarking scripts).

78 These actions are packaged as Docker containers. The action runner orchestrates the  
79 execution of these containers, providing them access to the necessary data from the data  
80 store using the client library or CLI.

- 81 ■ **Observability (Optional)** Monitoring and logging system performance, resource usage,  
82 and task execution status are crucial for managing a scalable data system. Integration  
83 with observability tools, such as the Grafana Stack (Loki for logs, Prometheus for metrics,  
84 Tempo for traces, Grafana for dashboards) can provide insights into the system's health  
85 and the progress of the data processing task.

## Usecase

Kleinkram has been used internally at the Robotic Systems Lab at ETH Zurich over the past year. During this time, it has stored over 20 TB of data collected from various robotic systems, effectively replacing the lab's previous reliance on Google Drive for data storage. The largest project supported by Kleinkram was the **GrandTour dataset** (Frey et al., 2025), in which our legged robot **ANYmal** (Hutter et al., 2016), equipped with **Boxi** (Frey et al., 2025), a multi-sensor payload, was deployed across various locations in Switzerland.

Following each data collection mission, raw data — recorded in the form of ROSbags and MCAP files — was uploaded directly to Kleinkram via its command-line interface (CLI). The intuitive Docker-based action integration allowed us to easily define and execute data verification tests. These include, for example, checks to ensure that all sensor streams were recorded at the expected frequencies and correct time synchronization was established during the distributed recordings, as well as common sense checking for validity of data (e.g. images are not black or the IMUs measure the gravity vector).

Beyond data verification, Kleinkram enabled us to run full SLAM pipelines retrospectively, automatically producing standard Absolute and Relative Trajectory Error (ATE/RTE) metrics. This compute integration was critical for development, benchmarking, and evaluation.

Equally important was Kleinkram's user-friendly CLI, which provided quick access to summary statistics such as dataset counts, durations, and other key metrics — many of which were directly used in associated publications. Given that for our use case, data has to be mainly accessed within the ETH network infrastructure; datasets can be pulled on demand and deleted afterwards, fully utilizing the fast on-premise interconnect infrastructure without relying on external servers.

Throughout the project, Kleinkram also enforced metadata submission during upload. Users were required to include a YAML file describing the mission, which captured essential information such as the robot operator, specific hardware configuration, location, and links to related resources (e.g., associated Google Drive folders for images). This structured metadata was essential for organizing and retrieving data at scale.

## Acknowledgements

This work was primarily supported by the Open Research Data Grant at ETH Zurich. Jonas Frey is supported by the Max Planck ETH Center for Learning Systems.

This work was supported and partially funded by Leica Geosystems, which is part of Hexagon. In addition, this work was supported by the National Centre of Competence in Research Robotics (NCCR Robotics), the ETH RobotX research grant funded through the ETH Zurich Foundation, the European Union's Horizon 2020 research and innovation program under grant agreement No 101016970, No 101070405, and No 101070596, and an ETH Zurich Research Grant No. 21-1 ETH-27.

We extend our sincere appreciation to Noel Kampus for the initial design of the web interface, to Lars Leuthold and Marvin Lichtsteiner for their valuable contributions, and to William Talbot and Turcan Tuna for their efforts in the internal testing of Kleinkram.

## References

- Frey, J., Tuna, T., Fu, L. F. T., Weibel, C., Patterson, K., Krummenacher, B., Müller, M., Nubert, J., Fallon, M., Cadena, C., & Hutter, M. (2025). Boxi: Design decisions in the context of algorithmic performance for robotics. *Proceedings of Robotics: Science and Systems*. <https://doi.org/10.48550/arXiv.2504.18500>

- 131 Hutter, M., Gehring, C., Jud, D., Lauber, A., Bellicoso, C. D., Tsounis, V., Hwangbo, J., Bodie,  
132 K., Fankhauser, P., Bloesch, M., Diethelm, R., Bachmann, S., Melzer, A., & Hoepflinger,  
133 M. (2016). ANYmal - a highly mobile and dynamic quadrupedal robot [Conference Paper].  
134 *2016 IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, 38–44.  
135 <https://doi.org/10.3929/ethz-a-010686165>

DRAFT