

Bamojax: Bayesian Modelling with JAX

Max Hinne  ^{1*}

¹ Radboud University, Nijmegen, The Netherlands * These authors contributed equally.

DOI: [10.21105/joss.08642](https://doi.org/10.21105/joss.08642)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: Kanishka B. Narayan 

Reviewers:

- [@matt-graham](#)
- [@alexlytle](#)

Submitted: 23 April 2025

Published: 27 October 2025

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/)).

Summary

Bayesian statistics offers a principled and elegant framework for inferring hidden causes from observed effects. It also provides a rigorous approach to hypothesis testing (model comparison), with advantages such as built-in complexity penalties, and the ability to quantify evidence in favour of the null hypothesis.

However, exact Bayesian inference is computationally intractable in all but the simplest of cases, and requires *approximate inference* techniques, such as Markov chain Monte Carlo (MCMC) and variational inference. Recent advances in the Python JAX ([Bradbury et al., 2018](#)) framework have enabled highly efficient implementations of these algorithms, due to features such as automated differentiation and GPU acceleration. These developments have the potential to greatly increase the efficiency of statistical modelling pipelines.

bamojax ('Bayesian Modelling in Jax') is a probabilistic programming language (PPL) that combines ease-of-use with access to advanced inference algorithms implemented in the Jax ecosystem.

Statement of need

Bamojax is a Bayesian modelling tool based on Python & JAX ([Bradbury et al., 2018](#)). It provides an intuitive, intermediate-level interface between defining a Bayesian statistical model conceptually, and performing efficient inference using the BlackJAX package ([Cabezas et al., 2024](#)).

Existing probabilistic programming languages, such as PyMC ([Abril-Pla O, 2023](#)), can export (the logarithm of) a probability density function that enables BlackJAX-based inference. However, this has two limitations:

1. It does not support Gibbs sampling, where variables are updated individually using their own MCMC kernels. For example, when approximating the posterior over a latent Gaussian process (GP) and its hyperparameters, elliptical slice sampling ([Murray et al., 2010](#)) for the GP can be more efficient than applying No-U-Turn Hamiltonian Monte Carlo (NUTS HMC; ([Hoffman & Gelman, 2014](#))) sampling to all variables jointly. This becomes even more important when embedding MCMC sampling in Sequential Monte Carlo ([Hinne, 2025](#)).
2. It makes it harder to apply tempered Sequential Monte Carlo methods that need separate prior and likelihood terms.

While users can circumvent these issues by manually implementing their models using BlackJAX, this is a labor-intensive and error-prone process. **Bamojax** addresses this gap by providing a user-friendly interface for model construction and Gibbs sampling on top of BlackJAX.

In **Bamojax**, users can define a probabilistic model by specifying variables as well as their associated distributions and dependencies, structured using a directed acyclic graph (DAG).

Under the hood, **Bamojax** translates this DAG and collection of probability distributions to the probability densities used in the approximate inference, leveraging the probability definitions defined in NumPyro (Phan et al., 2019). This abstraction allows users to focus on the conceptual model formulation, rather than the mathematical or inference details, leading to a more intuitive, less error-prone, and more efficient development workflow.

Bamojax is designed for researchers, students, and practitioners that want to make use of the extremely fast approximate inference offered by BlackJAX, but want to focus on model development instead of implementation.

Comparison with existing tools

Bamojax is comparable to existing modern probabilistic programming languages, like NumPyro (Phan et al., 2019), Oryx (The Oryx Authors, 2022), and PyMC (Abril-Pla O, 2023), that use the fast JAX backend for efficient computation of Bayesian inference. The aim of **Bamojax** is to give users full control in defining probability densities and transformations, while at the same time providing some level of abstraction, by automatically deriving densities from NumPyro primitives and the DAG structure of a Bayesian model. This allows for convenient integration of new methods with existing tools.

Bamojax easily admits Gibbs sampling, where individual model parameters are updated in turn, which in practice can lead to efficiency gains and sampling of discrete variables. These benefits increase further when Gibbs sampling is used with Sequential Monte Carlo (Hinne, 2025), which is straightforward to set up here.

Furthermore, **Bamojax** provides a convenient interface to the different sampling algorithms that are available in BlackJAX, giving the user fine-grained control over the inference strategy of their models. This enables users to mix-and-match BlackJAX MCMC kernels with elements of their probabilistic model, while maintaining the efficiency of JAX-based inference.

References

- Abril-Pla O, C. C., Andreani V. (2023). PyMC: A modern and comprehensive probabilistic programming framework in Python. *PeerJ Computer Science*, 9(e1516). <https://doi.org/10.7717/peerj-cs.1516>
- Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., Necula, G., Paszke, A., VanderPlas, J., Wanderman-Milne, S., & Zhang, Q. (2018). JAX: Composable transformations of Python+NumPy programs (Version 0.4.35). <http://github.com/jax-ml/jax>
- Cabezas, A., Corenflos, A., Lao, J., & Louf, R. (2024). BlackJAX: Composable Bayesian inference in JAX. <https://arxiv.org/abs/2402.10797>
- Hinne, M. (2025). An introduction to Sequential Monte Carlo for Bayesian inference and model comparison—with examples for psychology and behavioral science. *Behavior Research Methods*, 57(25). <https://doi.org/10.3758/s13428-025-02642-1>
- Hoffman, M. D., & Gelman, A. (2014). The No-U-Turn Sampler: Adaptively setting path lengths in Hamiltonian Monte Carlo. *Journal of Machine Learning Research*, 15(47), 1593–1623. <http://jmlr.org/papers/v15/hoffman14a.html>
- Murray, I., Adams, R., & MacKay, D. (2010). Elliptical slice sampling. In Y. W. Teh & M. Titterington (Eds.), *Proceedings of the thirteenth international conference on artificial intelligence and statistics* (Vol. 9, pp. 541–548). PMLR. <https://proceedings.mlr.press/v9/murray10a.html>
- Phan, D., Pradhan, N., & Jankowiak, M. (2019). Composable effects for flexible and accelerated probabilistic programming in NumPyro. *arXiv Preprint arXiv:1912.11554*.

The Oryx Authors. (2022). *Oryx: Probabilistic programming and deep learning in JAX.* <https://github.com/jax-ml/oryx>. <https://github.com/jax-ml/oryx>