

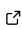


UnRaphael: A Python pipeline tool for analysing the similarity between paintings and their reproductions

Thijs Vroegh^{1*}, Stef Smeets^{1*}, and Lisandra S. Costiner^{2*}

¹ Netherlands eScience Center, The Netherlands ² Utrecht University, The Netherlands ¶ Corresponding author * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Arfon Smith](#) 

Reviewers:

- [@Atif-Anwer](#)
- [@tuelwer](#)

Submitted: 17 March 2025

Published: unpublished

License

Authors of papers retain copyright, and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Taking its name from the Renaissance artist Raphael and the popular practice of reproducing paintings, Unraphael(Vroegh & Smeets, 2024) is a Python workflow tool designed for art historical research. It aims to assist in the comparison of paintings, in particular of originals and their copies. It is designed to streamline the preprocessing, background removal, image alignment, outline extraction, and clustering of photographs of these paintings, grouping the images into similarity-based clusters. The tool specifically targets the challenge of comparing near-identical images by offering a comprehensive pipeline that includes state-of-the-art similarity measures, clustering techniques, together with several options for close visual inspection. Unraphael facilitates easy quantitative comparison of extracted figure-outlines across images of painting reproductions. Integrated with a user-friendly Streamlit interface, the package provides an efficient, flexible and accessible solution for art historians, museum professionals, digital humanities scholars and the general public to experiment, analyse and interpret visual data.

Statement of need

The comparison of designs or figures across paintings is of key interests to art historians, museum professionals and specialists, and is traditionally carried out by manually copying the originals using tracing paper (Holmes, 2004) [humfrey2022]. Among its many functions, it assists in uncovering the relationship between paintings, in understanding the practices used by an artist in the production and reproduction of a particular composition, in tracing the dissemination of designs within and beyond an artistic workshop, and in exploring creative and collaborative working methods (Bambach, 1999). The process of producing manual tracings comes with many challenges, including access to original artworks which may be on museum display or in situ, as well as concerns about working with and handling fragile original artworks. Recently a range of computational approaches have been added to the art historian's resources for studying paintings (Stork, 2023), although these have not focused specifically on the comparison of figures or compositions.

Unraphael addresses this challenge by providing an integrated pipeline that automates the essential steps of comparing elements within paintings. This relies on photographs of paintings alone and provides an intuitive and easy-to-use interface, accessible without specialised knowledge. This includes the steps of image pre-processing, segmentation, alignment, and clustering analysis based on structural similarities. In contrast to recent advancements in deep - and machine learning (Smith et al., 2024) (Ugail et al., 2023) as well as GPT-based agents (Tang et al., 2024), which have influenced the study of art and its attribution, Unraphael avoids these black box methods to ensure that the applied techniques can be interpreted and better understood by researchers from the art-historical domain. It integrates third-party algorithms

from established Python libraries in the computer-vision field (e.g., OpenCV and scikit-image), allowing researchers to focus on interpreting analytical results rather than managing complex image processing tasks. It is open-source and available on Github (Vroegh & Smeets, n.d.) under the Apache 2.0 license.

Overview

Unraphael integrates several image-processing functionalities into a cohesive pipeline, which can be used in sequence or independently from each other. It is supported by a [streamlit application](#) which offers an interactive web-based environment for conducting and visualising the analyses. Its user-friendly interface guides researchers through the entire process from data ingestion to result interpretation, enhancing the overall analytical experience. Several tutorials are available on the accompanying [documentation page](#), explaining how Unraphael can be applied using a small dataset of reproductions of paintings. The first **preprocessing** page allows users to upload one or multiple images (in *jpg* or *png*) and enables automatic standardisation through resizing, colour normalisation and enhancement to ensure consistency across images (see [Figure 1](#)).

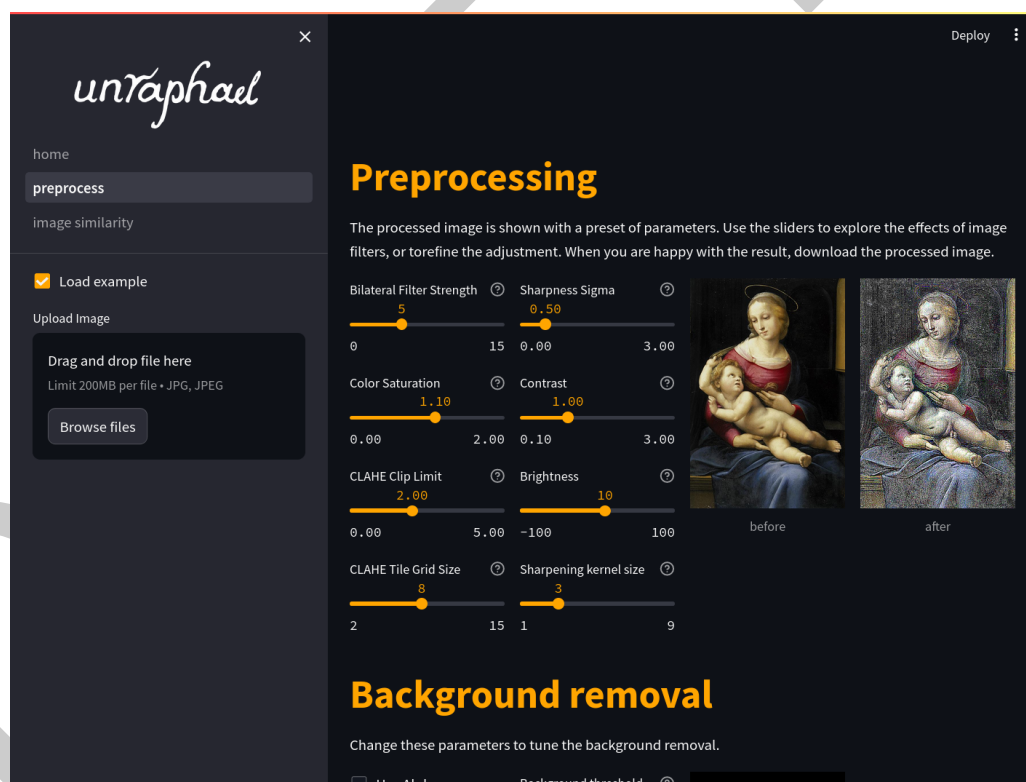


Figure 1: Visualisation of the preprocessing page

In particular, the pre-processing consists of the following features:

- **Bilateral Filter Strength:** This filter smooths the image while preserving edges, useful for reducing noise while keeping important details intact.
- **Color Saturation:** Adjusts the vibrancy of colours in the image, enhancing saturation which makes features stand out more clearly.
- **CLAHE (Contrast Limited Adaptive Histogram Equalization):** Enhances local contrast in the image, making details in darker or lighter regions more visible.

- Sharpness Sigma: Controls the sharpness of the image by applying a sharpening filter, which can help highlight fine details.
- Gamma: Adjusts the brightness and contrast of the image.
- Gain: Similar to gamma, but specifically enhances the overall brightness of the image.
- Sharpening Radius: Defines the extent of the sharpening effect applied to the image.

In addition, several options allow users to control and optimise the background removal process. The extracted figures in the foreground can be saved and used for subsequent analysis.

The **segmentation** page allows the isolation of key objects within each image such as individual figures, for further in-depth analysis. Additionally, analysing the pose structure enables a first visual comparison between the images. Next, the **alignment** page enables the alignment of multiple images to a selected base image. This feature is essential for comparing images against a consistent frame of reference. If desired, users can also adjust features such as brightness, contrast, sharpness, and colour to match the base image, ensuring uniformity before alignment and allowing for optimal (visual) comparison. Images can be aligned using:

- Feature-based alignment: Aligns images using detected features based on SIFT, SURF or ORB.
- Enhanced correlation coefficient maximization: Maximizes correlation between images.
- Fourier Mellin transform: Aligns images based on frequency content.
- FFT phase correlation: Aligns images using phase correlation in the Fourier domain.
- Rotational alignment: Aligns images by correcting rotational differences.

In addition, several motion models are available (see [Figure 2](#)). These models are used to describe the transformation between the base image and the other images.




Figure 2: Available motion models.

To facilitate detailed visual analysis, a number of ways of visualising the information are provided. An animation can be produced that slowly fades one of the images into the second one to aid comparison. A slider can be used to compare the base and aligned image. A third feature, produces an overlay of the outlines of the regions of interest (eg., figures in the painting). This is accompanied by the computations of similarity metrics such as *Frechet*-, *Procrustes*- and *Hausdorff* distances. These different features empower the user to more closely analyse and compare the results from different perspectives.

The **cluster** page of Unraphael is specifically designed to cluster images based on what in computer science is termed structural similarities, while in art history formal similarities, while being resilient to noise, outliers, or variations in the data (Yang & Wu, 2004). First, an option is presented to equalise brightness, sharpness, and contrast to reduce variations due to lighting varying photographic conditions. For instance, copies of paintings may have been photographed under different lighting conditions, or might have been printed or created with different sharpness and contrast levels. Equalising these factors ensures that the clustering

100 focuses on the actual content of the images and not on extraneous factors introduced during
101 photographing, such as lighting or image quality.

102 Next, aligning images to their mean ensures that the key elements of each image, typically the
103 elements to be compared, are in the same spatial location across images. This is critical for
104 structural similarity metrics, which assume that the images being compared are geometrically
105 aligned. Clustering takes place on either the foregrounded figures in the image (i.e., with
106 background removed) or the outer outlines of these figures. In the former case, clustering
107 is based on a) the extracted features from the figures or b) a similarity matrix containing
108 pairwise structural similarity indices. These indices include the *Structural Similarity Index*
109 (Wang et al., 2004), *Complex Wavelet Structural Similarity Index* (Sampat et al., 2009) and a
110 metric for *brushstroke signatures* based on a combination of edge detection techniques (Ugail
111 et al., 2023). *Spectral clustering*, *density-based spatial clustering of applications with noise*
112 (DBSCAN) and *affinity propagation* are among the available clustering methods attuned to
113 these similarity matrices. To assess the effectiveness of the clustering process, performance
114 metrics such as the *Silhouette*-*Davies-Bouldin*- and *Calinski Harabasz score* are displayed.
115 These metrics provide valuable insights into the cohesion and separation of clusters, ensuring
116 that the groupings reflect meaningful similarities. The process is complemented by several
117 visualizations. The first is a dendrogram which visualises relationships between the paintings.
118 The second and third are a silhouette- with a scatterplot that map how images group together
119 into clusters using a 2D projection based on Principal Component Analysis (see Figure 3).

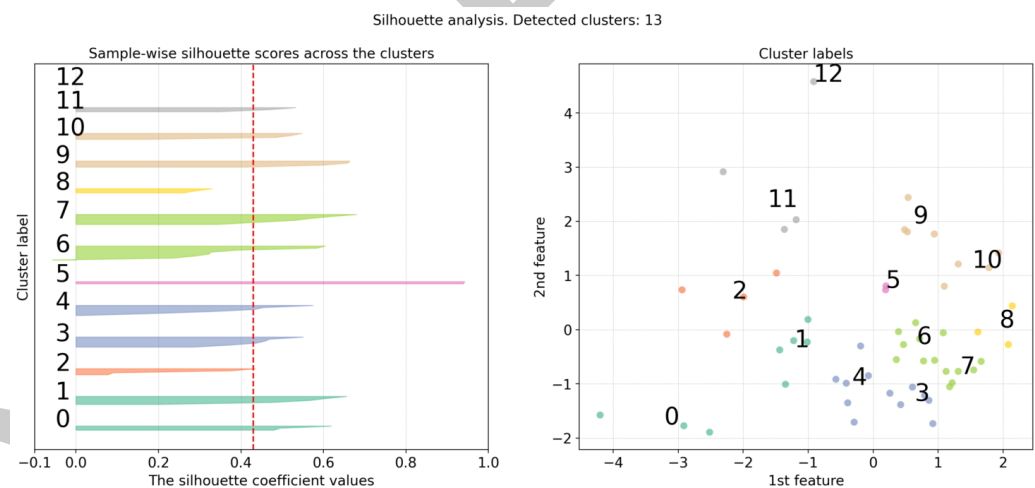


Figure 3: (left) Example of silhouette plot and (right) 2D PCA plot of image clustering.

Conclusion

120 Unraphael is a powerful tool that offers a user-friendly way of examining the visual relationships
121 between original paintings and their copies or derivatives. This new computational approach
122 not only aids in investigating copying methods—such as the use of stencils or cartoons—but
123 also contributes to broader discussions in art history about the transmission of compositions
124 and techniques across time and geography. Unraphael represents a significant advancement in
125 computational art analysis that can be integrated into the toolbox of art historians, museum
126 professionals, and art enthusiasts to further the understanding of paintings, their methods of
127 production and histories.

Acknowledgements

We acknowledge contributions from Christiaan Meijer and Elena Rangelova for scientific input and helpful discussions guiding the development of Unraphael. The development of this project was generously supported by the Netherlands eScience Center, under grant number NLESC.SSIH.2022a.SSIH016.

References

- Bambach, C. (1999). *Drawing and painting in the italian renaissance workshop: Theory and practice, 1300-1600*. Cambridge University Press.
- Holmes, M. (2004). Copying practices and marketing strategies in a fifteenth-century florentine painter's workshop. *Artistic Exchange and Cultural Translation in the Italian Renaissance City*.
- Sampat, M. P., Wang, Z., Gupta, S., Bovik, A. C., & Markey, M. K. (2009). Complex wavelet structural similarity: A new image similarity index. *IEEE Transactions on Image Processing*, 18(11), 2385–2401. <https://doi.org/10.1109/tip.2009.2025923>
- Smith, J. H., Holt, C., Smith, N. H., & Taylor, R. P. (2024). Using machine learning to distinguish between authentic and imitation jackson pollock poured paintings: A tile-driven approach to computer vision. *PIOS ONE*, 19(6), e0302962. <https://doi.org/10.1371/journal.pone.0302962>
- Stork, D. G. (2023). *Pixels & paintings: Foundations of computer-assisted connoisseurship*. John Wiley & Sons.
- Tang, T., Wu, Y., Gao, J., Ruan, K., Zhang, Y., Ye, S., Wu, Y., & Chen, X. (2024). ArtEyer: Enriching GPT-based agents with contextual data visualizations for fine art authentication. *Visual Informatics*, 8(4), 48–59. <https://doi.org/10.1016/j.visinf.2024.11.001>
- Ugail, H., Stork, D. G., Edwards, H., Seward, S. C., & Brooke, C. (2023). Deep transfer learning for visual analysis and attribution of paintings by raphael. *Heritage Science*, 11(1), 1–15. <https://doi.org/10.1186/s40494-023-01094-0>
- Vroegh, T., & Smeets, S. (n.d.). *Unraphael* (Version 0.1.0). <https://doi.org/10.5281/zenodo.11633433>
- Vroegh, T., & Smeets, S. (2024). *Unraphael* (Version 0.2.1). Zenodo. <https://doi.org/10.5281/zenodo.11633433>
- Wang, Z., Bovik, A. C., Sheikh, H. R., & Simoncelli, E. P. (2004). Image quality assessment: From error visibility to structural similarity. *IEEE Transactions on Image Processing*, 13(4), 600–612. <https://doi.org/10.1109/TIP.2003.819861>
- Yang, M.-S., & Wu, K.-L. (2004). A similarity-based robust clustering method. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 26(4), 434–448. <https://doi.org/10.1109/TPAMI.2004.1265860>