

# Bijx: Bijections and normalizing flows with JAX/NNX

Mathis Gerdes<sup>1</sup> and Miranda C. N. Cheng<sup>1,2</sup>

<sup>1</sup> University of Amsterdam, Amsterdam, The Netherlands <sup>2</sup> Academia Sinica, Taipei, Taiwan

DOI: 10.xxxxxx/draft

## Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [✉](#)

Submitted: 07 September 2025

Published: unpublished

## License

Authors of papers retain copyright  
and release the work under a

Creative Commons Attribution 4.0  
International License ([CC BY 4.0](#))

## Summary

Normalizing flows (NF) are a powerful class of generative models that allow for both efficient sampling and exact likelihood evaluation (Kobyzev et al., 2021; Papamakarios et al., 2021; Rezende & Mohamed, 2015). Here we present bijx, a JAX library (Bradbury et al., 2018) built on Flax NNX (Heek et al., 2024) that combines a general normalizing flow API with specialized tools for physics, especially in the context of lattice quantum field theory. The API mirrors standard patterns (e.g., Bijection, Distribution, Chain, Transformed) while making every component an nnx.Module for coherent parameter/state management and seamless use with modern JAX ML tools. Beyond this general core, bijx provides first-class support for continuous normalizing flows (via diffrax), structure-preserving Crouch–Grossmann integrators for matrix Lie groups, symmetry-aware vector fields for lattice data, and utilities for Fourier and Lie algebra valued data. The package is available via pip, with source code on GitHub (<https://github.com/mathisgerdes/bijx>). Each capability can be used independently: the general NF API, continuous normalizing flow (CNF) interface, Lie-group integrators, and physics utilities are modular and separable. General ML users can adopt the standard flow API and CNF interface without the physics utilities, and vice versa.

## Statement of need

The JAX ecosystem includes several flow libraries such as distrax (Babuschkin et al., 2020) and flowjax (Ward, 2025), but these primarily target discrete flows on Euclidean spaces and general ML applications. Many physics problems call for manifold-aware integration, continuous-time flows, and strict symmetry handling, which are not first-class in these packages.

bijx pursues an nnx.Module-centric design, a consistent runtime shape inference convention tailored to physics use-cases, and manifold-preserving ODE solvers as core abstractions. Packaging the physics capabilities inside a general, familiar flow API lowers barriers for users: one can prototype with standard components and opt into the specialized tools when needed. At the same time, the latter are also available independently and can be used in other contexts.

bijx consolidates, generalizes, and streamlines research code from prior work on equivariant models for lattice field theory and continuous flows for gauge theories (Gerdes et al., 2023, 2024). While the package itself is newly organized, the specialized components reflect methods validated in these research contexts and are actively used in ongoing research in this domain.

Primary contributions beyond existing JAX libraries:

- **Continuous flows:** A unified interface for CNFs (Chen et al., 2018) leveraging diffrax (Kidger, 2021), inheriting flexible adjoints and solver choices from the latter.
- **Lie-group integration:** Structure-preserving Crouch–Grossmann integrators (Crouch & Grossman, 1993) for matrix Lie groups, differentiable via adjoint sensitivity.
- **Symmetry-aware vector fields:** Convolutional CNF architectures for lattice data, applicable to other grid-like domains.
- **Physics-domain tools:** Fourier degrees of freedom handling, matrix Lie group operations, and symmetry-aware data transforms.

### 43 Design and Functionality

44 `bijx` favors composable primitives over monolithic training pipelines, as we anticipate its use in  
 45 (physics) research contexts. Users can assemble flows from small parts to match their problem,  
 46 e.g. following one of the examples outlined in the documentation.

47 Core abstractions follow standard patterns: Bijection and Distribution compose via Chain  
 48 and Transformed. All components are `nnx.Modules`, giving consistent parameter/state han-  
 49 dling, `jit/vmap` compatibility, and clean integration with modern tooling. Runtime shape  
 50 inference across batch/event/channel dimensions and an optional auto-vectorization utility  
 51 enable flexible shape handling.

52 Time-dependent vector fields  $f(t, x; \theta)$ , typically represented by neural networks, can be turned  
 53 into continuous flows via the CNF interface. This leverages `diffax`, exposing multiple  
 54 integrators and adjoint methods. For matrix Lie groups and Lie algebra-valued states, spe-  
 55 cialized Crouch–Grossmann integrators yield structure-preserving dynamics and remain fully  
 56 differentiable.

57 The library includes building blocks such as a highly general `GeneralCouplingLayer` that  
 58 reconstructs arbitrary bijection templates from parameters (automatic parameter extraction),  
 59 supports indexing or multiplicative masks, and offers optional auto-vectorization. It also  
 60 provides symmetry-aware convolutions for lattice data and utilities for vectorization and  
 61 batching. Extensive tests and tutorials support both standard and physics-focused workflows.

62 The library encourages extension by implementing minimal Bijection and Distribution inter-  
 63 faces, with a consistent runtime shape convention and `nnx.Module` state handling. Performance-  
 64 critical paths are JIT-compiled and `vmap`-friendly, minimizing Python overhead. A continuous  
 65 integration pipeline runs unit tests and doctests across configurations, and versioned documen-  
 66 tation provides API reference and tutorials.

67 An optional `bijx.flowjax` module provides a lightweight adapter that enables bidirectional  
 68 reuse of bijections and distributions between `bijx` and `FlowJAX`. This exemplifies how `bijx`  
 69 complements, rather than replaces, existing tools, and how similar patterns can be implemented  
 70 for other flow libraries.

### 71 Illustrative Example

72 The documentation (<https://mathisgerdes.github.io/bijx/>) includes basic 2D and higher-  
 73 dimensional examples using the general API. Following a very basic linear flow to showcase  
 74 the API, here we highlight a physics use-case from the tutorials to illustrate the specialized  
 75 capabilities.

76 A typical workflow composes bijections and wraps a base distribution.


```
import jax, jax.numpy as jnp
import bijx

prior = bijx.IndependentNormal(event_shape=(2,))
flow = bijx.Chain(
    bijx.Shift(jnp.array([1.0, -1.0])),
    bijx.Scaling(jnp.array([0.5, 2.0])),
)
dist = bijx.Transformed(prior, flow)
y, logp = dist.sample(batch_shape=(1024,), rng=jax.random.key(0))
```

77 Moving on to a physics use-case, we train a continuous flow on the  $SU(3)$  Lie group to learn a  
 78 Wilson-like target density,

$$p(U) \propto \exp \left( -\beta \sum_{i=1}^3 c_i \operatorname{Re}[\operatorname{tr}(U^i)] \right).$$

Here we choose  $c_1 = 0.17$ ,  $c_2 = -0.65$ , and  $c_3 = 1.22$ . Conjugation-invariance is enforced by building the vector field as the gradient of a “potential” defined in terms of invariant features, and dynamics are integrated with the Crouch–Grossmann solver to remain on the manifold. Starting from the Haar measure, the model is optimized to match the target by minimizing the reverse Kullback–Leibler divergence. As shown in Figure 1, the resulting density captures the target’s shape and symmetries.



**Figure 1:** Uniform Haar density on  $SU(3)$  (left), the learned  $SU(3)$  density (middle), and the target density (right), shown in eigenvalue-angle space.

## Acknowledgements

We thank collaborators who used this software and provided early feedback on the design of various components of this library.

## References

- Babuschkin, I., Baumli, K., Bell, A., Bhupatiraju, S., Bruce, J., Buchlovsky, P., Budden, D., Cai, T., Clark, A., Danihelka, I., Fantacci, C., Godwin, J., Jones, C., Hemsley, R., Hennigan, T., Hessel, M., Hou, S., Kapturowski, S., Kemaev, I., ... Viola, F. (2020). *The DeepMind JAX Ecosystem*.
- Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., Necula, G., Paszke, A., VanderPlas, J., Wanderman-Milne, S., & Zhang, Q. (2018). *JAX: Composable transformations of Python+NumPy programs* (Version 0.3.13). <http://github.com/jax-ml/jax>
- Chen, R. T. Q., Rubanova, Y., Bettencourt, J., & Duvenaud, D. (2018). Neural Ordinary Differential Equations. *Advances in Neural Information Processing Systems*, 31.
- Crouch, P. E., & Grossman, R. (1993). Numerical integration of ordinary differential equations on manifolds. *Journal of Nonlinear Science*, 3(1), 1–33. <https://doi.org/10.1007/BF02429858>
- Gerdes, M., de Haan, P., Rainone, C., Bondesan, R., & Cheng, M. C. N. (2023). Learning Lattice Quantum Field Theories with Equivariant Continuous Flows. *SciPost Phys.*, 15(6), 238. <https://doi.org/10.21468/SciPostPhys.15.6.238>
- Gerdes, M., Haan, P. de, Bondesan, R., & Cheng, M. C. N. (2024). *Continuous normalizing flows for lattice gauge theories* (No. arXiv:2410.13161). arXiv. <https://doi.org/10.48550/arXiv.2410.13161>
- Heek, J., Levskaya, A., Oliver, A., Ritter, M., Rondepierre, B., Steiner, A., & Zee, M. van. (2024). *Flax: A neural network library and ecosystem for JAX* (Version 0.11.1). <http://github.com/google/flax>

- 111 Kidger, P. (2021). *On Neural Differential Equations* [PhD thesis]. University of Oxford.
- 112 Kobyzev, I., Prince, S. J. D., & Brubaker, M. A. (2021). Normalizing flows: An introduction  
113 and review of current methods. *IEEE Transactions on Pattern Analysis and Machine*  
114 *Intelligence*, 43(11), 3964–3979. <https://doi.org/10.1109/TPAMI.2020.2992934>
- 115 Papamakarios, G., Nalisnick, E., Rezende, D. J., Mohamed, S., & Lakshminarayanan, B.  
116 (2021). Normalizing flows for probabilistic modeling and inference. *Journal of Machine*  
117 *Learning Research*, 22(57), 1–64.
- 118 Rezende, D., & Mohamed, S. (2015). Variational inference with normalizing flows. *Proceedings*  
119 *of the 32nd International Conference on Machine Learning*, 1530–1538. [https://arxiv.org/](https://arxiv.org/abs/1505.05770)  
120 [abs/1505.05770](https://arxiv.org/abs/1505.05770)
- 121 Ward, D. (2025). *FlowJAX: Distributions and normalizing flows in jax* (Version 17.2.0).  
122 <https://doi.org/10.5281/zenodo.10402073>

DRAFT