




# 1 aimz: Scalable probabilistic impact modeling

2 Eunseop Kim <sup>1</sup>

3 <sup>1</sup> Eli Lilly and Company, United States

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: 

Submitted: 05 October 2025

Published: unpublished

## License

Authors of papers retain copyright  
and release the work under a  
Creative Commons Attribution 4.0  
International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/))

## 4 Summary

5 aimz is a Python library for scalable probabilistic impact modeling, enabling assessment of  
6 intervention effects on outcomes while providing an intuitive interface for fitting Bayesian  
7 models, drawing posterior samples, generating large-scale posterior predictive simulations,  
8 and estimating interventional effects with minimal boilerplate. It combines the usability of  
9 general machine learning APIs with the flexibility of probabilistic programming through a single  
10 high-level object (ImpactModel). Built atop JAX ([Bradbury et al., 2018](#)) and NumPyro ([Phan  
11 et al., 2019](#)), it supports (minibatch) stochastic variational inference (SVI) and Markov chain  
12 Monte Carlo sampling, just-in-time (JIT)-compiled parallel predictive streaming to chunked Zarr  
([Miles et al., 2020](#)) stores exposed through Xarray ([Hoyer & Hamman, 2017](#)), and first-class  
13 intervention handling for effect estimation. Integrated MLflow ([Zaharia et al., 2018](#)) support  
14 enables experiment tracking and model lineage. These design choices reduce bespoke glue  
15 code and enable reproducible, high-throughput analyses on large datasets, while supporting  
16 rapid iteration and experimentation.

## 18 Statement of need

19 Applied analytics workflows often require: (1) fitting probabilistic models to large scale datasets,  
20 (2) generating posterior and posterior predictive samples for calibrated uncertainty, and (3)  
21 estimating intervention effects under explicit structural modifications. While core probabilistic  
22 programming frameworks (e.g., NumPyro, PyMC ([Oriol et al., 2023](#)), Stan ([Carpenter et al.,  
23 2017](#))) offer mature inference algorithms, recurring engineering tasks—such as streaming large  
24 predictive draws to disk, structuring outputs, coordinating intervention scenarios, managing  
25 device resources, and logging experiments—are typically reimplemented in an ad hoc manner.  
26 General machine learning libraries (e.g., scikit-learn ([Pedregosa et al., 2011](#))) lack native  
27 Bayesian sampling or causal intervention semantics, while many causal inference toolkits  
28 emphasize causal graph discovery or fixed-form treatment effect routines rather than scalable  
29 sampling workflows.

30 aimz consolidates these infrastructural concerns within a single object (ImpactModel) that  
31 provides: probabilistic model tracing and argument binding; SVI or MCMC with automatic  
32 posterior sample management; JIT-compiled, sharded posterior predictive sampling with  
33 concurrent streaming to Zarr stores surfaced as Xarray objects; structured intervention (“do-  
34 operation”) application; and optional MLflow integration for experiment tracking. The familiar  
35 “fit / predict / sample” interface further eases integration with machine learning tooling,  
36 emerging Model Context Protocol pipelines, and AI agents that work with simple estimator-like  
37 semantics. This unification reduces redundant glue code and minimizes potential failure points  
38 in applications such as marketing mix modeling, policy evaluation, and attribution of program  
39 impacts.

40 Existing impact or uplift modeling libraries (e.g., domain-specific frameworks such as Meridian  
41 ([Google Meridian Marketing Mix Modeling Team, 2025](#)), Robyn ([Zhou et al., 2024](#)), or PyMC-  
42 Marketing ([PyMC Labs, 2025](#))) typically standardize on a constrained family of time-series

or marketing response models and a fixed inference stack, making it difficult to deviate from their built-in assumptions without forking code or re-implementing infrastructure. aimz instead pursues generality—accepting arbitrary NumPyro model functions and multiple inference strategies—while still avoiding boilerplate through streamed predictive simulation, structured outputs, and intervention orchestration. By elevating scalable posterior predictive simulation and intervention effect estimation to first-class capabilities, aimz lowers the barrier between exploratory probabilistic modeling and production-grade Bayesian impact analysis.

## References

- Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., Necula, G., Paszke, A., VanderPlas, J., Wanderman-Milne, S., & Zhang, Q. (2018). *JAX: Composable transformations of Python+NumPy programs* (Version 0.3.13). <http://github.com/jax-ml/jax>
- Carpenter, B., Gelman, A., Hoffman, M. D., Lee, D., Goodrich, B., Betancourt, M., Brubaker, M., Guo, J., Li, P., & Riddell, A. (2017). Stan: A probabilistic programming language. *Journal of Statistical Software*, 76, 1–32.
- Google Meridian Marketing Mix Modeling Team. (2025). *Meridian: Marketing mix modeling* (Version 1.2.1). <https://github.com/google/meridian>
- Hoyer, S., & Hamman, J. (2017). Xarray: N-D labeled arrays and datasets in Python. *Journal of Open Research Software*, 5(1). <https://doi.org/10.5334/jors.148>
- Miles, A., Kirkham, J., Durant, M., Bourbeau, J., Onalan, T., Hamman, J., Patel, Z., shikharsg, Rocklin, M., dussin, raphael, Schut, V., Andrade, E. S. de, Abernathey, R., Noyes, C., sbalmer, bot, pyup.io, Tran, T., Saalfeld, S., Swaney, J., ... Banihirwe, A. (2020). *Zarr-developers/zarr-python: v2.4.0* (Version v2.4.0). Zenodo. <https://doi.org/10.5281/zenodo.3773450>
- Oriol, A.-P., Virgile, A., Colin, C., Larry, D., J., F. C., Maxim, K., Ravin, K., Jupeng, L., C., L. C., A., M. O., Michael, O., Ricardo, V., Thomas, W., & Robert, Z. (2023). PyMC: A modern and comprehensive probabilistic programming framework in python. *PeerJ Computer Science*, 9, e1516. <https://doi.org/10.7717/peerj-cs.1516>
- Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D., Brucher, M., Perrot, M., & Duchesnay, É. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830. <https://jmlr.org/papers/v12/pedregosa11a.html>
- Phan, D., Pradhan, N., & Jankowiak, M. (2019). Composable effects for flexible and accelerated probabilistic programming in NumPyro. *arXiv Preprint arXiv:1912.11554*.
- PyMC Labs. (2025). *Marketing statistical models in PyMC* (Version 0.16.0). <https://github.com/pymc-labs/pymc-marketing>
- Zaharia, M. A., Chen, A., Davidson, A., Ghodsi, A., Hong, S. A., Konwinski, A., Murching, S., Nykodym, T., Ogilvie, P., Parkhe, M., Xie, F., & Zumar, C. (2018). Accelerating the Machine Learning Lifecycle with MLflow. *IEEE Data Eng. Bull.*, 41, 39–45. <https://api.semanticscholar.org/CorpusID:83459546>
- Zhou, G., Lares, B., Skokan, I., & Sentana, L. (2024). *Robyn: Semi-automated marketing mix modeling (MMM) from meta marketing science* (Version 3.12.0). <https://github.com/facebookexperimental/Robyn>