

# <sup>1</sup> ETLForge: A unified framework for synthetic test-data generation and ETL validation

<sup>3</sup> Kyriakos Kartas  <sup>1</sup>

<sup>4</sup> 1 Independent Researcher

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

---

Editor: [Yuanqing Wang](#) 

## Reviewers:

- [@gwbischof](#)
- [@Nelly-Barret](#)

Submitted: 18 June 2025

Published: unpublished

## License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))<sup>3</sup>

## <sup>5</sup> Summary

<sup>6</sup> ETLForge is a Python package (Python  $\geq 3.9$ ) for tabular ETL testing workflows. It uses <sup>7</sup> one external YAML/JSON schema to drive both synthetic data generation and validation <sup>8</sup> for pandas DataFrames. The schema can describe field types, ranges, uniqueness, nullability <sup>9</sup> and optional Faker templates ([Faker Contributors, 2024](#)). ETLForge provides DataGenerator <sup>10</sup> (creates tabular test datasets) and DataValidator (checks datasets and returns row-level error <sup>11</sup> reports). A Click-based CLI ([The Pallets Projects, 2023](#)) and a Python API expose the same <sup>12</sup> core workflow, enabling local development and CI/CD automation ([Fowler & Foemmel, 2013](#)).

## <sup>13</sup> Statement of need

<sup>14</sup> Extract-Transform-Load (ETL) processes are critical for data-driven organizations, but testing <sup>15</sup> these pipelines remains challenging ([Kimball & Ross, 2013](#); [Kleppmann, 2017](#)). Teams typically <sup>16</sup> need both representative synthetic inputs and deterministic validation checks. Keeping those <sup>17</sup> artifacts aligned over time can be labor-intensive and can introduce drift between test setup <sup>18</sup> and quality checks ([Dasu & Johnson, 2003](#); [Loshin, 2010](#); [Redman, 2016](#)).

<sup>19</sup> This space already has mature tools. Faker focuses on synthetic value generation ([Faker <sup>20</sup> Contributors, 2024](#)). Great Expectations and Cerberus focus on data validation ([Cerberus <sup>21</sup> Contributors, 2024](#); [The Great Expectations Team, 2023](#)). pandera supports schema-based <sup>22</sup> validation and can also synthesize example data from schemas ([The pandera development <sup>23</sup> team, 2023](#)). ETLForge is positioned as a configuration-first tool for tabular workflows: one <sup>24</sup> external schema file is consumed by both generator and validator, and the same workflow is <sup>25</sup> available through CLI and Python interfaces.

## <sup>26</sup> State of the field

<sup>27</sup> The data-quality ecosystem is mature and diverse, with different projects optimizing for <sup>28</sup> different workflow styles. Tools such as Great Expectations, pandera and Cerberus offer strong <sup>29</sup> validation functionality, while Faker is commonly used for synthetic value generation ([Cerberus <sup>30</sup> Contributors, 2024](#); [Faker Contributors, 2024](#); [The Great Expectations Team, 2023](#); [The <sup>31</sup> pandera development team, 2023](#)).

<sup>32</sup> ETLForge does not attempt to replace these frameworks. Its scope is narrower: tabular <sup>33</sup> pandas DataFrames, declarative schema files (YAML/JSON), synthetic data generation and <sup>34</sup> constraint-based validation through both CLI and Python entry points. The contribution <sup>35</sup> is therefore a lightweight, configuration-first workflow that keeps generation and validation <sup>36</sup> aligned under one external schema.

## <sup>37</sup> Software description

<sup>38</sup> ETLForge implements a dual-purpose architecture where a single YAML/JSON schema drives <sup>39</sup> both data generation and validation processes for **tabular data** (pandas DataFrames). The

40 schema format supports common data types (integer, float, string, date, category), constraints  
41 (ranges, uniqueness, nullability) and realistic data generation via Faker integration.

42 **Schema standards support:** ETLForge includes adapters that detect and convert selected  
43 Frictionless Table Schema and JSON Schema definitions into ETLForge's internal tabular  
44 schema format:

- 45     ▪ **Frictionless Table Schema:** The widely-adopted standard for describing tabular data  
46         (<https://specs.frictionlessdata.io/table-schema/>)
- 47     ▪ **JSON Schema:** The popular standard for describing JSON data structures (<https://json-schema.org/>)

49 Adapter support is intentionally limited to flat/tabular field definitions. Nested object or  
50 array structures are not supported.

51 **Core components:** - DataGenerator: Creates synthetic tabular datasets (pandas DataFrames)  
52 - DataValidator: Validates tabular data (pandas DataFrames) against schema rules, returning  
53 detailed error reports - SchemaAdapter: Handles automatic detection and conversion of  
54 Frictionless Table Schema and JSON Schema formats - CLI interface: Enables command-line  
55 automation via Click ([The Pallets Projects, 2023](#))

56 **Typical workflow:** ETLForge is designed to support the following pipeline:

- 57     1. **Schema definition:** Define data structure and constraints based on target system  
58         requirements
- 59     2. **Test data generation:** Generate synthetic datasets for initial ETL pipeline development  
60         and unit testing
- 61     3. **Pipeline validation:** Use the same schema to validate production or external data after  
62         ETL transformations
- 63     4. **Quality assurance:** Identify discrepancies between expected schema constraints and  
64         actual data quality

65 This workflow demonstrates that while ETLForge generates synthetic test data, its primary  
66 value proposition is in the validation phase where real production data is checked against  
67 expected constraints. The generation capability primarily serves to create controlled test  
68 datasets for unit testing ETL transformations before production data becomes available.

69 **Integration approach:** Rather than replacing existing tools, ETLForge complements ETL  
70 testing workflows by ensuring test data and validation rules remain synchronized. It integrates  
71 with pandas-based pipelines and exports to common formats (CSV, Excel via openpyxl). The  
72 framework targets tabular data structures.

## 73 Software methodology

74 **Data generation algorithm:** The DataGenerator component parses the schema specification  
75 and creates pandas DataFrames by iterating through field definitions. For each field type, it  
76 applies the appropriate generation strategy:

- 77     ▪ **Numeric fields (int, float):** Uses Python's random module with specified ranges and  
78         precision constraints
- 79     ▪ **String fields:** Generates random strings or invokes Faker methods when faker\_template  
80         is specified
- 81     ▪ **Date fields:** Samples dates uniformly within specified ranges using Python datetime  
82         utilities
- 83     ▪ **Category fields:** Samples from predefined value sets with uniform distribution
- 84     ▪ **Uniqueness constraints:** Maintains sets of generated values to ensure uniqueness when  
85         required
- 86     ▪ **Nullability:** Applies configurable null rates to nullable fields using random sampling

87     **Validation algorithm:** The DataValidator component performs multi-pass validation on input  
88     datasets:

- 89       1. **Schema conformance:** Verifies all required columns exist and reports extra columns
- 90       2. **Type checking:** Validates each cell's data type matches schema specifications
- 91       3. **Constraint validation:** Checks range constraints, uniqueness requirements and categorical  
92       value memberships
- 93       4. **Null validation:** Ensures null values only appear in nullable fields
- 94       5. **Error aggregation:** Collects all validation failures with row and column identifiers for  
95       detailed reporting

96     The validation process runs all checks and aggregates all detected issues into a single report,  
97     prioritizing complete feedback over early termination.

98     **Quality control:** GitHub Actions run checks on Python 3.9-3.11, including unit tests, static  
99     analysis via flake8 and mypy, and integration tests through end-to-end workflows.

## 100    Performance characteristics

101    Benchmark results included in the repository (`benchmark_results.csv`) show near-linear scaling  
102    across the tested range (1,000 to 5,000,000 rows). These benchmarks were conducted using a  
103    representative schema containing 8 fields with varying complexity levels:

- 104       ■ 2 integer fields with range constraints (id: 1-10000000, age: 18-80)
- 105       ■ 1 float field with range constraints (30000.0-150000.0) and precision specifications
- 106       ■ 3 string fields, including two with Faker template integration (name, email) and one  
107       with length constraints
- 108       ■ 1 categorical field with 5 predefined values (Engineering, Marketing, Sales, HR and  
109       Finance)
- 110       ■ 1 date field with range constraints (2020-01-01 to 2024-12-31)

111    Performance scales approximately linearly with the number of rows and fields. Complex  
112    constraints such as uniqueness checking and Faker integration introduce additional overhead.  
113    The complete benchmark schema is available in the repository as `benchmark_schema.yaml` for  
114    reproducibility.

115    These results indicate that ETLForge can be integrated into CI/CD checks for tabular datasets  
116    while focusing on constraint-based validation rather than advanced statistical profiling.

## 117    Discussion

118    ETLForge unifies data generation and validation under one schema for tabular workflows, but it  
119    makes deliberate trade-offs compared with broader validation platforms. Compared with tools  
120    such as Great Expectations and pandera ([The Great Expectations Team, 2023](#); [The pandera  
121    development team, 2023](#)), ETLForge emphasizes a smaller feature set centered on schema  
122    conformance, type checks, basic constraints and row-level error reporting. This narrower scope  
123    prioritizes straightforward configuration and reproducible test workflows over advanced profiling  
124    and statistical quality analysis.

125    The framework currently has several technical limitations that constrain its applicability:

- 126       ■ **Dataset size:** Large datasets exceeding one million rows may require memory optimization  
127       strategies, as the current implementation loads entire datasets into pandas DataFrames  
128       during validation.
- 129       ■ **Nested structures:** Complex nested data structures are not supported. This limitation  
130       exists because ETLForge specifically targets tabular data formats (CSV and Excel) which  
131       are inherently flat.

- 132     ▪ **Statistical validation:** Advanced statistical validations (distribution testing, anomaly
- 133       detection and correlation analysis) require integration with specialized tools. ETLForge
- 134       provides constraint-based validation rather than statistical analysis.
- 135     ▪ **Custom validation logic:** While the framework validates against schema-defined
- 136       constraints, it does not currently support user-defined validation functions, limiting
- 137       extensibility for domain-specific validation rules.

## 138 Availability

139     The ETLForge source code is available on GitHub at <https://github.com/kkartas/ETLForge>  
140     under the MIT license. The latest release (v1.1.0) can be installed from the Python  
141     Package Index using `pip install etl-forge`, with optional extras `etl-forge[faker]` and  
142     `etl-forge[excel]`. Complete documentation is hosted at <https://etlforge.readthedocs.io/>.  
143     The software supports Linux, macOS and Windows operating systems and is compatible with  
144     Python versions 3.9 through 3.11.

## 145 References

- 146     Cerberus Contributors. (2024). *Cerberus: Lightweight, extensible data validation library for*  
147       *python* (Version 1.3.5). <https://docs.python-cerberus.org/>
- 148     Dasu, T., & Johnson, T. (2003). *Exploratory data mining and data cleaning*. John Wiley &  
149       Sons.
- 150     Faker Contributors. (2024). *Faker: A python package that generates fake data for you* (Version  
151       24.4.0). <https://doi.org/10.5281/zenodo.3823398>
- 152     Fowler, M., & Foemmel, M. (2013). Continuous integration. *Thought-Works*. <https:////martinfowler.com/articles/continuousIntegration.html>
- 154     Kimball, R., & Ross, M. (2013). *The data warehouse toolkit: The definitive guide to*  
155       *dimensional modeling* (3rd ed.). John wiley & sons.
- 156     Kleppmann, M. (2017). *Designing data-intensive applications: The big ideas behind reliable,*  
157       *scalable, and maintainable systems*. O'Reilly Media.
- 158     Loshin, D. (2010). *The practitioner's guide to data quality improvement*. Morgan Kaufmann.
- 159     Redman, T. C. (2016). *Getting in front on data: Who does what*. Harvard Business Review  
160       Press.
- 161     The Great Expectations Team. (2023). *Great Expectations: Always know what to expect from*  
162       *your data* (Version 0.18.0). <https://doi.org/10.5281/zenodo.1497623>
- 163     The Pallets Projects. (2023). *Click: Command Line Interface Creation Kit* (Version 8.1.3).  
164       <https://doi.org/10.5281/zenodo.5123732>
- 165     The pandera development team. (2023). *pandera: Data validation for scientists, engineers,*  
166       *and analysts* (Version 0.17.0). <https://doi.org/10.25080/gerudo-f2bc6f59-010>