

diffsol: Rust crate for solving differential equations

Martin Robinson   and Alex Allmont  

1 Oxford Research Software Engineering Group, Doctoral Training Centre, University of Oxford, Oxford, UK  Corresponding author

DOI: [10.21105/joss.09384](https://doi.org/10.21105/joss.09384)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: Neea Rusch  

Reviewers:

- [@DiogoRibeiro7](#)
- [@ArmaVica](#)
- [@jeertmans](#)

Submitted: 06 November 2025

Published: 23 January 2026

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Ordinary Differential Equations (ODEs) are powerful tools for modelling a wide range of physical systems. Unlike purely data-driven models, ODEs can be based on the underlying physics, biology, or chemistry of the system being modelled, making them particularly useful for predicting the behaviour of a system under conditions that have not been observed. ODEs can be used to model everything from the motion of planets to the spread of infectious diseases.

`diffsol` is a Rust crate for solving ordinary differential equations (ODEs) or semi-explicit differential algebraic equations (DAEs). It can solve equations in the following form:

$$M \frac{dy}{dt} = f(t, y, p)$$

where y is the state of the system, p are a set of parameters, t is time, $f(t, y, p)$ is a function that describes how the state of the system changes over time and M is an optional and possibly singular mass matrix. The solution to an ODE is a function $y(t)$ that satisfies the ODE and any initial conditions.

The equations (e.g. $f(t, y, p)$) can be provided by the user either using Rust code, or a custom Domain Specific Language (DSL) called DiffSL. DiffSL uses automatic differentiation using Enzyme ([Moses et al., 2022](#)) to calculate the necessary gradients, and JIT compilation (using either LLVM ([Lattner & Adve, 2004](#)) or Cranelift ([Bytecode Alliance, 2025](#))) to generate efficient native code at runtime. The DSL is ideal for using `diffsol` from a higher-level language like Python or R while still maintaining similar performance to pure rust. `Diffsol` currently provides Python bindings through the `pydiffsol` package, with further language bindings planned.

ODE solvers require linear algebra containers (e.g. vectors, matrices), operators and linear solvers. `diffsol` allows users to choose both dense and sparse matrices and solvers from the nalgebra ([Dimforge, 2025](#)) or faer ([Kazdadi, 2025](#)) crates, and uses a trait-based approach to allow other linear algebra libraries to be added at a later date.

Statement of need

ODE solvers have a long history in scientific computing, and many libraries currently exist. Some notable examples include `scipy.integrate.odeint` ([Virtanen et al., 2020](#)) in Python, `ode45` ([Shampine & Reichelt, 1997](#)) in MATLAB, and the Sundials suite of solvers ([Gardner et al., 2022](#)) in C. Rust is a systems programming language that is gaining popularity in the scientific computing community due to its performance, safety, and ease of use. There is currently no ODE solver library written in Rust that provides the same level of functionality as these other libraries, and this is the gap that `diffsol` aims to fill.

ODE solvers written in lower-level languages like C, Fortran or Rust offer significant performance benefits. However, these solvers are often more difficult to wrap and use in higher-level languages like Python or MATLAB, primarily because users must supply their equations in the language of the solver. `diffsol` solves this issue by providing its own custom DiffSL DSL which is JIT compiled to efficient native code at run-time, meaning that users from a higher-level language like Python or R can specify their equations using a simple string-based format while still maintaining similar performance to pure Rust. Two other popular ODE solvers that take advantage of JIT compilation are `DifferentialEquations.jl` ([Rackauckas & Nie, 2017](#)) in Julia, and `diffrafx` ([Kidger, 2021](#)) in Python. However, both these packages compile the entire solver as well as the equations, which is a significant amount of code. `diffSol` only compiles the equations, meaning that it has a significantly smaller “time-to-first-plot” for users. Another popular differential equations solver package utilising a DSL is `OpenModelica` ([Fritzson et al., 2020](#)). Wrappers to this package in higher-level languages like Python rely on messaging to a separate `OpenModelica` server, which can be slow and more complicated to set up. In contrast, `diffsol` can be integrated directly into higher-level languages using language bindings and linking to a single shared library, see for example the `pydiffsol` Python bindings discussed below.

Features

The following solvers are available in `diffsol`:

1. A variable order Backwards Difference Formulae (BDF) solver, suitable for stiff problems and singular mass matrices. The basic algorithm is derived in ([Byrne & Hindmarsh, 1975](#)), however this particular implementation follows that implemented in the MATLAB routine `ode15s` ([Shampine & Reichelt, 1997](#)) and the SciPy implementation ([Virtanen et al., 2020](#)), which features the NDF formulas for improved stability.
2. A Singly Diagonally Implicit Runge-Kutta (SDIRK or ESDIRK) solver, suitable for moderately stiff problems and singular mass matrices. Two different butcher tableau are provided, TR-BDF2 ([Bank et al., 1985](#); [Hosea & Shampine, 1996](#)) and ESDIRK34 ([Jørgensen et al., 2018](#)), or users can supply their own.
3. A variable order Explicit Runge-Kutta (ERK) solver, suitable for non-stiff problems. One butcher tableau is provided, the 4th order TSIT45 ([Tsitouras, 2011](#)), or users can supply their own.

All solvers feature:

- Linear algebra containers and linear solvers from the `nalgebra` or `faer` crates, including both dense and sparse matrix support.
- Adaptive step-size control to given relative and absolute tolerances. Tolerances can be set separately for the main equations, quadrature of the output function, and sensitivity analysis.
- Dense output, interpolating to times provided by the user.
- Event handling, stopping when a given condition $g_e(t, y, p)$ is met or at a specific time.
- Numerical quadrature of an optional output $g_o(t, y, p)$ function over time.
- Forward sensitivity analysis, calculating the gradient of an output function or the solver states y with respect to the parameters p .
- Adjoint sensitivity analysis, calculating the gradient of cost function $G(p)$ with respect to the parameters p . The cost function can be the integral of a continuous output function $g(t, y, p)$ or a sum of a set of discrete functions $g_i(t_i, y, p)$ at time points t_i .

Bindings in higher-level languages

`pydiffsol` provides Python bindings to `diffsol` using the PyO3 ([PyO3 Project and Contributors, 2025](#)) crate. It allows users to define ODEs in Python using the DiffSL DSL, and solve them

using the Rust `diffsol` library. `pydiffsol` aims to provide a simple and easy-to-use interface for solving ODEs in Python, while still maintaining the performance benefits of using Rust under the hood.

The goal is to develop further [bindings to other higher-level languages](#), including R, JAX and JavaScript, exploiting the DiffSL DSL to provide high performance while maintaining ease of use and positioning the core `diffsol` library as a widely used cross-language, cross-platform, high-performance ODE solver.

Acknowledgements

We gratefully acknowledge the support of all the past and future contributors to the `diffsol` project, for their advice, enthusiasm, bug reports and code. In particular, we would like to thank the authors of the `pharmsol` crate, Julian Otalvaro and Markus Hovd.

References

- Bank, R. E., Coughran, W. M., Fichtner, W., Grosse, E. H., Rose, D. J., & Smith, R. K. (1985). Transient simulation of silicon devices and circuits. *IEEE Transactions on Computer-Aided Design of Integrated Circuits and Systems*, 4(4), 436–451. <https://doi.org/10.1109/T-ED.1985.22232>
- Byrne, G. D., & Hindmarsh, A. C. (1975). A polyalgorithm for the numerical solution of ordinary differential equations. *ACM Transactions on Mathematical Software (TOMS)*, 1(1), 71–96. <https://doi.org/10.1145/355626.355636>
- Bytecode Alliance. (2025). *Craneflift* (Version 0.115.1). <https://wasmtime.dev>
- Dimforge. (2025). *Nalgebra* (Version 0.33.2). <https://github.com/dimforge/nalgebra>
- Fritzson, P., Pop, A., Abdelhak, K., Ashgar, A., Bachmann, B., Braun, W., Bouskela, D., Braun, R., Buffoni, L., Casella, F., Castro, R., Franke, R., Fritzson, D., Gebremedhin, M., Heuermann, A., Lie, B., Mengist, A., Mikelsons, L., Moudgalya, K., ... Östlund, P. (2020). The OpenModelica Integrated Environment for Modeling, Simulation, and Model-Based Development. *Modeling, Identification and Control*, 41(4), 241–295. <https://doi.org/10.4173/mic.2020.4.1>
- Gardner, D. J., Reynolds, D. R., Woodward, C. S., & Balos, C. J. (2022). Enabling new flexibility in the SUNDIALS suite of nonlinear and differential/algebraic equation solvers. *ACM Transactions on Mathematical Software (TOMS)*, 48(3), 1–24. <https://doi.org/10.1145/3539801>
- Hosea, M., & Shampine, L. (1996). Analysis and implementation of TR-BDF2. *Applied Numerical Mathematics*, 20(1-2), 21–37. [https://doi.org/10.1016/0168-9274\(95\)00115-8](https://doi.org/10.1016/0168-9274(95)00115-8)
- Jørgensen, J. B., Kristensen, M. R., & Thomsen, P. G. (2018). A family of ESDIRK integration methods. *arXiv Preprint arXiv:1803.01613*. <https://doi.org/10.48550/arXiv.1803.01613>
- Kazdadi, S. E. (2025). *Faer-rs* (Version 0.1.0). <https://faer-rs.github.io>
- Kidger, P. (2021). *On Neural Differential Equations* [PhD thesis, University of Oxford]. <https://doi.org/h10.48550/arXiv.2202.02435>
- Lattner, C., & Adve, V. (2004). LLVM: A compilation framework for lifelong program analysis & transformation. *International Symposium on Code Generation and Optimization, 2004. CGO 2004.*, 75–86. <https://doi.org/10.1109/CGO.2004.1281665>
- Moses, W. S., Narayanan, S. H. K., Paehler, L., Churavy, V., Schanen, M., Hückerheim, J., Doerfert, J., & Hovland, P. (2022). Scalable automatic differentiation of multiple

parallel paradigms through compiler augmentation. *Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis*. <https://doi.org/10.1109/SC41404.2022.00065>

PyO3 Project and Contributors. (2025). PyO3 (Version v0.27.2). <https://github.com/PyO3/pyo3>

Rackauckas, C., & Nie, Q. (2017). DifferentialEquations.jl – a performant and feature-rich ecosystem for solving differential equations in julia. *The Journal of Open Research Software*, 5(1). <https://doi.org/10.5334/jors.151>

Shampine, L. F., & Reichelt, M. W. (1997). The MATLAB ODE suite. *SIAM Journal on Scientific Computing*, 18(1), 1–22. <https://doi.org/10.1137/s1064827594276424>

Tsitouras, C. (2011). Runge–kutta pairs of order 5 (4) satisfying only the first column simplifying assumption. *Computers & Mathematics with Applications*, 62(2), 770–775. <https://doi.org/10.1016/j.camwa.2011.06.002>

Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D., Burovski, E., Peterson, P., Weckesser, W., Bright, J., & others. (2020). SciPy 1.0: Fundamental algorithms for scientific computing in python. *Nature Methods*, 17(3), 261–272. <https://doi.org/10.1038/s41592-019-0686-2>