

# <sup>1</sup> Maui: A Python Visualization Framework for Ecoacoustics Data

<sup>3</sup> Caio Ferreira Bernardo  <sup>1</sup> and Maria Cristina Ferreira de Oliveira 

<sup>4</sup> 1 Instituto de Ciências Matemáticas e de Computação - Universidade de São Paulo (USP), São Carlos,  
<sup>5</sup> Brazil

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Nick Golding](#)  

## Reviewers:

- [@markur4](#)
- [@louisfh](#)

Submitted: 10 September 2025

Published: unpublished

## License

Authors of papers retain copyright  
and release the work under a  
Creative Commons Attribution 4.0  
International License ([CC BY 4.0](#)).  


## <sup>6</sup> Summary

<sup>7</sup> Passive Acoustic Monitoring (PAM) technology generates vast volumes of environmental audio  
<sup>8</sup> recordings. There is a diverse range of tools, either to compute mathematical features (acoustic  
<sup>9</sup> indices) or to perform machine learning tasks such as classification. However, researchers  
<sup>10</sup> still lack unified tools for visually exploring soundscape repositories. We introduce Maui, an  
<sup>11</sup> open-source Python framework designed to simplify exploratory analysis of large ecoacoustic  
<sup>12</sup> datasets by placing visualization directly into the PAM workflow. Maui provides modules  
<sup>13</sup> for data ingestion and metadata parsing, supports flexible incorporation of user-computed  
<sup>14</sup> acoustic indices, and offers a suite of high-level plotting functions. It includes methods for  
<sup>15</sup> creating false-color spectrograms ([M. Towsey et al., 2014](#)), customizable Diel plots, and radar,  
<sup>16</sup> violin, and parallel plots for multivariate index comparison. These visualizations are useful  
<sup>17</sup> to reveal temporal, spatial, and taxonomic patterns at scale. With a modular architecture,  
<sup>18</sup> Maui leverages existing Python libraries for bioacoustic data processing, while focusing on  
<sup>19</sup> interactive, publication-quality graphics that facilitate hypothesis generation and large-scale  
<sup>20</sup> soundscape synthesis. By filling a gap in ecoacoustic visual data analytics and bringing together  
<sup>21</sup> visualizations drawn from acoustic ecology literature ([Phillips et al., 2018](#); [M. Towsey et al., 2014](#);  
<sup>22</sup> [M. W. Towsey et al., 2015](#)), Maui helps researchers study soundscape dynamics.

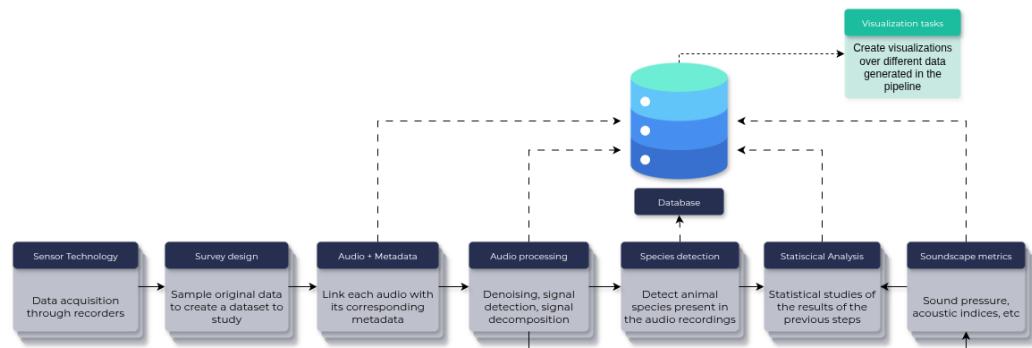
## <sup>23</sup> Statement of need

<sup>24</sup> The sounds originating from anthropogenic, biophonic, and geophonic sources in a landscape  
<sup>25</sup> define its soundscape ([Pijanowski et al., 2011](#)). Substantial amounts of acoustic data, particu-  
<sup>26</sup> larly from biophonic sources, can be captured using low-cost autonomous recorders deployed  
<sup>27</sup> for Passive Acoustic Monitoring (PAM) ([Browning et al., 2017](#)). An active research line in  
<sup>28</sup> environmental ecology, acoustic ecology addresses the study of natural soundscapes ([Grinfeder](#)  
<sup>29</sup> [et al., 2022](#)). The discipline relies heavily on computational methods for audio data processing  
<sup>30</sup> and analysis ([Napier et al., 2024](#); [Pijanowski et al., 2024](#)).

<sup>31</sup> A diversity of PAM data processing pipelines are described in the literature. In the particular  
<sup>32</sup> context of detecting the presence of animal species in the recordings, ([Gibb et al., 2019](#)) defines  
<sup>33</sup> a seven-step pipeline (see 1). Departing from data acquisition, sampling the recordings before  
<sup>34</sup> analysis is often necessary, given the large data volumes collected. Researchers extract the  
<sup>35</sup> associated metadata from the resulting subset of audio files, including location, time, climate  
<sup>36</sup> conditions, and recorder type. A fourth step involves preprocessing the audio files, e.g., to  
<sup>37</sup> reduce noise and emphasize relevant signals. Researchers can then perform acoustic event  
<sup>38</sup> detection and labeling, e.g., to identify animal species. This typically involves multiple iterations  
<sup>39</sup> of computing metrics such as acoustic features, ecological indices, and conducting statistical  
<sup>40</sup> analyses of intermediate results to gain a comprehensive understanding of the soundscape.

<sup>41</sup> As PAM pipelines are instantiated multiple times (see Figure 1), researchers typically accumulate

vast datasets of audio recordings collected at multiple sites over extended periods. These large repositories of environmental acoustic recordings are a valuable source of knowledge when analyzed from a global perspective. Knowledge extraction requires practical software tools and libraries to streamline data exploration and analysis. Analysts need flexibility to investigate their accumulated data from multiple perspectives, considering different data and metadata properties. Conducting global investigations can uncover insights beyond previous soundscape analyses, help identify potential improvements in existing practices and methodologies, and support large-scale PAM data analysis in the long term ([Napier et al., 2024](#)).



**Figure 1:** Multiple executions of the PAM pipeline generate soundscape repositories amenable to exploration with visualization methods.

Data visualization is a powerful tool for extracting insights in this inherently exploratory context. It enables the representation of metadata and acoustic features across time and locations, providing comprehensive overviews of the data repositories from multiple perspectives. Compelling visualizations can enhance data exploration and summarization beyond the standard processing pipeline. We consider an extended PAM pipeline that integrates visualization into a framework to promote exploratory analysis of the data accumulated in acoustic repositories resulting from multiple instantiations of the standard pipeline, as illustrated in [Figure 1](#). The dissemination of acoustic ecology practices has motivated many open-source tools and libraries to facilitate tasks such as computing acoustic features and applying machine learning algorithms. Following this trend, we introduce Maui, a Python package built on top of Plotly to support visual exploratory tasks on ecoacoustic data repositories.

## 61 Modules architecture and useflow

62 Maui implements methods that focuses primarily on creating visualizations that require the  
 63 computation of acoustic features. We assume users already have their preferred tools for this  
 64 computation. A brief description of each module follows.

65 **File Metadata:** a helper module, it provides methods to decode the relevant metadata values  
 66 (e.g., location, date, time) encoded in audio file names. It is common practice to adopt some file  
 67 naming template to encode information; e.g., a file named “LEEC02\_20161202\_050100\_br”  
 68 refers to an acoustic recording obtained in a landscape identified as “LEEC02” on December  
 69 2, 2016, recording capture started at time 05h:01m from a device placed at an environment  
 70 identified as “br”. This module provides a method for users to specify how metadata must be  
 71 decoded from a given naming template. Once the encoding policy is informed, the method  
 72 parses the file names to extract the corresponding metadata values.

73 **IO:** implements multiple input and output methods, e.g., to load a single file or an entire  
 74 dataset consisting of multiple audio files and create a Python data frame that incorporates  
 75 the extracted metadata, as per the policy defined by the previous module. This module exists  
 76 so that users can focus on understanding the data without being concerned with low-level

<sup>77</sup> operations, such as parsing metadata from file names to obtain the data frame. The remaining  
<sup>78</sup> methods from this and other modules will operate on the resulting data frame.

<sup>79</sup> **Samples:** a utility module to retrieve a small sample dataset already embedded in Maui for  
<sup>80</sup> demonstration purposes.

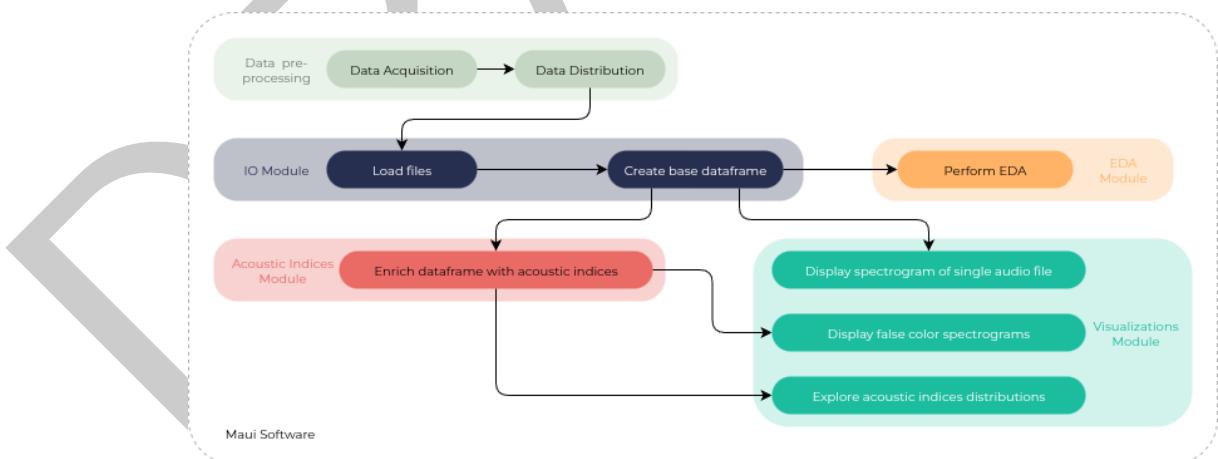
<sup>81</sup> **EDA:** facilitates creating visualizations that convey overviews of the dataset stored in the  
<sup>82</sup> Python data frame, depicting data sample distributions across multiple user-defined dimensions,  
<sup>83</sup> such as date, time, and location. It includes methods to generate summary reports, duration  
<sup>84</sup> analysis views, daily distribution views, heatmaps, and histograms.

<sup>85</sup> **Acoustic Indices:** Maui does not include modules or methods for acoustic index computation or  
<sup>86</sup> feature extraction. Instead, this module provides an interface to incorporate into the working  
<sup>87</sup> data frame the audio features obtained using some user-defined method or external tools. As  
<sup>88</sup> feature computation on large datasets can be computationally demanding, we considered it  
<sup>89</sup> necessary to streamline the acoustic feature computation task.

<sup>90</sup> **Visualizations:** a core module that incorporates methods to create visualizations of audio data  
<sup>91</sup> with a few lines of code, simplifying data exploration tasks.

<sup>92</sup> **Utils:** another utility module that implements methods for data preprocessing operations,  
<sup>93</sup> such as audio segmentation and data preparation steps required, e.g., to create false color  
<sup>94</sup> spectrograms.

<sup>95</sup> Figure [Figure 2](#) shows the different modules, their relationships, and the major tasks they  
<sup>96</sup> implement. Each module focuses on a specific task and operates independently from the others.  
<sup>97</sup> Still, they interact as a user executes data processing tasks and creates data visualizations. As  
<sup>98</sup> such, they together implement a complete data visualization solution. The **IO** module is central  
<sup>99</sup> to Maui because it provides methods that simplify the data loading process. Nonetheless, it  
<sup>100</sup> is not required to load a dataset, as long as the user provides the required data frame. A  
<sup>101</sup> complete example of each method and resulting visualizations created from real world datasets  
<sup>102</sup> are available at example notebooks hosted on [GitHub](#)<sup>1</sup>.



**Figure 2:** Flow of Maui software usage: each activity is represented within the respective module where it is performed.

## Related Work

<sup>103</sup> There are several softwares solutions that address acoustic ecology challenges, but, as far as we  
<sup>104</sup> are concerned, Maui is the first visualization framework developed specifically to meet the needs

<sup>105</sup><sup>1</sup><https://github.com/maui-software/maui-software-examples>

106 of acoustic ecology. In the realm of open source Python packages, Scikit-maad ([Ulloa et al., 2021](#)) and Open Soundscape ([Lapp et al., 2023](#)) are tools that complement Maui. Scikit-maad  
107 encompasses a complete workflow to load, preprocess, and transform data, find regions of  
108 interest, compute temporal and spectral acoustic indices, estimate sound pressure levels, and  
109 calculate the distance from an audio source to the recording device. Open Soundscape is  
110 focused on data classification and spatial localization of acoustic events. The package includes  
111 methods for training convolutional neural networks (CNNs), performing data augmentation,  
112 and other utility functions required to execute machine learning tasks.  
113

## 114 Acknowledgements

115 This project was supported by grants from the State of São Paulo Research Foundation  
116 (FAPESP 2021/08322-3) and the Brazilian Research Council (CNPq 301847/2017-7). We  
117 thank Dr. Milton Cezar Ribeiro, from LEEC, for insightful discussions.

## 118 References

- 119 Browning, E., Gibb, R., Glover-Kapfer, P., & Jones, K. E. (2017). *Passive acoustic monitoring*  
120 *in ecology and conservation*.
- 121 Gibb, R., Browning, E., Glover-Kapfer, P., & Jones, K. E. (2019). Emerging opportunities  
122 and challenges for passive acoustics in ecological assessment and monitoring. *Methods in*  
123 *Ecology and Evolution*, 10(2), 169–185.
- 124 Grinfeder, E., Lorenzi, C., Haupert, S., & Sueur, J. (2022). What do we mean by “soundscape”?  
125 A functional description. *Frontiers in Ecology and Evolution*, 10, 894232.
- 126 Lapp, S., Rhinehart, T., Freeland-Haynes, L., Khilnani, J., Syunkova, A., & Kitzes, J. (2023).  
127 OpenSoundscape: An open-source bioacoustics analysis package for Python. *Methods in*  
128 *Ecology and Evolution*, 14(9), 2321–2328.
- 129 Napier, T., Ahn, E., Allen-Ankins, S., Schwarzkopf, L., & Lee, I. (2024). Advancements in  
130 preprocessing, detection and classification techniques for ecoacoustic data: A comprehensive  
131 review for large-scale passive acoustic monitoring. *Expert Systems with Applications*, 252,  
132 124220. <https://doi.org/https://doi.org/10.1016/j.eswa.2024.124220>
- 133 Phillips, Y. F., Towsey, M., & Roe, P. (2018). Revealing the ecological content of long-duration  
134 audio-recordings of the environment through clustering and visualisation. *PLoS One*, 13(3),  
135 e0193345.
- 136 Pijanowski, B. C., Fuenzalida, F. R., Banerjee, S., Minghim, R., Lima, S. L., Bowers-Sword, R.,  
137 Guzman, S. R., Revuelta-Acosta, J., Adeniji, A. E., Grimes, S. E., Sarker, S. K., Hossain,  
138 Md. R., Anika, T. T., & Savage, D. (2024). Soundscape analytics: A new frontier of  
139 knowledge discovery in soundscape data. *Current Landscape Ecology Reports*, 9(4), 88–107.  
140 <https://doi.org/10.1007/s40823-024-00101-9>
- 141 Pijanowski, B. C., Villanueva-Rivera, L. J., Dumyahn, S. L., Farina, A., Krause, B. L.,  
142 Napoletano, B. M., Gage, S. H., & Pieretti, N. (2011). Soundscape ecology: The science  
143 of sound in the landscape. *BioScience*, 61(3), 203–216.
- 144 Towsey, M. W., Truskinger, A. M., & Roe, P. (2015). The navigation and visualisation of  
145 environmental audio using zooming spectrograms. *2015 IEEE International Conference on*  
146 *Data Mining Workshop (ICDMW)*, 788–797.
- 147 Towsey, M., Zhang, L., Cottman-Fields, M., Wimmer, J., Zhang, J., & Roe, P. (2014).  
148 Visualization of long-duration acoustic recordings of the environment. *Procedia Computer*  
149 *Science*, 29, 703–712.

- 150 Ulloa, J. S., Haupert, S., Latorre, J. F., Aubin, T., & Sueur, J. (2021). Scikit-maad: An  
151 open-source and modular toolbox for quantitative soundscape analysis in Python. *Methods*  
152 in Ecology and Evolution, 12(12), 2334–2340.

DRAFT