

PatientProfiles: An R package to identify patient characteristics in data mapped to the OMOP common data model

Martí Català ¹¶, Mike Du ¹, Yuchen Guo ¹, Kim López-Güell ¹, Núria Mercadé-Besora ¹, Xihang Chen ¹, Marta Alcalde-Herraiz ¹, Xintong Li ¹, Daniel Prieto-Alhambra ^{1,2}, and Edward Burn ¹

¹ Health Data Sciences Group, Nuffield Department of Orthopaedics, Rheumatology and Musculoskeletal Sciences, University of Oxford, United Kingdom ² Department of Medical Informatics, Erasmus University Medical Center, Rotterdam, The Netherlands ¶ Corresponding author

DOI: [10.21105/joss.09654](https://doi.org/10.21105/joss.09654)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Nick Golding](#)  

Reviewers:

- [@javier-gracia-tabuenca-tuni](#)
- [@pacoramon](#)

Submitted: 09 July 2025

Published: 23 September 2026

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

Real-world data (RWD) mapped to the Observational Medical Outcomes Partnership Common Data Model (OMOP CDM) offers a standardised framework for conducting observational health research across diverse data sources. However, identifying and summarising patient-level characteristics within this model often requires custom code, limiting efficiency and reproducibility. To address this, we developed the open-source PatientProfiles R package. This package streamlines the process of extracting demographic characteristics, computing intersections between cohorts and clinical events, and generating standard summaries of patient populations in OMOP CDM datasets.

Built on the tidyverse ([Wickham et al., 2019](#)) and omopgenerics infrastructure, PatientProfiles supports SQL translation for scalable database operations and includes comprehensive test coverage across multiple database systems. It provides a suite of functions grouped into demographics, intersections, summaries, utility, and mock data generation. The package is designed for transparency, modularity, and reusability in epidemiological workflows and is available via CRAN and GitHub, along with documentation and vignettes to support users.

Statement of need

Real-world data (RWD), routinely collected health data such as GP records, hospital data, and insurance claims data are valuable resources for conducting epidemiological research studies. However, with such data typically not collected primarily for research, different RWD sources can vary substantially in format and clinical coding systems. To overcome this difficulty a common data model (CDM) is often used. A CDM helps standardising data structures across various sources, enhancing data consistency, quality, and interoperability. A particularly popular data model is the Observational Medical Outcomes Partnership (OMOP) CDM, with more than 800 million patients' health care data transformed into this format ([Overhage et al., 2011](#)).

The OMOP CDM is a person-centric relational data model. Patients' data is spread across various tables related to different clinical domains with, for example, the *condition occurrence* table containing diagnoses while the *drug exposure* table contains drug prescriptions. These different clinical tables are all linked back to the *person* table which contains a unique identifier for each individual along with some key demographic data such as their date of birth.

Meanwhile, records in the *observation period* table define the period of calendar time over which an individual is followed up. (Blacketer, 2025)

One of the principal benefits of mapping data to a CDM is that it allows for the same analytic code to be run across different datasets. Developing well-tested and easy-to-use software for common analytic tasks can therefore bring significant benefits, both improving the speed at which analyses can be performed and improving quality by reducing the amount of study-specific bespoke code needing to be written.

Obtaining the characteristics of individuals is one of the most common first tasks when working with patient-level data. In almost all analyses specific characteristics of individuals will need to be identified, after which groups of individuals who share some specific common condition or characteristic need to be identified and relationships between these groups are described (for example the time between a given diagnosis and a health outcome of interest).

We created the PatientProfiles R package to support identifying patient characteristics in data mapped to the OMOP CDM. It provides functionality to obtain demographic information (such as age, sex, prior observation time, future observation time, and so on), describe intersections between different groups of patients, and summarise the results in a standard output format.

State of the field

PatientProfiles is part of the tidyverse-oriented OMOP R ecosystem described in *Tidy R programming with the OMOP Common Data Model* (Burn & Català, 2025), alongside packages such as CDMConnector, CodelistGenerator, CohortCharacteristics, CohortConstructor, OmopSketch, omock, and omopgenerics. These packages are examples from a growing ecosystem of OMOP R tools that supports different analytical needs. Together, they form a modular stack: CDMConnector and omopgenerics support OMOP CDM references, omock provides mock data, CohortConstructor creates and manipulates cohorts, PatientProfiles adds demographics and intersections, and CohortCharacteristics summarises and visualises cohort characteristics. The components can be combined as needed to build different analysis pipelines.

HADES is another established approach, providing a broad suite for end-to-end OMOP analytics, including population characterisation, causal effect estimation, and patient-level prediction (Schuemie et al., 2024). HADES emphasises specifying and executing large-scale observational analyses and exploring their standardised results through integrated workflows and user interfaces. PatientProfiles instead contributes a focused, composable layer for patient characteristics, intersections, and standard summaries, providing a complementary modular approach aligned with tidyverse-style R programming. Community projects such as HadesExtras (FinnGen, n.d.) illustrate efforts to extend HADES workflows; the two ecosystems are complementary and fit different needs rather than competing alternatives.

Design principles

PatientProfiles was designed to adhere to the tidyverse tidy design principles. The tidyverse is a collection of R packages designed for data science, offering a cohesive and consistent syntax for data manipulation, and analysis (Wickham et al., 2019). The dplyr package defines multiple methods that can be implemented to many different sources of data. Of particular relevance to working with OMOP CDM data which is typically stored in a database, the dbplyr package provides translations of dplyr methods to SQL.

The core dependency of PatientProfiles is the omopgenerics package (Català & Burn, 2026), which provides methods, classes and basic operations for packages working with data in the OMOP CDM format. It defines a central object, a `cdm_reference`, that provides a central reference to all the different OMOP CDM tables, along with various other S3 classes and methods that facilitate working with the data contained in this reference.

Development of the PatientProfiles R package

PatientProfiles was developed in accordance with best practices for R packages with the devtools and usethis R packages used for common development tasks. The core, general dependencies of the package include dplyr and tidyr for common data manipulations and dbplyr which provides translations to SQL. In addition the core dependency related to OMOP CDM data is the omopgenerics package which provides core classes and methods specific to this data format.

The PatientProfiles package includes functionality to create its own mock data in the OMOP CDM format. This mock data is used to test the package using the testthat framework (Wickham, 2011). Every line of the package is tested multiple times trying to account for various edge cases. Currently, the package is tested iteratively against different database management systems: PostgreSQL, SQL Server, Amazon Redshift, and DuckDB. In addition to unit tests, end-to-end integration tests of the package have been conducted to ensure the face validity of results.

The package is open-source and released via CRAN: <https://CRAN.R-project.org/package=PatientProfiles> (Català et al., 2026) (version 1.6.1, released 5 August 2026) and also available on GitHub: <https://github.com/darwin-eu/PatientProfiles> with its own website with more documentation and vignettes that cover the content of the package in more depth.

Overview of the PatientProfiles R package

PatientProfiles contains three main groups of functions (Figure 1). **Demographics** functions are used to add information contained in person and observation period tables to other tables or objects of interest. **Intersections** are used to intersect a table with an object of interest (it can be another table, a cohort of patients or a particular clinical concept). The **summarise** functions are used to create standard objects that summarise the content of a table of interest. Finally, the package also contains some complementary utility functions to, for example, create mock data.

PatientProfiles works with valid OMOP cohort tables; apart from generating mock data for examples and tests, cohort creation and modification are outside the scope of the package. Dedicated OMOP packages such as CohortConstructor (Burn et al., 2026), CDMConnector (Inberg et al., 2026), CohortGenerator (Sena et al., 2026), and other cohort-generation tools can be used to create and modify cohorts. PatientProfiles then adds characteristics, identifies intersections, and summarises the resulting cohort tables.




Figure 1: PatientProfiles functions blocks. Note that each demographic function has its own analogous *query* function to only add a query to the data, e.g. `addAge()` and `addAgeQuery()`.

Mock data

A reference to an OMOP CDM instance is needed to use PatientProfiles. In this simple tutorial we will use mock toy data produced by the same package, using the `omock` mock-data infrastructure (Du et al., 2025), and request that it is copied into an in-process duckdb database.

```
library(PatientProfiles)
library(CohortConstructor)
cdm <- mockPatientProfiles(numberIndividuals = 1000, source = "duckdb")
# use CohortConstructor to customise cohorts
cdm$my_flu_cohort <- cdm$cohort1 |>
  subsetCohorts(cohortId = 1, name = "my_flu_cohort") |>
  renameCohort(newCohortName = "flu")
cdm$target <- cdm$cohort2 |>
  renameCohort(
    cohortId = c(1L, 2L, 3L),
    newCohortName = c("covid_test", "flu_test", "asthma")
  ) |>
  dplyr::compute(name = "target")
```

Demographics

`addDemographics()` is used to characterise the demographics of a table. The table must be part of a `cdm_reference` object and contain a person identifier column (either `person_id` or `subject_id`). There are multiple columns that can be added with this function:

- *age*: the age at a certain `indexDate`. You can also add an *age group* column grouping individuals into different age ranges.
- *sex*: the sex of the individual.
- *prior observation*: the number of days between start of observation and `indexDate`.
- *future observation*: the number of days between `indexDate` and end of observation.
- *date of birth*: the birth date of the individual.

An example to add the demographics to a mock cohort table is:

```
cdm$my_flu_cohort |>
  addDemographics(
    indexDate = "cohort_start_date",
    ageGroup = list("children" = c(0, 17), "adult" = c(18, Inf))
  ) |>
  dplyr::glimpse()

## Rows: ??
## Columns: 10
## Database: DuckDB v1.0.0 [root@Darwin 23.4.0:R 4.4.1/:memory:]
## $ cohort_definition_id <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
## $ subject_id           <int> 117, 818, 634, 729, 886, 245, 761, 385, 597, ...
## $ cohort_start_date    <date> 1959-06-03, 1936-09-03, 1979-08-13, 2010-03-...
## $ cohort_end_date      <date> 1960-11-29, 1981-03-23, 2053-12-02, 2044-04-...
## $ age                  <int> 27, 2, 16, 36, 58, 20, 38, 82, 94, 77, 18, 5,...
## $ age_group            <chr> "adult", "children", "children", "adult", "ad...
## $ sex                  <chr> "Female", "Female", "Female", "Male", "Male",...
## $ prior_observation     <int> 10015, 976, 6068, 13221, 21371, 7397, 13961, ...
## $ future_observation    <int> 768, 41462, 32133, 40592, 20967, 10212, 20892...
## $ date_of_birth        <date> 1932-01-01, 1934-01-01, 1963-01-01, 1974-01-...
```

For each functionality, there is an individual function: `addAge()`, `addSex()`, `addPriorObservation()`,

`addFutureObservation()` or `addDateOfBirth()`.

Observation period id

The *observation_period* contains the period of time that an individual in the database is in observation. There might be multiple individual periods per person, but they cannot overlap each other. When doing analysis it can be of interest knowing if a certain date is in observation, whether the individual will be in observation after a certain time, and from which observation period is an observation. To do so we have two functions:

- `addInObservation()` to identify if an individual is in observation in a certain *window* with respect to an *indexDate*.
- `addObservationPeriodId()` to identify the observation period ordinal for that date.

```
cdm$my_flu_cohort |>
  addInObservation(
    indexDate = "cohort_start_date",
    window = list("obs_index_date" = c(0, 0), "in_1_year" = c(365, 365)),
    nameStyle = "{window_name}"
  ) |>
  addObservationPeriodId() |>
  dplyr::glimpse()

## Rows: ??
## Columns: 7
## Database: DuckDB v1.0.0 [root@Darwin 23.4.0:R 4.4.1/:memory:]
## $ cohort_definition_id <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
## $ subject_id           <int> 962, 1158, 4462, 351, 3556, 320, 1965, 2105, ...
## $ cohort_start_date    <date> 1995-07-09, 2016-12-27, 1990-10-23, 2018-06-...
## $ cohort_end_date      <date> 2019-06-14, 2017-02-15, 2018-04-27, 2018-06-...
## $ obs_index_date       <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
## $ in_1_year            <int> 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, ...
## $ observation_period_id <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
```

Query functions

Usually OMOP CDM instances are stored in SQL databases. The functions that we have seen take the original table add the new columns and save the result into a new table (name argument). Each function has its own homologous one (same name terminated with 'Query') that instead of saving the result to a table only returns a query to generate the table. For local instances both functions provide exactly the same result.

Intersections

PatientProfiles has 15 functions that are used to compute intersections between tables, cohorts, and concepts. Common function parameters are:

- *indexDate*: Name of the column that contains the date that will be the origin time of our calculations.
- *sensorDate*: Name of the column that contains the date to censor the observation window.
- *window*: Window of time with respect to the index date that we will consider relevant events.

There are 5 different function types:

- *Flag*: Creates a new integer column that can have three possible values: 1 if an event of interest is observed; 0 if the event is not observed; and NA if the individual is not in observation within that window.
- *Count*: Creates a new integer column with the number of observed events; NA is reported if the individual is not in observation in that window.
- *Date*: Creates a new date column containing the date of a certain event; NA is reported if the event is not observed or the individual is not in observation in that window.
- *Days*: Creates a new integer column containing the time difference to a certain event; NA is reported if the event is not observed or the individual is not in observation in that window.
- *Field*: Creates a new column containing a selected field from the first or last matching event in the window; NA is reported if no event is observed or the individual is not in observation in that window.

For the *Flag* and *Count* functions there are two extra parameters:

- *targetStartDate*: Name of the column that identifies the start of the event.
- *targetEndDate*: Name of the column identifying the end of the episode. If NULL, the event is considered to start and end on *targetStartDate*.

With the following code you can add the number of visits recorded in the prior year (*number_visits*) and a flag to see if there is a record of asthma at any time before the index date.

```
cdm$my_flu_cohort |>
  addTableIntersectCount(
    tableName = "visit_occurrence",
    window = c(-365, 0),
    nameStyle = "number_visits"
  ) |>
  addCohortIntersectFlag(
    targetCohortTable = "target",
    targetCohortId = 3,
    window = c(-Inf, 0),
    nameStyle = "prior_asthma"
  ) |>
  dplyr::glimpse()

## Rows: ??
## Columns: 6
## Database: DuckDB v1.0.0 [root@Darwin 23.4.0:R 4.4.1/:memory:]
## $ cohort_definition_id <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
## $ subject_id           <int> 117, 818, 634, 886, 245, 761, 597, 53, 124, 2...
## $ cohort_start_date    <date> 1959-06-03, 1936-09-03, 1979-08-13, 1996-07-...
## $ cohort_end_date      <date> 1960-11-29, 1981-03-23, 2053-12-02, 2048-10-...
## $ number_visits        <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
## $ prior_asthma         <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, ...
```

For the *Date*, *Days*, and *Field* functions there are two extra parameters:

- *targetDate*: Name of the column that contains the event of interest.
- *order*: Whether we are interested in the “first” or “last” event in the window.

The *Field* functions also require *field*, the name of the field to return from the selected event.

With the following code you would add the date of the next test (flu or COVID) after the index date:

```
cdm$my_flu_cohort |>
  addCohortIntersectDate(
    targetCohortTable = "target",
    targetCohortId = c(1, 2),
    window = c(1, Inf)
  ) |>
  dplyr::glimpse()

## Rows: ??
## Columns: 6
## Database: DuckDB v1.0.0 [root@Darwin 23.4.0:R 4.4.1/:memory:]
## $ cohort_definition_id <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
## $ subject_id           <int> 761, 597, 124, 285, 157, 348, 297, 919, 830, ...
## $ cohort_start_date    <date> 2037-03-23, 2093-08-24, 1981-12-13, 1993-01-...
## $ cohort_end_date      <date> 2058-09-28, 2141-02-14, 2020-03-08, 2006-09-...
## $ covid_test_1_to_inf  <date> 2077-08-29, NA, NA, NA, 2043-06-27, NA, NA, ...
## $ flu_test_1_to_inf    <date> NA, 2194-01-24, 2058-08-07, 2093-09-06, NA, ...
```

Note that each function has some arguments related to the intersecting target (cohort, concept or clinical table).

Summarise data

`summariseResult()` allows users to summarise multiple columns into multiple estimates (see `availableEstimates()`) in a standard output format, as shown below:

```
cdm$my_flu_cohort |>
  addDemographics() |>
  addTableIntersectCount(
    tableName = "visit_occurrence",
    window = c(-365, 0),
    nameStyle = "number_visits"
  ) |>
  addCohortIntersectFlag(
    targetCohortTable = "target",
    targetCohortId = 3,
    window = c(-Inf, 0),
    nameStyle = "prior_asthma"
  ) |>
  addCohortIntersectDate(
    targetCohortTable = "target",
    targetCohortId = c(1, 2),
    window = c(1, Inf)
  ) |>
  addCohortName() |>
  summariseResult(
    group = "cohort_name",
    strata = list("sex", c("sex", "prior_asthma")),
    variables = list(
      c("number_visits", "age"),
      c("covid_test_1_to_inf", "flu_test_1_to_inf")
    ),
    estimates = list(
      c("median", "q25", "q75"),
      c("min", "max")
    )
  )
```



```
) |>
dplyr::glimpse()

## Rows: 84
## Columns: 13
## $ result_id      <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
## $ cdm_name       <chr> "PP MOCK", "PP MOCK", "PP MOCK", "PP MOCK", "PP_M...
## $ group_name     <chr> "cohort_name", "cohort_name", "cohort_name", "coh...
## $ group_level    <chr> "flu", "flu", "flu", "flu", "flu", "flu", "flu", ...
## $ strata_name    <chr> "overall", "overall", "overall", "overall", "over...
## $ strata_level   <chr> "overall", "overall", "overall", "overall", "over...
## $ variable_name  <chr> "number records", "number subjects", "age", "age"...
## $ variable_level <chr> NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, NA, N...
## $ estimate_name  <chr> "count", "count", "median", "q25", "q75", "median...
## $ estimate_type  <chr> "integer", "integer", "integer", "integer", "inte...
## $ estimate_value <chr> "350", "350", "49", "23", "87", "0", "0", "0", "1...
## $ additional_name <chr> "overall", "overall", "overall", "overall", "over...
## $ additional_level <chr> "overall", "overall", "overall", "overall", "over..."
```

Conclusions

The PatientProfiles R package provides functionality to assist researchers working with data mapped to the OMOP CDM format. By basing the package around this data model which has a known structure the package could be developed with simple interfaces yet deep functionality. The package has already been used in published studies (Català et al., 2024; Mercadé-Besora et al., 2024) and is freely available to be used in future research.

Funding information

Development of the PatientProfiles R package was funded by the European Medicines Agency as part of the Data Analysis and Real World Interrogation Network (DARWIN EU®). This manuscript represents the views of the DARWIN EU® Coordination Centre only and cannot be interpreted as reflecting the views of the European Medicines Agency or the European Medicines Regulatory Network.

References

- Blacketer, C. (2025). *Definition and DDLs for the OMOP common data model (CDM)*.
- Burn, E., & Català, M. (2025). *Tidy r programming with the OMOP common data model*. Zenodo. <https://doi.org/10.5281/zenodo.17532124>
- Burn, E., Català, M., Mercadé-Besora, N., Alcalde-Herraiz, M., Du, M., Guo, Y., Chen, X., Lopez-Guell, K., & Rowlands, E. (2026). *CohortConstructor: Build and manipulate study cohorts using a common data model*. <https://doi.org/10.32614/CRAN.package.CohortConstructor>
- Català, M., & Burn, E. (2026). *Omopgenerics: Methods and classes for the OMOP common data model*. <https://doi.org/10.32614/CRAN.package.omopgenerics>
- Català, M., Guo, Y., Du, M., Lopez-Guell, K., Burn, E., & Mercade-Besora, N. (2026). *PatientProfiles: Identify characteristics of patients in the OMOP common data model*. <https://doi.org/10.32614/CRAN.package.PatientProfiles>
- Català, M., Mercadé-Besora, N., Kolde, R., Trinh, N. T. H., Roel, E., Burn, E., Rathod-Mistry, T., Kostka, K., Man, W. Y., Delmestri, A., Nordeng, H. M. E., Uusküla, A.,

- Duarte-Salles, T., Prieto-Alhambra, D., & Jödicke, A. M. (2024). The effectiveness of COVID-19 vaccines to prevent long COVID symptoms: Staggered cohort study of data from the UK, Spain, and Estonia. *The Lancet Respiratory Medicine*, 12, 226–240. [https://doi.org/10.1016/S2213-2600\(23\)00414-9](https://doi.org/10.1016/S2213-2600(23)00414-9)
- Du, M., Mercadé-Besora, N., Alcalde-Herraiz, M., Chen, X., Guo, Y., López-Güell, K., Burn, E., & Català, M. (2025). Omock: A R package for mock data generation for the observational medical outcomes partnership common data model. *Journal of Open Source Software*, 10(113), 8178. <https://doi.org/10.21105/joss.08178>
- FinnGen. (n.d.). *HadesExtras*. Retrieved September 9, 2026, from <https://github.com/FINNGEN/HadesExtras>
- Inberg, G., Black, A., Gorbachev, A., Burn, E., Català, M., & Nika, I. (2026). *CDMConnector: Connect to an OMOP common data model*. <https://doi.org/10.32614/CRAN.package.CDMConnector>
- Mercadé-Besora, N., Li, X., Kolde, R., Trinh, N. T., Sanchez-Santos, M. T., Man, W. Y., Roel, E., Reyes, C., Delmestri, A., Nordeng, H. M. E., Uusküla, A., Duarte-Salles, T., Prats, C., Prieto-Alhambra, D., Jödicke, A. M., & Català, M. (2024). The role of COVID-19 vaccines in preventing post-COVID-19 thromboembolic and cardiovascular complications. *Heart*, 110(9), 635–643. <https://doi.org/10.1136/heartjnl-2023-323483>
- Overhage, J. M., Ryan, P. B., Reich, C. G., Hartzema, A. G., & Stang, P. E. (2011). Validation of a common data model for active safety surveillance research. *Journal of the American Medical Informatics Association*, 19(1), 54–60. <https://doi.org/10.1136/amiajnl-2011-000376>
- Schuemie, M. J., Reps, J. M., Black, A., DeFalco, F., Evans, L., Fridgeirsson, E., Gilbert, J. P., Knoll, C., Lavalley, M., Rao, G., Rijnbeek, P., Sadowski, K., Sena, A., Swerdel, J., Williams, R. D., & Suchard, M. A. (2024). Health-analytics data to evidence suite (HADES): Open-source software for observational research. *Studies in Health Technology and Informatics*, 310, 966–970. <https://doi.org/10.3233/SHTI231108>
- Sena, A., Gilbert, J., Rao, G., Cruz, F. A., Schuemie, M., Science, O. H. D., & Informatics. (2026). *CohortGenerator: Cohort generation for the OMOP common data model*. <https://doi.org/10.32614/CRAN.package.CohortGenerator>
- Wickham, H. (2011). Testthat: Get started with testing. *The R Journal*, 3, 5–10. <https://doi.org/10.32614/rj-2011-002>
- Wickham, H., Averick, M., Bryan, J., Chang, W., McGowan, L. D., François, R., Golemund, G., Hayes, A., Henry, L., Hester, J., Kuhn, M., Pedersen, T. L., Miller, E., Bache, S. M., Müller, K., Ooms, J., Robinson, D., Seidel, D. P., Spinu, V., ... Yutani, H. (2019). Welcome to the tidyverse. *Journal of Open Source Software*, 4(43), 1686. <https://doi.org/10.21105/joss.01686>