

¹ PatientProfiles: An R package to identify patient characteristics in data mapped to the OMOP common data model

⁴ Martí Català  ¹¶, Mike Du  ¹, Yuchen Guo  ¹, Kim Lopez-Guell  ¹, Núria Mercadé-Besora  ¹, Xihang Chen  ¹, Marta Alcalde-Herraiz  ¹, Xintong Li  ¹, Daniel Prieto-Alhambra  ^{1,2}, and Edward Burn  ¹

⁷ 1 Health Data Sciences Group, Nuffield Department of Orthopaedics, Rheumatology and
⁸ Musculoskeletal Sciences, University of Oxford, United Kingdom 2 Department of Medical Informatics,
⁹ Erasmus University Medical Center, Rotterdam, The Netherlands ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Nick Golding](#)  

Reviewers:

- [@javier-gracia-tabuenac-tuni](#)
- [@pacoramón](#)

Submitted: 09 July 2025

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).
22
23
24

¹⁰ Summary

¹¹ Real-world data (RWD) mapped to the Observational Medical Outcomes Partnership Common
¹² Data Model (OMOP CDM) offers a standardised framework for conducting observational
¹³ health research across diverse data sources. However, identifying and summarising patient-
¹⁴ level characteristics within this model often requires custom code, limiting efficiency and
¹⁵ reproducibility. To address this, we developed the open-source PatientProfiles R package.
¹⁶ This package streamlines the process of extracting demographic characteristics, computing
¹⁷ intersections between cohorts and clinical events, and generating standard summaries of patient
¹⁸ populations in OMOP CDM datasets.

¹⁹ Built on the tidyverse and omopgenetics infrastructure, PatientProfiles supports SQL translation
²⁰ for scalable database operations and includes comprehensive test coverage across multiple
²¹ database systems. It provides a suite of functions grouped into demographics, intersections,
²² summaries, utility, and mock data generation. The package is designed for transparency,
²³ modularity, and reusability in epidemiological workflows and is available via CRAN and GitHub,
²⁴ along with documentation and vignettes to support users.
25

Statement of need

²⁶ Real-world data (RWD), routinely collected health data such as GP records, hospital data, and
²⁷ insurance claims data are valuable resources for conducting epidemiological research studies.
²⁸ However, with such data typically not collected primarily for research, different RWD sources
²⁹ can vary substantially in format and clinical coding systems. To overcome this difficulty a
³⁰ common data model (CDM) is often used. A CDM helps standardising data structures across
³¹ various sources, enhancing data consistency, quality, and interoperability. A particularly popular
³² data model is the Observational Medical Outcomes Partnership (OMOP) CDM, with more
³³ than 800 million patients' health care data transformed into this format ([Overhage et al., 2011](#)).
34

³⁵ The OMOP CDM is a person-centric relational data model. Patients' data is spread across
³⁶ various tables related to different clinical domains with, for example, the *condition occurrence*
³⁷ table containing diagnoses while the *drug exposure* table contains drug prescriptions. These
³⁸ different clinical tables are all linked back to the *person* table which contains a unique
³⁹ identifier for each individual along with some key demographic data such as their date of birth.

Meanwhile, records in the *observation period* table define the period of calendar time over which an individual is followed-up.([Blacketer, 2025](#))

One of the principal benefits of mapping data to a CDM is that it allows for the same analytic code to be run across different datasets. Developing well-tested and easy to use software for common analytic tasks can therefore bring significant benefits, both improving the speed in which analyses can be performed and improved quality by reducing the amount of study-specific bespoke code needing to be written.

Obtaining the characteristics of individuals is one of the most common first tasks when working with patient-level data. In almost all analyses specific characteristics of individuals will need to be identified, after which groups of individuals who share some specific common condition or characteristic need to be identified and relationships between these groups are described (for example the time between a given diagnosis and a health outcome of interest).

We created the PatientProfiles R package to support identifying patient characteristics in data mapped to the OMOP CDM. It provides functionality to obtain demographic information (such as age, sex, prior observation time, future observation time, and so on), describe intersections between different groups of patients, and summarise the results in a standard output format.

Design principles

PatientProfiles was designed to adhere to the tidyverse tidy design principles. The tidyverse is a collection of R packages designed for data science, offering a cohesive and consistent syntax for data manipulation, and analysis ([Wickham et al., 2019](#)). The dplyr package defines multiple methods that can be implemented to many different sources of data. Of particular relevance to working with OMOP CDM data which is typically stored in a database, the dbplyr package provides translations of dplyr methods to SQL.

The core dependency of PatientProfiles is the omopgenerics package ([Català & Burn, 2024](#)), which provides methods, classes and basic operations for packages working with data in the OMOP CDM format. It defines a central object, a `cdm_reference`, that provides a central reference to all the different OMOP CDM tables, along with various other S3 classes and methods that facilitate working with the data contained in this reference.

Development of the PatientProfiles R package

PatientProfiles was developed in accordance with best practices for R packages with the devtools and usethis R packages used for common development tasks. The core, general dependencies of the package include dplyr and tidyr for common data manipulations and dbplyr which provides translations to SQL. In addition the core dependency related to OMOP CDM data is the omopgenerics package which provides core classes and methods specific to this data format.

The PatientProfiles package includes functionality to create its own mock data in the OMOP CDM format. This mock data is used to test the package using the testthat framework ([Wickham, 2011](#)). Every line of the packages is tested multiple times trying to account for various edge cases. Currently, the package is tested iteratively against different database management systems: PostgreSQL, SQL Server, Amazon Redshift, and DuckDB. In addition to unit tests, end-to-end integration tests of the package have been conducted to ensure the face validity of results.

The package is open-source and released via CRAN: <https://CRAN.R-project.org/package=PatientProfiles> ([Català et al., 2025](#)) (version 1.4.1 as of 9th July 2025) and also available on github: <https://github.com/darwin-eu/PatientProfiles> with its own website with more documentation and vignettes that cover the content of the package more in depth.

86 Overview of the PatientProfiles R package

87 PatientProfiles contains three main groups of functions (Figure 1). **Demographics** functions
 88 are used to add information contained in person and observation period tables to other tables
 89 or objects of interest. **Intersections** are used to intersect a table with an object of interest
 90 (it can be another table, a cohort of patients or a paritciar clinical concept). The **summarise**
 91 functions are used to create standard objects that summarise the content of a table of interest.
 92 Finally, the package also contains some complementary utility functions to for example create
 93 mock data.

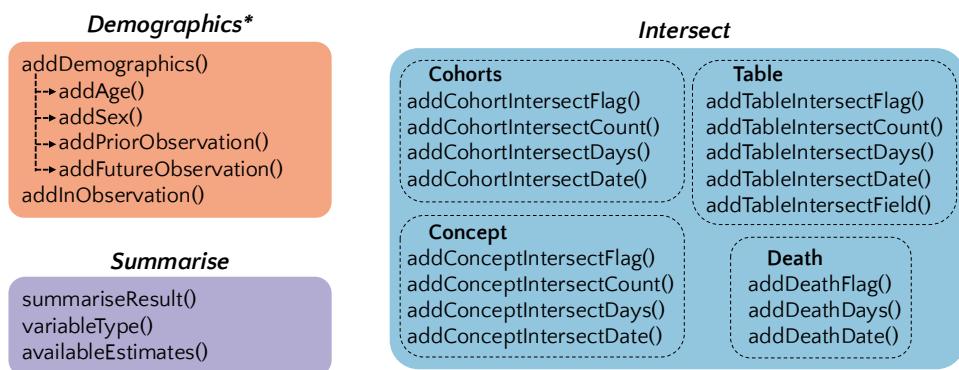


Figure 1: PatientProfiles functions blocks. Note that each demographic function has its own analogous query function to only add a query to the data, e.g. addAge() -> addAgeQuery().

94 Mock data

95 A reference to an OMOP CDM instance is needed to use PatientProfiles. In this simple tutorial
 96 we will use mock toy data produced by the same package. By default this toy data is copied
 97 into an in-process duckdb database.

```

98 library(PatientProfiles)
99 cdm <- mockPatientProfiles(numberIndividuals = 1000)
100 # to customise cohorts
101 cdm$my_flu_cohort <- cdm$cohort1 |>
102   dplyr::filter(cohort_definition_id == 1L) |>
103   omopgenerics::newCohortTable(cohortSetRef = dplyr::tibble(
104     cohort_definition_id = 1L, cohort_name = "flu"
105   ))
106 cdm$target <- cdm$cohort2 |>
107   omopgenerics::newCohortTable(cohortSetRef = dplyr::tibble(
108     cohort_definition_id = c(1L, 2L, 3L),
109     cohort_name = c("covid_test", "flu_test", "asthma")
110   ))

```

111 Demographics

112 addDemographics() is used to characterise the demographics of a table. The table is needed to
 113 be part of a `cdm_reference` object and to contain a person identifier column (either `person_id`
 114 or `subject_id`). There are multiple columns that can be added with this function:

- 115 ▪ `age`: the age at a certain `indexDate`. You can also add an `age group` column grouping
 116 individuals for different age group ranges.

- 117 ▪ *sex*: the sex of the individual.
- 118 ▪ *prior observation*: the number of days between start of observation and `indexDate`.
- 119 ▪ *future observation*: the number of days between `indexDate` and end of observation.
- 120 ▪ *date of birth*: the birth date of the individual.

121 An example to add the demographics to a mock cohort table is:

```

122 cdm$my_flu_cohort |>
123   addDemographics(
124     indexDate = "cohort_start_date",
125     ageGroup = list("children" = c(0, 17), "adult" = c(18, Inf))
126   ) |>
127   dplyr::glimpse()

128 ## Rows: ???
129 ## Columns: 10
130 ## Database: DuckDB v1.0.0 [root@Darwin 23.4.0:R 4.4.1/:memory:]
131 ## $ cohort_definition_id <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
132 ## $ subject_id           <int> 117, 818, 634, 729, 886, 245, 761, 385, 597, 53, ...
133 ## $ cohort_start_date    <date> 1959-06-03, 1936-09-03, 1979-08-13, 2010-03-
134 14, ...
135 ## $ cohort_end_date      <date> 1960-11-29, 1981-03-23, 2053-12-02, 2044-04-
136 28, ...
137 ## $ age                  <int> 27, 2, 16, 36, 58, 20, 38, 82, 94, 77, 18, 5, 5, ...
138 ## $ age_group            <chr> "adult", "children", "children", "adult", "adult"...
139 ## $ sex                  <chr> "Female", "Female", "Female", "Male", "Male", "Ma...
140 ## $ prior_observation    <int> 10015, 976, 6068, 13221, 21371, 7397, 13961, 3005...
141 ## $ future_observation   <int> 768, 41462, 32133, 40592, 20967, 10212, 20892, 19...
142 ## $ date_of_birth         <date> 1932-01-01, 1934-01-01, 1963-01-01, 1974-01-
143 01, ...
144 For each one of the functionalities there exist individual functions: addAge(), addSex(),
145 addPriorObservation(), addFutureObservation() and addDateOfBirth().

```

146 Observation period id

147 The *observation_period* contains the period of time that an individual in the database is in
148 observation. There might be multiple individual periods per person, but they can not overlap
149 each other. When doing analysis it can be of interest knowing if a certain date is in observation,
150 whether the individual will be in observation after a certain time, and from which observation
151 period is an observation. To do so we have two functions:

- 152 ▪ `addInObservation()` to identify if an individual is in observation in a certain *window*
153 respect an *indexDate*.
- 154 ▪ `addObservationPeriodId()` to identify in which observation period ordinal is that date
155 from.

```

cdm$gibleed |>
  addInObservation(
    indexDate = "cohort_start_date",
    window = list("obs_index_date" = c(0, 0), "in_1_year" = c(365, 365)),
    nameStyle = "{window_name}"
  ) |>
  addObservationPeriodId() |>
  dplyr::glimpse()

## Rows: ???
## Columns: 7

```

```

158 ## Database: DuckDB v1.0.0 [root@Darwin 23.4.0:R 4.4.1/:memory:]
159 ## $ cohort_definition_id <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
160 ## $ subject_id <int> 962, 1158, 4462, 351, 3556, 320, 1965, 2105, 259...
161 ## $ cohort_start_date <date> 1995-07-09, 2016-12-27, 1990-10-23, 2018-06-
162 28, ...
163 ## $ cohort_end_date <date> 2019-06-14, 2017-02-15, 2018-04-27, 2018-06-
164 29, ...
165 ## $ obs_index_date <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
166 ## $ in_1_year <int> 1, 0, 1, 0, 1, 1, 1, 1, 1, 1, 1, 1, 1, 0, 1, 1, 1, ...
167 ## $ observation_period_id <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...

```

168 Query functions

169 Usually OMOP CDM instances are stored in SQL databases. The functions that we have
 170 seen take the original table add the new columns and save the result into a new table (name
 171 argument). Each function has its own homologous one (same name terminated with 'Query')
 172 that instead of saving the result to a table only returns a query to generate the table. For
 173 local instances both functions provide exactly the same result.

174 Intersections

175 PatientProfiles has 15 functions that are used to compute intersections between tables.
 176 Common functions parameters:

- 177 ▪ indexDate Name of the column that contains the date that will be the origin time of
 178 our calculations.
- 179 ▪ censorDate Name of the column that contains the date to censor the observation
 180 window.
- 181 ▪ window Window of time respect to the index date that we will consider relevant events
 182 on.

183 There exist 4 different function types:

- 184 ▪ *Flag*: It creates a new integer column that can have 3 possible values: 1 whether an
 185 event of interest is observed; 0 if the event is not observed; NA if the individual is not in
 186 observation within that window.
- 187 ▪ *Count*: It creates a new integer column with the number of observed events, NA is
 188 reported if the individual is not in observation in that window.
- 189 ▪ *Date*: It creates a new date column that contains the date of a certain event, NA is
 190 reported if the event is not observed or the individual is not in observation in that
 191 window.
- 192 ▪ *Days*: It creates a new integer date with the time difference with a certain event, NA
 193 is reported if the event is not observed or the individual is not in observation in that
 194 window.

195 For the *Flag* and *Count* functions there are 2 extra parameters: - targetStartDate Nome of
 196 the column that identifies the start of the event. - targetEndDate Nome of the end of the
 197 episode, if NULL the event is considered to start and end on the targetStartDate.

198 With the following code you can add the number of visits recorded in the prior year
 199 (number_visits) and a flag to see if there is a record of a asthma test any time prior to the
 200 index date.

```

201 cdm$cohort1 |>
202   addTableIntersectCount(
203     tableName = "visit_occurrence",
204     window = c(-365, 0),
205     nameStyle = "number_visits"

```

```

206     ) |>
207     addCohortIntersectFlag(
208       cohortTableName = "cohort2",
209       cohortId = 3,
210       window = c(-Inf, 0),
211       nameStyle = "prior_asthma"
212     ) |>
213     dplyr::glimpse()

214 ## Rows: ???
215 ## Columns: 6
216 ## Database: DuckDB v1.0.0 [root@Darwin 23.4.0:R 4.4.1/:memory:]
217 ## $ cohort_definition_id <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
218 ## $ subject_id <int> 117, 818, 634, 886, 245, 761, 597, 53, 124, 285, ...
219 ## $ cohort_start_date <date> 1959-06-03, 1936-09-03, 1979-08-13, 1996-07-
220 06, ...
221 ## $ cohort_end_date <date> 1960-11-29, 1981-03-23, 2053-12-02, 2048-10-
222 18, ...
223 ## $ number_visits <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0...
224 ## $ prior_asthma <dbl> 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0...

225 For the Date and Days functions there are 2 extra parameters: - targetDate Name of the
226 column that contains the event of interest. - order Whether we are interested with the “first”
227 or “last” event in the window.

228 With the following code you would add which is the date of the next test (of flu or covid) after
229 the index date:

cdm$cohort1 |>
  addCohortIntersectDate(
    targetCohortTable = "cohort2",
    targetCohortId = c(1, 2),
    window = c(1, Inf)
  ) |>
  dplyr::glimpse()

230 ## Rows: ???
231 ## Columns: 6
232 ## Database: DuckDB v1.0.0 [root@Darwin 23.4.0:R 4.4.1/:memory:]
233 ## $ cohort_definition_id <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1...
234 ## $ subject_id <int> 761, 597, 124, 285, 157, 348, 297, 919, 830, 741, ...
235 ## $ cohort_start_date <date> 2037-03-23, 2093-08-24, 1981-12-13, 1993-01-
236 28, ...
237 ## $ cohort_end_date <date> 2058-09-28, 2141-02-14, 2020-03-08, 2006-09-
238 29, ...
239 ## $ covid_test_1_to_inf <date> 2077-08-29, NA, NA, NA, 2043-06-27, NA, NA, 2045...
240 ## $ flu_test_1_to_inf <date> NA, 2194-01-24, 2058-08-07, 2093-09-06, NA, 2041...

241 NOTE that each function has some arguments related to the intersecting target (cohort,
242 concept or clinical table).

```

243 Summarise data

244 summariseResult() is a function that allow the user to summarise multiple columns into
 245 multiple estimates (see availableEstimates()) into a standard format output, see the below
 246 example:

```
247 cdm$cohort1 |>
```

```

248   addCohortName() |>
249     summariseResult(
250       group = "cohort_name",
251       strata = list("sex", c("sex", "prior_asthma")),
252       variables = list(c("number_visits", "age"), c("covid_test_1_to_inf", "flu_test_1_to_"
253       estimates = list(c("median", "q25", "q75"), c("min", "max"))
254     ) |>
255     dplyr::glimpse()

256   ## Rows: 84
257   ## Columns: 13
258   ## $ result_id      <int> 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, 1, ...
259   ## $ cdm_name       <chr> "PP_MOCK", "PP_MOCK", "PP_MOCK", "PP_MOCK", ...
260   ## $ group_name     <chr> "cohort_name", "cohort_name", "cohort_name", ...
261   ## $ group_level    <chr> "flu", "flu", "flu", "flu", "flu", "flu", ...
262   ## $ strata_name    <chr> "overall", "overall", "overall", "overall", ...
263   ## $ strata_level   <chr> "overall", "overall", "overall", "overall", ...
264   ## $ variable_name  <chr> "number records", "number subjects", "age", ...
265   ## $ variable_level <chr> NA, ...
266   ## $ estimate_name  <chr> "count", "median", "q25", "q75", "median", ...
267   ## $ estimate_type   <chr> "integer", "integer", "integer", "integer", ...
268   ## $ estimate_value  <chr> "350", "350", "49", "23", "87", "0", "0", "0", ...
269   ...
270   ## $ additional_name <chr> "overall", "overall", "overall", "overall", ...
271   ## $ additional_level <chr> "overall", "overall", "overall", "overall", ...

```

272 Conclusions

273 The PatientProfiles R package provides functionality to assist researchers working with data
274 mapped to the OMOP CDM format. By basing the package around this data model which has
275 a known structure the package could be developed with simple interfaces yet deep functionality.
276 The package has already been used in published studies ([Català et al., 2024](#); [Mercadé-Besora](#)
277 et al., 2024) and is freely available to be used in future research.

278 Funding information

279 Development of the PatientProfiles R package was funded by the European Medicines Agency
280 as part of the Data Analysis and Real World Interrogation Network (DARWIN EU®). This
281 manuscript represents the views of the DARWIN EU® Coordination Centre only and cannot
282 be interpreted as reflecting the views of the European Medicines Agency or the European
283 Medicines Regulatory Network.

284 References

- 285 Blacketer, C. (2025). *Definition and DDLs for the OMOP common data model (CDM)*.
- 286 Català, M., & Burn, E. (2024). *Omopgenerics: Methods and classes for the OMOP com-*
287 *mon data model.* <https://darwin-eu.github.io/omopgenerics/>, <https://github.com/darwin-eu/omopgenerics>
- 288 Català, M., Guo, Y., Du, M., Lopez-Guell, K., Burn, E., & Mercade-Besora, N. (2025).
289 *PatientProfiles: Identify characteristics of patients in the OMOP common data model.*
290 <https://darwin-eu.github.io/PatientProfiles/>, <https://github.com/darwin-eu/PatientPro->
291 [files](#)

- 293 Català, M., Mercadé-Besora, N., Kolde, R., Trinh, N. T. H., Roel, E., Burn, E., Rathod-
294 Mistry, T., Kostka, K., Man, W. Y., Delmestri, A., Nordeng, H. M. E., Uusküla, A.,
295 Duarte-Salles, T., Prieto-Alhambra, D., & Jödicke, A. M. (2024). The effectiveness
296 of COVID-19 vaccines to prevent long COVID symptoms: Staggered cohort study of
297 data from the UK, spain, and estonia. *The Lancet Respiratory Medicine*, 12, 226–240.
298 [https://doi.org/10.1016/S2213-2600\(23\)00414-9](https://doi.org/10.1016/S2213-2600(23)00414-9)
- 299 Mercadé-Besora, N., Li, X., Kolde, R., Trinh, N. T., Sanchez-Santos, M. T., Man, W. Y., Roel,
300 E., Reyes, C., Delmestri, A., Nordeng, H. M. E., Uusküla, A., Duarte-Salles, T., Prats,
301 C., Prieto-Alhambra, D., Jödicke, A. M., & Català, M. (2024). The role of COVID-19
302 vaccines in preventing post-COVID-19 thromboembolic and cardiovascular complications.
303 *Heart*, 110(9), 635–643. <https://doi.org/10.1136/heartjnl-2023-323483>
- 304 Overhage, J. M., Ryan, P. B., Reich, C. G., Hartzema, A. G., & Stang, P. E. (2011).
305 Validation of a common data model for active safety surveillance research. *Journal of
306 the American Medical Informatics Association*, 19(1), 54–60. <https://doi.org/10.1136/amiajnl-2011-000376>
- 308 Wickham, H. (2011). Testthat: Get started with testing. *The R Journal*, 3, 5–10. https://journal.r-project.org/archive/2011-1/RJournal_2011-1_Wickham.pdf
- 310 Wickham, H., Averick, M., Bryan, J., Chang, W., McGowan, L. D., François, R., Grolemund,
311 Hayes, A., Henry, L., Hester, J., Kuhn, M., Pedersen, T. L., Miller, E., Bache, S. M.,
312 Müller, K., Ooms, J., Robinson, D., Seidel, D. P., Spinu, V., ... Yutani, H. (2019). Welcome
313 to the tidyverse. *Journal of Open Source Software*, 4(43), 1686. <https://doi.org/10.21105/joss.01686>
- 314