

¹ CD-Dynamax: A JAX-based Python package for continuous-discrete probabilistic state space modeling and inference

⁴ Matthew Levine  ^{1,2}¶ and Iñigo Urteaga  ^{3,4}¶

⁵ 1 Broad Institute of MIT and Harvard, USA ² Basis Research Institute, USA ³ BCAM – Basque Center
⁶ for Applied Mathematics, Spain ⁴ IKERBASQUE — Basque Foundation for Science, Spain ¶
⁷ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: ↗

Submitted: 09 December 2025

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).¹⁹

⁸ Summary

⁹ Dynamical systems, often modeled as stochastic differential equations (SDEs), are widely used mathematical tools to describe complex phenomena in various scientific fields, including engineering, economics, neuroscience, ecology, and climate science. In most real-world scenarios, observations of these systems are collected, subject to noise, at discrete and irregular time intervals, requiring nuanced modeling approaches.

¹⁰ Continuous-discrete state space models (CD-SSMs) provide a powerful probabilistic framework for modeling such systems ([Särkkä & Svensson, 2023](#)). These models describe the latent state evolution continuously over time according to an SDE, while noisy observations are obtained at specific, discrete time instants.

¹¹ Mathematically, a CD-SSM is described according to:

- A (possibly unknown) stochastic dynamical system, i.e.,

$$\frac{dx(t)}{dt} = f(x(t), t)dt + L(x(t), t)dw(t) ,$$

²⁰ where:

- $x \in \mathbb{R}^{d_x}$ and $x(0) \sim P(x_0)$,
- f is a (possibly time-dependent) drift function,
- L is a (possibly state and/or time-dependent) diffusion coefficient, and
- dw is the derivative of a d_x -dimensional Brownian motion with a covariance Q .

- Data is observed at arbitrary times $\{t_k\}_{k=1}^K$ via a measurement process

$$y(t) = h(x(t)) + \eta(t) ,$$

²⁶ where:

- $h : \mathbb{R}^{d_x} \rightarrow \mathbb{R}^{d_y}$; i.e., h transforms the d_x -dimensional state of the dynamical system $x(t)$ (a realization of the SDE) to a d_y -dimensional observation, and
- $\eta(t)$ is an independent and identically distributed noise process that corrupts the observations.

- The collection of CD-SSM parameters is denoted with θ , which may include parameters governing the latent dynamics (e.g., parameters of f and L) and/or parameters of the observation model (e.g., parameters of h and, in the Gaussian observation noise case, its covariance matrix R).

³⁵ This mathematical framework describes ***continuous (dynamics) - discrete (observation) state space models***. Under this formulation, CD-SSMs enable accurate modeling of dynamical systems where noisy data are collected at irregular intervals and the underlying processes

³⁸ evolve continuously over time. Note that a CD-SSM may also include inputs (i.e., controls),
³⁹ u_1, \dots, u_K , also occurring at times $\{t_k\}_{k=1}^K$ to steer the latent state dynamics and influence
⁴⁰ the observations.

⁴¹ When constructing a CD-SSM for a specific application, the modeler must define the functional
⁴² forms of the latent dynamics and observation models. Namely, there are two key design choices
⁴³ to make:

- ⁴⁴ 1. How do the latent states evolve over time? E.g., are the latent dynamics linear or
⁴⁵ nonlinear? What is the form of the drift governing the latent state evolution? How
⁴⁶ random is the evolution, i.e., what is the diffusion coefficient? How is the randomness
⁴⁷ structured, i.e., what is the covariance of the possibly multi-dimensional driving Brownian
⁴⁸ motion?
- ⁴⁹ 2. How are the observations related to the latent states? E.g., is the observation model
⁵⁰ linear or nonlinear? How noisy are the observations? Is the observation noise Gaussian
⁵¹ or non-Gaussian?

⁵² Due to the range and combination of choices available to the modeler when defining a CD-SSM,
⁵³ these can be tailored to capture the specific characteristics of the system being modeled,
⁵⁴ making CD-SSMs highly versatile and applicable across a wide range of domains.

⁵⁵ However, the flexibility and expressiveness of CD-SSMs come at the cost of increased complexity
⁵⁶ (in implementation and in computational resources) for state inference and parameter estimation
⁵⁷ tasks. Hence, efficient and robust tools for CD-SSM modeling, inference and learning are
⁵⁸ crucial to researchers in both theoretical and applied domains.

⁵⁹ In general, for a given set of observations $Y_K = [y(t_1), \dots, y(t_K)]$ of a CD-SSM, the main
⁶⁰ objectives of interest to theorists and practitioners are:

- ⁶¹ ▪ **Filtering**: to estimate the distribution of $x(t_K) | Y_K, \theta$.
- ⁶² ▪ **Smoothing**: to estimate the distribution of $\{x(t)\}_{t \leq t_K} | Y_K, \theta$.
- ⁶³ ▪ **Forecasting**: to estimate the distribution of $\{x(t)\}_{t > t_K} | Y_K, \theta$.
- ⁶⁴ ▪ **Parameter learning**: to estimate $\theta | Y_K$, either point-wise or in distribution.

⁶⁵ With this context in mind, we present cd-dynamax: a **CD-SSM modeling framework, with**
⁶⁶ **inference and learning algorithms** for the tasks outlined above.

⁶⁷ cd-dynamax is a JAX-based open-source Python package for continuous-discrete state space
⁶⁸ modeling and inference:

- ⁶⁹ ▪ cd-dynamax not only supports canonical CD-SSMs, e.g., the continuous-discrete linear
⁷⁰ dynamical system (CD-LGSSM), but allows for easy construction, modeling and inference
⁷¹ of flexible CD-SSMs as needed: the practitioner is only required to specify the drift
⁷² function f , the diffusion coefficient L of the latent SDE, and the observation function h
⁷³ for each specific model of interest.
- ⁷⁴ ▪ cd-dynamax's flexibility with respect to model definition means that users can define and
⁷⁵ work with a wide range of custom CD-SSMs that include mechanistic and/or flexible
⁷⁶ (e.g., neural network) components for the latent dynamics and observation models.
- ⁷⁷ ▪ cd-dynamax provides robust implementations of several, state-of-the-art continuous-
⁷⁸ discrete inference algorithms in an efficient, autodifferentiable framework, enabling the
⁷⁹ use of modern general-purpose libraries for parameter inference (e.g., stochastic gradient
⁸⁰ descent, Hamiltonian Monte Carlo). cd-dynamax is designed to allow users to flexibly
⁸¹ choose among a host of algorithms for their specific CD-SSM model and application.

⁸² Statement of need

⁸³ cd-dynamax is a JAX-based (Bradbury et al., 2018) open-source Python package for continuous-
⁸⁴ discrete state space modeling, where observations are made at specified discrete times (rather

85 than at regular intervals) and are driven by latent SDEs.

86 Other Python libraries exist for state space modeling (Aicher et al., 2025; Corenflos & Särkkä,
87 2021; Johnson, 2020; Lee et al., 2023; S. Linderman et al., 2020; Weiss et al., 2024), which are
88 primarily focused on Hidden Markov Models and discrete-time state space models, with their
89 corresponding Bayesian inference algorithms. Amongst the JAX-native libraries, dynamax (S.
90 W. Linderman et al., 2025) provides a comprehensive framework for discrete-time state space
91 modeling and inference, while (Lysy, 2023) offers particle filtering capabilities for discrete-time
92 state space models. The rodeo (Wu & Lysy, 2025) library provides JAX-based probabilistic
93 numerics tools for approximating likelihoods of noisy partially observed data under deterministic
94 continuous-time systems (i.e., ordinary differential equations) but, crucially, does not address
95 state-stochasticity.

96 To the best of our knowledge, there is no existing Python-based library that provides a
97 comprehensive framework for continuous-discrete state space modeling and inference.

98 cd-dynamax fills this gap by providing a user-friendly interface for defining CD-SSMs, along with
99 efficient implementations of state-of-the-art filtering, smoothing, forecasting, and parameter
100 learning algorithms specifically designed for continuous-discrete dynamical systems:

- 101 ▪ cd-dynamax extends the dynamax (S. W. Linderman et al., 2025) library by exploiting
102 diffraex (Kidger, 2021) —a JAX-based library providing numerical differential equation
103 solvers— to enable accurate and efficient simulation of continuous-time dynamics, as well
104 as gradient-based backpropagation through automatic differentiation. By relying on JAX,
105 cd-dynamax supports automatic autodifferentiation and just-in-time (JIT) compilation
106 for hardware acceleration on CPU, GPU, and TPU machines.
- 107 ▪ cd-dynamax is particularly suited for domains where continuous-time dynamics are
108 prevalent, and observations are collected at irregular intervals. Its internal structure
109 is designed to interact with any CD-SSM model (linear or nonlinear) in a unified way
110 (rather than being treated separately) for model definition, state-inference and system-
111 identification, producing a consistently structured library.
- 112 ▪ cd-dynamax is developed for both methodological researchers (interested in advancing
113 state space modeling and inference algorithms) and practitioners (interested in applying
114 CD-SSMs to real-world problems in fields such as systems biology, neuroscience, finance,
115 and engineering).

116 On the importance of continuous-time modeling

117 While continuous-time SSMs can be represented as discrete-time SSMs when sampling at fixed
118 intervals, there remain fundamental differences between these two modeling paradigms: the
119 former cannot be perfectly translated into the latter without loss of information or introduction
120 of artifacts.

121 Succinctly put, the relationship between the discrete and continuous frameworks is one of
122 approximation —a mapping that may involve significant information loss: while it is possible to
123 derive a discrete-time model from a continuous-time model through discretization, the reverse
124 process of obtaining a continuous-time model from a discrete-time model is generally ill-posed
125 and non-unique.

126 There are two fundamental issues introduced by discretization:

- 127 ▪ **Information Loss:** Sampling inevitably obscures the system's true dynamics, distorting
128 the signal in a process known as aliasing. Discretization results in the loss of inter-sample
129 behavior, and hence, a system can appear stable at the sampling points while actually
130 experiencing oscillations between them.
- 131 ▪ **Artifact Creation:** The choice of a discrete-time representation of a model, along
132 with the definition of its sampling interval, can create non-physical, artificial dynamics.

133 Discretization choices can introduce entirely new behaviors not present in the original
134 continuous-time system. For instance, naive sampling can induce the emergence (or
135 destruction) of chaos in simple discrete maps (entirely absent, or assured, in their
136 stable continuous-time counterparts) or instability of control-systems (where a stable
137 continuous-time system can be rendered catastrophically unstable by choosing incorrect
138 sampling intervals).

139 There are significant benefits of a continuous-time treatment of dynamical systems:

- 140 ▪ *Data agnosticism*: continuous-time models are inherently suited to handle real-world,
141 irregularly-spaced, and missing data: they model the underlying process, not the measure-
142 ment grid. Thus, continuous-time models naturally generalize to arbitrary observation
143 time grids without retraining or modification.
- 144 ▪ *Discretize at the end, not at the beginning*: a continuous-time framing allows for
145 discretization choices to be deferred until the final stages of analysis, enabling the
146 use of adaptive solvers and multi-rate sampling strategies that can better capture the
147 system's dynamics. A history of successes in numerical analysis has shown that delaying
148 discretization until the final stages of computation often leads to more accurate and
149 stable results.
- 150 ▪ *Physical interpretability*: continuous-time model parameters represent fundamental,
151 invariant physical properties of the system (e.g., reaction rates, physical constants,
152 clearance rates), whereas discrete-time parameters are a conflation of physical properties
153 and the choices of sampling intervals. In physics-aware modeling, prior knowledge is
154 often most naturally expressed in a continuous-time formulation.
- 155 ▪ *First-principles-based theory*: continuous-time models, expressed as differential equations,
156 are the “first principles” foundation for many physical and life sciences. The discrete-
157 time model is most accurately viewed as a subsequent numerical implementation or
158 approximation of this theoretical truth.

159 cd-dynamax aims to facilitate the adoption of continuous-discrete state space models in various
160 scientific domains, removing the user’s burden of implementing continuous-time dynamical
161 system models, solvers, and inference algorithms from scratch. Instead, users can focus on
162 defining their inductive modeling biases, their priors and high-level inference choices.

163 Additionally, it enables methodological research in continuous-time state space modeling and
164 inference by providing a single, flexible framework for experimentation, development, and
165 benchmarking of new algorithms.

166 CD-dynamax modeling and inference framework

167 For researchers and practitioners interested in continuous-time modeling, cd-dynamax provides
168 a robust, efficient, and user-friendly framework for CD-SSM modeling, inference, and learning.
169 cd-dynamax provides (i) continuous-discrete linear and nonlinear state space model definitions,
170 (ii) state-of-the-art filtering and smoothing algorithm implementations, and (iii) flexible tools
171 for system identification and model parameter estimation.

172 Currently, cd-dynamax offers:

- 173 1. A set of modular definitions of CD-SSM models, capturing both linear (CD-LGSSM) and
174 nonlinear (CD-NLGSSM) dynamics and observation functions, seamlessly incorporating
175 non-regular, noisy observation time instants. More information about state space
176 modeling can be found in the textbooks by Murphy (2023) and Särkkä & Svensson
177 (2023).
- 178 2. Low-level, probabilistic inference algorithms for filtering and smoothing. There exist many
179 algorithms for state inference and parameter estimation in CD-SSMs (Särkkä & Svensson,

180 2023), e.g., Extended Kalman Filter/Smoother, Unscented Kalman Filter/Smoother,
181 Particle Filter/Smoother. Specifically, cd-dynamax provides JAX implementations for:
182 ▪ Kalman filtering and smoothing for linear Gaussian CD-SSMs,
183 ▪ Extended Kalman filtering and smoothing for nonlinear CD-SSMs,
184 ▪ Unscented and Ensemble Kalman filtering for nonlinear CD-SSMs.
185 3. A high-level interface for constructing and fitting probabilistic SSMs. We provide readily
186 usable functions for:
187 ▪ point-estimation of model parameters via gradient-based or black-box optimization
188 ([Scipy](#), [Scipy-jaxopt](#)) of the (approximate) marginal log-likelihood
189 ▪ Bayesian posterior parameter estimation, i.e., Markov Chain Monte Carlo via the
190 [BlackJAX](#) library.
191 The publicly available cd-dynamax documentation and demos provide informative resources
192 describing the use of cd-dynamax for CD-SSM modeling, filtering, smoothing, forecasting and
193 fitting to data, for experts and newcomers alike.

194 Acknowledgements

195 Iñigo Urteaga acknowledges the support of "la Caixa" foundation's LCF/BQ/PI22/11910028
196 award, as well as funds by MICIU/AEI/10.13039/501100011033 and the BERC 2022-2025
197 program funded by the Basque Government. Matthew Levine acknowledges support by Basis
198 Research Institute and the Eric and Wendy Schmidt Center at the Broad Institute of MIT and
199 Harvard.

200 References

- 201 Aicher, C., Putcha, S., Nemeth, C., Fearnhead, P., & Fox, E. (2025). Stochastic gradient
202 MCMC for nonlinear state space models. *Bayesian Analysis*, 20(1). <https://doi.org/10.1214/23-BA1395>
- 204 Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., Necula, G.,
205 Paszke, A., VanderPlas, J., Wanderman-Milne, S., & Zhang, Q. (2018). JAX: Composable
206 transformations of Python+NumPy programs (Version 0.3.13). <http://github.com/google/jax>
- 208 Corenflos, A., & Särkkä, S. (2021). *Code companion for Bayesian Filtering and Smoothing*
209 (Version 1.0). <https://github.com/EEA-sensors/Bayesian-Filtering-and-Smoothing>
- 210 Johnson, M. J. (2020). PyHSMM: Bayesian inference in HSMMs and HMMs (Version 0.0.0).
211 <https://github.com/mattjj/pyhsmm>
- 212 Kidger, P. (2021). *On Neural Differential Equations* [PhD thesis, University of Oxford].
213 <https://docs.kidger.site/diffrax/>
- 214 Lee, M. A., Yi, B., Martín-Martín, R., Savarese, S., & Bohg, J. (2023). TorchFilter: Differentiable
215 particle filtering in PyTorch. <https://github.com/stanford-iprl-lab/torchfilter>
- 216 Linderman, S. W., Chang, P., Harper-Donnelly, G., Kara, A., Li, X., Duran-Martin, G., &
217 Murphy, K. (2025). Dynamax: A python package for probabilistic state space modeling
218 with JAX. *Journal of Open Source Software*, 10(108), 7069. <https://doi.org/10.21105/joss.07069>
- 220 Linderman, S., Antin, B., Zoltowski, D., & Glaser, J. (2020). SSM: Bayesian Learning and
221 Inference for State Space Models (Version 0.0.1). <https://github.com/lindermanlab/ssm>
- 222 Lysy, M. (2023). PFJax: Particle filtering in JAX. <https://github.com/mlsy/pfjax>

- 223 Murphy, K. P. (2023). *Probabilistic machine learning: Advanced topics*. MIT Press. <http://probml.github.io/book2>
- 224 Särkkä, S., & Svensson, L. (2023). *Bayesian filtering and smoothing* (Vol. 17). Cambridge University Press. <https://doi.org/10.1017/CBO9781139344203>
- 225 Weiss, R., Du, S., Grobler, J., Cournapeau, D., Pedregosa, F., Varoquaux, G., Mueller, A., Thirion, B., Nouri, D., Louppe, G., Vanderplas, J., Benediktsson, J., Buitinck, L., Korobov, M., McGibbon, R., Lattarini, S., Niculae, V., Gramfort, A., Lebedev, S., ... Rockhill, A. (2024). *Hmmlearn* (Version 0.3.2). <https://github.com/hmmlearn/hmmlearn>
- 226
- 227
- 228
- 229
- 230
- 231 Wu, M., & Lysy, M. (2025). *Rodeo: Probabilistic methods of parameter inference for ordinary differential equations*. <https://arxiv.org/abs/2506.21776>
- 232

DRAFT