

# <sup>1</sup> PyTupli: Enabling Collaboration in Offline Reinforcement Learning

<sup>3</sup> **Hannah Markgraf**  <sup>1</sup>¶, **Michael Eichelbeck**  <sup>1</sup>, **Daria Cappey** <sup>1</sup>, **Selin Demirtürk** <sup>1</sup>, **Yara Schattschneider** <sup>1</sup>, and **Matthias Althoff**  <sup>1</sup>

<sup>5</sup> **1** Technical University of Munich, Germany ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: ↗

Submitted: 10 December 2025

Published: unpublished

## License

Authors of papers retain copyright<sup>15</sup> and release the work under a<sup>16</sup> Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/))<sup>18</sup>

## <sup>6</sup> Summary

<sup>7</sup> Offline reinforcement learning (RL) offers a powerful way to derive effective decision-making policies for control problems using pre-collected data. However, managing and sharing such datasets in collaborative research settings can be challenging, as proper versioning, filtering, and access control are required. PyTupli is a free, Python-based toolkit designed to support this workflow by providing an easy-to-use yet specialized framework that remains independent of external service providers. Unlike centralized platforms, PyTupli is designed for self-hosted deployment, giving research teams full control over their data infrastructure and access policies. Its client library allows users to serialize and store control problems, upload new data, and retrieve precisely the subsets they need through flexible and expressive filters. Built-in metrics help evaluate dataset coverage and utility, informing both dataset selection and algorithm design for offline RL. To ensure secure deployment, a container-based server component offers authentication, role-based access control, and automated certificate provisioning. Together, these capabilities enable researchers to create, manage, exchange, and analyze datasets for offline RL in a robust and accessible manner.

## <sup>21</sup> Statement of need

<sup>22</sup> Reinforcement learning (RL) provides powerful methods for decision-making under uncertainty, but training RL agents typically requires extensive interaction with the underlying system or a computationally expensive simulator. Offline RL has emerged as a paradigm that alleviates this requirement by training agents solely on previously collected data ([Lange et al., 2012](#)). Such datasets contain tuples consisting of *state*, *action*, *next state*, and *reward*, obtained from recordings of real systems or generated through simulators.

<sup>28</sup> Effectively managing, sharing, and curating these datasets is essential for collaborative offline RL research, yet existing tools provide only partial support. Platforms like Zenodo or the free version of HuggingFace allow users to share finalized datasets with a general public but are not suitable for ongoing or private collaborations. Furthermore, they lack mechanisms for tracking internal dataset structure or performing efficient, fine-grained queries. The same applies to version control systems such as GitHub. Traditional databases are better suited for this purpose but require substantial expertise to design and maintain robust workflows.

<sup>35</sup> PyTupli addresses this gap by providing a dedicated Python toolkit for creating, storing, and sharing tuple datasets for custom environments. We provide a Docker container that each research group or collaboration can deploy independently to maintain complete control over their data. Each dataset is associated to a benchmark. Benchmarks are stored as JSON-serialized objects and can be linked to related artifacts, including time-series data, algorithm hyperparameters, or trained policies, allowing multiple benchmarks to reference

41 shared resources. Because it is built for scalable collaboration, PyTupli includes integrated  
 42 user and access management features.

43 Since the performance of offline RL algorithms often depends critically on the quality of the  
 44 underlying dataset (Asadulaev et al., 2025; Schweighofer et al., 2022; Suttle et al., 2025),  
 45 PyTupli offers extensive filtering capabilities. Whereas established offline RL datasets primarily  
 46 support filtering by entire episodes (Liu et al., 2023; Younis et al., 2024), PyTupli enables  
 47 tuple-level filtering as well. This can be used, for example, to rebalance datasets with sparse  
 48 rewards or selectively include transitions from specific regions of the state space. In addition,  
 49 PyTupli provides a suite of metrics for assessing dataset coverage and reward characteristics,  
 50 which can serve as predictors of offline RL performance (Asadulaev et al., 2025; Schweighofer  
 51 et al., 2022) and help guide algorithm selection.

## 52 Related Software

53 Publicly available tuple datasets have been essential for advancing offline RL algorithms  
 54 (Kostrikov et al., 2021; Kumar et al., 2020). These curated collections span various domains,  
 55 such as robotics and games (Formanek et al., 2023; Fu et al., 2020; Gulcehre et al., 2020;  
 56 Younis et al., 2024), power system control (Qin et al., 2022), and autonomous driving (Lee et  
 57 al., 2024; Liu et al., 2023), and are typically designed to support the development of improved  
 58 offline RL methods. As offline RL techniques mature, they are being applied to increasingly  
 59 diverse and task-specific control problems. Yet, to the best of our knowledge, no existing  
 60 toolbox supports the collaborative creation, management, and sharing of datasets for custom  
 61 environments. Minari (Younis et al., 2024) is the closest related tool, offering a repository of  
 62 standardized datasets along with functionality for filtering, environment reconstruction, and  
 63 recording new interactions. Its focus, however, remains on distributing datasets for established  
 64 benchmarks.

65 PyTupli instead targets collaborative workflows for custom control tasks. It enables researchers  
 66 to share datasets and benchmarks directly within project teams without dependence on a  
 67 central public server. Although Minari permits users to request publication on its official  
 68 platform, this approach is often unsuitable for ongoing or proprietary work or for datasets that  
 69 evolve over time. In addition, Minari does not provide tools for assessing dataset quality or  
 70 coverage, which PyTupli includes to support informed dataset curation and algorithm selection  
 71 in offline RL research.




Figure 1: Overview of the core functionalities of PyTupli.

## 72 Core Functionalities

73 We briefly describe the core functionalities of PyTupli which are illustrated in Figure 1.

74 **Benchmark and Artifact Management:** PyTupli enables users to store any control task defined  
 75 as a gymnasium environment. Environments may include parameterizable configurations that

76 produce variations in task dynamics. A fully specified environment is stored as a benchmark,  
 77 providing a unique reference for reproducible evaluation of controllers. Benchmarks can  
 78 reference additional data, such as exogenous inputs, time-series data, or pre-trained models.  
 79 These external units, referred to as artifacts, are stored independently and linked to multiple  
 80 benchmarks to avoid duplication.

81 **Data Management:** PyTupli supports ingesting, storing, and querying structured datasets (RL  
 82 tuples), including their relation to existing benchmark problems and any relevant metadata.  
 83 MongoDB serves as the backend, providing scalable storage and efficient retrieval for large  
 84 datasets. [Table 1](#) shows how ingestion and retrieval times scale with dataset size.

85 **Multi-User Collaboration and Access Control:** PyTupli facilitates collaborative workflows  
 86 through private, group, and public scopes. Based on their assigned role, users can store,  
 87 retrieve, delete, and publish objects. A server-side backend with FastAPI provides a REST  
 88 interface for secure, programmatic access, while token-based authentication ensures secure  
 89 sharing across teams or organizations.

90 **Integration with Existing Offline RL Infrastructure:** An interface to the gymnasium framework  
 91 enables users to record interactions with gymnasium environments as RL tuples. Furthermore,  
 92 retrieved tuple datasets are made available in a form that can easily be converted into the  
 93 dataset formats used by existing offline RL libraries such as d3rlpy ([Seno & Imai, 2022](#)).

94 **Assessment of Dataset Quality:** PyTupli implements metrics to evaluate dataset quality in  
 95 terms of coverage and expected returns. These metrics inform dataset selection and can  
 96 provide guidance for algorithm choice. Detailed formulations are provided in the following  
 97 section.

**Table 1:** Upload and download times for established datasets averaged over 10 runs. We chose two examples with low, medium, and high dataset size from the Minari collection. However, not only the size, but also the nature of observations has a strong influence on processing times.

Dataset	Size	M	N	Upload (s)	Download (s)
<b>D4RL</b>					
door/human-v2	3.5MB	25	7K	0.83	0.24
hammer/human-v2	6.2MB	25	11K	1.11	0.40
antmaze/medium-play-v1	605.2MB	1K	1M	173.26	126.62
<b>Atari</b>					
pitfall/expert-v0	351.7MB	10	65K	18.56	16.51
<b>Mujoco</b>					
ant/expert-v0	1.92GB	2K	2M	64.17	29.65
humanoid/expert-v0	2.95GB	1K	999K	96.61	55.32

## 98 Quality Metrics

### 99 Return-Based Metrics

100 Return-based metrics, such as trajectory quality (TQ) ([Schweighofer et al., 2022](#)) or the average  
 101 Q-value ([Asadulaev et al., 2025](#)) can inform algorithm decision. For example, Schweighofer et  
 102 al. ([2022](#)) show that behavioral cloning performs well despite its simplicity if the dataset has  
 103 a high TQ. For datasets with low TQ, algorithms from the deep Q-network family perform  
 104 well in their experiments as they do not constrain the learned policy towards the distribution  
 105 of the behavioral policy. TQ normalizes the average return of a dataset with respect to  
 106 the returns obtained by a minimal performant and an expert policy. To provide similar  
 107 insights without relying on such additional information, estimated relative return improvement  
 108 ([Swazinna et al., 2021](#)) relates the maximum trajectory return in the dataset to its average  
 109 return. While estimated return improvement and TQ operate on a trajectory level, average

110 Q-value estimation offers insights on the tuple level, making it a better predictor of offline RL  
111 performance (Asadulaev et al., 2025). It requires fitting a Q-function using Bellman updates,  
112 which is closely related to the objectives used in offline RL training. However, for continuous  
113 action spaces, the user needs to provide an evaluation policy to predict the next actions in the  
114 Bellman target. Such a policy can, for example, be obtained using behavioral cloning on the  
115 dataset.

### 116 Coverage-Based Metrics

117 An important question when assessing a dataset is whether the behavioral policy (or policies)  
118 used to generate it did explore the state and action space well enough to learn a meaningful  
119 target policy from the data. A common approach for quantifying explorativeness is to  
120 approximate the entropy of transition probabilities for the behavior policy. For discrete state  
121 and action spaces, Schweighofer et al. (2022) suggest to approximate this by counting unique  
122 state-action pairs. Optionally, this value can be normalized using a reference dataset  $\mathcal{D}_{\text{ref}}$  of  
123 same size, for example, the replay buffer collected during online training. Schweighofer et al.  
124 (2022) show that low state-action coverage values hinder performance of a large variety of  
125 algorithms. While counting unique state-action pairs aims at estimating the Shannon entropy  
126 of the transition probabilities, Suttle et al. (2025) suggest that datasets that maximize their  
127 proposed behavioral entropy metric support better offline RL performance. They suggest  
128 a  $k$ -nearest-neighbor estimator of the true behavioral entropy that relies on density-based  
129 weighting of different regions in the state-action space.

## 130 Acknowledgements

131 This work was partially supported by the German Research Foundation (AL 1185/9-1).

## 132 References

- 133 Asadulaev, A., Karray, F., & Takac, M. (2025). Expert or not? Assessing data quality in  
134 offline reinforcement learning. *arXiv Preprint arXiv:2510.12638*. <https://doi.org/10.48550/arXiv.2510.12638>
- 136 Formanek, C., Jeewa, A., Shock, J., & Pretorius, A. (2023). Off-the-grid MARL: Datasets  
137 and baselines for offline multi-agent reinforcement learning. 2442–2444. <https://doi.org/10.5555/3545946.3598961>
- 139 Fu, J., Kumar, A., Nachum, O., Tucker, G., & Levine, S. (2020). D4rl: Datasets for deep  
140 data-driven reinforcement learning. *arXiv Preprint arXiv:2004.07219*. <https://doi.org/https://doi.org/10.48550/arXiv.2004.07219>
- 142 Gulcehre, C., Wang, Z., Novikov, A., Paine, T., Gómez, S., Zolna, K., Agarwal, R., Merel, J. S.,  
143 Mankowitz, D. J., Paduraru, C., & others. (2020). RI unplugged: A suite of benchmarks  
144 for offline reinforcement learning. *Advances in Neural Information Processing Systems*, 33,  
145 7248–7259. <https://doi.org/10.5555/3495724.3496332>
- 146 Kostrikov, I., Nair, A., & Levine, S. (2021). Offline reinforcement learning with implicit  
147 q-learning. *arXiv Preprint arXiv:2110.06169*. <https://doi.org/10.48550/arXiv.2110.06169>
- 148 Kumar, A., Zhou, A., Tucker, G., & Levine, S. (2020). Conservative q-learning for offline  
149 reinforcement learning. *Advances in Neural Information Processing Systems*, 33, 1179–1191.  
150 <https://doi.org/10.5555/3495724.3495824>
- 151 Lange, S., Gabel, T., & Riedmiller, M. (2012). Batch reinforcement learning. In *Reinforcement  
152 learning: State-of-the-art* (pp. 45–73). Springer. [https://doi.org/10.1007/978-3-642-27645-3\\_2](https://doi.org/10.1007/978-3-642-27645-3_2)

- 154 Lee, D., Eom, C., & Kwon, M. (2024). AD4RL: Autonomous driving benchmarks for offline  
155 reinforcement learning with value-based dataset. *2024 IEEE International Conference on*  
156 *Robotics and Automation (ICRA)*, 8239–8245. <https://doi.org/10.1109/ICRA57147.2024.10610308>
- 158 Liu, Z., Guo, Z., Lin, H., Yao, Y., Zhu, J., Cen, Z., Hu, H., Yu, W., Zhang, T., Tan, J., &  
159 others. (2023). Datasets and benchmarks for offline safe reinforcement learning. *arXiv*  
160 *Preprint arXiv:2306.09303*. <https://doi.org/10.48550/arXiv.2306.09303>
- 161 Qin, R.-J., Zhang, X., Gao, S., Chen, X.-H., Li, Z., Zhang, W., & Yu, Y. (2022). NeoRL: A near  
162 real-world benchmark for offline reinforcement learning. *Advances in Neural Information*  
163 *Processing Systems*, 35, 24753–24765. <https://doi.org/10.5555/3600270.3602065>
- 164 Schweigofer, K., Dinu, M., Radler, A., Hofmarcher, M., Patil, V. P., Bitto-Nemling, A.,  
165 Eghbal-Zadeh, H., & Hochreiter, S. (2022). A dataset perspective on offline reinforcement  
166 learning. *Conference on Lifelong Learning Agents*, 470–517. <https://doi.org/10.48550/arXiv.2111.04714>
- 168 Seno, T., & Imai, M. (2022). d3rlpy: An offline deep reinforcement learning library. *Journal of*  
169 *Machine Learning Research*, 23(315), 1–20. <https://doi.org/10.5555/3586589.3586904>
- 170 Suttle, W. A., Suresh, A., & Nieto-Granda, C. (2025). Behavioral entropy-guided dataset  
171 generation for offline reinforcement learning. *arXiv Preprint arXiv:2502.04141*. <https://doi.org/10.48550/arXiv.2502.04141>
- 173 Swazinna, P., Udluft, S., & Runkler, T. (2021). Measuring data quality for dataset selection in  
174 offline reinforcement learning. *2021 IEEE Symposium Series on Computational Intelligence*  
175 (*SSCI*), 1–8. <https://doi.org/10.1109/SSCI50451.2021.9660006>
- 176 Younis, O. G., Perez-Vicente, R., Balis, J. U., Dudley, W., Davey, A., & Terry, J. K. (2024).  
177 *Minari* (Version 0.5.0). Zenodo. <https://doi.org/10.5281/zenodo.13767625>