

# <sup>1</sup> sklearn-migrator: Cross-version migration of <sup>2</sup> scikit-learn models for reproducible MLOps

<sup>3</sup> Alberto Andres Valdes Gonzalez  <sup>1</sup>

## **4 1 Independent Researcher (Chile)**

DOI: [10.xxxxxxx/draft](https://doi.org/10.xxxxxxx/draft)

Software

- [Review ↗](#)
  - [Repository ↗](#)
  - [Archive ↗](#)

**Editor:** 

**Submitted:** 13 December 2025

**Submitted:** 15 December

## License

Authors of papers retain copyright<sup>14</sup> and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).<sup>15,16</sup>

## Summary

sklearn-migrator is a Python library for **safely migrating scikit-learn models across versions** while preserving inference behavior and remaining robust to internal attribute changes. scikit-learn is among the most widely used machine learning libraries in both research and industry, and its estimators are commonly deployed in tabular-data domains such as finance, risk, operations, and marketing ([Kaggle, n.d.](#); [Pedregosa et al., 2011](#)). In these settings, model upgrades often coincide with dependency upgrades, container base-image updates, or security patching cycles—making version-to-version portability a practical requirement for production MI Ops.

The core problem is that standard persistence mechanisms (pickle/joblib) are **version-fragile**: models saved under one scikit-learn release may fail to load—or may load with altered behavior—under another. This limitation is explicitly cautioned in the official documentation and has been observed in empirical and experiential work ([Fitzpatrick & Manning, 2024](#); [Parida et al., 2025](#); [scikit-learn developers, n.d.](#)). sklearn-migrator addresses this gap by exporting supported estimators into **portable, JSON-compatible Python dictionaries** and reconstructing them in a different environment running a target scikit-learn version. The resulting payloads are readable, inspectable, and transportable across environments and teams, enabling long-term reproducibility and governance.

The migration proceeds in two stages. **Stage 1 (parity)**: the library captures a minimal set of prediction-critical attributes (constructor parameters and other parity-relevant settings) that guarantee parity of outputs across versions; these are used to reconstruct an equivalent estimator in the target version and validated under a strict tolerance (e.g.,  $\max |y_{in} - y_{out}| < 1e-2$ ). **Stage 2 (compatibility)**: remaining attributes are serialized with a version-aware policy that gracefully handles additions, removals, and renames so deserialization does not break across releases. This strategy is designed around the practical reality that estimator internals, defaults, and attribute names shift over time—even when the public API remains stable (Parida et al., 2025; scikit-learn developers, n.d.).

32 Concretely, consider version 0.21.3 exposing param\_1, param\_2, param\_3, param\_4 and  
33 version 1.7.0 exposing param\_1, param\_2, param\_3, param\_5. We use param\_1 and param\_2  
34 to enforce identical predictions across versions. Since param\_3 exists in both versions, it is  
35 stored in the payload and reassigned on load. For version-specific attributes, the serializer  
36 records values for both param\_4 and param\_5: in 0.21.3, the payload stores the real value of  
37 param\_4 and the **default** value that param\_5 would take in newer versions; in 1.7.0, it stores  
38 the real value of param\_5 and the default value that param\_4 would take in older versions.  
39 During deserialization, attributes are assigned using a simple try/except (or hasattr) so only  
40 valid attributes for the target version are set, and missing ones are safely skipped. This design  
41 makes migrations resilient to evolutionary changes such as renamed fields (e.g., affinity →  
42 metric), reorganized tree/boosting internals, and added default parameters across releases.

<sup>43</sup> This submission supports 21 models across supervised and unsupervised learning:

44     ▪ **Classification (7):** DecisionTreeClassifier, RandomForestClassifier, GradientBoostingClass  
45       LogisticRegression, KNeighborsClassifier, SVC, MLPClassifier  
46     ▪ **Regression (10):** DecisionTreeRegressor, RandomForestRegressor, GradientBoostingRegressor  
47       LinearRegression, Ridge, Lasso, KNeighborsRegressor, SVR, AdaBoostRegressor,  
48       MLPRegressor  
49     ▪ **Clustering (3):** AgglomerativeClustering, KMeans, MiniBatchKMeans  
50     ▪ **Dimensionality reduction (1):** PCA

51     Collectively, these estimators represent a large fraction of classical ML model families used  
52     in practice and commonly reported in practitioner surveys and applied workflows ([Kaggle, n.d.](#)). sklearn-migrator has been validated across 32 scikit-learn versions (0.21.3 → 1.7.2),  
53     covering 1,024 migration pairs, with automated, environment-isolated testing and strict parity  
54     checks. Continuous integration, unit tests, and an MIT license support reproducibility and  
55     adoption in production MLOps workflows.

## 57 Statement of need

58     Persisted scikit-learn models frequently break across library upgrades because internal attributes,  
59     defaults, and serialization details change over time. Standard persistence mechanisms (e.g.,  
60     pickle/joblib) are **version-fragile**: a model saved under one release may fail to load—or load with  
61     altered behavior—under another. This is explicitly cautioned in the official documentation and  
62     has been observed in empirical and experiential studies ([Fitzpatrick & Manning, 2024](#); [Parida et al., 2025](#); [scikit-learn developers, n.d.](#)). The resulting brittleness complicates production  
63     upgrades, environment migrations, cross-team sharing, and long-term reproducibility—especially  
64     in regulated or audit-heavy contexts where models must remain verifiable over time.

65     In practice, this fragility creates failure modes that are costly and hard to debug. A dependency  
66     upgrade may break deserialization of a mission-critical model artifact. Conversely, pinning  
67     old versions indefinitely increases security risk and operational burden. Teams often face  
68     an uncomfortable trade-off: **upgrade safely** (but risk breaking legacy models) or **freeze environments**  
69     (but accumulate technical debt). This is particularly acute in organizations  
70     that operate many model-serving services, notebooks, and batch pipelines—each with slightly  
71     different dependency constraints.

72     Given scikit-learn's broad adoption in research and industry, there is strong practical value  
73     in keeping legacy models usable across releases ([Kaggle, n.d.](#); [Pedregosa et al., 2011](#)). The  
74     estimator families supported in this submission—tree ensembles, linear/logistic models, nearest  
75     neighbors, support vector machines, and MLPs—are among the most commonly taught,  
76     prototyped, and deployed methods in applied machine learning. Unsupervised components such  
77     as k-means clustering and PCA are similarly ubiquitous in feature engineering and exploratory  
78     analysis pipelines.

79     sklearn-migrator addresses this need with a two-stage, version-aware (de)serialization strategy.  
80     First, it captures the minimal, prediction-critical attributes to guarantee parity of outputs  
81     between versions. Second, it serializes remaining attributes with explicit, version-conditioned  
82     defaults so that parameters added, removed, or renamed across releases do not break deserializa-  
83     tion. Unlike pickle/joblib, the library uses portable, JSON-compatible Python dictionaries,  
84     enabling safe transport, inspection, and storage independent of the original runtime. This design  
85     aligns with modern reproducibility needs and with ecosystem efforts focused on lightweight,  
86     inspectable artifacts and reproducible computational environments ([Fitzpatrick & Manning,](#)  
87     [2024](#); [Parida et al., 2025](#)).

88     The library targets practitioners and MLOps teams who must migrate or reproduce models  
89     across heterogeneous environments. It supports forward and backward migration and has been  
90     exercised across 32 scikit-learn releases (0.21.3 → 1.7.2), covering 1,024 version pairs with  
91     unit tests and environment-isolated validation. This foundation reduces upgrade risk today  
92     while remaining extensible to additional estimators and components in future releases.

## <sup>94</sup> Design and validation

### <sup>95</sup> Serialization format

<sup>96</sup> Each supported estimator is serialized into a Python dictionary containing:

- <sup>97</sup> 1. **Metadata**: source version, estimator type, and migration-relevant flags.
- <sup>98</sup> 2. **Parity-critical reconstruction parameters**: the minimal set of fields required to reconstruct
- <sup>99</sup> an estimator that produces matching predictions under strict tolerance.
- <sup>100</sup> 3. **Compatibility attributes**: additional learned attributes and internal fields stored with
- <sup>101</sup> version-aware rules, including explicit defaults for fields that exist only in some versions.

<sup>102</sup> To keep payloads portable, the library restricts values to JSON-encodable primitives (numbers,  
<sup>103</sup> strings, booleans, lists, dicts) and encodes arrays using standard Python lists where necessary.  
<sup>104</sup> This enables storage in plain JSON files, object storage, databases, or artifact registries, and  
<sup>105</sup> supports inspection and debugging without executing arbitrary code (a common concern with  
<sup>106</sup> pickle).

### <sup>107</sup> Validation methodology

<sup>108</sup> sklearn-migrator validates migrations through:

- <sup>109</sup> ▪ **Environment isolation** (e.g., containers) to ensure `version_in` and `version_out` represent
- <sup>110</sup> real installations.
- <sup>111</sup> ▪ **Fixed synthetic datasets** for deterministic evaluation.
- <sup>112</sup> ▪ **Strict parity checks** comparing source and migrated predictions under a tolerance (e.g.,  
<sup>113</sup> `1e-4`).

<sup>114</sup> The library has been tested across a full  $32 \times 32$  version compatibility matrix, totaling **1,024**  
<sup>115</sup> **migration pairs**, and across all supported estimators. This automated validation provides  
<sup>116</sup> confidence that the two-stage strategy behaves consistently across a large portion of the  
<sup>117</sup> modern scikit-learn release history.

## <sup>118</sup> Example

<sup>119</sup> The example below trains a `RandomForestRegressor`, serializes it to a portable (JSON-  
<sup>120</sup> compatible) dictionary, deserializes it, and checks prediction parity within a strict tolerance. In  
<sup>121</sup> practice, the `version_out` may correspond to a different scikit-learn installation (e.g., upgrading  
<sup>122</sup> from 1.5.x to 1.7.x).

```
import json
import numpy as np
import sklearn
from sklearn.datasets import make_regression
from sklearn.ensemble import RandomForestRegressor

from sklearn_migrator.regression.random_forest_reg import (
    serialize_random_forest_reg,
    deserialize_random_forest_reg,
)

# Train in the "source" environment
X, y = make_regression(n_samples=200, n_features=10, random_state=0)
src_version = sklearn.__version__

src_model = RandomForestRegressor(
    n_estimators=50, random_state=0
```

```

).fit(X, y)

# Serialize to a portable payload (JSON-compatible dict)
payload = serialize_random_forest_reg(src_model, version_in=src_version)

# (Optional) store as JSON
with open("model.json", "w") as f:
    json.dump(payload, f)

# --- In a different environment, load and deserialize ---
# (For illustration we reuse the same environment; in practice, version_out may differ.)
tgt_version = sklearn.__version__
with open("model.json") as f:
    payload_loaded = json.load(f)

tgt_model = deserialize_random_forest_reg(payload_loaded, version_out=tgt_version)

# Prediction parity check
y_src = src_model.predict(X)
y_tgt = tgt_model.predict(X)
assert np.max(np.abs(y_src - y_tgt)) < 1e-4
print("Prediction parity verified.")

```

123 Analogous functions exist for all covered estimators (classification, regression, clustering, and  
 124 dimensionality reduction). In a true two-environment workflow, the serialization code runs in the  
 125 source environment (e.g., 0.24.2) and the deserialization code runs in the target environment  
 126 (e.g., 1.7.2); the assertion above remains the same.

## Limitations

- 128 ▪ **Two environments required.** End-to-end validation relies on two isolated environments: a  
 129 source environment (`version_in`) to serialize and a target environment (`version_out`) to  
 130 deserialize and verify prediction parity. While this can be simulated on one machine, truly  
 131 isolated setups (e.g., Docker images) are recommended to avoid dependency leakage.
- 132 ▪ **Partial scikit-learn coverage.** scikit-learn contains many estimators and pipeline compo-  
 133 nents beyond the 21 currently supported. Additional models (e.g., pipelines, transformers,  
 134 and further estimators) are not yet supported but are planned for future releases. We  
 135 actively welcome community contributions to expand coverage. **Parity tolerance depends**  
 136 **on model family.** Some model families may be sensitive to floating-point or solver  
 137 differences across versions; the library uses strict tolerances by default but these can be  
 138 adjusted depending on operational requirements.

## Acknowledgements

- 140 We thank the scikit-learn core developers and contributors for their open-source work and  
 141 documentation, as well as the broader NumPy/SciPy/joblib ecosystems on which this project  
 142 depends. We are grateful to colleagues and early adopters for testing and feedback that shaped  
 143 the design, and to the JOSS editors and reviewers for guidance during the review process.
- 144 Fitzpatrick, P. C., & Manning, J. R. (2024). Davos: A python package “smuggler” for  
 145 constructing lightweight reproducible notebooks. *SoftwareX*, 25, 101614. <https://doi.org/10.1016/j.softx.2023.101614>
- 146 Kaggle. (n.d.). *Kaggle's state of machine learning and data science 2021*. <https://storage.googleapis.com/kaggle-media/surveys/Kaggle%27s%20State%20of%20Machine%20Learning%20and%20Data%20Science%202021.pdf>

- 149        [20Machine%20Learning%20and%20Data%20Science%202021.pdf](#).
- 150      Parida, S. K., Gerostathopoulos, I., & Bogner, J. (2025). *How do model export formats*  
151      *impact the development of ML-enabled systems? A case study on model integration.*  
152      <https://doi.org/10.48550/arXiv.2502.00429>
- 153      Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M.,  
154      Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau, D.,  
155      Brucher, M., Perrot, M., & Duchesnay, É. (2011). Scikit-learn: Machine learning in python.  
156      *Journal of Machine Learning Research*, 12, 2825–2830.
- 157      scikit-learn developers. (n.d.). *Model persistence*. [https://scikit-learn.org/stable/model\\_persistence.html](https://scikit-learn.org/stable/model_persistence.html).
- 158

DRAFT