

pyfive: A pure-Python HDF5 reader

Bryan N. Lawrence  ^{1,2}, Ezequiel Cimadevilla  ³, Wout De Nolf  ⁴, David Hassell  ^{1,2}, Jonathan Helmus⁵, Benjamin Hodel⁸, Brian Maranville  ⁶, Kai Mühlbauer  ⁷, and Valeriu Predoi  ^{1,2}

¹ National Center for Atmospheric Science (NCAS), United Kingdom.  ² Department of Meteorology, University of Reading, Reading, United Kingdom.  ³ Instituto de Física de Cantabria (IFCA), CSIC-Universidad de Cantabria, Santander, Spain. ⁴ European Synchrotron Radiation Facility (ESRF), Grenoble, France. ⁵ Astral Software Inc., USA. ⁶ NIST Center for Neutron Research ⁷ Institute of Geosciences, Meteorology Section, University of Bonn, Germany. ⁸ Independent Researcher, USA.

DOI: [10.21105/joss.09688](https://doi.org/10.21105/joss.09688)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: Richard Liu 

Reviewers:

- [@Zeitsperre](#)
- [@ayenpure](#)
- [@dostuffthatmatters](#)

Submitted: 15 December 2025

Published: 12 February 2026

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

pyfive is an open-source and thread-safe pure Python package for reading data stored in HDF5. While it is not a complete implementation of all the specifications and capabilities of HDF5, it includes all the core functionality necessary to read gridded datasets, whether stored contiguously or with chunks (with or without standard compression options). All data access is fully lazy as the data is only read from storage when the numpy data arrays are manipulated. Originally developed some years ago, the package has recently been expanded to support lazy data access, and to add missing features necessary for handling all the HDF5-based environmental data known to the authors. It is now a realistic option for production data access in environmental science and more broadly across other domains. The API is based on that of h5py (<https://github.com/h5py/h5py>, a Python shimmy over the HDF5 C-library which itself is not thread-safe), with some API extensions to help optimise remote access. With these extensions, coupled with thread safety, many of the limitations precluding the efficient use of HDF5 (and netCDF4) on cloud storage have been removed.

Statement of need

HDF5¹ (Folk et al., 2011) is arguably the most important data format in physical science. It is of particular importance in the environmental sciences that rely on the netCDF4² (Rew et al., 2006) data format, which itself uses the HDF data format underneath. From satellite missions to climate models and radar systems, the default binary format has been HDF5 for decades. While newer data formats are starting to get mindshare, there are petabytes, if not exabytes, of existing HDF5, and there remain many good use cases for creating new data in the HDF5 format today. However, despite its historical importance, there are few libraries available for reading HDF5 file data that do not depend on the official HDF5 library maintained by the HDF Group. In particular, apart from pyfive, there are no Python HDF5 libraries that address the data access needs of environmental science. While the HDF5 C library is reliable and performant, and battle-tested over decades, there are some caveats to depending upon it. Firstly, it is not thread-safe. Secondly, the underlying code is large and complex, and should anything happen to the financial stability of the HDF Group, it is not obvious it could be maintained. Finally, the code complexity also means that it is not suitable for developing bespoke code for data recovery in the case of partially corrupt data. From a long-term curation perspective these last two constraints present a major concern.

¹<https://www.hdfgroup.org/solutions/hdf5/>

²<https://www.unidata.ucar.edu/software/netcdf>

Reliance on a complex codebase controlled by a single private company presents significant challenges for long-term data access. Addressing these challenges requires well-documented data formats, the use of only those documented features, and the existence of publicly available code that can be sustainably maintained. The HDF Group have provided good documentation for the HDF5 format, but while there are communities of developers beyond those of the HDF Group, recent events suggest that given most of those developers and their existing funding are based in the USA, some spreading of risk would be desirable. To that end, a pure Python code covering the core HDF5 features of interest to the target scientific community, which is relatively small and maintained by an international constituency, provides some assurance that the community can maintain HDF5 access for the foreseeable future. A pure Python code also makes it easier to develop scripts that can work around data and metadata corruption should they occur, and has the additional advantage of being able to be deployed in resource or operating-system constrained environments (such as on mobile).

Current Status of pyfive

The original implementation of pyfive (by JH), which included all the low-level functionality to deal with the internals of an HDF5 file, was developed with POSIX access in mind. The recent upgrades were developed with the use cases of performant remote access to curated data as the primary motivation - including full support for lazy loading only the relevant parts of chunked datasets as they are needed.

Thread safety has become a concern given the wide use of Dask³ in Python-based analysis workflows, and this, coupled with a lack of user knowledge about how to efficiently use HDF5, has led to a community perception that HDF5 is not fit for remote access (especially on cloud storage). pyfive addresses thread safety by bypassing the underlying HDF5 C library. It addresses some of the issues with remote access by supporting the determination of whether or not a given file is cloud-optimised, and by optimising access to internal file metadata (in particular, the chunk indexes).

To improve internal metadata access, pyfive supports several levels of laziness for instantiating chunked datasets (variables). The default method preloads internal indices to make parallelism more efficient, but a completely lazy option without index loading is also possible. Neither method loads data until it is requested.

To be fully cloud-optimised, files need sensible chunking, and variables need contiguous indices. Chunking information has always been easy to determine. pyfive now also provides simple methods to expose information about internal file layout - both in API extensions, and via a new p5dump utility packaged with the pyfive library⁴. Either method allows one to determine whether the key internal “b-tree” indices are contiguous in storage, and to determine the parameters necessary to rewrite the data with contiguous indices. While pyfive itself cannot rewrite files to address chunking or layout, tools such as the HDF5 `repack` utility can do this very efficiently ([Hassell & Cimadevilla Alvarez, 2025](#)).

With the use of pyfive, suitably repacked and rechunked HDF5 data can now be considered “cloud-optimised”, insofar as with lazy loading, improved index handling, and thread-safety, there are no “format-induced” constraints on performance during remote access.

Acknowledgements

Most of the recent developments outlined have been supported by the UK Met Office and UKRI via 1) UK Excalibur Exascale programme (ExcaliWork), 2) the UKRI Digital Research Infrastructure programme (WacaSoft), and 3) the national capability funding of the UK

³<https://www.dask.org/>

⁴<https://pyfive.readthedocs.io/>

National Center for Atmospheric Science (NCAS). Ongoing maintenance of pyfive is expected to continue with NCAS national capability funding.

References

- Folk, M., Heber, G., Koziol, Q., Pourmal, E., & Robinson, D. (2011). An overview of the HDF5 technology suite and its applications. *Proceedings of the EDBT/ICDT 2011 Workshop on Array Databases*, 36–47. <https://doi.org/10.1145/1966895.1966900>
- Hassell, D., & Cimadevilla Alvarez, E. (2025). *Cmip7repack: Repack CMIP7 netCDF-4 datasets*. Zenodo. <https://doi.org/10.5281/zenodo.17550920>
- Rew, R., Hartnett, E., & Caron, J. (2006). NetCDF-4: Software implementing an enhanced data model for the geosciences. *22nd International Conference on Interactive Information Processing Systems for Meteorology, Oceanography and Hydrology*.