

¹ MFGLib: A Library for Mean-Field Games

² Xin Guo¹, Anran Hu³, Matteo Santamaria¹, Mahan Tajrobehkar¹, and Junzi
³ Zhang⁴

⁴ 1 University of California, Berkeley 3 Columbia University 4 Citadel Securities

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- Review ↗
 - Repository ↗
 - Archive ↗

Editor:

Submitted: 04 January 2026

Submitted: 31 January

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).¹⁵

Summary

Mean-field games (MFGs) provide scalable models for large-population strategic interactions. They approximate an N -player game by analyzing the limiting regime as $N \rightarrow \infty$, replacing explicit multi-agent interactions with the interaction between a representative agent and the population distribution (Lasry & Lions, 2007). It has been shown that the Nash equilibrium policy of the mean-field game is an ϵ -Nash equilibrium of the N -player game with $\epsilon = O(1/\sqrt{N})$ (Huang et al., 2006) and in practice even games with small N on the order of tens can be well-approximated by MFGs (Cabannes et al., 2021; Guo et al., 2019; Kizilkale et al., 2019). Due to their tractability, MFG models have become widely used in applications such as digital advertising, high-frequency trading, dynamic pricing, transportation, and behavioral modeling. Despite the rapid growth of the MFG literature, however, researchers and practitioners lack a unified, open-source software package for defining and solving their own MFG problems.

¹⁷ MGGLib is an open-source Python library that addresses this gap by providing:

- A modular and extensible API for defining arbitrary discrete-time finite-state MFGs
 - Implementations of state-of-the-art algorithms for computing (approximate) Nash equilibria
 - A collection of customizable benchmark environments drawn from the literature
 - Tight integration with Optuna ([Akiba et al., 2019](#)) to provide automatic hyperparameter selection
 - Clear documentation and examples, facilitating both research and industry use

²⁵ The library is implemented in Python, maintained on GitHub, and can be installed via pip
²⁶ install mfglib. Full documentation, tutorials, and example notebooks are available at
²⁷ <https://mfglib.readthedocs.io/en/latest/>.

Statement of Need

29 Large-population games, where a massive number of agents interact, are ubiquitous in fields such
 30 as economic modeling, crowd dynamics, and smart-grid management. However, as the number
 31 of agents N increases, traditional game-theoretic methods face exponential computational
 32 growth. MFGs provide a tractable approximation by modeling the collective behavior of
 33 agents in the infinite-population limit ($N \rightarrow \infty$). Despite the mathematical maturity of MFG
 34 theory, the research ecosystem lacks a standardized software framework. This has led to a
 35 fragmented landscape where researchers must frequently re-implement population dynamics
 36 and equilibrium solvers from scratch for every publication, a process that is time-intensive and
 37 prone to implementation errors.

³⁸ MFGLib is designed to provide a unified, modular foundation for MFG research. By offering a
³⁹ standardized API that decouples the environment definition from the equilibrium solver, the
⁴⁰ library lowers the barrier to entry for new researchers, enables rapid experimentation, and offers

⁴¹ practitioners a way to prototype MFG-based models without requiring deep expertise in game
⁴² theory or optimal control.

⁴³ MFGLib is for both researchers and practitioners from a broad range of backgrounds working
⁴⁴ on large-population strategic interactions who need a standardized, extensible toolkit to model,
⁴⁵ solve, and experiment with mean-field games.

⁴⁶ State of the Field

⁴⁷ The current landscape of game-theoretic software is bifurcated between general-purpose N -
⁴⁸ player frameworks and specialized research scripts. N -player libraries like QuantEcon ([Batista et al., 2024](#)), Nashpy ([Knight & Campbell, 2018](#)), and ilqgames ([Fridovich-Keil et al., 2020](#)) are
⁴⁹ restricted to small N and lack the mean-field approximations necessary to handle the complexity
⁵⁰ of large-scale games. Conversely, MFG-specific repositories such as gmfg-learning ([Cui & Koepll, 2022](#)) and entropic-mfg ([Benamou et al., 2019](#)) are typically “static” artifacts designed
⁵¹ for single papers; they lack the unit testing, extensible interfaces, and documentation required
⁵² for broader community adoption. Among the very few existing MFG libraries, OpenSpiel
⁵³ ([Lanctot et al., 2019](#)), a collection of environments and algorithms for research in reinforcement
⁵⁴ learning and planning in games, is the closest one to MFGLib. OpenSpiel has dedicated
⁵⁵ a module to MFGs implementing several environments and algorithms. However, it lacks
⁵⁶ customizability and a user-friendly API. In fact, according to its documentation, their code is
⁵⁷ still experimental and is only recommended for internal use.
⁵⁸

⁶⁰ Software Design

⁶¹ User-Friendly Environment Creation

⁶² Users can define custom MFG environments by providing reward functions, transition functions,
⁶³ and basic problem parameters (time horizon, state/action space sizes, initial distribution). The
⁶⁴ reward and transition functions are simple callables that map time and population distribution
⁶⁵ to tensors, allowing users to create environments with minimal code while maintaining
⁶⁶ mathematical clarity.

⁶⁷ Pre-Implemented Algorithms

⁶⁸ MFGLib implements several widely used algorithms, including **Online Mirror Descent** ([Perolat et al., 2021](#)), **Fictitious Play** ([Perrin et al., 2020](#)), **MFOMO** ([Guo et al., 2024](#)), **MFOMI** ([Hu & Zhang, 2025](#)), and **Prior Descent** ([Cui & Koepll, 2021](#)). These algorithms encompass many
⁶⁹ other existing methods as special cases, such as fixed point iteration and **GMF-V** ([Guo et al., 2019](#)). The unified solver interface returns policy iterates, exploitability scores (which evaluate
⁷⁰ closeness to Nash equilibrium), and cumulative runtimes, with optional real-time logging to
⁷¹ monitor convergence.
⁷²

⁷⁵ Automatic Hyperparameter Tuning

⁷⁶ Every algorithm requires hyperparameters that can drastically influence convergence properties.
⁷⁷ MFGLib provides a built-in tuner based on Optuna ([Akiba et al., 2019](#)) to automatically select
⁷⁸ optimal hyperparameters. The tuner can optimize across single instances or environment suites
⁷⁹ with multiple policy initializations and customizable metrics (e.g., shifted geometric mean of
⁸⁰ exploitability). Users can also implement their own metrics with minimal effort.

⁸¹ High-Dimensional Representation

⁸² MFGLib uses PyTorch tensors to represent policies, mean-fields, and rewards while preserving
⁸³ the original structure of state and action spaces. Rather than flattening high-dimensional spaces

⁸⁴ into one-dimensional representations, the library maintains their natural structure, providing
⁸⁵ higher interpretability and more flexible user interactions.

⁸⁶ Example

⁸⁷ We demonstrate the usage of our library with a brief example: a mean-field variant of Rock-
⁸⁸ Paper-Scissors introduced by Cui & Koepll (2021). Each agent chooses rock, paper or scissors,
⁸⁹ and receives a reward proportional to twice the number of beaten agents minus the number of
⁹⁰ agents that beat them.

⁹¹ Formally, the state and action spaces are $\mathcal{S} = \{0, R, P, S\}$ and $\mathcal{A} = \mathcal{S} \setminus \{0\}$, respectively.
⁹² The initial state distribution is fixed at $\mu_0(0) = 1$, and the game occurs over timesteps
⁹³ $\mathcal{T} = \{0, 1, \dots, T\}$. Agent rewards are specified by

$$\begin{aligned} r(R, a, \mu_t) &= 2 \cdot \mu_t(S) - 1 \cdot \mu_t(P) \\ r(P, a, \mu_t) &= 4 \cdot \mu_t(R) - 2 \cdot \mu_t(S) \\ r(S, a, \mu_t) &= 6 \cdot \mu_t(P) - 3 \cdot \mu_t(R) \end{aligned}$$

⁹⁴ The transition function allows agents to pick their next state directly and independently of the
⁹⁵ population; that is, for all $s, s' \in \mathcal{S}$ and $a \in \mathcal{A}$,

$$\Pr(s_{t+1} = s' \mid s_t = s, a_t = a) = \mathbf{1}_{\{s'\}}(a)$$

⁹⁶ We tune two algorithms – Online Mirror Descent (Perolat et al., 2021) and Occupation Measure
⁹⁷ Inclusion (Hu & Zhang, 2025) – on this environment.

```
from mfglib.env import Environment
from mfglib.alg import OnlineMirrorDescent, OccupationMeasureInclusion
from mfglib.tuning import GeometricMean

solve_kwargs = {"atol": None, "rtol": None, "max_iter": 300}

env = Environment.rock_paper_scissors(T=20)

omd_orig = OnlineMirrorDescent()
omi_orig = OccupationMeasureInclusion()

# Compute exploitability traces for untuned algorithms
_, omd_expls_orig, _ = omd_orig.solve(env, **solve_kwargs)
_, omi_expls_orig, _ = omi_orig.solve(env, **solve_kwargs)

metric = GeometricMean(shift=0.1)

# Optimize algorithms over hyperparameters
N_TRIALS = 50
omd_study = omd_orig.tune(
    metric=metric, envs=[env], n_trials=N_TRIALS, solve_kwargs=solve_kwargs
)
omi_study = omi_orig.tune(
    metric=metric, envs=[env], n_trials=N_TRIALS, solve_kwargs=solve_kwargs
)

# Initialize new algorithms objects from the tuning results
omd_tuned = omd_orig.from_study(omd_study)
omi_tuned = omi_orig.from_study(omi_study)
```

```

# Compute exploitability traces for tuned algorithms
_, omd_expls_tuned, _ = omd_tuned.solve(env, **solve_kwargs)
_, omi_expls_tuned, _ = omi_tuned.solve(env, **solve_kwargs)

98 Plotting the exploitability scores of the two algorithms, before and after tuning, we observe
99 that tuning significantly improves performance by achieving faster exploitability reduction.

```

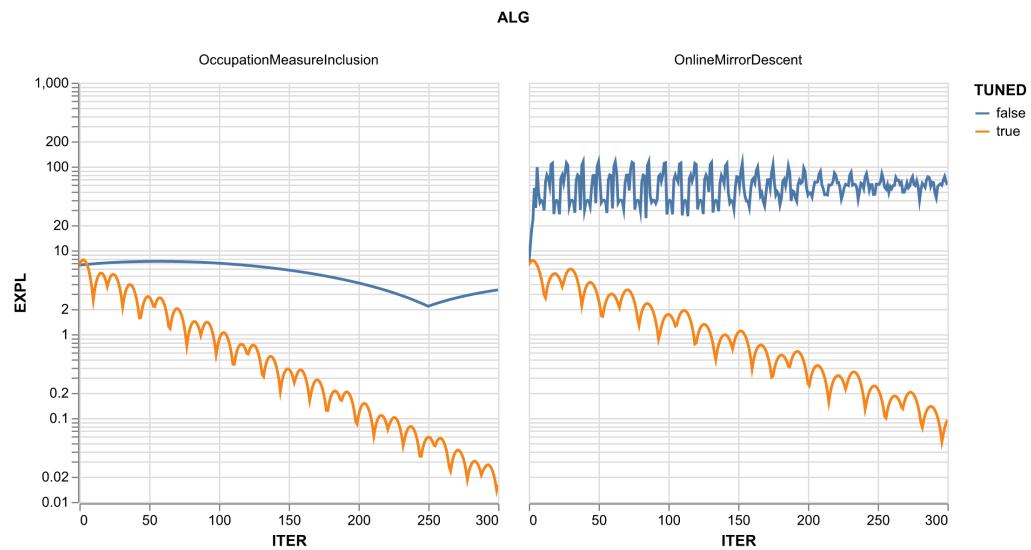


Figure 1: Exploitability curves before and after hyperparameter tuning

100 Research Impact Statement

101 The library has developed an active user community, including researchers and practitioners who
 102 have already used MFGLib in their work, contributed issues and pull requests on GitHub, and
 103 engaged with the tutorials and documentation. This activity demonstrates that the package
 104 is both useful to the community and actively maintained. Since its release, MFGLib has
 105 supported the development of several cutting-edge algorithms for MFGs, such as MF-OMO
 106 ([Guo et al., 2024](#)), MF-OMI ([Hu & Zhang, 2025](#)), and MESOB ([Guo et al., 2023](#)), as well
 107 as new models including MFG-MCDM ([Becherer et al., 2025](#)) and (α, β) -symmetric games
 108 ([Yardim & He, 2024](#)). It has also been used internally by Amazon Advertising for both research
 109 and production purposes. We believe it will continue to serve as an important building block
 110 for researchers in both academia and industry.

111 Acknowledgments

112 We thank the Amazon research teams for feedback and stress-testing the library in production
 113 settings. The authors would especially like to thank Sareh Nabi, Rabih Salhab, and Lihong Li of
 114 Amazon, Xiaoyang Liu of Columbia University, and Zhaoran Wang of Northwestern University
 115 for their valuable comments.

116 AI Usage Disclosure

117 The MFGLib library was developed without AI assistance; for the manuscript, Gemini 3 Flash
 118 was utilized for grammatical refinement and copy-editing. All AI-generated content was
 119 reviewed and edited for accuracy by the authors.

120 References

- 121 Akiba, T., Sano, S., Yanase, T., Ohta, T., & Koyama, M. (2019). Optuna: A next-
122 generation hyperparameter optimization framework. *Proceedings of the 25th ACM SIGKDD
123 International Conference on Knowledge Discovery & Data Mining*, 2623–2631.
- 124 Batista, Q., Coleman, C., Furusawa, Y., Hu, S., Lunagariya, S., Lyon, S., McKay, M., Oyama,
125 D., Sargent, T. J., Shi, Z., Stachurski, J., Winant, P., Watkins, N., Yang, Z., & Zhang,
126 H. (2024). QuantEcon.py: A community based python library for quantitative economics.
127 *Journal of Open Source Software*, 9(93), 5585. <https://doi.org/10.21105/joss.05585>
- 128 Becherer, D., Reisinger, C., & Tam, J. (2025). *Mean-field games of speedy information access
129 with observation costs*. <https://arxiv.org/abs/2309.07877>
- 130 Benamou, J.-D., Carlier, G., Di Marino, S., & Nenna, L. (2019). An entropy minimization
131 approach to second-order variational mean-field games. *Mathematical Models and Methods
132 in Applied Sciences*, 29(08), 1553–1583. <https://doi.org/10.1142/S0218202519500283>
- 133 Cabannes, T., Lauriere, M., Perolat, J., Marinier, R., Girgin, S., Perrin, S., Pietquin, O., Bayen,
134 A. M., Goubault, E., & Elie, R. (2021). Solving N-player dynamic routing games with
135 congestion: A mean field approach. *arXiv Preprint arXiv:2110.11943*.
- 136 Cui, K., & Koepll, H. (2021). Approximately solving mean field games via entropy-regularized
137 deep reinforcement learning. *International Conference on Artificial Intelligence and Statistics*,
138 1909–1917.
- 139 Cui, K., & Koepll, H. (2022). *Learning graphon mean field games and approximate nash
140 equilibria*. <https://arxiv.org/abs/2112.01280>
- 141 Fridovich-Keil, D., Ratner, E., Peters, L., Dragan, A. D., & Tomlin, C. J. (2020). Efficient
142 iterative linear-quadratic approximations for nonlinear multi-player general-sum differential
143 games. *2020 IEEE International Conference on Robotics and Automation (ICRA)*,
144 1475–1481. <https://doi.org/10.1109/ICRA40945.2020.9197129>
- 145 Guo, X., Hu, A., Xu, R., & Zhang, J. (2019). Learning mean-field games. *Advances in Neural
146 Information Processing Systems*, 32.
- 147 Guo, X., Hu, A., & Zhang, J. (2024). MF-OMO: An optimization formulation of mean-field
148 games. *SIAM Journal on Control and Optimization*, 62(1), 243–270. <https://doi.org/10.1137/22M1524084>
- 149 Guo, X., Li, L., Nabi, S., Salhab, R., & Zhang, J. (2023). MESOB: Balancing equilibria &
150 social optimality. <https://arxiv.org/abs/2307.07911>
- 151 Hu, A., & Zhang, J. (2025). MF-OML: Online mean-field reinforcement learning with
152 occupation measures for large population games. *Appl. Math. Optim.*, 92(3). <https://doi.org/10.1007/s00245-025-10328-5>
- 152 Huang, M., Malhamé, R. P., & Caines, P. E. (2006). Large population stochastic dynamic
153 games: closed-loop McKean-Vlasov systems and the Nash certainty equivalence principle.
154 *Communications in Information & Systems*, 6(3), 221–252. <https://doi.org/10.4310/CIS.2006.v6.n3.a5>
- 155 Kizilkale, A. C., Salhab, R., & Malhamé, R. P. (2019). An integral control formulation of
156 mean field game based large scale coordination of loads in smart grids. *Automatica*, 100,
157 312–322. <https://doi.org/10.1016/j.automatica.2018.11.029>
- 158 Knight, V., & Campbell, J. (2018). Nashpy: A python library for the computation of nash
159 equilibria. *Journal of Open Source Software*, 3(30), 904. <https://doi.org/10.21105/joss.00904>
- 160 Lanctot, M., Lockhart, E., Lespiau, J.-B., Zambaldi, V., Upadhyay, S., Pérolat, J., Srinivasan,
161
- 162
- 163
- 164
- 165

- 166 S., Finbarr Timbers, ', Tuyls, K., Omidshafiei, S., Hennes, D., Morrill, D., Muller, P.,
167 Ewalds, T., Faulkner, R., Kramár, J., Vylder, B. D., Saeta, B., Bradbury, J., ... Ryan-
168 Davis, J. (2019). OpenSpiel: A framework for reinforcement learning in games. *CoRR*,
169 [abs/1908.09453](https://arxiv.org/abs/1908.09453).[http://arxiv.org/abs/1908.09453](https://arxiv.org/abs/1908.09453)
- 170 Lasry, J.-M., & Lions, P.-L. (2007). Mean field games. *Japanese Journal of Mathematics*,
171 2(1), 229–260.
- 172 Perolat, J., Perrin, S., Elie, R., Laurière, M., Piliouras, G., Geist, M., Tuyls, K., & Pietquin,
173 O. (2021). Scaling up mean field games with online mirror descent. *arXiv Preprint*
174 [arXiv:2103.00623](https://arxiv.org/abs/2103.00623).
- 175 Perrin, S., Pérolat, J., Laurière, M., Geist, M., Elie, R., & Pietquin, O. (2020). Fictitious
176 play for mean field games: Continuous time analysis and applications. *Advances in Neural*
177 *Information Processing Systems*, 33, 13199–13213.
- 178 Yardim, B., & He, N. (2024). *Exploiting approximate symmetry for efficient multi-agent*
179 *reinforcement learning*.<https://arxiv.org/abs/2408.15173>

DRAFT