

¹ PipeOptz: A Python Library for Pipeline Optimization

² Nicolas Pengov  ^{1,2}, Théo Jaunet  ^{1,2}, and Romain Vuillemot  ^{1,2}

³ 1 Ecole Centrale de Lyon, France  2 LIRIS CNRS UMR 5205, France

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: Erik Whiting  

Reviewers:

- [@jqwang2373](#)
- [@gflofst](#)

Submitted: 21 November 2025

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)),

⁴ Summary

⁵ PipeOptz is a Python library for the building, visualizing, and fine-tuning of processing pipelines. It enables users to define a series of operations as a DAG (Directed Acyclic Graph) which parameters can then be optimized to achieve a desired outcome. The library is designed to be suitable for a wide range of applications, and is particularly suited for image processing where workflows can be dense and often require parameter tuning.

¹⁰ Statement of need

¹¹ In many scientific and engineering domains, complex data processing workflows are common. These workflows, i.e., pipelines, often consist of multiple steps, each with its own set of parameters and outputs. Finding the optimal set of parameters, and their individual influence, for a given task can be a tedious and time-consuming process which is often solved through manual trial and error. This is especially true in fields like image processing ([Walt et al., 2014](#)), where a sequence of filters and transformations is applied to an image, e.g., to find the best thresholding parameters. As opposed standard deep learning systems, pipelines, and their parameters have also the benefit to be more interpretable, through visualizations, and more easily reproducible.

²⁰ Existing tools for pipeline management often fall into two categories: heavy-weight workflow orchestration frameworks (e.g., Airflow ([Apache Airflow, n.d.](#)), Prefect ([Prefect, n.d.](#))) that are designed for large-scale data engineering tasks, or more specialized machine learning pipeline libraries (e.g., Scikit-learn pipelines ([Pedregosa et al., 2011](#))) that are focused on linear sequences of operations. From our experience, we found a need for a lightweight, flexible, and Pythonic library that is suited for the easy creation, visualization, and optimization of, non-linear pipelines directly within a Python script.

²⁷ PipeOptz addresses this need by providing an API for defining pipelines as Directed Acyclic Graphs (DAGs), with support for conditional branching and looping. In this graph, each node is a user-defined function in python, to ensure expressivity, and application to various domains. It integrates parameter optimization as a core feature, enabling users to define a search space for their pipeline's parameters and use various baseline optimization algorithms to find the best configuration.

³³ State of the field

³⁴ PipeOptz sits at the intersection of workflow orchestration, pipeline representation, and hyperparameter/black-box optimization. Workflow orchestrators such as Apache Airflow ([Apache Airflow, n.d.](#)) and Prefect ([Prefect, n.d.](#)) provide rich operational features (scheduling, monitoring, retries, deployments) and are well-suited for production batch workflows, but they are not designed as lightweight research libraries that expose the pipeline graph as a first-class object for iterative experimentation and optimization inside another tool. On the other end

40 of the spectrum, hyperparameter optimization (HPO) frameworks and Bayesian optimization
41 toolkits typically assume a user-written objective function and leave the internal structure of
42 the computational pipeline implicit in the user's code, which limits explicit control-flow nodes,
43 graph-level visualization, and step-wise traceability.

44 PipeOptz was created to support the needs of Descript, where we required (i) an expressive,
45 multi-step pipeline with explicit control-flow (loops and conditionals), (ii) heterogeneous tunable
46 parameters optimized against an application-specific loss function, and (iii) built-in graph
47 visualization and execution traceability for rapid research iteration. In principle, part of this
48 functionality could be implemented by extending an existing optimization toolkit with a custom
49 loss, but the core requirement here is the combination of "pipeline-as-a-graph" modeling,
50 control-flow nodes, and a clean separation between execution and optimization. Because
51 these constraints cut across the fundamental abstractions of existing orchestration/HPO tools,
52 we implemented a dedicated library designed as a reusable backend component for research
53 workflows rather than an operational orchestrator.

54 To clarify our position with respect to closely related optimization libraries, Bayesian optimization
55 frameworks such as BayesO (Kim & Choi, 2023) and pyGPGO (Jiménez & Ginebra, 2017)
56 focus primarily on sample-efficient search strategies for expensive black-box objectives. In these
57 systems the pipeline is usually encoded inside a single objective function, so the optimizer does
58 not directly represent intermediate steps or control flow. PipeOptz keeps the optimization
59 goal identical (minimize a user-defined loss), but makes the evaluation procedure explicit: the
60 workflow is represented as a graph of nodes with dependencies and control-flow constructs,
61 enabling node-level traceability and visualization while still treating the overall pipeline outcome
62 as the quantity to optimize.

63 AutoML frameworks such as NiaAML (Pečnik & Fister, 2021) also address "pipeline +
64 optimization", but they target the automated composition and tuning of machine-learning
65 pipelines within a predefined space of ML components and objectives. PipeOptz is intentionally
66 not ML-specific: it targets research workflows where the pipeline steps are arbitrary Python
67 functions and the loss can encode domain-specific criteria (e.g., balancing geometric accuracy
68 and the number of extracted targets), making it suitable as a backend for alternative approaches
69 beyond conventional ML pipelines.

70 Finally, some optimization problems are best addressed by algebraic modeling and solver-based
71 approaches. Linopy (Hofmann, 2023), for example, provides a modeling layer for linear and
72 mixed-integer optimization with labeled n-dimensional variables and solver backends. PipeOptz
73 is complementary: it targets workflows whose objective is evaluated by executing an end-to-end
74 pipeline and cannot be naturally expressed as a linear/mixed-integer model.

75 Software Design

76 PipeOptz is designed to make research pipelines explicit and optimizable while keeping them
77 lightweight and fully Python-native. The main design trade-off is to favor expressivity and
78 traceability over an "objective-function-only" interface: instead of hiding the workflow inside
79 a single function, PipeOptz represents it as a pipeline graph with explicit dependencies and
80 control-flow nodes. This matters in research workflows where debugging, profiling, and iterating
81 on multi-step processing chains is as important as finding good parameter values.

82 PipeOptz is built around the following core concepts:

83 **▪ Node:** The basic building block of a pipeline. A Node wraps a single Python function
84 and its parameters. To support non-linear workflows beyond simple DAG composition,
85 we provide dedicated control-flow nodes that embed sub-pipelines:

- 86 – NodeIf: for conditional branching (if/else).
87 – NodeFor: for 'for' loops.
88 – NodeWhile: for 'while' loops.

- 89 ▪ **Pipeline**: A Pipeline holds the entire workflow. Nodes are added to the pipeline with
90 their dependencies, forming a DAG. The pipeline manages execution by following a
91 topological order ([Kahn, 1962](#)). During execution, the library can cache node outputs
92 and record per-node execution time, supporting fine-grained inspection and iterative
93 refinement of complex pipelines.
- 94 ▪ **Parameter**: A Parameter defines the type and search space for a value to be optimized.
95 PipeOptz provides several types of parameters: IntParameter, FloatParameter,
96 ChoiceParameter, MultiChoiceParameter, and BoolParameter.
- 97 ▪ **PipelineOptimizer**: The optimization layer is separated from pipeline execution. It
98 takes a pipeline, a set of parameters to optimize, and a user-defined loss function to
99 minimize, and then evaluates candidate configurations by running the pipeline. PipeOptz
100 provides several baseline optimization strategies :
- 101 – Grid Search (GS),
 - 102 – Bayesian Optimization (BO) ([Shahriari et al., 2016; Snoek et al., 2012](#)),
 - 103 – Ant Colony Optimization (ACO) ([Dorigo & Gambardella, 1997](#)),
 - 104 – Simulated Annealing (SA) ([Kirkpatrick et al., 1983](#)),
 - 105 – Particle Swarm Optimization (PSO) ([Kennedy & Eberhart, 1995](#))
 - 106 – Genetic Algorithm (GA) ([Holland, 1975](#)).

107 The library also provides features for: - **Visualization**: Pipelines can be visualized as graphs
108 using the `to_dot` and `to_image` methods, which generate Graphviz dot files and PNG images
109 ([Gansner & North, 2000](#)). - **Serialization**: Pipelines can be saved to and loaded from JSON
110 files using the `to_json` and `from_json` methods, enabling reuse and sharing beyond a single
111 Python script (the `.dot` export is used for visualization and does not capture full pipeline
112 semantics).

113 Research Impact Statement

114 PipeOptz is currently used as a workflow-level optimization backend in ongoing applied research
115 prototypes where the target objective is not a standard machine-learning training loss, but an
116 application-specific loss computed by executing a multi-step processing workflow. In these
117 settings, practitioners need to iterate quickly over non-linear pipelines (including branching and
118 loops) and tune heterogeneous parameters while keeping the workflow explicit, inspectable,
119 and reproducible.

120 As this work is ongoing, we focus on credible near-term significance and reusability signals.
121 PipeOptz is distributed as a Python package via PyPI, released under an OSI-approved
122 license, and includes continuous integration, automated tests, and structured documentation
123 with runnable examples. The repository provides executable examples demonstrating core
124 capabilities (including control-flow pipelines and end-to-end optimization), and the runtime
125 interface exposes node-level outputs and execution timing to support debugging and profiling
126 of research workflows. These materials make PipeOptz reusable by other researchers who
127 need to define, visualize, and optimize DAG-based pipelines in a lightweight, Pythonic way,
128 especially in image-processing and related scientific workflows.

129 AI usage disclosure

130 We used generative-AI-assisted developer tools during software development and documentation
131 writing, but not for drafting the JOSS manuscript.

132 Code: GitHub Copilot (Visual Studio Code extension) (last version used : 1.104.1) and Gemini
133 Code Assist (Visual Studio Code extension) (last version used : 2.53) were used primarily for
134 code completion and small refactoring suggestions during implementation. No large, unverified

135 code blocks were accepted as-is; all AI-assisted edits were reviewed, tested, and integrated by
 136 the authors, who made the primary architectural and design decisions.

137 Documentation: Gemini Code Assist was used to accelerate writing of repetitive API
 138 documentation (docstrings, README sections). All generated text was manually reviewed
 139 and edited for correctness and consistency with the implemented behavior.

140 Manuscript: No generative AI tools were used to write paper.md.

141 Audience

142 PipeOptz is intended for researchers, data scientists, and engineers who need to build, visualize,
 143 and optimize data processing workflows in Python. It is particularly useful for those working in
 144 image processing, computer vision, and other scientific domains where pipeline-based workflows
 145 are common.

146 Example Usage

147 The following example demonstrates how to use PipeOptz to find the minimum of a simple
 148 function $f(x, y) = (x - 3)^2 + (y + 1)^2$ using Bayesian Optimization. The pipeline is constructed
 149 from multiple nodes to showcase the graph-based approach.

```
from pipeoptz import Pipeline, Node, FloatParameter, PipelineOptimizer

# Define the functions for the nodes
def squared_error(x, y):
    return (x - y)**2

def add(x, y):
    return x + y

# Create the pipeline
pipe = Pipeline("SimplePipeline")
pipe.add_node(Node("X", squared_error, fixed_params={"x": 0, "y": -3}))
pipe.add_node(Node("Y", squared_error, fixed_params={"x": -1, "y": 0}))
pipe.add_node(Node("Add", add), predecessors={"x": "X", "y": "Y"})

# The loss is the function's output, as we want to minimize it
def loss_func(result, _):
    return result

# Set up the optimizer with tunable parameters
optimizer = PipelineOptimizer(pipe, loss_function=loss_func)
optimizer.add_param(FloatParameter("X", "x", -5.0, 5.0))
optimizer.add_param(FloatParameter("Y", "y", -5.0, 5.0))

# Run the Bayesian Optimization
# We provide a dummy dataset ([{}]) and [0] as this example does not depend on external
best_params, loss_log = optimizer.optimize([{}], [0], method="BO", iterations=25, init_p
```

150 This script will search for the optimal values for X.x and Y.y that minimize the final output of
 151 the pipeline. The expected output will show the best parameters found, which should be close

¹⁵² to `{'X.x': 3.0, 'Y.y': -1.0}`, and a final loss close to 0.

¹⁵³ We can visualize the pipeline using Graphviz.

```
from PIL import Image
pipe.to_image("pipeline.png")
im = Image.open("pipeline.png")
im.show()
```

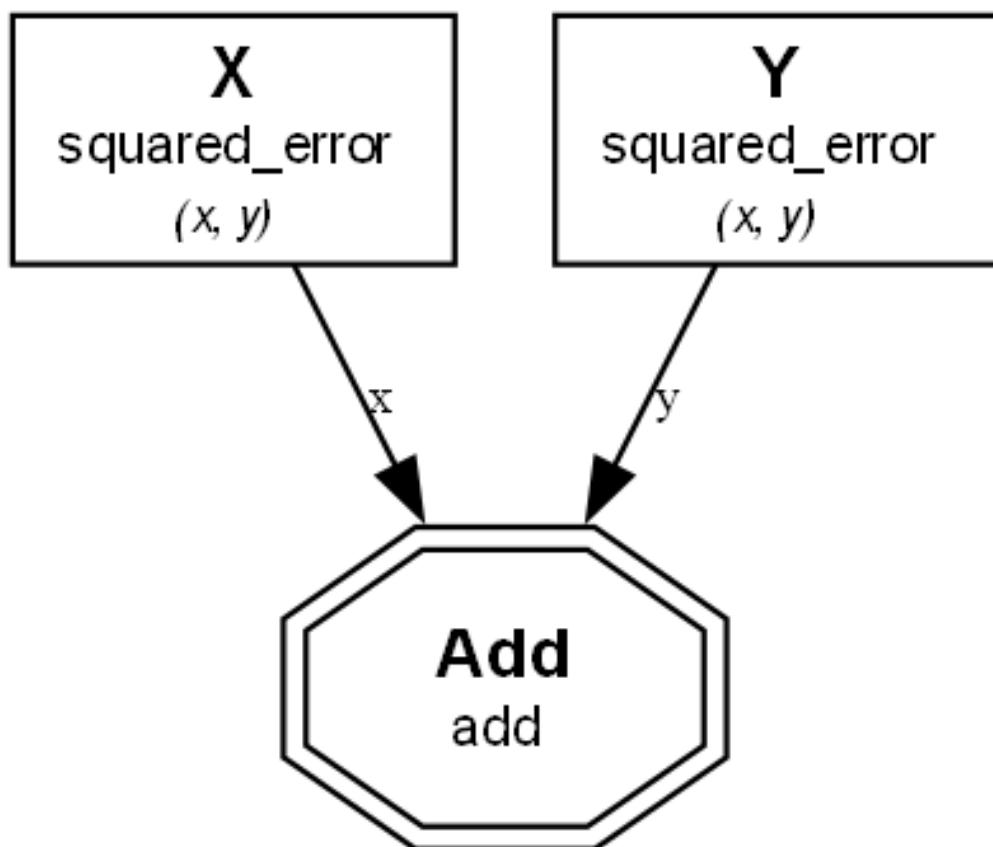


Figure 1: Visualization of the example pipeline.

¹⁵⁴ Citations

¹⁵⁵ PipeOptz complements lightweight helpers designed for algorithm evaluation (Küderle et al.,
¹⁵⁶ 2023), and relies on NumPy (Harris et al., 2020), SciPy (Virtanen et al., 2020), and scikit-
¹⁵⁷ learn (Pedregosa et al., 2011) for numerical computing and Gaussian-process-based Bayesian
¹⁵⁸ optimization (Rasmussen & Williams, 2006). Its optimization engine builds on Bayesian
¹⁵⁹ Optimization techniques (Shahriari et al., 2016; Snoek et al., 2012) alongside the broader
¹⁶⁰ family of metaheuristics surveyed in (Tomar et al., 2023).

¹⁶¹ Acknowledgements

¹⁶² This research is partially funded by ANR, the French National Research Agency with the
¹⁶³ GLACIS project (grant ANR-21-CE33-0002).

References

- 164 Apache airflow. (n.d.). Retrieved December 21, 2025, from <https://airflow.apache.org/>
- 165 Dorigo, M., & Gambardella, L. M. (1997). Ant colony system: A cooperative learning approach
166 to the traveling salesman problem. *IEEE Transactions on Evolutionary Computation*, 1(1),
167 53–66. <https://doi.org/10.1109/4235.585892>
- 168 Gansner, E. R., & North, S. C. (2000). An open graph visualization system and its applications
169 to software engineering. *Software: Practice and Experience*, 30(11), 1203–1233. [https://doi.org/10.1002/1097-024X\(200009\)30:11%3C1203::AID-SPE338%3E3.0.CO;2-N](https://doi.org/10.1002/1097-024X(200009)30:11%3C1203::AID-SPE338%3E3.0.CO;2-N)
- 170 Harris, C. R., Millman, K. J., Walt, S. J. van der, Gommers, R., Virtanen, P., Cournapeau, D.,
171 Wieser, E., Taylor, J., Berg, S., Smith, N. J., Kern, R., Picus, M., Hoyer, S., Kerkwijk,
172 M. H. van, Brett, M., Haldane, A., Río, J. F. del, Wiebe, M., Peterson, P., ... Oliphant,
173 T. E. (2020). Array programming with NumPy. *Nature*, 585(7825), 357–362. <https://doi.org/10.1038/s41586-020-2649-2>
- 174 Hofmann, F. (2023). Linopy: Linear optimization with n-dimensional labeled variables. *Journal
175 of Open Source Software*, 8(84), 4823. <https://doi.org/10.21105/joss.04823>
- 176 Holland, J. H. (1975). *Adaptation in natural and artificial systems*. University of Michigan
177 Press.
- 178 Jiménez, J., & Ginebra, J. (2017). pyGPGO: Bayesian optimization for python. *Journal of
179 Open Source Software*, 2(19), 431. <https://doi.org/10.21105/joss.00431>
- 180 Kahn, A. B. (1962). Topological sorting of large networks. *Communications of the ACM*,
181 5(11), 558–562. <https://doi.org/10.1145/368996.369025>
- 182 Kennedy, J., & Eberhart, R. (1995). Particle swarm optimization. *Proceedings of the IEEE
183 International Conference on Neural Networks*, 1942–1948. <https://doi.org/10.1109/ICNN.1995.488968>
- 184 Kim, J., & Choi, S. (2023). BayesO: A bayesian optimization framework in python. *Journal of
185 Open Source Software*, 8(90), 5320. <https://doi.org/10.21105/joss.05320>
- 186 Kirkpatrick, S., Gelatt, C. D., & Vecchi, M. P. (1983). Optimization by simulated annealing.
187 *Science*, 220(4598), 671–680. <https://doi.org/10.1126/science.220.4598.671>
- 188 Küderle, A., Richer, R., Sîmpetre, R. C., & Eskofier, B. M. (2023). Tpcp: Tiny pipelines for
189 complex problems - a set of framework independent helpers for algorithms development
190 and evaluation. *Journal of Open Source Software*, 8(82), 4953. <https://doi.org/10.21105/joss.04953>
- 191 Pečnik, L., & Fister, I. (2021). NiaAML: AutoML framework based on stochastic population-
192 based nature-inspired algorithms. *Journal of Open Source Software*, 6(61), 2949. <https://doi.org/10.21105/joss.02949>
- 193 Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel,
194 M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau,
195 D., Brucher, M., Perrot, M., & Duchesnay, É. (2011). Scikit-learn: Machine learning in
196 Python. *Journal of Machine Learning Research*, 12, 2825–2830. <http://jmlr.org/papers/v12/pedregosa11a.html>
- 197 Prefect. (n.d.). Retrieved December 21, 2025, from <https://www.prefect.io>
- 198 Rasmussen, C. E., & Williams, C. K. I. (2006). *Gaussian processes for machine learning*. MIT
199 Press. <http://www.gaussianprocess.org/gpml/>
- 200 Shahriari, B., Swersky, K., Wang, Z., Adams, R. P., & Freitas, N. de. (2016). Taking the
201 human out of the loop: A review of Bayesian optimization. *Proceedings of the IEEE*,
- 202 9871. 6
- 203 <https://doi.org/10.xxxxxx/draft>

- 209 104(1), 148–175. <https://doi.org/10.1109/JPROC.2015.2494218>
- 210 Snoek, J., Larochelle, H., & Adams, R. P. (2012). *Practical bayesian optimization of machine*
211 *learning algorithms*. <https://arxiv.org/abs/1206.2944>
- 212 Tomar, V., Bansal, M., & Singh, P. (2023). Metaheuristic algorithms for optimization: A brief
213 review. *Engineering Proceedings*, 59(1). <https://doi.org/10.3390/engproc2023059238>
- 214 Virtanen, P., Gommers, R., Oliphant, T. E., Haberland, M., Reddy, T., Cournapeau, D.,
215 Burovski, E., Peterson, P., Weckesser, W., Bright, J., Walt, S. J. van der, Brett, M.,
216 Wilson, J., Millman, K. J., Mayorov, N., Nelson, A. R. J., Jones, E., Kern, R., Larson, E.,
217 ... Mulbregt, P. van. (2020). SciPy 1.0: Fundamental algorithms for scientific computing
218 in Python. *Nature Methods*, 17(3), 261–272. <https://doi.org/10.1038/s41592-019-0686-2>
- 219 Walt, S. van der, Schönberger, J. L., Nunez-Iglesias, J., Boulogne, F., Warner, J. D., Yager,
220 N., Gouillart, E., Yu, T., & contributors, the scikit-image. (2014). Scikit-image: Image
221 processing in Python. *PeerJ*, 2, e453. <https://doi.org/10.7717/peerj.453>

DRAFT