

¹ TulipaEnergyModel.jl: A Modelling Framework ² Breaking the Tradeoff Between Fidelity and ³ Computational Load

⁴ Abel Soares Siqueira  ¹, Diego A. Tejada-Arango  ^{2,6}, Grigory
⁵ Neustroev  ³, Juha Kiviluoma  ⁴, Lauren Clisby  ², Maaike Elgersma³, Ni
⁶ Wang  ², Suvayu Ali¹, Zhi Gao  ⁵, Mathijs M. de Weerdt  ³, and Germán
⁷ Morales-España 

⁸ 1 Netherlands eScience Center, Netherlands  ² TNO - Netherlands Organisation for Applied
⁹ Scientific Research, Netherlands  ³ Delft University of Technology, Netherlands  ⁴ Nodal Tools,
¹⁰ Finland ⁵ Utrecht University, Netherlands  ⁶ Universidad Pontificia Comillas 

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: Marlin Arnz  

Reviewers:

- [@DavidPozo](#)
- [@trulsf](#)

Submitted: 20 October 2025

Published: unpublished

License

Authors of papers retain copyright²³ and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).²⁴²⁵²⁶

¹¹ Summary

¹² TulipaEnergyModel.jl is a modelling framework for analysing investment and operational decisions of future energy systems through capacity expansion and dispatch optimisation.
¹³ TulipaEnergyModel.jl is the main package of the Tulipa Energy ecosystem, and it is developed in [Julia](#) ([Bezanson et al., 2017](#)) using [JuMP.jl](#) ([Lubin et al., 2023](#)). As a framework, Tulipa formulates models entirely based on input data. This allows users to analyse virtually any system using the generalised building blocks – production, consumption, conversion, storage, and transport. TulipaEnergyModel.jl focuses on model quality and efficient implementation, allowing it to break the tradeoff between model fidelity and computational load through: tighter MIP formulations; exact LP reformulations with fewer constraints and variables; more accurate LP approximations; and flexible model fidelity across temporal, technological, and spatial dimensions.²¹²²

²⁷ Statement of Need

²⁸ Existing models and frameworks in Energy System Optimisation include [EnergyModelsX](#) ([Hellemo et al., 2024](#)), [PowerModels](#) ([Coffrin et al., 2018](#)), [SpineOpt](#) ([Ihleman et al., 2022](#)), [Sienna](#) ([Lara et al., 2021](#)), [GenX](#) ([Bonaldo et al., 2025](#)), [PyPSA](#) ([Brown et al., n.d.](#)), and [Calliope](#) ([Pfenninger & Pickering, 2018](#)). However, they run into computational limits when solving large-scale problems and must resort to (over)simplifying the model to reduce computational burden. The common misconception is that the only strategy to speed up solving times without sacrificing model fidelity is through faster solvers or computers. However, the strategy that is widely overlooked is improving the quality of the mathematical formulations, which increases model fidelity while simultaneously solving faster than standard formulations. This insight inspired the development of TulipaEnergyModel.jl, with the core philosophy of advancing the state-of-the-art in formulation quality by: 1) lowering computational cost while maintaining model fidelity, by reducing the problem size ([Tejada-Arango et al., 2025](#)), and by creating tighter mixed-integer programs (MIP) ([Morales-España et al., 2013](#)). 2) increasing model fidelity without extra computational cost, e.g., by developing more accurate linear programming (LP) approximations ([Elgersma et al., 2025](#); [Gentile et al., 2016](#); [Morales-España et al., 2022](#)). Finally, 3) balancing computational burden with adaptive/flexible model fidelity, i.e., having different levels of detail in various parts of the model, in the temporal ([Gao et al., 2025](#)), technological ([Morales-España et al., 2022](#)) and spatial dimensions. These modelling strategies offer significant computational benefits, especially when handling large-scale problems:³³³⁴³⁵³⁶³⁷³⁸³⁹⁴⁰⁴¹⁴²


43 covering a continent with multiple energy carriers, and optimising over decades while maintaining
 44 hourly resolution for key aspects (e.g., renewable generation). TulipaEnergyModel.jl had to be
 45 developed from scratch to be able to include all of these modelling breakthroughs, since they
 46 alter the foundation and structure of the model. Below, we present some of core modelling
 47 and software design innovations.

48 Modelling Innovations

49 Two of the main innovations include:

- 50 1. Fully flexible temporal resolution (Gao et al., 2025) for the assets/ flows,
- 51 2. Direct connections between assets (Tejada-Arango et al., 2025).

52 The following example illustrates these concepts:



53 **Figure 1:** Example of network with flexible resolution of assets and flows

54 For the fully flexible temporal resolution, consider the 6-hour duration of this system. The
 55 flow between “H2” and “ccgt” has a resolution of 6 hours, while the flow between “ccgt”
 56 and “balance” is 1 hour. The resolution from “wind” to “phs” is 3 hours, and from “phs”
 57 to “balance” is irregular, a 4-hour block followed by a 2-hour block. Tulipa allows different
 58 temporal resolutions throughout the model, thereby reducing the number of variables and
 59 constraints. This feature can drastically speed up solving with little loss in accuracy (Gao et
 al., 2025).

60 For the direct connection between assets, the storage “phs” is directly connected to the “wind”
 61 to charge, and to “balance” to discharge. This for direct connection between assets completely
 62 avoids intermediate elements (connections/nodes), thereby eliminating unnecessary variables
 63 and constraints. Thus, accelerating solving times without any loss of accuracy (Tejada-Arango
 64 et al., 2025).

65 TulipaEnergyModel.jl is fundamentally focused on high-quality mathematical formulations. The
 66 model also includes other key features such as seasonal storage modelling using representative
 67 periods (Neustroev et al., 2025; Tejada-Arango et al., 2018), tight and compact MIP
 68 formulations for storage (Elgersma et al., 2025), unit commitment (Morales-España et al.,

69 2013), and compact formulations for multi-year investment ([Wang & Morales-España, 2025](#);
 70 [Wang & Tejada-Arango, 2025](#)).

71 Software Design Innovations

72 To accommodate flexible temporal resolution for assets and flows, many variables have a “time
 73 block” component instead of a “time step”, and many variables and constraints have sparse
 74 indices. To explain how this is implemented, revisit the system from Figure 1. In an hourly
 75 implementation, there would be variables such as $f_{(H2,ccgt),t}$, $f_{(phs,balance),t}$ for $t = 1, \dots, 6$.
 76 Instead, looking at the time blocks of each flow, there are $f_{(H2,ccgt),1:6}$, $f_{(phs,balance),1:4}$, and
 77 $f_{(phs,balance),5:6}$. In other words, this flow variable could be defined as

$$f_{e,b} \quad \forall e \in E, b \in B(e),$$

78 where E is the set of edges of the graph, and $B(e)$ is the set of time blocks for each edge.

79 Similarly, some constraints and expressions have their own time blocks (e.g., the balance
 80 constraints), and obtaining the coefficients of the flows involves matching indices and computing
 81 the intersections between time blocks. Defining variables like this in JuMP leads to sparse
 82 indices, which are cumbersome and potentially slow. It can also create performance pitfalls,
 83 such as the “[sum-if problem](#)”, when trying to determine the coefficients of each flow in each
 84 balance constraint.

85 To improve the storage of these objects, Tulipa uses a tabular format to collect the indices of
 86 each variable, constraint, and expression. Each row of the table stores the corresponding set
 87 of indices for each JuMP object, and the JuMP objects themselves are stored in a linearised
 88 way in an array. Returning to the example system, the first three rows of the flow variable
 89 defined above are shown in Table 1.

Table 1: “Example of linearised tabular indices of the var_flow table”

id	from asset	to asset	time block start	time block end
1	H2	ccgt	1	6
2	phs	balance	1	4
3	phs	balance	5	6

90 Storing the indices in tables avoids the sparse storage of JuMP objects, which can instead be
 91 stored in an array with their positions matching the IDs in the index tables. Figure 2 illustrates
 92 this method of storage.




Figure 2: Representation of the storage of indices and JuMP objects

93 Another key design decision in Tulipa is handling data using a [DuckDB](#) ([Raasveldt & Muehleisen, n.d.](#)) connection - from processing input data, to generating internal tables for model creation, to storing outputs. This enables manipulating different data formats using DuckDB’s capabilities (instead of Julia’s), as well as decreasing data movement and duplication.

97 When creating variables and constraints, Tulipa can read the DuckDB tables row by row,
 98 reducing memory use compared to using an explicit Julia table, such as DataFrames.

99 Furthermore, using DuckDB for everything allows cleaner separation of the pipeline into (i)
 100 data ingestion and manipulation, (ii) indices creation, and (iii) model creation. Figure 3
 101 summarises how Tulipa interacts with the DuckDB connection.




Figure 3: Overview of Tulipa's integration with DuckDB

102 First, users can prepare their data using any tools they prefer before loading it into the DuckDB
 103 connection. Then, the indices are created from the data in the DuckDB connection and saved
 104 back into it. While this step currently occurs in Julia, many of the operations use SQL and
 105 could potentially be moved outside Julia. Finally, the JuMP pipeline begins by reading the
 106 indices table and complementing the data as necessary, then creates the JuMP objects and
 107 generates the complete model, which is sent to the solver. Tulipa is tested and benchmarked
 108 using the HiGHS (Huangfu & Hall, 2017) solver, but other MIP solvers accepted by JuMP can
 109 be used as well.

110 Acknowledgements

111 We thank Oscar Dowson for the initial suggestion to use DataFrames, which motivated us to
 112 use tabular storage for all indices.

113 This publication is part of the project NextGenOpt ESI.2019.008, which is financed by the
 114 Dutch Research Council (NWO) and supported by the Netherlands eScience Center NLeSC
 115 C 21.0226. This project was also supported by TNO's internal R&D. This research is also
 116 (partially) funded by the European Climate, Infrastructure and Environment Executive Agency
 117 under the European Union's HORIZON Research and Innovation Actions under grant agreement
 118 no. 101095998.

119 References

- 120 Bezanson, J., Edelman, A., Karpinski, S., & Shah, V. B. (2017). Julia: A fresh approach to
 121 numerical computing. *SIAM Review*, 59(1), 65–98. <https://doi.org/10.1137/141000671>
- 122 Bonaldo, L., Chakrabarti, S., Cheng, F., Ding, Y., Jenkins, J., Luo, Q., Macdonald, R.,
 123 Mallapragada, D., Manocha, A., Mantegna, G., Morris, J., Patankar, N., Pecci, F.,
 124 Schwartz, A., Schwartz, J., Schivley, G., Sepulveda, N., & Xu, Q. (2025). GenX (Version
 125 0.4.4). <https://doi.org/10.5281/zenodo.10846070>
- 126 Brown, T., Hörsch, J., Hofmann, F., Neumann, F., Zeyen, L., Syranidis, C., Frysztacki, M.,
 127 Schlachtberger, D., Glaum, P., & Parzen, M. (n.d.). PyPSA: Python for Power System
 128 Analysis (Version 0.35.1). <https://doi.org/10.5334/jors.188>

- 129 Coffrin, C., Bent, R., Sundar, K., Ng, Y., & Lubin, M. (2018). PowerModels.jl: An open-
 130 source framework for exploring power flow formulations. *2018 Power Systems Computation
 131 Conference (PSCC)*, 1–8. <https://doi.org/10.23919/PSCC.2018.8442948>
- 132 Elgersma, M. B., Morales-España, G., Aardal, K. I., Helistö, N., Kiviluoma, J., & Weerdt, M.
 133 M. de. (2025). *Tight MIP formulations for optimal operation and investment of storage
 134 including reserves*. <https://arxiv.org/abs/2411.17484>
- 135 Gao, Z., Gazzani, M., Tejada-Arango, D. A., Soares Siqueira, A., Wang, N., Gibescu, M., &
 136 Morales-España, G. (2025). Fully flexible temporal resolution for energy system optimization.
 137 *Applied Energy*, 396, 126267. <https://doi.org/10.1016/j.apenergy.2025.126267>
- 138 Gentile, C., Morales-España, G., & Ramos, A. (2016). A tight MIP formulation of the
 139 unit commitment problem with start-up and shut-down constraints. *EURO Journal on
 140 Computational Optimization*, 5(1), 177–201. <https://doi.org/10.1007/s13675-016-0066-y>
- 141 Hellemo, L., Bødal, E. F., Holm, S. E., Pinel, D., & Straus, J. (2024). EnergyModelsX: Flexible
 142 energy systems modelling with multiple dispatch. *Journal of Open Source Software*, 9(97),
 143 6619. <https://doi.org/10.21105/joss.06619>
- 144 Huangfu, Q., & Hall, J. A. J. (2017). Parallelizing the dual revised simplex method.
 145 *Mathematical Programming Computation*, 10(1), 119–142. <https://doi.org/10.1007/s12532-017-0130-5>
- 147 Ihlemann, M., Kouveliotis-Lysikatos, I., Huang, J., Dillon, J., O'Dwyer, C., Rasku, T., Marin,
 148 M., Poncelet, K., & Kiviluoma, J. (2022). SpineOpt: A flexible open-source energy system
 149 modelling framework. *Energy Strategy Reviews*, 43, 100902. <https://doi.org/10.1016/j.esr.2022.100902>
- 151 Lara, J. D., Barrows, C., Thom, D., Krishnamurthy, D., & Callaway, D. (2021). PowerSystems.jl
 152 — A power system data management package for large scale modeling. *SoftwareX*, 15.
 153 <https://doi.org/10.1016/j.softx.2021.100747>
- 154 Lubin, M., Dowson, O., Dias Garcia, J., Huchette, J., Legat, B., & Vielma, J. P. (2023).
 155 JuMP 1.0: Recent improvements to a modeling language for mathematical optimization.
 156 *Mathematical Programming Computation*. <https://doi.org/10.1007/s12532-023-00239-3>
- 157 Morales-España, G., Latorre, J. M., & Ramos, A. (2013). Tight and compact MILP formulation
 158 for the thermal unit commitment problem. *IEEE Transactions on Power Systems*, 28(4),
 159 4897–4908. <https://doi.org/10.1109/TPWRS.2013.2251373>
- 160 Morales-España, G., Martínez-Gordón, R., & Sijm, J. (2022). Classifying and modelling demand
 161 response in power systems. *Energy*, 242, 122544. <https://doi.org/10.1016/j.energy.2021.122544>
- 163 Neustroev, G., Tejada-Arango, D. A., Morales-Espana, G., & Weerdt, M. M. de. (2025). *Hull
 164 Clustering with Blended Representative Periods for Energy System Optimization Models*.
 165 arXiv. <https://doi.org/10.48550/arXiv.2508.21641>
- 166 Pfenninger, S., & Pickering, B. (2018). Calliope: A multi-scale energy systems modelling
 167 framework. *Journal of Open Source Software*, 3(29), 825. <https://doi.org/10.21105/joss.00825>
- 169 Raasveldt, M., & Muehleisen, H. (n.d.). DuckDB. <https://github.com/duckdb/duckdb>
- 170 Tejada-Arango, D. A., Domeshek, M., Wogrin, S., & Centeno, E. (2018). Enhanced
 171 representative days and system states modeling for energy storage investment analysis. *IEEE
 172 Transactions on Power Systems*, 33(6), 6534–6544. <https://doi.org/10.1109/TPWRS.2018.2819578>
- 174 Tejada-Arango, D. A., Kiviluoma, J., & Morales-España, G. (2025). Debunking the speed-
 175 fidelity trade-off: Speeding-up large-scale energy models while keeping fidelity. *International*

¹⁷⁶ *Journal of Electrical Power & Energy Systems*, 168, 110674. <https://doi.org/10.1016/j.ijepes.2025.110674>

¹⁷⁷ Wang, N., & Morales-España, G. (2025). *Vintage-based formulations in multi-year investment modelling for energy systems*. <https://arxiv.org/abs/2505.00379>

¹⁷⁸ Wang, N., & Tejada-Arango, D. A. (2025). *Discounting approaches in multi-year investment modelling for energy systems*. <https://arxiv.org/abs/2504.21709>

DRAFT