

¹ **Nexarag: Democratizing Reproducible Knowledge Graph Contexts for LLM Research**

³ **Thomas J. Kerby¹, Benjamin N. Fuller³, and Kevin R. Moon²**

⁴ 1 Brigham Young University, Provo, UT 2 Utah State University, Logan, UT 3 Independent Researcher

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review ↗](#)
- [Repository ↗](#)
- [Archive ↗](#)

Editor: [Wentao Ye](#) ↗ ⓘ

Reviewers:

- [@pebabion](#)
- [@aslan-ng](#)
- [@gsquared94](#)
- [@r0hin](#)

Submitted: 01 December 2025

Published: unpublished

⁵ Summary

⁶ Large language models (LLMs) are widely used in research workflows but struggle with
⁷ hallucinations, short context windows, and weak reproducibility in literature reviews ([Huang et al., 2025; Ji et al., 2023](#)). Nexarag is a modular, open-source platform that lets researchers
⁸ curate, visualize, and share custom knowledge graphs (KGs) from academic sources stored in
⁹ Neo4j ([Neo4j Inc., 2024](#)). Through native support for the Model Context Protocol (MCP),
¹⁰ any MCP-compatible LLM can access these curated KGs for controllable, reproducible context
¹¹ injection ([Anthropic, 2024; Model Context Protocol Contributors, 2024](#))—including fully private,
¹² air-gapped deployments via containers ([Boettiger, 2015](#))—so teams can explore literature more
¹³ deeply and transparently. Nexarag provides interactive graph/semantic visualizations using
¹⁴ Cytoscape.js and D3 ([Bostock et al., 2011; Franz et al., 2023](#)).

¹⁶ Statement of need

License

Authors of papers retain copyright¹⁷ and release the work under a¹⁸ Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))¹⁹. This weakens controllability, auditability, and reproducibility for complex research tasks.

Knowledge graphs provide a structured and interpretable alternative by modeling entities and relations explicitly ([Reinanda et al., 2020](#)). However, existing KG-powered tooling is either proprietary and expensive or technically demanding to deploy and maintain. Nexarag fills this gap with a researcher-friendly platform that automates KG creation from academic inputs, supports semantic exploration, and standardizes LLM access via MCP ([Anthropic, 2024](#)), enabling reproducible literature synthesis across local and cloud settings.

RAG improves access to external knowledge and has become a standard strategy for knowledge-intensive NLP ([Gao et al., 2023; Guu et al., 2020; Lewis et al., 2020](#)), yet it struggles with long contexts and multi-step reasoning when similarity search is the only primitive ([Wang et al., 2024](#)). KGs address this by enabling path-based queries and explicit relation reasoning while preserving transparency and updatability ([Reinanda et al., 2020; Sahlab et al., 2022; Xu et al., 2024](#)). Nexarag's contribution is to operationalize these advantages in a package that researchers can run locally, share with collaborators, and connect to a wide range of LLM hosts through MCP ([Anthropic, 2024](#)).

³⁶ State of the field

³⁷ The open-source ecosystem around retrieval-augmented generation (RAG) has matured quickly,
³⁸ with widely used orchestration frameworks such as LangChain, LlamaIndex, and Haystack
³⁹ providing reusable components for ingesting documents, building indexes, and implementing

40 retrieve-then-generate pipelines ([deepset Haystack, n.d.](#); [LangChain, n.d.-a](#); [LlamalIndex, n.d.-a](#)).
41 These frameworks are well suited for building bespoke RAG applications, but they are primarily
42 developer libraries: users typically assemble pipelines in code, and the resulting “context
43 construction” logic is often tightly coupled to a specific project’s embeddings, chunking choices,
44 runtime configuration, and LLM host.

45 Graph-augmented retrieval approaches have emerged to address limitations of similarity-
46 only retrieval for multi-document reasoning. For example, LlamalIndex provides a
47 KnowledgeGraphIndex that supports automated knowledge graph construction from
48 unstructured text and entity-based querying ([LlamalIndex, n.d.-b](#)). LangChain offers graph
49 question-answering utilities, including GraphCypherQACChain for translating natural-language
50 questions into Cypher queries over Neo4j ([LangChain, n.d.-b](#)). Dedicated GraphRAG toolkits
51 such as Microsoft GraphRAG and the Neo4j GraphRAG package for Python provide pipelines
52 for extracting structured representations from text and using graph-aware retrieval strategies
53 during LLM inference ([Microsoft, n.d.](#); [Neo4j, n.d.](#)). These projects demonstrate that
54 knowledge graphs can improve retrieval structure and interpretability, but they are generally
55 positioned as building blocks for developers rather than as reproducible, end-to-end research
56 applications for literature-centric workflows.

57 In parallel, literature discovery tools such as Connected Papers, ResearchRabbit, and Litmaps
58 offer citation-network visualizations and recommendations to help researchers explore related
59 work and track research areas over time ([Connected Papers, n.d.](#); [Litmaps, n.d.](#); [ResearchRabbit,](#)
60 [n.d.](#)). While valuable for interactive discovery, these tools are not designed to serve as a
61 researcher-owned, queryable knowledge graph substrate that can be versioned, shared, and
62 integrated as a controllable context source for LLM experiments across local and secure
63 environments.

64 Nexarag fills the gap between these two tool families by packaging knowledge-graph-based
65 context construction as a reproducible, shareable research artifact and workflow. Instead of
66 treating graph retrieval as a code-level feature, Nexarag provides a self-hostable platform
67 centered on a persistent Neo4j knowledge graph, interactive graph curation and visualization,
68 and standardized LLM access through the Model Context Protocol (MCP) ([Anthropic,](#)
69 [2024](#); [Model Context Protocol Contributors, 2024](#)). This design enables researchers to
70 hold the underlying graph context and retrieval tools fixed while varying models, prompts, and
71 deployment settings, supporting more transparent and reproducible LLM-assisted literature
72 synthesis.

73 We chose to build Nexarag rather than contribute directly to an existing RAG or GraphRAG
74 library because our core contribution requires application-level infrastructure that is outside the
75 scope of most libraries’ design constraints: a literature-oriented ingestion workflow (BibTeX
76 and paper lists, citation expansion, and persistent graph storage), a collaborative and visual
77 curation interface, containerized offline deployment, and an MCP-first tool surface for model-
78 agnostic experimentation. Nexarag therefore complements, rather than replaces, existing RAG
79 frameworks by providing a researcher-facing system for producing and reusing reproducible
80 knowledge graph contexts in LLM research.

81 Software Design

82 In designing Nexarag we focused on four principles: (1) ease-of-use (in deployment and practical
83 application), (2) flexibility (in model selection and configuration), (3) modularity (to promote
84 scale and independent contribution), and (4) privacy and security. For ease-of-use, we chose
85 popular frontend technologies such as Angular, D3.js, and Cytoscape.js and leveraged existing
86 component libraries to provide a simple and familiar frontend user interface with intuitive
87 tools for building knowledge graphs, searching for relevant papers, conversing with LLMs,
88 and visualizing data. Nexarag is also fully containerized, with release artifacts produced by
89 automated build pipelines readily available for local deployment through Docker. For flexibility,

90 we integrate with Ollama and support any embedding model and LLM that is also supported
 91 by the user's hardware, making it simple to switch between models for different tasks and
 92 as new models are released. Users can also plug in their preferred LLM or coding agent of
 93 choice using the built-in MCP server. Additionally, Nexarag features a highly modular design,
 94 with distinct services for the REST API, the neo4j knowledge graph, the MCP server, and the
 95 frontend application, all bound together with a RabbitMQ messaging backbone. This allows
 96 components to scale horizontally, and to minimize the blast radius of contributions in a single
 97 service. Finally, Nexarag supports on-premises, air-gapped deployments, offering privacy and
 98 security that is not available in cloud-based applications.

99 Software overview

100 **Core capabilities.** Nexarag provides: (i) automated KG construction from BibTeX, paper lists,
 101 search queries, and citation expansion (Semantic Scholar integration) ([Kinney et al., 2023](#);
 102 [Semantic Scholar, 2024](#)); (ii) Neo4j-backed storage and Cypher querying ([Neo4j Inc., 2024](#));
 103 (iii) interactive graph and semantic visualizations (Cytoscape.js and D3.js) ([Bostock et al.,](#)
 104 [2011](#); [Franz et al., 2023](#)); and (iv) an AI "Talk To Your Data" interface that supports both
 105 simple retrieve-and-generate and ReAct-style agentic workflows ([Yao et al., 2022](#)).

106 **Architecture.** The system uses a containerized, microservices design orchestrated with Docker
 107 Compose ([Docker, Inc., 2024](#); [Merkel, 2014](#)). Primary services include: a FastAPI service for
 108 HTTP coordination ([FastAPI Contributors, 2024](#)), a Neo4j database for graph storage ([Neo4j](#)
 109 [Inc., 2024](#)), and a Knowledge Graph service for document processing/embeddings/AI tasks.
 110 Services communicate asynchronously via RabbitMQ, enabling horizontal scaling ([VMware,](#)
 111 [Inc., 2024](#)).

112 **MCP integration.** Nexarag ships an MCP-compatible server that exposes graph querying,
 113 semantic search over embedded content, and external search via Semantic Scholar to any
 114 MCP-enabled LLM (local via Ollama or remote via hosted providers) ([Anthropic, 2024](#); [Model](#)
 115 [Context Protocol Contributors, 2024](#); [Ollama Team, 2024](#); [OpenAI, 2023](#)). This standardizes
 116 context delivery and promotes reproducible prompt-driven research workflows.

117 **Install & minimal run.** (see repository docs for full instructions)

```
# CPU example
docker compose -f docker-compose.cpu.yml up -d
# or on macOS
docker compose -f docker-compose.macos.yml up -d
```

118 Optionally pull local models for embedding/LLM integration with Ollama ([Ollama Team,](#)
 119 [2024](#)); for example, a long-context text embedder like Nomic Embed ([Nussbaum et al., 2025](#)):

```
# inside the Ollama container or on macOS host
ollama pull nomic-embed-text:v1.5
ollama pull gemma3:1b
```

120 **Repository:** <https://github.com/REPLACE-WITH-YOUR-ORG/nexarag>

121 **License:** GNU General Public License v3.0.

122 Use cases

- 123 ■ **Reproducible literature reviews.** Build a KG from a seed set (e.g., via BibTeX), expand
 124 by citations, and generate a structured review through the MCP interface ([Sahlab et al.,](#)
 125 [2022](#)).
- 126 ■ **Private research contexts.** Run entirely offline (air-gapped) with local LLMs for sensitive
 127 domains (e.g., healthcare, legal, proprietary research) ([Boettiger, 2015](#)).

- 128 ▪ **Collaborative curation.** Share/export/import graphs across teams to support longitudinal
129 projects.

130 **Quality control**

131 Nexarag emphasizes verifiable operation through containerized deployment and a guided quick
132 start (Boettiger, 2015). Reviewers can launch the full stack with Docker Compose, query/persist
133 KGs in Neo4j, and exercise end-to-end flows (semantic search, citation expansion, MCP tools).
134 A worked MCP chat transcript and an automatically generated literature review illustrate
135 that the system's graph building, retrieval, and reporting features execute as described. The
136 repository includes example datasets/notebooks and scripts for running tests where applicable,
137 supporting broader reproducibility goals in research practice (Rothacher et al., 2023).

138 **Research impact statement**

139 Nexarag addresses a growing need in LLM-assisted research for transparent, reproducible, and
140 inspectable context construction beyond embedding-only retrieval. While many RAG systems
141 remain opaque and difficult to reproduce, Nexarag operationalizes knowledge-graph-based
142 context building in a form that researchers can deploy locally, inspect visually, and share across
143 projects. By combining Neo4j-backed knowledge graphs with standardized access through
144 the Model Context Protocol (MCP), the software provides a reproducible bridge between
145 structured scholarly knowledge and LLM-driven analysis.

146 Although Nexarag is a relatively new project and has not yet accumulated extensive downstream
147 citations, it demonstrates credible near-term research impact through its design, documentation,
148 and reproducible reference materials. The repository includes end-to-end examples that
149 reproduce literature expansion, graph construction, semantic querying, and LLM-mediated
150 synthesis from fixed inputs, allowing independent researchers to verify behavior and compare
151 results across models and deployment environments. Containerized deployment and air-gapped
152 operation further support use in domains where reproducibility, auditability, or data sensitivity
153 are critical.

154 Nexarag is positioned to serve as shared research infrastructure for studies on retrieval-
155 augmented generation, knowledge-graph-augmented reasoning, and AI-assisted literature
156 review workflows. Its model-agnostic design, enabled by MCP, allows researchers to interchange
157 local or API-hosted LLMs while holding the underlying knowledge graph and retrieval logic fixed.
158 This supports a direct comparison of LLM behavior under identical, graph-derived contexts,
159 facilitating methodological research on controllability, hallucination reduction, and long-context
160 reasoning. By lowering the technical barrier to building, inspecting, and sharing reproducible
161 knowledge graph contexts, Nexarag enables researchers to move beyond ad hoc, model-coupled
162 RAG pipelines toward more transparent and portable AI-assisted research practices.

163 **AI usage disclosure**

164 Generative AI tools were used in the development of the software, supporting code reviews,
165 providing minor features in the frontend, and identifying and fixing bugs. Generative AI tools
166 were also used to generate some of the documentation, and assisted with paper authoring. We
167 primarily used:

- 168 ▪ ChatGPT with the GPT-4o model for writing tasks
169 ▪ Claude Code with the Sonnet 4 model for coding tasks

170 All AI-generated material was explicitly reviewed by at least one author, and all major design
171 decisions were formalized by multiple authors.

172 Acknowledgements

173 We acknowledge the open-source ecosystems behind Neo4j, Cytoscape.js, D3.js, RabbitMQ,
174 FastAPI, Ollama, and the Model Context Protocol, as well as contributors and users who
175 provided feedback during development. This research was supported in part by the NSF under
176 Grant 221235.

177 References

- 178 Anthropic. (2024). *Introducing the model context protocol*. <https://www.anthropic.com/news/model-context-protocol> <https://www.anthropic.com/news/model-context-protocol>
- 180 Boettiger, C. (2015). An introduction to docker for reproducible research. *ACM SIGOPS Operating Systems Review*, 49(1), 71–79. <https://doi.org/10.1145/2723872.2723882>
- 182 Bostock, M., Ogievetsky, V., & Heer, J. (2011). D3: Data-driven documents. *IEEE Transactions on Visualization and Computer Graphics*, 17(12), 2301–2309. <https://doi.org/10.1109/TVCG.2011.185>
- 185 Connected Papers. (n.d.). *Connected papers*. <https://www.connectedpapers.com/>.
- 186 deepset Haystack. (n.d.). *Get started (haystack documentation)*. <https://docs.haystack.deepset.ai/docs/get-started>.
- 188 Docker, Inc. (2024). *Docker: Accelerated container application development*. <https://www.docker.com/>. <https://www.docker.com/>
- 190 FastAPI Contributors. (2024). *FastAPI: Modern, fast (high-performance), web framework for building APIs with python 3.6+ based on standard python type hints*. <https://fastapi.tiangolo.com/>. <https://fastapi.tiangolo.com/>
- 193 Franz, M., Lopes, C. T., Fong, D., Kucera, M., Cheung, M., Siper, M. C., Huck, G.,
194 Dong, Y., Sumer, O., & Bader, G. D. (2023). Cytoscape.js 2023 update: A graph
195 theory library for visualization and analysis. *Bioinformatics*, 39(1), btad031. <https://doi.org/10.1093/bioinformatics/btad031>
- 197 Gao, Y., Xiong, Y., Gao, X., Jia, K., Pan, J., Bi, Y., Dai, Y., Sun, J., Wang, H., & Wang,
198 H. (2023). Retrieval-augmented generation for large language models: A survey. *arXiv Preprint arXiv:2312.10997*, 2.
- 200 Guu, K., Lee, K., Tung, Z., Pasupat, P., & Chang, M.-W. (2020). REALM: Retrieval-
201 augmented language model pre-training. *Proceedings of the 37th International Conference
202 on Machine Learning*.
- 203 Huang, L., Yu, W., Ma, W., Zhong, W., Feng, Z., Wang, H., Chen, Q., Peng, W., Feng,
204 X., Qin, B., & Liu, T. (2025). A survey on hallucination in large language models:
205 Principles, taxonomy, challenges, and open questions. *ACM Trans. Inf. Syst.*, 43(2).
206 <https://doi.org/10.1145/3703155>
- 207 Ji, Z., Lee, N., Frieske, R., Yu, T., Su, D., Xu, Y., Ishii, E., Bang, Y. J., Madotto, A., & Fung,
208 P. (2023). Survey of hallucination in natural language generation. *ACM Comput. Surv.*,
209 55(12). <https://doi.org/10.1145/3571730>
- 210 Kinney, R. M., Anastasiades, C., Authur, R., Beltagy, I., Bragg, J., Buraczynski, A., Cachola,
211 I., Candra, S., Chandrasekhar, Y., Cohan, A., Crawford, M., Downey, D., Dunkelberger,
212 J., Etzioni, O., Evans, R., Feldman, S., Gorney, J., Graham, D. W., Hu, F. Q., ...
213 Weld, D. S. (2023). The semantic scholar open data platform. *ArXiv*, *abs/2301.10140*.
214 <https://api.semanticscholar.org/CorpusID:256194545>
- 215 LangChain. (n.d.-a). *Build a RAG agent with LangChain*. <https://docs.langchain.com/oss/>

- 216 [python/langchain/rag](https://github.com/llamaindex/llamaindex/blob/main/python/langchain/rag.py).
- 217 LangChain. (n.d.-b). *Neo4j (LangChain integration documentation)*. <https://docs.langchain.com/oss/python/integrations/providers/neo4j>.
- 218
- 219 Lewis, P., Perez, E., Piktus, A., Petroni, F., Karpukhin, V., Goyal, N., Küttler, H., Lewis, M.,
220 Yih, W., Rocktäschel, T., Riedel, S., & Kiela, D. (2020). Retrieval-augmented generation
221 for knowledge-intensive NLP tasks. *Proceedings of the 34th International Conference on
222 Neural Information Processing Systems*.
- 223 Litmaps. (n.d.). *Litmaps*. <https://www.litmaps.com/>.
- 224 Llamaindex. (n.d.-a). *Introduction to RAG (retrieval-augmented generation)*. <https://developers.llamaindex.ai/python/framework/understanding/rag/>.
- 225
- 226 Llamaindex. (n.d.-b). *Knowledge graph index (Llamaindex python documentation)*.
227 https://developers.llamaindex.ai/python/examples/index_structs/knowledge_graph/
228 knowledgegraphdemo/.
- 229 Merkel, D. (2014). Docker: Lightweight linux containers for consistent development and
230 deployment. *Linux Journal*, 2014(239), 2.
- 231 Microsoft. (n.d.). *Microsoft/graphrag: A modular graph-based retrieval-augmented generation
232 (RAG) system*. <https://github.com/microsoft/graphrag>.
- 233 Model Context Protocol Contributors. (2024). *Model context protocol: A protocol for seamless
234 integration between LLM applications and external data sources*. [https://github.com/modelcontextprotocol](https://github.com/
235 modelcontextprotocol)
- 236 Neo4j. (n.d.). *neo4j/neo4j-graphrag-python: Neo4j GraphRAG package for python*. <https://github.com/neo4j/neo4j-graphrag-python>.
- 237
- 238 Neo4j Inc. (2024). *Neo4j graph database platform*. <https://neo4j.com/product/neo4j-graph-database/>. <https://neo4j.com/product/neo4j-graph-database/>
- 239
- 240 Nussbaum, Z., Morris, J. X., Mulyar, A., & Duderstadt, B. (2025). Nomic embed: Training
241 a reproducible long context text embedder. *Transactions on Machine Learning Research*.
242 <https://openreview.net/forum?id=IPmzyQSjQE>
- 243 Ollama Team. (2024). *Ollama: Get up and running with llama 3.2, mistral, gemma 2, and
244 other large language models*. <https://ollama.com/>. <https://ollama.com/>
- 245 OpenAI. (2023). *OpenAI API*.
- 246 Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using siamese
247 BERT-networks. *Proceedings of the 2019 Conference on Empirical Methods in Natural
248 Language Processing*, 3973–3983. <http://arxiv.org/abs/1908.10084>
- 249 Reinanda, R., Meij, E., & Rijke, M. de. (2020). Knowledge graphs: An information retrieval
250 perspective. *Foundations and Trends in Information Retrieval*, 14(2-3), 289–444. <https://doi.org/10.1561/1500000063>
- 251
- 252 ResearchRabbit. (n.d.). *ResearchRabbit*. <https://www.researchrabbit.ai/>.
- 253 Rothacher, L., Engel, C., Garbeva, P., Grüning, B., Goble, C., Gundersen, S., Ioannidis, J. P.,
254 Marwick, B., Parsons, S., Pronk, T. E., & others. (2023). Eleven strategies for making
255 reproducible research and open science training the norm at research institutions. *eLife*, 12,
256 RP89736. <https://doi.org/10.7554/eLife.89736>
- 257 Sahlab, N., Kahoul, H., Jazdi, N., & Weyrich, M. (2022). A knowledge graph-based method
258 for automating systematic literature reviews. *arXiv Preprint arXiv:2208.02334*. <https://doi.org/10.48550/arXiv.2208.02334>
- 259
- 260 Semantic Scholar. (2024). *Semantic scholar academic graph API*. <https://www.semanticscholar.org/api>.

- 261 [semanticscholar.org/product/api](https://www.semanticscholar.org/product/api). <https://www.semanticscholar.org/product/api>
- 262 VMware, Inc. (2024). *RabbitMQ: The most widely deployed open source message broker*.
- 263 <https://www.rabbitmq.com/>. <https://www.rabbitmq.com/>
- 264 Wang, M., Chen, L., Cheng, F., Liao, S., Zhang, X., Wu, B., Yu, H., Xu, N., Zhang, L.,
265 Luo, R., Li, Y., Yang, M., Huang, F., & Li, Y. (2024). Leave no document behind:
266 Benchmarking long-context LLMs with extended multi-doc QA. *Proceedings of the 2024*
267 *Conference on Empirical Methods in Natural Language Processing*, 5627–5646. <https://doi.org/10.18653/v1/2024.emnlp-main.322>
- 268
- 269 Xu, Z., Cruz, M. J., Guevara, M., Wang, T., Deshpande, M., Wang, X., & Li, Z. (2024).
270 Retrieval-augmented generation with knowledge graphs for customer service question
271 answering. *Proceedings of the 47th International ACM SIGIR Conference on Research*
272 *and Development in Information Retrieval*, 2905–2909. <https://doi.org/10.1145/3626772.3661370>
- 273
- 274 Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2022). ReAct:
275 Synergizing reasoning and acting in language models. *arXiv Preprint arXiv:2210.03629*.
276 <https://doi.org/10.48550/arXiv.2210.03629>