

# <sup>1</sup> graph\_framework: A Domain Specific Compiler for Building Physics Applications

<sup>3</sup> M. Cianciosa  <sup>1</sup>, D. Batchelor  <sup>2</sup>, and W. Elwasif  <sup>1</sup>

<sup>4</sup> 1 Oak Ridge National Laboratory 2 Diditco, Oak Ridge TN 37831

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

## Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: 

Submitted: 03 December 2025

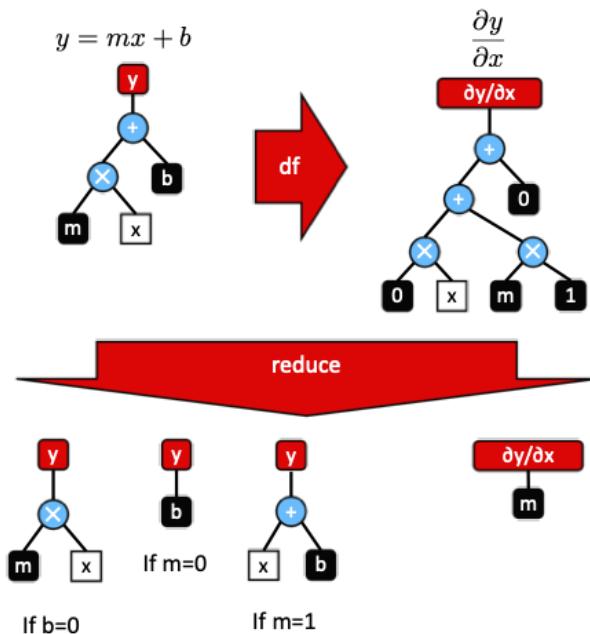
Published: unpublished

## <sup>5</sup> Summary<sup>1</sup>

The graph\_framework is a domain specific compiler which enables domain scientists to create optimized kernels that can operate on Graphics Processing Units (GPUs) or central processing unit (CPUs). This framework works by first building data structures of the operations making up a physics equations. Algebraic simplifications are applied to the graphs to reduce them to simpler forms. Auto differentiation is supported by traversing existing graphs and creating new graphs by applying the chain rule. These graphs can be Just-In-Time (JIT) compiled to central processing unit (CPUs), Apple GPUs, NVidia GPUs, and initial support for AMD GPUs.

## License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).



**Figure 1:** Mathematical operations are defined as a tree of operations. A df method transforms the tree by applying the derivative chain rule to each node. A reduce method applies algebraic rules removing nodes from the graph.

<sup>1</sup>Notice of Copyright This manuscript has been authored by UT-Battelle, LLC under Contract No. DE-AC05-00OR22725 with the U.S. Department of Energy. The United States Government retains and the publisher, by accepting the article for publication, acknowledges that the United States Government retains a non-exclusive, paidup, irrevocable, world-wide license to publish or reproduce the published form of this manuscript, or allow others to do so, for United States Government purposes. The Department of Energy will provide public access to these results of federally sponsored research in accordance with the DOE Public Access Plan (<http://energy.gov/downloads/doe-public-access-plan>).

13 This framework focuses on the domain of physics problems where a the same physics is being  
14 applied to large ensemble of particles or rays. Applications have been developed for tracing large  
15 numbers of Radio Frequency (RF) rays in fusion devices and particle tracing for understanding  
16 how particles distributions are lost or evolve over time. The exploitation of GPU resources  
17 afforded by this framework allows high fidelity simulations at low computational cost.

## 18 Statement of need

19 Modern supercomputers are increasingly relying on Graphic Processing Units (GPUs) and other  
20 accelerators to achieve exa-scale performance at reasonable energy usage. A major challenge of  
21 exploiting these accelerators is the incompatibility between different vendors. A scientific code  
22 written using CUDA will not operate on a AMD gpu. Frameworks that can abstract the physics  
23 from the accelerator kernel code are needed to exploit the current and future hardware. In the  
24 world of machine learning, several auto differentiation frameworks have been developed that  
25 have the promise of abstracting the math from the compute hardware. However in practice,  
26 these framework often lag in supporting non-CUDA platforms. Their reliance on python makes  
27 them challenging to embed within non python based applications.

28 Fusion energy is a grand engineering challenge to make into a viable power source. Beyond the  
29 technical challenges towards making it work in the first place, there is an economic challenge  
30 that it needs to be addressed. For fusion energy to be competitive in the market place.  
31 Addressing the economic challenge is tackled though design optimization. However, a barrier to  
32 optimization is the computational costs associated with exploring the different configurations.

33 Low fidelity models like systems codes(Kovari et al., 2014),(Kovari et al., 2016), can miss  
34 critical physics that enable optimized designs. High fidelity models, are too costly to run for  
35 multiple configurations. GPUs offer tremendous processing power that is largely untapped  
36 in codes developed by domain scientists. Due to the challenges of exploiting GPUs they are  
37 largely relegated to hero class codes which can use a large percentage of Exa-scale machines.

38 However, there is an intermediate scale of problems which individually can operate using  
39 modest computational requirements but become a challenge when generating large ensembles.  
40 These codes are typically CPU only due to the challenges of adopting GPUs. As more super  
41 computers are diminishing CPU capacity in favor of GPU support, we are losing the capacity  
42 computing necessary to explore large ensembles necessary for device optimization.

43 The goal of the graph\_framework is to lower the barrier of entry for adopting GPU code.  
44 While there are many different solutions to the problem of performance portable code, different  
45 solutions have different drawbacks or trade offs. With that in mind the graph\_framework was  
46 developed to address the specific capabilities of:

- 47     ■ Transparently support multiple CPUs and GPUs including Apple GPUs.
- 48     ■ Use an API that is as simple as writing equations.
- 49     ■ Allow easy embedding in legacy code (Doesn't rely on python).
- 50     ■ Enables automatic differentiation.

51 With these design goals in mind this framework is limited to the classes of problems which the  
52 same physics is applied to a large ensemble of particles. This limitation simplifies the complexity  
53 of this framework making future extensibility simpler as a need arises for a new problem domain.  
54 In this paper will describe the frameworks design and capabilities. Demonstrate applications  
55 to problems in radio frequency (RF) heating and particle tracing, and show its performance  
56 scaling.

## 57 Background

Table 1: Overview of GPU capable frameworks.

Framework	Language	Cuda Support	Metal Support	RocM Support	Auto Differentiation
graph_framework	C++, C, Fortran	Official	Official	Preliminary	Yes
Cuda Metal	C Objective C, Swift	Official None	None Official	None Deprecated	No No
Kokkos OpenACC	C++ C, C++, Fortran	Official Official	None None	Official None	No No
OpenMP	C, C++, Fortran	Compiler Dependent	None	Compiler Dependent	No
OpenCL	C	Official	Deprecated	Official	No
Vulcan	C	Official	Unofficial	Official	No
HIP	C	Official	None	Official	No
TensorFlow	Python, C++	Official	Unofficial/Incomplete	Unofficial	Yes
JAX PyTorch	Python Python, C++, Java	Official Official	Unofficial/Incomplete Official	Official Official	Yes Yes
mlx	Python, C++, Swift	Official	Official	Experimental	Yes

58 Standardized programming languages such as Fortran([Backus & Heising, 1964](#)), C([Ritchie, 1993](#)), C++([Stroustrup, 2013](#)), have simplified the development of cross platform programs.

59 60 Scientific codes have relied on the ability to write source code which can operate on multiple  
61 processor architectures and operating systems (OSs) with no or little changes given an  
62 appropriate compiler. However, modern super computers rely on graphical processing units  
63 (GPUs) to achieve exa-scale performance([Hines, 2018](#)),([Yang & Deslippe, 2020](#)),([Schneider, 2022](#)) with reasonable energy usage. Unlike central processing units (CPUs), the instruction  
64 sets of GPUs are proprietary information. Additionally, since accelerators typically are hardware  
65 accessories, an OS requires device drivers which are also proprietary. NVidia GPUs are best  
66 programmed using CUDA([Cuda Documentation, n.d.](#)) while Apple GPUs use Metal([Metal Documentation, n.d.](#)) and AMD GPUs use HIP([Hip Documentation, n.d.](#)).

67 68  
69 There are many potential solutions to cross performance portable support. Low level cross  
70 platform frameworks general purpose GPU (GPGPU) programming frameworks such as  
71 OpenCL([Munshi et al., 2011](#)) and Vulkan([Vulkan Specification, n.d.](#)) requires direct vendor  
72 support. HIP can support NVidia GPUs by abstracting the driver API and rewriting kernel  
73 code. However these frameworks are the lowest level and require GPU programming expertise  
74 to utilize them effectively that a domain scientist may not have. A higher level approach  
75 used in OpenACC([Farber, 2016](#)) and OpenMP([OpenMP Specification, n.d.](#)) use source  
76 code annotation to transform loops and code blocks into GPU kernels. The drawback of  
77 this approach is that source code written for CPUs can result in poor GPU performance.  
78 Kokkos([Edwards et al., 2011](#)) is a collection of performance portable array operations for  
79 building device agnostic applications. However, the framework only support AMD and Nvidia  
80 GPUs and doesn't have out of box support for auto differentiation.

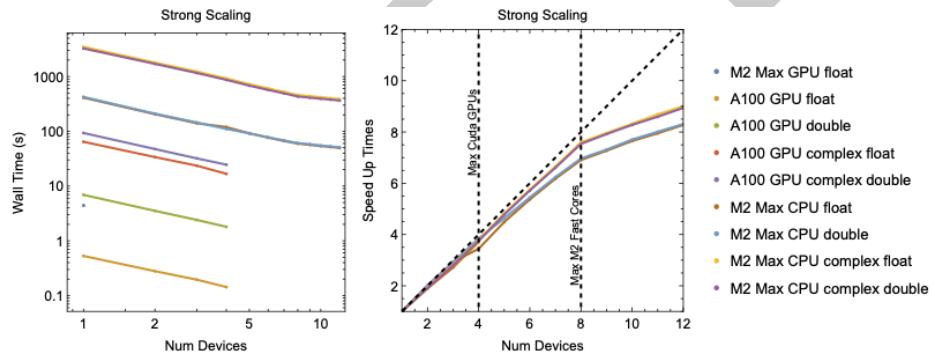
81 82 With the advent of Machine learning, several machine learning frameworks have been created  
83 such as TensorFlow([Abadi et al., 2015](#)), JAX([Bradbury et al., 2018](#)), PyTorch([Paszke et  
84 al., 2017](#)), and MLX([Hannun et al., 2023](#)). These frameworks build a graph representation  
operations that can be auto-differentiated and compiled to GPUs. These frameworks are

intended to be used through a python interface which lowers one barrier to using but also introduces new barriers. For instance, it's not straight forward to embed these frameworks in non-python codes and their non-python API's don't always support all the features or are as well documented as their python API's. Additionally performance is not guaranteed. It is not always straight forward to understand what the framework is doing. Additionally cross platform support is often unofficial and can be incomplete. Table 1 shows an overview of these frameworks.

## Performance

To demonstrate the performance of the optimized kernels created using this framework we measured the strong scaling using the the RF ray tracing problem in a realistic tokamak geometry. To compare against other frameworks we benchmarked the achieved throughput for simulating gyro motion in a uniform magnetic field.

### Strong Scaling



**Figure 2:** Left: Strong scaling wall time for 100000 Rays traced in a realistic tokamak equilibrium. Right: Strong scaling speedup normalized to the wall time for a single device or core. The dashed diagonal line references the best possible scaling. The M2 Max has 8 fast performance cores and 4 slower energy efficiency cores resulting drop off in improvement beyond 8 cores.

To benchmark code performance we traced  $10^6$  rays for  $10^3$  time steps using the cold plasma dispersion relation in a realistic tokamak equilibrium. A benchmarking application is available in the git repository. The figure above shows the strong scaling of wall time as the number of GPU and CPU devices are increased. The figure above shows the strong scaling speed up

$$SpeedUp = \frac{time(1)}{time(n)}$$

Benchmarking was prepared on two different setups. The first set up as a Mac Studio with an Apple M2 Max chip. The M2 chip contains a 12 core CPU where 8 cores are faster performance cores and the remaining 4 are slower efficiency cores. The M2 Max also contains a single 38-core GPU which only support single precision operations. The second setup is a server with 4 Nvidia A100 GPUs. Benchmarking measures the time to trace  $10^6$  rays but does not include the setup and JIT times.

Figure 2 shows the advantage even a single GPU has over CPU execution. In single precision, the M2's GPU is almost  $100\times$  faster than single CPU core while the a single A100 has a nearly  $800\times$  advantage. An interesting thing to note is the M2 Max CPU show no advantage between single and double precision execution.

112 For large problem sizes the framework is expected to show good scaling with number of devices  
 113 as the problems we are applying are embarrassingly parallel in nature. The figure above shows  
 114 the strong scaling speed up with the number of devices. The framework shows good strong  
 115 scaling as the problem is split among more devices. The architecture of the M2 Chip contains  
 116 8 fast performance cores and 4 slower energy efficiency cores. This produces a noticeable knee  
 117 in the scaling after 8 core are used. Overall, the framework demonstrates good scaling across  
 118 CPU and GPU devices.

### 119 Comparison to other frameworks

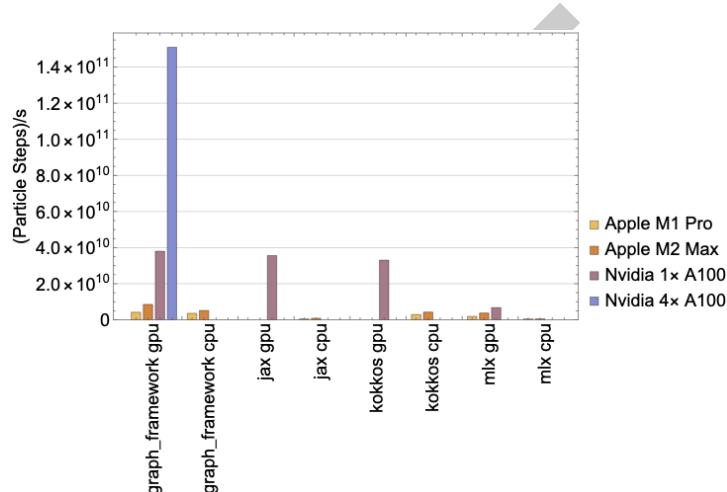


Figure 3: Particle throughput for graph framework compared to MLX and JAX.

120 To benchmark against other frameworks we will look at the simple case of gyro motion in a  
 121 uniform magnetic field  $\vec{B} = B_0\hat{z}$ .

$$\frac{\partial \vec{v}}{\partial t} = dt\vec{v} \times \vec{B}$$

$$\frac{\partial \vec{x}}{\partial t} = dt\vec{v}$$

122  
 123 We compared the graph framework against the MLX framework since it supports Apple GPUs  
 124 and JAX due to its popularity. Source codes for this benchmark case is available in the  
 125 graph\_framework documentation. Figure 3 shows the throughput of pushing  $10^8$  particles for  
 126  $10^3$  time steps. The graph\_framework consistently shows the best throughput on both CPUs  
 127 and GPUs. Note MLX CPU throughput could be improved by splitting the problem to multiple  
 128 threads.

### 129 Acknowledgements

- 130 The authors would like to thank Dr. Yashika Ghai, Dr. Rhea Barnett, and Dr. David Green  
 131 for their valuable insights when setting up test cases for the RF-Ray Tracing.
- 132 Abadi, M., Agarwal, A., Barham, P., Brevdo, E., Chen, Z., Citro, C., Corrado, G. S., Davis,  
 133 A., Dean, J., Devin, M., Ghemawat, S., Goodfellow, I., Harp, A., Irving, G., Isard, M., Jia,  
 134 Y., Jozefowicz, R., Kaiser, L., Kudlur, M., ... Zheng, X. (2015). *TensorFlow: Large-scale  
 135 machine learning on heterogeneous systems*. <https://www.tensorflow.org/>
- 136 Backus, J. W., & Heising, W. P. (1964). Fortran. *IEEE Transactions on Electronic Computers*,  
 137 EC-13(4), 382–385. <https://doi.org/10.1109/PGEC.1964.263818>

- 138 Bradbury, J., Frostig, R., Hawkins, P., Johnson, M. J., Leary, C., Maclaurin, D., Necula, G.,  
139 Paszke, A., VanderPlas, J., Wanderman-Milne, S., & Zhang, Q. (2018). *JAX: Composable*  
140 *transformations of Python+NumPy programs* (Version 0.3.13). <http://github.com/jax-ml/jax>
- 141
- 142 *Cuda documentation*. (n.d.). <https://docs.nvidia.com/cuda/>.
- 143 Edwards, H. C., Sunderland, D., Amsler, C., & Mish, S. (2011). Multicore/GPGPU portable  
144 computational kernels via multidimensional arrays. *2011 IEEE International Conference on*  
145 *Cluster Computing*, 363–370. <https://doi.org/10.1109/CLUSTER.2011.47>
- 146 Farber, R. (2016). *Parallel programming with OpenACC*. Newnes.
- 147 Hannun, A., Digani, J., Katharopoulos, A., & Collobert, R. (2023). *MLX: Efficient and flexible*  
148 *machine learning on apple silicon* (Version 0.0). <https://github.com/ml-explore>
- 149 Hines, J. (2018). Stepping up to summit. *Computing in Science & Engineering*, 20(2), 78–82.  
150 <https://doi.org/10.1109/MCSE.2018.021651341>
- 151 *Hip documentation*. (n.d.). <https://rocm.docs.amd.com/projects/HIP/en/latest/index.html>.
- 152 Kovari, M., Fox, F., Harrington, C., Kembleton, R., Knight, P., Lux, H., & Morris, J. (2016).  
153 “PROCESS”: A systems code for fusion power plants – part 2: engineering. *Fusion*  
154 *Engineering and Design*, 104, 9–20. <https://doi.org/10.1016/j.fusengdes.2016.01.007>
- 155
- 156 Kovari, M., Kemp, R., Lux, H., Knight, P., Morris, J., & Ward, D. J. (2014). “PROCESS”: A  
157 systems code for fusion power plants—part 1: physics. *Fusion Engineering and Design*,  
158 89(12), 3054–3069. <https://doi.org/10.1016/j.fusengdes.2014.09.018>
- 159 *Metal documentation*. (n.d.). <https://developer.apple.com/metal/>.
- 160 Munshi, A., Gaster, B., Mattson, T. G., & Ginsburg, D. (2011). *OpenCL programming guide*.  
161 Pearson Education.
- 162 *OpenMP specification*. (n.d.). <https://www.openmp.org/wp-content/uploads/OpenMP-API-Specification-6-0.pdf>.
- 163
- 164 Paszke, A., Gross, S., Chintala, S., Chanan, G., Yang, E., DeVito, Z., Lin, Z., Desmaison, A.,  
165 Antiga, L., & Lerer, A. (2017). *Automatic differentiation in PyTorch*.
- 166 Ritchie, D. M. (1993). The development of the c language. *ACM Sigplan Notices*, 28(3),  
167 201–208.
- 168 Schneider, D. (2022). The exascale era is upon us: The frontier supercomputer may be the  
169 first to reach 1,000,000,000,000,000,000 operations per second. *IEEE Spectrum*, 59(1),  
170 34–35. <https://doi.org/10.1109/MSPEC.2022.9676353>
- 171 Stroustrup, B. (2013). *The c++ programming language*. Pearson Education.
- 172 *Vulkan specification*. (n.d.). <https://docs.vulkan.org/spec/latest/index.html>.
- 173 Yang, C., & Deslippe, J. (2020). Accelerate science on perlmutter with NERSC. *Bulletin of*  
174 *the American Physical Society*, 65.