

¹ PointCloudCrafter: Tools for Handling, Manipulating and Analyzing of Point Clouds

³ **Dominik Kulmer**  ^{1*} and **Maximilian Leitenstern**  ^{1*}

⁴ ¹ Institute of Automotive Technology, Department of Mobility System Engineering, School of
⁵ Engineering and Design, Technical University of Munich, Germany. * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: 

Submitted: 30 January 2026

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#))¹⁵

⁶ Summary

⁷ PointCloudCrafter is a C++ command-line interface (CLI) toolkit for the processing, manipulation, and analysis of three-dimensional point cloud data. The software provides a collection of compiled executables for integration into automated research pipelines, enabling common point cloud operations without requiring custom code.

¹⁰ The toolkit supports the conversion and processing of data from multiple acquisition and storage formats, including ROS2 bag recordings, binary files, plain-text files, and standard OBJ, PLY and PCD files. A central design feature of PointCloudCrafter is its schema-agnostic handling of per-point attributes. Arbitrary scalar fields associated with individual points are preserved throughout conversion, transformation, filtering, and aggregation steps, ensuring that auxiliary metadata remains available for downstream analysis. Furthermore, it allows the accumulation of multiple point clouds based on their timestamps and transforms. By combining support for robotic middleware data with batch-oriented file processing, PointCloudCrafter facilitates reproducible workflows across a range of scientific and engineering applications that rely on large-scale three-dimensional point cloud data.

²¹ Statement of need

²² Three-dimensional point cloud data is widely used in robotics, autonomous driving, remote sensing, and scientific data analysis. In practice, such data is commonly encountered in two distinct representations. During acquisition and experimentation, point clouds are often recorded as time-stamped data streams using robotic middleware such as ROS2. For benchmarking, archival, and offline analysis, the same data is typically stored and distributed as static files in dataset-specific or standardized formats. Especially, the use of automated pipelines for neural network training often requires the input data to be static files with a specific naming scheme. Transitioning between these representations, or performing large-scale batch operations on point cloud collections, remains a recurring challenge.

³¹ The Point Cloud Library (PCL) ([Rusu & Cousins, 2011](#)) provides a comprehensive set of data structures and algorithms for three-dimensional processing and serves as the technical foundation of PointCloudCrafter, but its use generally requires direct integration into custom software projects. Graphical tools such as CloudCompare are well-suited for interactive inspection but are not designed for automated, headless processing of large datasets. Python-based libraries such as Open3D ([Zhou et al., 2018](#)) and Pyoints ([Lamprecht, 2019](#)) offer accessible interfaces for point cloud manipulation, but performance and input-output overhead can become limiting factors when processing large directories of high-resolution scans. PDAL ([Contributors, 2022](#)) is a mature framework for large-scale geospatial point cloud processing and provides a flexible pipeline abstraction for working with formats such as LAS and LAZ. While PDAL is well-suited for geospatial analysis and national lidar datasets, it does not natively integrate with

42 robotic middleware or support time-resolved transformations based on sensor pose information.
43 Recent toolkits such as PointClouds.jl (Schmid et al., 2025) provide efficient and flexible
44 programmatic interfaces for point cloud processing within the Julia ecosystem, with a focus on
45 interactive development and integration with geospatial data sources. While well-suited for
46 algorithm development and exploratory analysis, such approaches require users to implement
47 custom processing logic and are not designed as standalone command-line tools for automated
48 conversion of robotic middleware data. Libraries such as libpointmatcher (Pomerleau et
49 al., 2013) focus on modular implementations of point cloud registration algorithms and are
50 typically embedded within state estimation or SLAM systems rather than used for dataset-level
51 preprocessing and format conversion.

52 A further limitation of many existing tools is the implicit assumption of fixed-point schemas.
53 Conversion utilities and parsers often restrict data to spatial coordinates and optional color
54 information, discarding non-standard scalar fields. However, many scientific applications depend
55 on additional per-point attributes. Automotive LiDAR sensors provide channels such as intensity,
56 timestamp, and ring index, which are required for calibration, motion correction, and mapping.
57 Furthermore, novel imaging RaDAR sensors export point cloud data with metadata, such as
58 Doppler velocity, RaDAR-cross-section, or signal-to-noise-ratio. In other domains, including
59 structural biology, point clouds may encode physicochemical properties such as hydrophobicity
60 or electrostatic potential (Bazzano et al., 2025; Parigger et al., 2024). Preserving these
61 attributes across processing steps is essential for reproducibility and downstream analysis.

62 PointCloudCrafter addresses these needs by providing a lightweight CLI toolkit focused on
63 batch-oriented point cloud conversion and processing. The primary software contribution is a
64 schema-agnostic, command-line-driven framework that unifies middleware-based and file-based
65 point cloud processing while preserving arbitrary per-point attributes. Arbitrary scalar fields
66 refer to dynamically discovered per-point attributes whose names, data types, and memory
67 layout are preserved without enforcing a predefined schema.

68 Software design

69 PointCloudCrafter is implemented in C++ and is designed primarily as a command-line toolkit
70 composed of standalone executables for point cloud processing. The software builds on the
71 Point Cloud Library (PCL) (Rusu & Cousins, 2011), which provides the underlying data
72 structures and algorithms for ROS2 point cloud streaming and for commonly used file formats
73 such as PCD. PointCloudCrafter does not reimplement PCL functionality; instead, it exposes
74 selected operations and extends them through a unified, reproducible command-line interface.

75 The software is developed using modern C++ design practices, including a clear separation
76 between executable logic and shared processing components. This design improves
77 maintainability and allows new command-line tools to be added in the future without altering
78 existing workflows. Documentation of the codebase and available commands is generated with
79 Doxygen and published as a static website on GitHub.

80 The toolkit consists of two main executables that target distinct workflows. These executables
81 share a common internal library that implements functionality for loading, saving, analyzing,
82 transforming, filtering, and merging point clouds. Its purpose is to ensure consistent behavior
83 and reduce code duplication across the executables, while keeping the user-facing interface
84 focused on command-line usage.

85 Internally, all point cloud data is represented using the PCLPointCloud2 data structure. This
86 representation enables schema-agnostic handling of point clouds by allowing arbitrary per-point
87 attributes to be discovered and preserved at runtime without enforcing a predefined schema.

88 PointCloudCrafter provides two primary execution modes, *rosbag* and *file*, for streaming
89 and static data workflows, respectively. The *rosbag* executable enables extraction of
90 `sensor_msgs::msg::PointCloud2` topics from ROS2 bag recordings. This mode integrates
91 with the TF2 library (Foote, 2013) to perform time-resolved pose lookups based on the

recorded transform tree. Each point cloud can be transformed from the sensor coordinate frame into a user-specified fixed reference frame using the pose corresponding to the exact acquisition timestamp. This enables motion-consistent export of point cloud data from dynamic sensor platforms and a timestamp-based fusion of different point cloud sensor topics. The `file` executable focuses on batch-oriented processing of static point cloud files. Supported formats include common text and binary formats such as PLY, PCD, OBJ, and TXT, as well as dataset-specific binary formats used in KITTI (Geiger et al., 2012) and nuScenes (Caesar et al., 2020).

Both executables support a common set of processing operations implemented in the shared internal codebase. It allows users to apply rigid-body transformations, merge multiple point clouds into a single representation, analyse timestamps, and perform filtering operations. Filters include geometric cropping using box, sphere, or cylindrical regions, voxel grid downsampling, and statistical or radius-based outlier removal. All processing steps are fully parameterized via command-line options, enabling deterministic and reproducible execution across datasets. Due to the schema-agnostic internal representation, per-point attributes, such as sensor-specific metadata or domain-specific properties, are preserved throughout processing whenever supported by the input and output formats.

A comprehensive overview of the available commands, parameters, and usage examples is provided in the online documentation at <https://tumftm.github.io/PointCloudCrafter/usage.html>.

Research impact statement

PointCloudCrafter has been developed to support reproducible preprocessing of large-scale point cloud datasets in research environments. Early versions of the toolkit have been used in student research projects, scientific publications (Kulmer, Leitenstern, et al., 2025; Leitenstern et al., 2025), and the creation of publicly available datasets (Kulmer, Tahiraj, et al., 2025). It has been used in a range of internal research workflows related to autonomous driving and robotic perception. These applications include analyzing sensor drivers by inspecting timestamp consistency and packet loss, which manifest as incomplete or irregular point clouds, as well as fusing multiple ROS2 LiDAR topics into a unified point cloud representation using a user-specified static reference frame. The toolkit has been further applied to the creation of custom datasets for evaluating state-of-the-art open-source algorithms in pseudo-labeling, object detection, localization, and mapping. In addition, PointCloudCrafter has been integrated into automated data processing pipelines that convert logged ROS2 data into established dataset formats, such as KITTI, to enable systematic evaluation and labeling for subsequent machine learning training. The software has also been used to generate large-scale, high-resolution point cloud maps by aggregating temporally indexed scans using externally estimated sensor poses.

By providing a middleware-aware, schema-agnostic CLI toolkit, PointCloudCrafter lowers the barrier for converting raw sensor data into analysis-ready datasets. Its focus on ease of use, reproducibility, and preservation of per-point metadata makes it applicable across multiple domains that rely on three-dimensional point cloud data.

Availability

The source code is available at <https://github.com/TUMFTM/PointCloudCrafter>.
The Python package is available at <https://pypi.org/project/pointcloudcrafter>.
The Docker is available at <https://ghcr.io/tumftm/pointcloudcrafter:latest>.
The documentation is available at <https://tumftm.github.io/PointCloudCrafter>.
The license is Apache-2.0.

¹³⁹ AI usage disclosure

¹⁴⁰ Artificial intelligence tools were used in a limited and transparent manner during the development
¹⁴¹ of this work. Claude Opus 4.5 was used to assist with code structuring and to generate code
¹⁴² comments, as well as test cases. GitHub Copilot was used to review Pull Request. The
¹⁴³ manuscript text was reviewed for grammar and clarity using OpenAI GPT-5.2 and Grammarly,
¹⁴⁴ and selected passages were rephrased to improve readability and flow. All scientific content,
¹⁴⁵ software design decisions, and final wording were reviewed and approved by the authors, who
¹⁴⁶ retain full responsibility for the work.

¹⁴⁷ Acknowledgements

¹⁴⁸ The authors acknowledge the developers and maintainers of the Point Cloud Library (PCL)
¹⁴⁹ ([Rusu & Cousins, 2011](#)), upon which PointCloudCrafter is built. We also acknowledge the
¹⁵⁰ contributions of student researchers at the Institute of Automotive Technology at the Technical
¹⁵¹ University of Munich who supported the development and evaluation of this software.
¹⁵² This work was funded by basic research funds of the Institute of Automotive Technology at
¹⁵³ the Technical University of Munich.

¹⁵⁴ References

- ¹⁵⁵ Bazzano, C. F., Alves, L. F. G., Telles, G. P., & Trivella, D. B. B. (2025). Labeled dataset of
¹⁵⁶ x-ray protein ligand images in 3D point cloud and validated deep learning models. *Scientific
Data*, 12, 1726. <https://doi.org/10.1038/s41597-025-06002-8>
- ¹⁵⁸ Caesar, H., Bankiti, V., Lang, A., Vora, S., Liong, V., Xu, Q., Krishnan, A., Pan, Y., Baldan,
¹⁵⁹ G., & Beijbom, O. (2020). *nuScenes: A multimodal dataset for autonomous driving* (pp.
¹⁶⁰ 11621–11631) [Data set].
- ¹⁶¹ Contributors, P. (2022). *PDAL point data abstraction library*. <https://doi.org/10.5281/zenodo.2616780>
- ¹⁶³ Foote, T. (2013). Tf: The transform library. *2013 IEEE Conference on Technologies
for Practical Robot Applications (TePRA)*, 1–6. <https://doi.org/10.1109/TePRA.2013.6556373>
- ¹⁶⁶ Geiger, A., Lenz, P., & Urtasun, R. (2012). Are we ready for autonomous driving? The KITTI
¹⁶⁷ vision benchmark suite. *Conference on Computer Vision and Pattern Recognition (CVPR)*.
- ¹⁶⁸ Kulmer, D., Leitenstern, M., Weinmann, M., & Lienkamp, M. (2025). *OpenLiDARMap:
Zero-drift point cloud mapping using map priors* (pp. 178–188). INSTICC; SciTePress.
<https://doi.org/10.5220/0013405400003941>
- ¹⁷¹ Kulmer, D., Tahiraj, I., & Lienkamp, M. (2025). *Autonomous Driver Licence (ADL) dataset
[Forschungsdaten]*. <https://doi.org/10.14459/2025mp1785105>
- ¹⁷³ Lamprecht, S. (2019). Pyoints: A python package for point cloud, voxel and raster processing.
¹⁷⁴ *Journal of Open Source Software*, 4, 990. <https://doi.org/10.21105/joss.00990>
- ¹⁷⁵ Leitenstern, M., Alten, M., Bolea-Schaser, C., Kulmer, D., Weinmann, M., & Lienkamp, M.
¹⁷⁶ (2025). *FlexCloud: Direct, modular georeferencing and drift-correction of point cloud maps*
¹⁷⁷ (pp. 157–165). INSTICC; SciTePress. <https://doi.org/10.5220/0013359600003941>
- ¹⁷⁸ Parigger, L., Krassnigg, A., Hetmann, M., Hofmann, A., Gruber, K., Steinkellner, G., &
¹⁷⁹ Gruber, C. C. (2024). CavitOmiX drug discovery: Engineering antivirals with enhanced
¹⁸⁰ spectrum and reduced side effects for arboviral diseases. *Viruses*, 16(8), 1186. <https://doi.org/10.3390/v16081186>

- ¹⁸² Pomerleau, F., Colas, F., Siegwart, R., & Magnenat, S. (2013). Comparing ICP Variants on
¹⁸³ Real-World Data Sets. *Autonomous Robots*, 34(3), 133–148.
- ¹⁸⁴ Rusu, R. B., & Cousins, S. (2011). 3D is here: Point cloud library (PCL). *IEEE International
185 Conference on Robotics and Automation (ICRA)*, 1–4. <https://doi.org/10.1109/ICRA.2011.5980567>
- ¹⁸⁷ Schmid, M. F., Massey, J. D., & Giometto, M. G. (2025). PointClouds.jl: Fast & flexible
¹⁸⁸ processing of lidar data. *Journal of Open Source Software*, 10, 7952. <https://doi.org/10.21105/joss.07952>
- ¹⁹⁰ Zhou, Q.-Y., Park, J., & Koltun, V. (2018). Open3D: A modern library for 3D data processing.
¹⁹¹ [arXiv:1801.09847](https://arxiv.org/abs/1801.09847).

DRAFT