

¹ AutoEmulate: A PyTorch tool for end-to-end emulation workflows

³ Radka Jersakova  ^{1*}, Sam F. Greenbury  ^{1*}, Ed Chalstrey  ¹, Edwin Brown  ², Marjan Famili  ¹, Chris Sprague  ¹, Paolo Conti  ¹, Camila Rangel Smith  ¹, Martin A. Stoffel  ¹, Bryan M. Li  ^{1,5}, Kalle Westerling  ¹, Sophie Arana  ¹, Max Balmus  ^{1,3}, Eric Daub  ¹, Steve Niederer  ^{1,3}, Andrew B. Duncan  ³, and Jason D. McEwen  ^{1,4}

⁸ 1 The Alan Turing Institute, London, United Kingdom ² University of Sheffield, Sheffield, United
⁹ Kingdom ³ Imperial College London, London, United Kingdom ⁴ University College London, London,
¹⁰ United Kingdom ⁵ University of Edinburgh, Edinburgh, United Kingdom ¶ Corresponding author *
¹¹ These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: Evan Spotte-Smith  ¹⁶

Reviewers:

- [@stefanradev93](#)
- [@pmeier](#)

Submitted: 25 September 2025

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/)).

¹² Summary

¹³ Computational simulations lie at the heart of modern science and engineering, but they are often slow and computationally costly. A common solution is to use emulators: fast, cheap ¹⁴ models trained to approximate the simulator. However, constructing these requires substantial expertise. AutoEmulate ([Stoffel et al., 2025](#)) is a low-code Python package for emulation ¹⁵ workflows, making it easy to replace simulations with fast, accurate emulators. AutoEmulate has now been fully refactored to use PyTorch as a backend, enabling GPU acceleration, automatic differentiation, and seamless integration with the broader PyTorch ecosystem (released in v1.0). The toolkit has also been extended with easy-to-use interfaces for common emulation tasks, including model calibration (determining which input values are most likely to have generated real-world observations) and active learning (where simulations are chosen to improve emulator performance at minimal computational cost). Together these updates make AutoEmulate uniquely suited to running performant end-to-end emulation workflows.

²⁴ Statement of need

²⁶ Physical systems are often modelled using computer simulations. Depending on the complexity of the system, these simulations can be computationally expensive and time-consuming. This ²⁷ bottleneck can be resolved by approximating simulations with emulators, which can be orders ²⁸ of magnitudes faster ([Kennedy & O'Hagan, 2000](#)).

³⁰ Emulation requires significant expertise in machine learning as well as familiarity with a broad ³¹ and evolving ecosystem of tools. This creates a barrier to entry for domain researchers ³² whose focus is on the underlying scientific problem. AutoEmulate ([Stoffel et al., 2025](#)) ³³ lowers the barrier to entry by automating the entire emulator construction process (training, ³⁴ hyperparameter tuning and model selection). This makes emulation accessible to non-specialists ³⁵ while also offering a reference set of emulators for benchmarking to experienced users.

³⁶ AutoEmulate was originally built on scikit-learn, which is well suited for traditional machine ³⁷ learning but less flexible for complex workflows. AutoEmulate v1.0 introduced a PyTorch ³⁸ ([Paszke et al., 2019](#)) backend that provides GPU acceleration for faster training and inference ³⁹ and automatic differentiation via PyTorch's autograd system. It also made AutoEmulate easy ⁴⁰ to integrate with other PyTorch-based tools. For example, the PyTorch refactor enabled fast

⁴¹ Bayesian model calibration (identifying input values most likely to have generated real-world
⁴² observations) using gradient-based inference methods such as Hamiltonian Monte Carlo exposed
⁴³ through Pyro ([Bingham et al., 2018](#)).

⁴⁴ The latest version of AutoEmulate now also supports direct integration of custom simulators and
⁴⁵ active learning, in which the tool adaptively selects informative simulations to run to improve
⁴⁶ emulator performance at minimal computational cost. Additionally, the AutoEmulate refactor
⁴⁷ improved support for high-dimensional data through dimensionality reduction techniques such
⁴⁸ as principal component analysis (PCA) and variational autoencoders (VAEs).

⁴⁹ State of the field

⁵⁰ This paper describes an extensive contribution to an existing package. We felt that AutoEmulate
⁵¹ ([Stoffel et al., 2025](#)) already filled a unique gap in the ecosystem by focusing on making
⁵² emulation accessible to domain researchers unfamiliar with ML. However, in its reliance on
⁵³ scikit-learn as a backend we could not extend it to handle use cases that we were targeting.
⁵⁴ Refactoring the backend to be PyTorch-first allowed us to leverage the wider PyTorch ecosystem
⁵⁵ as well as the benefits of having end-to-end automatically differentiable emulators and GPU
⁵⁶ acceleration. This has resulted in a tool that uniquely brings together a wide range of emulation
⁵⁷ capabilities (e.g., sensitivity analysis, calibration, active learning) and translated to a significant
⁵⁸ growth of the user base and package contributors. We have also retained support for some of
⁵⁹ the non-PyTorch features following discussions with the community (e.g., the users can still
⁶⁰ opt in to fit classic ML models such as SVMs although this results in loss of compatibility with
⁶¹ some of the more advanced features).

⁶² Software Design

⁶³ AutoEmulate design is centered around (i) low-code mode, (ii) modularity and (iii) integrating
⁶⁴ with the wider ecosystem wherever possible. The design has now been updated from being
⁶⁵ scikit-learn oriented to PyTorch-first.

⁶⁶ AutoEmulate primarily targets users who are simulation but not ML experts, aiming to make
⁶⁷ it as easy as possible to fit emulators to their simulated data. We also offer flexibility to
⁶⁸ advanced users by exposing customizable parameters through our APIs (set to sensible defaults
⁶⁹ to abstract complexity away from novice users).

⁷⁰ The software's modular design is built on base classes for each component, enabling users to
⁷¹ easily add new emulators and functionality. Our documentation showcases how to do this,
⁷² which has already encouraged community contributions to the software. We chose PyTorch
⁷³ as the backend because of its autodiff and GPU capabilities as well as the mature ecosystem
⁷⁴ that we could integrate with. For example, both GPyTorch ([Gardner et al., 2021](#)) and Pyro
⁷⁵ ([Bingham et al., 2018](#)) are extensively utilised within the package.

⁷⁶ Example usage

⁷⁷ The AutoEmulate documentation provides a comprehensive set of [tutorials](#) showcasing all
⁷⁸ functionality. We are also collecting [case studies](#) demonstrating how to use AutoEmulate for
⁷⁹ real-world problems and complex workflows. Below we provide a brief overview of the main
⁸⁰ features.

⁸¹ The core use case for AutoEmulate is emulator construction. AutoEmulate takes as input
⁸² variables x , y . The variable x is a 2D array with columns corresponding to simulation parameters
⁸³ and rows corresponding to parameter sets. The variable y is an array of one or more simulation
⁸⁴ outputs corresponding to each set of parameters. From this data, AutoEmulate constructs an
⁸⁵ emulator in just a few lines of code:

```
from autoemulate import AutoEmulate
```

```
ae = AutoEmulate(x, y)
```

```
result = ae.best_result()
```

```
emulator = result.model
```

86 This simple script runs a search over a library of emulator models, performs hyperparameter
 87 tuning and compares models using cross validation. Each model is stored along with metadata
 88 in a Results object. The user can then easily extract the best performing emulator.

89 AutoEmulate can additionally search over different data preprocessing methods, such as
 90 normalization or dimensionality reduction techniques (PCA, VAEs). Any Transform from
 91 PyTorch distributions can also be used. The transforms are passed as a list to permit the
 92 user to define a sequence of transforms to apply to the data. For example, the following
 93 code standardizes the input data and compares three different output transformations: no
 94 transformation, PCA with 16 components, and PCA with 32 components in combination with
 95 the default set of emulators:

```
from autoemulate.transforms import PCATransform, StandardizeTransform
```

```
ae = AutoEmulate(  
    x,  
    y,  
    x_transforms_list=[[StandardizeTransform]]  
    y_transforms_list=[  
        [],  
        [PCATransform(n_components=16)],  
        [PCATransform(n_components=32)]  
    ],  
)
```

96 The result in this case will return the best combination of model and output transform. The
 97 returned emulator and transforms are wrapped together in a TransformedEmulator class,
 98 which outputs predictions in the original data space. The figure below shows an example
 99 result of fitting a Gaussian Process emulator in combination with PCA to a reaction-diffusion
 100 simulation (see the full [tutorial](#) for a detailed overview).




Figure 1: GP with PCA emulator prediction for a reaction diffusion simulation compared to the ground truth.

101 Once an emulator has been trained it can generate fast predictions for new input values,
 102 enabling [downstream tasks](#) such as [sensitivity analysis](#) or [model calibration](#). For example, to
 103 run Sobol sensitivity analysis one only needs to pass the trained emulator and some information
 104 about the data. Below is a dummy example assuming a simulation with two input parameters
 105 param1 and param2, each with a plausible range of values, and two outputs output1 and
 106 output2:

```

from autoemulate.core.sensitivity_analysis import SensitivityAnalysis

input_parameters_ranges = {
    'param1': (0, 1),
    'param2': (0, 10),
}

problem = {
    'num_vars': 2,
    'names': ["param1", "param2"],
    'bounds': input_parameters_ranges.values(),
    'output_names': ["output1", "output2"],
}

sa = SensitivityAnalysis(emulator, problem=problem)
sobol_df = sa.run()

107 AutoEmulate also provides a simple interface for calibration given a trained emulator, input
108 parameter ranges (same as in the sensitivity analysis example), and real-world observations:
109
110 from autoemulate.calibration.bayes import BayesianCalibration

111 observations = {'output1': 0.5, 'output2': 7.2}

112 bc = BayesianCalibration(
113     emulator,
114     input_parameters_ranges,
115     observations,
116 )
117 mcmc = bc.run()

118 Lastly, AutoEmulate makes it easy to integrate custom simulators through subclassing. This
119 enables simulator-in-the-loop workflows like active learning, which selects the most informative
120 simulations to improve emulator performance at minimal computational cost.

```

112 Research Impact Statement

113 In the last year, we have worked with around 10 collaborators across diverse domains including
114 biomedicine and materials science. This has led to academic outputs such as a poster at
115 OFEME2025. Additionally, our collaborations have driven software development through
116 numerous feature requests and bug reports that we have addressed. For example, we have
117 implemented a full end-to-end calibration workflow used by our collaborators in cardiac
118 modelling and demonstrated how to use AutoEmulate in their pipelines in one of our [case](#)
119 [studies](#). We have also had contributions outside the core development team. This has included
120 external contributors responding to existing issues as well as users adapting the tool for their
121 own use cases (e.g., contributing new types of emulators).

122 AI usage disclosure

123 Human authors have made all the core design decisions and authored much of the code and
124 documentation. Generative AI tools have been used to assist with code and documentation
125 writing. Specifically, the team uses GitHub Copilot in auto mode or selects one of the available
126 versions of Claude, GPT and Gemini. We confirm that human authors have reviewed, edited
127 and validated all AI-assisted outputs. We have also added a section on the use of generative
128 AI tools in our contributing guidelines.

129 References

- 130 Bingham, E., Chen, J. P., Jankowiak, M., Obermeyer, F., Pradhan, N., Karaletsos, T., Singh,
131 R., Szerlip, P. A., Horsfall, P., & Goodman, N. D. (2018). Pyro: Deep universal probabilistic
132 programming. *CoRR*, *abs/1810.09538*. <http://arxiv.org/abs/1810.09538>
- 133 Gardner, J. R., Pleiss, G., Bindel, D., Weinberger, K. Q., & Wilson, A. G. (2021). *GPyTorch:*
134 *Blackbox matrix-matrix gaussian process inference with GPU acceleration*. <https://arxiv.org/abs/1809.11165>
- 135 Kennedy, M. C., & O'Hagan, A. (2000). Predicting the output from a complex computer code
136 when fast approximations are available. *Biometrika*, *87*(1), 1–13. <http://www.jstor.org/stable/2673557>
- 137 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z.,
138 Gimelshein, N., Antiga, L., Desmaison, A., Köpf, A., Yang, E. Z., DeVito, Z., Raison,
139 M., Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., ... Chintala, S. (2019). PyTorch:
140 An imperative style, high-performance deep learning library. *CoRR*, *abs/1912.01703*.
141 <http://arxiv.org/abs/1912.01703>
- 142 Stoffel, M. A., Li, B. M., Westerling, K., Arana, S., Balmus, M., Daub, E., & Niederer, S.
143 (2025). AutoEmulate: A python package for semi-automated emulation. *Journal of Open
144 Source Software*, *10*(107), 7626. <https://doi.org/10.21105/joss.07626>
- 145
- 146