

¹ Uno: a composable framework for nonlinearly constrained optimization

³ Charlie Vanaret  ^{1¶} and Alexis Montoison  ²

⁴ 1 Applied Optimization Department, Zuse-Institut Berlin, Germany ² Mathematics and Computer Science Division, Argonne National Laboratory, USA ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: 

Submitted: 20 February 2026

Published: unpublished

License

Authors of papers retain copyright

and release the work under a

Creative Commons Attribution 4.0

International License ([CC BY 4.0](#))

⁶ Summary

⁷ Uno is a composable software framework for nonlinearly constrained optimization written
⁸ in modern C++. It unifies the workflows of Lagrange-Newton methods, i.e., gradient-based
⁹ algorithms that iteratively solve the KKT optimality conditions using Newton's method. As
¹⁰ of February 2026, Uno supports sequential (convex and nonconvex) quadratic programming,
¹¹ interior-point (barrier) methods, and sequential linear programming.

¹² Uno breaks down optimization algorithms into reusable modular components such as step
¹³ computation, constraint reformulation, globalization techniques, and acceptance criteria. This
¹⁴ allows classical and hybrid methods to be configured and compared within a single framework.
For full mathematical details of the algorithms implemented in Uno, see ([Vanaret & Leyffer, 2024](#)).

¹⁵ The core C++ code of Uno is organized into modular, object-oriented components that
¹⁶ separate the mathematical logic of the algorithms from implementation details such as
¹⁷ memory management, data structures, and computational routines. Uno provides interfaces to
¹⁸ Julia, Python, C, Fortran, and AMPL, enabling use across scientific computing environments.
¹⁹ Precompiled artifacts are available on GitHub, and the solver can be accessed directly via
²⁰ UnoSolver.jl in Julia or unopy in Python.

²³ Statement of need

²⁴ Nonlinearly constrained optimization is central to engineering, optimal control, machine learning,
²⁵ and scientific modeling ([Nocedal & Wright, 2006](#)). It also plays a key role in mixed-integer
²⁶ nonlinear optimization ([Lee & Leyffer, 2011](#)).

²⁷ Typical nonlinear solvers share common building blocks such as step computation, constraint
²⁸ handling, and globalization strategies. These components can be found across major paradigms
²⁹ such as sequential quadratic programming, interior-point methods, and augmented Lagrangian
³⁰ methods, but these are usually developed as separate solver families. Algorithmic ideas are
³¹ typically tested at the level of complete solvers, rather than individual components.

³² These observations motivate the development of a unified and composable framework in which
³³ algorithmic components are made explicit.

³⁴ State of the field

³⁵ Popular solvers such as IPOPT ([Wächter & Biegler, 2006](#)), KNITRO ([Byrd et al., 2006](#)), and
³⁶ SNOPT ([Gill et al., 2005](#)) are robust and efficient, but they are typically monolithic: parameter
³⁷ tuning is exposed while internal components remain rigid. These solvers are primarily designed

38 for end users rather than algorithmic experimentation. Testing new algorithmic variants usually
 39 requires intrusive modification or reimplementation.

40 Unification framework

41 Uno enables experimentation at the component level by providing a composable optimization
 42 framework in which algorithms are constructed from modular components that reflect
 43 mathematical concepts such as step computation, constraint reformulation, and globalization
 44 techniques.

45 Within this framework, strategies are organized into a coherent hierarchy, illustrated in the
 46 wheel of strategies (Figure 1): the outer ring represents high-level layers, the middle ring
 47 represents algorithmic ingredients that are combined automatically, and the inner ring lists
 48 possible strategies for each ingredient. The strategies currently implemented in Uno are
 49 highlighted in green.



Figure 1: Unification framework: wheel of strategies.

50 Software design

51 In Uno's object-oriented architecture, the ingredients of Figure 1 are abstract classes whose
 52 interfaces should be implemented by subclasses (the strategies). Uno's simplified UML diagram
 53 is shown in Figure 2. Inheritance is represented as dashed lines with white arrows, while
 54 composition is represented as solid lines with black diamonds. Abstract classes are written in
 55 italic, while subclasses are written in bold. This modular architecture offers a clear separation
 56 between the mathematical logic of the optimization algorithm (the reformulation of the problem,
 57 the definition of the subproblem, and the globalization techniques) and the computational
 58 aspects (evaluating the model's functions, and solving the subproblems).

59 Uno implements a generic Lagrange-Newton method in which the abstract classes interact with

one another and exchange data, while being agnostic of the underlying strategies. Strategies are picked by the user at runtime via options. Uno also implements presets, that is particular combinations of strategies (and sets of hyperparameters) that correspond to state-of-the-art solvers. Uno currently implements an ipopt preset that mimics the IPOPT solver ([Wächter & Biegler, 2006](#)) (a *line-search restoration filter interior-point method with exact Hessian and primal-dual inertia correction*), and a filtersqp preset that mimics the filterSQP solver ([Fletcher & Leyffer, 1998](#)) (a *trust-region restoration filter SQP method with exact Hessian and no inertia correction*). While, in theory, all combinations of strategies can be generated, some are not supported yet (e.g., interior-point method with a trust-region constraint).

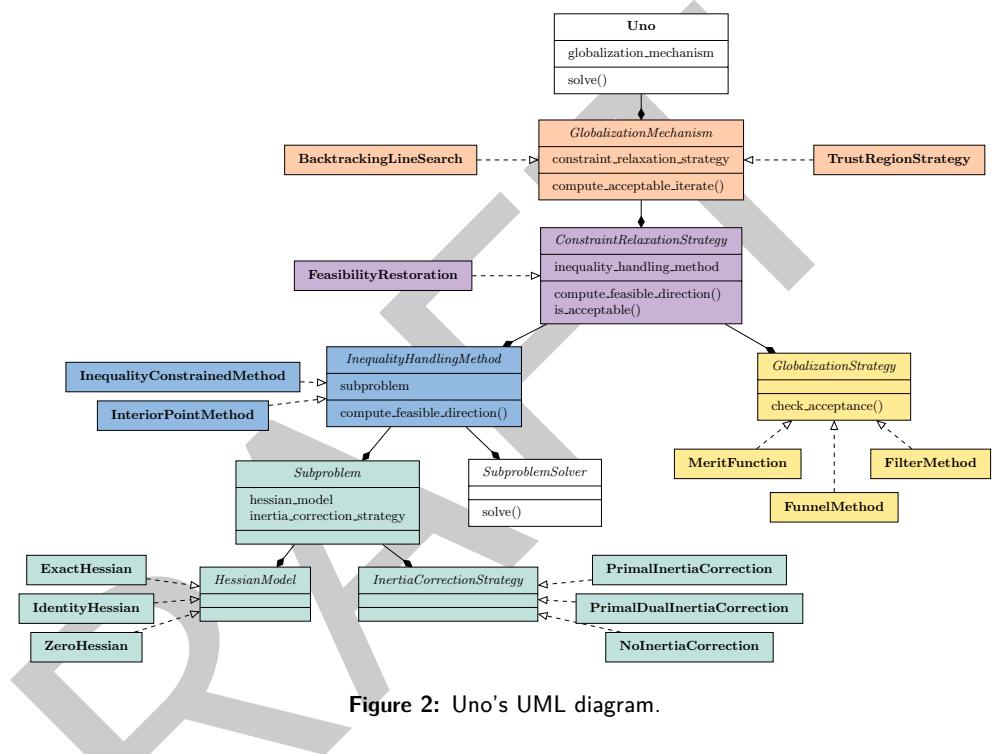


Figure 2: Uno's UML diagram.

Interfaces

- To make Uno accessible to a wide range of users, we provide multiple language interfaces.
- The AMPL Solver Library ([Gay, 1997](#)) interface gives access to the AMPL ([Fourer et al., 1990](#)) modeling language for optimization. It is distributed as a binary that takes a compiled AMPL model (.nl file) as input, allowing users to solve problems without directly interacting with the C++ core.
- The C interface provides direct access to Uno's core functionality while maximizing interoperability with other programming languages and tools. The optimization process is expressed as the interaction between two independent opaque objects: a model describing the optimization problem and a solver holding the algorithmic configuration and execution state. The interface is therefore centered around two main structures:
- **Model:** represents an optimization problem and stores information about variables, bounds, constraints, the objective function, and derivative information. Users create a model and set its components (objective, constraints, derivatives, initial primal-dual point).
 - **Solver:** represents the algorithm used to solve a given model. Users set solver options, attach callbacks, and access results such as primal and dual solutions, residuals, iteration counts, and performance metrics.

87 The Fortran interface provides access to Uno through the C API using `iso_c_binding`. It
88 is split into two files: `uno_c.f90` for low-level C bindings, and `uno_fortran.f90` for Fortran-
89 friendly wrappers that handle string arguments. The interface can be included directly in a
90 source file or wrapped in a module for cleaner use statements. It is provided as source rather
91 than a precompiled Fortran module (`.mod`) to maximize portability and interoperability across
92 compilers and platforms.

93 The Julia interface is distributed as the registered package `UnoSolver.jl`. It integrates Uno
94 into the Julia optimization ecosystem through:

- 95 ▪ a thin wrapper around the full C API,
96 ▪ an interface to `NLPModels.jl` for solving problems following the `NLPModels` API, such
97 as `ADNLPModels.jl` or `ExaModels.jl`,
98 ▪ an interface to `MathOptInterface.jl` for handling JuMP models.

99 Precompiled artifacts are downloaded automatically, making the package plug-and-play without
100 requiring user compilation.

101 The Python interface is registered on PyPI as `unopy`, providing Uno bindings via precompiled
102 wheels on most platforms.

103 A MATLAB interface is also under development, further expanding Uno's accessibility.

104 Research impact statement

105 Uno's `ipopt` and `filtersqp` presets currently perform on a par with the state-of-the-art solvers
106 IPOPT (Uno is slightly less robust) and filterSQP (Uno is slightly more robust) in terms
107 of function evaluations on a set of 429 small CUTE instances (Bongartz et al., 1995). An
108 up-to-date performance profile is maintained on Uno's GitHub README page.

109 Uno's ongoing developments were presented at several international conferences over the years
110 (ISMP 2018, SIAM OP 2021, ICCOPT 2022, ISMP 2024, and ICCOPT 2025), and were the
111 subject of invited talks at Zuse-Institut Berlin (2020), Argonne National Laboratory (2022),
112 and KU Leuven (2025). This resulted in scientific cooperations with the MECO team at KU
113 Leuven (Kiessling et al., 2024, 2025) and the HiGHS team in Edinburgh.

114 Uno is currently used as a nonlinear optimization solver in:

- 115 ▪ the JuMP.jl ecosystem,
116 ▪ DNLP, an extension of CVXPY to general nonlinear programming,
117 ▪ Vecchia.jl, a package for Gaussian processes approximation,
118 ▪ FelooPy, a user-friendly tool for coding, modeling, and solving decision problems,
119 ▪ IMPL © /IMPL-DATA © by Industrial Algorithms Limited, a modeling and solving
120 platform used in the process industries especially suited for economic, efficiency and
121 emissions optimization and estimation.

122 Ongoing discussions and community interest indicate potential future integrations in CasADI,
123 Pyomo, pyOptSparse, Minotaur, and the NEOS Server.

124 Acknowledgments

125 This work was supported by the Applied Mathematics activity of the U.S. Department of
126 Energy Office of Science, Advanced Scientific Computing Research, under Contract No. DE-
127 AC02-06CH11357. This work was also supported in part by NSF CSSI Grant No. 2104068.

¹²⁸ AI usage disclosure

¹²⁹ No generative AI was used in the creation of the software, interfaces, implementation of
¹³⁰ algorithms, or documentation. ChatGPT was used to check spelling, grammar, and clarity of
¹³¹ the English text in this paper.

¹³² References

- ¹³³ Bongartz, I., Conn, A. R., Gould, N., & Toint, P. L. (1995). CUTE: Constrained and
¹³⁴ unconstrained testing environment. *ACM Transactions on Mathematical Software (TOMS)*,
¹³⁵ 21(1), 123–160. <https://doi.org/10.1145/200979.201043>
- ¹³⁶ Byrd, R. H., Nocedal, J., & Waltz, R. A. (2006). KNITRO: An integrated package for
¹³⁷ nonlinear optimization. In *Large-scale nonlinear optimization* (pp. 35–59). Springer.
¹³⁸ https://doi.org/10.1007/0-387-30065-1_4
- ¹³⁹ Fletcher, R., & Leyffer, S. (1998). User manual for filterSQP. *Numerical Analysis Report*
¹⁴⁰ NA/181, Department of Mathematics, University of Dundee, Dundee, Scotland, 35.
- ¹⁴¹ Fourer, R., Gay, D. M., & Kernighan, B. W. (1990). AMPL: A mathematical programming
¹⁴² language. *Management Science*, 36(5), 519–554. https://doi.org/10.1007/978-3-642-83724-1_12
- ¹⁴⁴ Gay, D. M. (1997). *Hooking your solver to AMPL*. Technical Report 93-10, AT&T Bell
¹⁴⁵ Laboratories, Murray Hill, NJ, 1993, revised.
- ¹⁴⁶ Gill, P. E., Murray, W., & Saunders, M. A. (2005). SNOPT: An SQP algorithm for large-
¹⁴⁷ scale constrained optimization. *SIAM Review*, 47(1), 99–131. <https://doi.org/10.1137/S0036144504446096>
- ¹⁴⁹ Kiessling, D., Leyffer, S., & Vanaret, C. (2025). A unified funnel restoration SQP algorithm:
¹⁵⁰ D. Kiessling et al. *Mathematical Programming*, 1–45. <https://doi.org/10.1007/s10107-025-02284-3>
- ¹⁵² Kiessling, D., Vanaret, C., Astudillo, A., Decré, W., & Swevers, J. (2024). An almost feasible
¹⁵³ sequential linear programming algorithm. *2024 European Control Conference (ECC)*,
¹⁵⁴ 2328–2335. <https://doi.org/10.23919/ECC64448.2024.10590864>
- ¹⁵⁵ Lee, J., & Leyffer, S. (2011). *Mixed integer nonlinear programming*. Springer Science &
¹⁵⁶ Business Media. <https://doi.org/10.1007/978-1-4614-1927-3>
- ¹⁵⁷ Nocedal, J., & Wright, S. J. (2006). *Numerical optimization*. Springer. <https://doi.org/10.1007/978-0-387-40065-5>
- ¹⁵⁹ Vanaret, C., & Leyffer, S. (2024). *Implementing a unified solver for nonlinearly constrained*
¹⁶⁰ *optimization*.
- ¹⁶¹ Wächter, A., & Biegler, L. T. (2006). On the implementation of an interior-point filter
¹⁶² line-search algorithm for large-scale nonlinear programming. *Mathematical Programming*,
¹⁶³ 106(1), 25–57. <https://doi.org/10.1007/s10107-004-0559-y>