

Speechmetryflow: A reproducible Nextflow-based pipeline for the automatic extraction of text and acoustic features from speech

Christophe Bedetti^{1,2}, Natasha Clarke^{1,2}, B  reng  re Houz  ^{1,2},
Antoine Slegers³, Sophie Pellerin^{1,2}, Pierre-Briac M  tayer-Mariotti^{1,2},
and Simona Maria Brambati^{1,2}

¹ Department of Psychology, University of Montr  al, Montr  al, Canada ² Research Centre of the Institut
Universitaire de G  riatrie de Montr  al, Montr  al, Canada ³ Institut universitaire en sant   mentale de
Montr  al, Montr  al, Canada  Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: 

Submitted: 06 March 2026

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/))

Summary

The automatic extraction of linguistic and acoustic features from speech is essential for analysing speech in neuroscience and clinical research. Speechmetryflow is an open-source pipeline, built using Python, Nextflow ([Di Tommaso et al., 2017](#)) and Apptainer ([Kurtzer et al., 2021](#)), to enable large-scale, offline, reproducible speech preprocessing and feature extraction. The pipeline currently supports English and French, and extracts general features applicable to any mono-speaker speech task, depending on the length of the sample (for example, features that calculate over a window size of tokens). Features specific to two picture description tasks commonly used in clinical speech evaluation - the Cookie Theft and the Picnic Scene - relating to predetermined information content units specific to each image can also be extracted. For multi-speaker speech samples, we recommended that users first conduct diarisation to isolate utterances from a single speaker. The extracted features can be easily used in downstream tasks such as classification. Speechmetryflow features have already been reported in several peer-reviewed publications ([Slegers et al., 2021](#)) ([Pellerin et al., 2025](#)).

Statement of need

Speech analysis has long been used across various fields of human sciences, including sociolinguistics, education and literacy, psychology and mental health, and political science. A growing body of evidence suggests that it has strong potential as an early biomarker for different clinical populations. However, its adoption in clinical settings remains limited due to the lack of automated analysis pipelines. The role of speech in diagnosing and monitoring disease has led to a demand for tools that automate the extraction of linguistic and acoustic features at scale. This need will continue to increase with advances in automatic speech recognition (ASR) that should enable faster, more accurate transcription of speech tasks. The majority of existing feature extraction pipelines either focus on a narrow subset of features or are not freely available, leading to more complex workflows (utilising multiple tools) that lack reproducibility. Efforts are often duplicated across researchers and teams, using more resources and time. This is compounded by the complexities of speech analyses, in which careful consideration must be given at different steps of the pipeline (for example, lemmatisation versus stemming). Speechmetryflow combines extensive text features from multiple linguistic domains, with established acoustic features implemented in the same workflow, offering a user-friendly “one-stop shop” for comprehensive speech analysis that works out-of-the-box, using the command line. Text features have been implemented according to current best

42 practice in neuroscience and clinical research, ensuring consistency and reproducibility across
43 studies.

44 Workflow overview

45 To enable the user to run the pipeline without an internet connection, the first step is to build
46 the different containers using the provided Apptainer recipes.

47 Subsequently, input data may be in audio, text, or a combination of both formats. It is not
48 necessary to provide all formats for each participant. It is also possible to provide multiple
49 audio or text files for the same participant.

50 The acoustic pipeline processes WAV audio files and extracts features using Parselmouth
51 (Jadoul et al., 2018), the Python interface of Praat (Boersma & Weenink, 2025), including
52 duration, pitch, jitter, shimmer, formant and mel frequency cepstral coefficient features, as
53 well as implementing the Praat toolbox “uhm-o-meter” (Jong et al., 2021) to assess fluency
54 speed and breakdown, measuring silent pauses and filled pauses. Using the Python interface of
55 OpenSmile (Eyben et al., 2010), Speechmetryflow extracts the Geneva Minimalistic Acoustic
56 Parameter Set (Eyben et al., 2016) a minimal set of voice parameters selected based on their
57 potential to index affective physiological changes in voice production, and their proven value
58 in former studies. Each set of parameters results in a CSV file containing the features. Audio
59 files are also automatically transcribed with whisperX (Bain et al., 2023), an implementation
60 of whisper (Radford et al., 2022).

61 The text processing pipeline accepts, in parallel, TXT files provided by the user and the
62 automatic transcriptions obtained from the acoustic pipeline. It implements tokenization,
63 part-of-speech tagging, syntactic parsing, lemmatization and feature extraction, outputting a
64 csv file including measures of speech production, lexical diversity, lexico-semantics, syntactic
65 complexity, coherence and psycholinguistics. Text features are extracted with the Python
66 libraries NLTK (Bird et al., 2009) and TextDescriptives (Hansen et al., 2023), which use
67 spaCy (Montani et al., 2023). The largest spaCy pre-trained statistical model is utilised:
68 en_core_web_lg (343k unique vectors in 300 dimensions) for English and fr_core_news_lg
69 (500k unique vectors in 300 dimensions) for French. Features describing feelings and emotions
70 are also calculated using models available on Hugging Face (Barbieri et al., 2022) (Raffel et
71 al., 2023). An overview of the workflow can be seen in Figure 1.

72 The pipeline uses nextflow, which enables the provision of detailed execution metrics, including
73 CPU usage, memory consumption, runtime duration, and disk I/O statistics for each process,
74 as well as error messages and debugging information.

General overview of speechmetryflow.

Figure 1: General overview of speechmetryflow.

75 Acknowledgements

76 The study is funded by the Natural Sciences and Engineering Research Council of Canada
77 (NSERC) and the Chaire Courtois en recherche fondamentale III (neuroscience) de l'Université
78 de Montréal. NC is funded by a Postdoctoral Fellowship from the Canadian Consortium on
79 Neurodegeneration in Aging Sex & Gender Hub.

80 Availability

81 Speechmetryflow is distributed as open-source software. Its source code, documentation, and
82 usage examples are available at <https://github.com/lingualab/speechmetryflow>. Contributions

as well as community feedback are encouraged.

References

- Bain, M., Huh, J., Han, T., & Zisserman, A. (2023). *WhisperX: Time-accurate speech transcription of long-form audio*. 4489–4493. <https://doi.org/10.21437/Interspeech.2023-78>
- Barbieri, F., Anke, L. E., & Camacho-Collados, J. (2022). *XLM-T: Multilingual Language Models in Twitter for Sentiment Analysis and Beyond*. arXiv. <https://doi.org/10.48550/arXiv.2104.12250>
- Bird, S., Edward, L., & Ewan, K. (2009). *Natural Language Processing with Python*. <https://tjzhifei.github.io/resources/NLTK.pdf>
- Boersma, P., & Weenink, D. (2025). *Praat: Doing phonetics by computer [Computer program]* (Version v6.4.47). <https://praat.org>
- Di Tommaso, P., Chatzou, M., Floden, E. W., Barja, P. P., Palumbo, E., & Notredame, C. (2017). Nextflow enables reproducible computational workflows. *Nature Biotechnology*, 35(4), 316–319. <https://doi.org/10.1038/nbt.3820>
- Eyben, F., Scherer, K. R., Schuller, B. W., Sundberg, J., André, E., Busso, C., Devillers, L. Y., Epps, J., Laukka, P., Narayanan, S. S., & Truong, K. P. (2016). The geneva minimalistic acoustic parameter set (GeMAPS) for voice research and affective computing. *IEEE Transactions on Affective Computing*, 7(2), 190–202. <https://doi.org/10.1109/TAFFC.2015.2457417>
- Eyben, F., Wöllmer, M., & Schuller, B. (2010). Opensmile: The munich versatile and fast open-source audio feature extractor. *Proceedings of the 18th ACM International Conference on Multimedia*, 1459–1462. <https://doi.org/10.1145/1873951.1874246>
- Hansen, L., Olsen, L. R., & Enevoldsen, K. (2023). TextDescriptives: A python package for calculating a large variety of metrics from text. *Journal of Open Source Software*, 8(84), 5153. <https://doi.org/10.21105/joss.05153>
- Jadoul, Y., Thompson, B., & Boer, B. de. (2018). Introducing Parselmouth: A Python interface to Praat. *Journal of Phonetics*, 71, 1–15. <https://doi.org/10.1016/j.jwocn.2018.07.001>
- Jong, N. H. de, Pacilly, J., & Heeren, W. (2021). PRAAT scripts to measure speed fluency and breakdown fluency in speech automatically. *Assessment in Education: Principles, Policy & Practice*, 28(4), 456–476. <https://doi.org/10.1080/0969594X.2021.1951162>
- Kurtzer, G. M., cclerget, Bauer, M., Kaneshiro, I., Trudgian, D., & Godlove, D. (2021). *Hpcng/singularity: Singularity 3.7.3*. Zenodo. <https://doi.org/10.5281/zenodo.4667718>
- Montani, I., Honnibal, M., Boyd, A., Landeghem, S. V., & Peters, H. (2023). *Explosion/spaCy: v3.7.2: Fixes for APIs and requirements*. Zenodo. <https://doi.org/10.5281/ZENODO.1212303>
- Pellerin, S., Houzé, B., Bedetti, C., Phillips, N., & Brambati, S. M. (2025). Connected speech profiles in mild cognitive impairment reflect global cognition. *Journal of Alzheimer's Disease*, 105(2), 582–597. <https://doi.org/10.1177/13872877251329615>
- Radford, A., Kim, J. W., Xu, T., Brockman, G., McLeavey, C., & Sutskever, I. (2022). *Robust Speech Recognition via Large-Scale Weak Supervision*. arXiv. <https://doi.org/10.48550/arXiv.2212.04356>
- Raffel, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., & Liu, P. J. (2023). *Exploring the Limits of Transfer Learning with a Unified Text-to-Text*

128 *Transformer*. arXiv. <https://doi.org/10.48550/arXiv.1910.10683>

129 Slegers, A., Chafouleas, G., Montembeault, M., Bedetti, C., Welch, A. E., Rabinovici, G.
130 D., Langlais, P., Gorno-Tempini, M. L., & Brambati, S. M. (2021). Connected speech
131 markers of amyloid burden in primary progressive aphasia. *Cortex*, 145, 160–168. [https:](https://doi.org/10.1016/j.cortex.2021.09.010)
132 [//doi.org/10.1016/j.cortex.2021.09.010](https://doi.org/10.1016/j.cortex.2021.09.010)

DRAFT