

Sentinel: Signal Validation and Anomaly Detection for Enterprise Log Data

José Manuel Vergara Álvarez¹, Nicolás Laverde Manotas¹, Juan Pablo Aguilar Calle¹, Jeisson Vicente Niño Castillo¹, Julián David Muñoz Pertuz¹, Daniel Monsalve Muñoz¹, and Sebastián Osorio Agudelo¹

¹ Bancolombia, Medellín, Colombia

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [↗](#)

Submitted: 27 March 2026

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/)).

Summary

Modern enterprise systems generate large volumes of operational logs that contain information about system behavior, performance, and failure conditions. These logs are frequently used for troubleshooting, monitoring, and anomaly detection. However, raw logs are typically unstructured and heterogeneous, making it difficult to determine whether they contain meaningful signals for analytical methods. Applying anomaly detection algorithms without first validating the quality and informational content of the data can lead to misleading results or unnecessary computational costs.

Sentinel is a Python library designed to address this challenge by introducing **early signal validation for enterprise logs prior to anomaly detection**. The library provides a modular pipeline that transforms unstructured logs into structured data, evaluates the presence of meaningful statistical signals, and applies anomaly detection algorithms only when the data meets predefined quality criteria. This fail-fast validation approach helps practitioners determine whether log datasets are suitable for anomaly analysis before investing effort in complex modeling pipelines.

The library supports common enterprise log formats and provides tools for ingestion, transformation, signal validation, anomaly detection, visualization, and streaming simulation. Sentinel is implemented in Python and builds upon widely used scientific computing libraries including pandas (McKinney, 2010) and scikit-learn (Pedregosa et al., 2011), with optional deep learning detectors implemented using PyTorch (Paszke et al., 2019).

Statement of Need

Enterprise infrastructures such as application servers, message brokers, security modules, and networking systems continuously produce operational logs. These logs often contain indicators of system failures, abnormal activity, or performance degradation. Detecting anomalies in log data is therefore an important task in operational monitoring, cybersecurity, and reliability engineering.

Despite the extensive literature on anomaly detection methods (Chandola et al., 2009), a practical challenge often arises in real-world environments: **not all log data contains analyzable signals**. Logs may contain excessive missing values, constant fields, inconsistent timestamps, or insufficient variability to support meaningful detection algorithms. In such situations, applying machine learning models can produce unstable or misleading outputs.

Most anomaly detection libraries focus on the modeling stage and assume that the input data already contains useful patterns. Libraries such as PyOD (Zhao et al., 2019), ADTK (Arundo Analytics, 2019), pySAD (Yilmaz & Kozat, 2021), TODS (Lai et al., 2021), and

40 Anomalib (Akçay et al., 2022) provide extensive implementations of detection algorithms, but
41 they generally rely on pre-processed, signal-rich datasets.

42 Sentinel addresses the earlier stage of the workflow: **determining whether a log dataset is**
43 **suitable for anomaly detection in the first place.** The library evaluates signal quality metrics
44 before applying detectors, enabling practitioners to avoid unnecessary modeling steps when the
45 underlying data lacks informative structure. This approach is particularly valuable for enterprise
46 logs, where large datasets may still contain limited analytical value.

47 State of the Field

48 The ecosystem of anomaly detection software has expanded significantly in recent years. PyOD
49 (Zhao et al., 2019) provides a unified interface for a wide range of outlier detection algorithms.
50 ADTK (Arundo Analytics, 2019) focuses on time-series anomaly detection using rule-based
51 and statistical techniques. pySAD (Yilmaz & Kozat, 2021) addresses streaming anomaly
52 detection scenarios, while TODS (Lai et al., 2021) introduces automated machine learning
53 pipelines for time-series outlier detection. Anomalib (Akçay et al., 2022) provides deep learning
54 architectures primarily targeting computer vision applications.

55 In the context of log analysis, research has also explored specialized anomaly detection
56 techniques such as DeepLog, which applies recurrent neural networks to learn sequential
57 patterns in system logs (Du et al., 2017). Public datasets such as LogHub provide standardized
58 benchmarks for evaluating log anomaly detection approaches (He et al., 2020).

59 These tools and datasets have significantly advanced anomaly detection research. However,
60 they typically assume that the dataset has already undergone preprocessing and validation. In
61 many operational environments, the primary difficulty lies in determining whether the available
62 logs contain sufficient information to justify anomaly detection.

63 Sentinel complements existing tools by introducing a **signal validation stage before detection.**
64 Rather than replacing existing anomaly detection libraries, Sentinel can be used alongside them
65 to assess data suitability and prepare structured log datasets for downstream analysis.

66 Software Design

67 Sentinel implements a modular architecture designed to support the complete lifecycle of
68 enterprise log analysis. The system is organized into six core modules that can be used
69 independently or combined into a unified pipeline.




Figure 1: Sentinel modular architecture for log analysis and anomaly detection.

Module	Purpose	Techniques and Components
Ingestion	Convert raw log files into structured tabular data	Parsers for WAS, HSM, HDC, IBMMQ and ZTNA logs; extensible parser interface
Transformer	Prepare logs for analysis	Time aggregation, rolling statistics, feature generation
Explorer	Validate signal presence before anomaly detection	IQR anomaly checks, variance thresholds, label validation, correlation analysis
Detectors	Apply anomaly detection algorithms	Isolation Forest (Liu et al., 2008), Robust Random Cut Forest (Guha et al., 2016), LSTM Autoencoder, Liquid Neural Networks (Hasani et al., 2022)
Visualization	Interpret anomaly detection outputs	Time-series anomaly visualization and SHAP-based explainability (?)
Simulation	Test streaming anomaly detection workflows	Real-time chunk processing and adaptive thresholds

70 The **Ingestion module** converts unstructured log files into structured tabular representations
71 using pandas. Sentinel includes built-in parsers for enterprise systems such as WebSphere

72 Application Server (WAS), Hardware Security Modules (HSM), High-Density Computing
73 (HDC), IBM Message Queue (IBMMQ), and ZTNA logs. A base parser interface allows users
74 to implement custom parsers for additional log formats.

75 The **Transformer module** prepares structured logs for analysis by aggregating events into time
76 windows and computing rolling statistics. These transformations allow event-based logs to be
77 converted into feature matrices suitable for statistical or machine learning methods.

78 The **Explorer module** is the central innovation of the library. It evaluates signal quality using
79 statistical metrics including interquartile range (IQR) anomaly checks, variance thresholds,
80 label presence validation, and correlation analysis. The module also evaluates simple predictive
81 models on individual features to estimate signal relevance. Based on these metrics, the module
82 provides a fail-fast decision indicating whether further analysis is justified.

83 The **Detectors module** provides a unified interface for multiple anomaly detection algorithms.
84 Implementations include Isolation Forest (Liu et al., 2008), Robust Random Cut Forest (Guha
85 et al., 2016), an LSTM-based autoencoder implemented with PyTorch (Paszke et al., 2019),
86 and Liquid Neural Networks (Hasani et al., 2022). The unified interface enables practitioners
87 to compare different algorithms within the same workflow.

88 The **Visualization module** provides tools for analyzing anomaly detection results. Anomaly
89 scores can be visualized as time-series plots with anomaly markers, and SHAP explanations
90 can be used to interpret feature contributions to anomaly predictions.

91 The **Simulation module** enables testing of anomaly detection pipelines in streaming scenarios.
92 Data can be processed in configurable chunks to emulate real-time ingestion pipelines, allowing
93 practitioners to evaluate detection strategies under operational conditions.

94 Research Impact Statement

95 Sentinel has demonstrated research and operational impact through its adoption in production
96 and research activities within the organization. The software has been used to assess the
97 analytical potential of operational data sources for service availability and anomaly detection
98 studies, helping teams identify log datasets that contain sufficient statistical signal before
99 investing effort in advanced analytics.

100 To date, Sentinel has been applied to seven operational data sources, including Cortex XDR logs
101 from three endpoint environments, as well as logs from HSM, HDC, WebSphere Application
102 Server (WAS), and IBM MQ platforms. These evaluations have supported the selection of
103 data sources suitable for observability and anomaly detection research.

104 The software has been operational in production for approximately one year and currently
105 supports analytical activities associated with five services. During this period, Sentinel has
106 enabled the validation of log quality and analytical readiness in real-world environments,
107 providing practical evidence of the applicability of its signal-validation approach.

108 Sentinel has been utilized by three different teams across research and operational contexts,
109 supporting investigations related to observability, service availability, and anomaly detection.
110 By standardizing preprocessing, signal validation, and detector integration workflows, the
111 software has facilitated reproducible experimentation while simultaneously supporting production
112 monitoring activities.

113 The combination of sustained production use, adoption across multiple teams, and application
114 to diverse enterprise log sources demonstrates Sentinel's value as both a research-enabling
115 platform and a practical operational analytics tool.

116 AI Usage Disclosure

117 Generative AI tools were used to assist with code scaffolding, documentation drafting, and
118 language refinement during the development of this software and manuscript. All generated
119 content was reviewed, validated, and modified as necessary by the authors, who assume full
120 responsibility for the accuracy and integrity of the work.

121 Acknowledgements

122 The authors thank the Technology Innovation team (ARQUITECTURA INNOVACION TI) at
123 Bancolombia for supporting the development of Sentinel and for providing feedback during the
124 design of the library.

125 References

- 126 Akcay, S., Ameln, D., Vaidya, A., Lakshmanan, B., Ahuja, N., & Genc, U. (2022). Anomalib:
127 A deep learning library for anomaly detection. *2022 IEEE International Conference on*
128 *Image Processing (ICIP)*, 1706–1710. <https://doi.org/10.1109/ICIP46576.2022.9897283>
- 129 Arundo Analytics. (2019). *ADTK: Anomaly detection toolkit*. GitHub. <https://github.com/arundo/adtk>
- 131 Chandola, V., Banerjee, A., & Kumar, V. (2009). Anomaly detection: A survey. *ACM*
132 *Computing Surveys*, 41(3), 1–58. <https://doi.org/10.1145/1541880.1541882>
- 133 Du, M., Li, F., Zheng, G., & Srikumar, V. (2017). DeepLog: Anomaly detection and diagnosis
134 from system logs through deep learning. *Proceedings of the 2017 ACM SIGSAC Conference*
135 *on Computer and Communications Security*, 1285–1298. <https://doi.org/10.1145/3133956.3134015>
- 137 Guha, S., Mishra, N., Roy, G., & Schrijvers, O. (2016). Robust random cut forest based
138 anomaly detection on streams. *Proceedings of the 33rd International Conference on Machine*
139 *Learning (ICML 2016)*, 2712–2721. <https://proceedings.mlr.press/v48/guha16.html>
- 140 Hasani, R., Lechner, M., Amini, A., Liebenwein, L., Ray, A., Tschaikowski, M., Teschl, G., &
141 Rus, D. (2022). Closed-form continuous-time neural networks. *Nature Machine Intelligence*,
142 4, 992–1003. <https://doi.org/10.1038/s42256-022-00556-7>
- 143 He, P., Zhu, J., He, S., Li, J., & Lyu, M. (2020). LogHub: A large collection of system log
144 datasets for AI-driven log analytics. *Proceedings of the IEEE International Symposium on*
145 *Software Reliability Engineering*, 1–11. <https://doi.org/10.1109/issre59848.2023.00071>
- 146 Lai, K.-H., Zha, D., Wang, G., Xu, J., Zhao, Y., Kumar, D., Chen, Y., Zumkhawaka, P., Wan,
147 M., Martinez, D., & Hu, X. (2021). TODS: An automated time series outlier detection
148 system. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35, 16060–16062.
149 <https://doi.org/10.1609/aaai.v35i18.18012>
- 150 Liu, F. T., Ting, K. M., & Zhou, Z.-H. (2008). Isolation forest. *2008 Eighth IEEE International*
151 *Conference on Data Mining*, 413–422. <https://doi.org/10.1109/ICDM.2008.17>
- 152 McKinney, W. (2010). Data structures for statistical computing in Python. *Proceedings of the*
153 *9th Python in Science Conference (SciPy 2010)*, 56–61. <https://doi.org/10.25080/Majora-92bf1922-00a>
- 155 Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z.,
156 Gimelshein, N., Antiga, L., Desmaison, A., Köpf, A., Yang, E., DeVito, Z., Raison, M.,
157 Tejani, A., Chilamkurthy, S., Steiner, B., Fang, L., ... Chintala, S. (2019). PyTorch: An
158 imperative style, high-performance deep learning library. *Advances in Neural Information*

- 159 *Processing Systems 32 (NeurIPS 2019)*, 8024–8035. [https://proceedings.neurips.cc/paper/](https://proceedings.neurips.cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html)
160 [2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html](https://proceedings.neurips.cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html)
- 161 Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel,
162 M., Prettenhofer, P., Weiss, R., Dubourg, V., Vanderplas, J., Passos, A., Cournapeau,
163 D., Brucher, M., Perrot, M., & Duchesnay, É. (2011). Scikit-learn: Machine learning in
164 Python. *Journal of Machine Learning Research*, 12, 2825–2830. [http://jmlr.org/papers/](http://jmlr.org/papers/v12/pedregosa11a.html)
165 [v12/pedregosa11a.html](http://jmlr.org/papers/v12/pedregosa11a.html)
- 166 Yilmaz, S. F., & Kozat, S. S. (2021). *pySAD: A streaming anomaly detection framework in*
167 *Python*. GitHub. <https://github.com/selimfirat/pysad>
- 168 Zhao, Y., Nasrullah, Z., & Li, Z. (2019). PyOD: A Python toolbox for scalable outlier detection.
169 *Journal of Machine Learning Research*, 20(96), 1–7. [http://jmlr.org/papers/v20/19-](http://jmlr.org/papers/v20/19-011.html)
170 [011.html](http://jmlr.org/papers/v20/19-011.html)

DRAFT