

The Irchel Geoparser: A Modular Python Library for Toponym Recognition and Resolution

Diego Gomes¹ and Ross S. Purves¹

¹ Department of Geography, University of Zurich, Switzerland  Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [James Gaboardi](#) 

Reviewers:

- [@SamHames](#)
- [@cforgaci](#)

Submitted: 28 November 2025

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).

Summary

The Irchel Geoparser is an open-source Python library for recognising and resolving place names (toponyms) in unstructured text. Based on a modular architecture, it provides an end-to-end geoparsing pipeline in which the components responsible for toponym recognition, toponym resolution, and the underlying gazetteer data can be freely adapted or exchanged. The library includes robust, configurable built-in modules and a GeoNames-based gazetteer configuration, offering strong performance for English text out of the box. Users can also choose to fine-tune the built-in modules or integrate their own custom recognisers and resolvers. Similarly, it is possible to include domain-specific gazetteers, allowing geoparsing workflows to be tailored to diverse research and practical contexts.

Statement of Need

Geoparsing is an essential step in linking text to geographical locations ([Purves et al., 2018](#)), enabling diverse applications including mapping news coverage, exploring historical documents, or analysing discourse across regions ([Chesnokova et al., 2019](#); [Grover et al., 2010](#); [Teitler et al., 2008](#)). Typically, geoparsing is a two-stage process involving firstly toponym recognition and subsequently toponym resolution. Toponyms, or place names, are first identified in text, a special case of the more general problem of named entity recognition. Toponym recognition must differentiate between instances of words being used in other contexts - for example to describe generic objects (as for instance in the case of bath (a place to wash oneself and an English city)) or to name people and firms (for example (George) Washington or Zurich (Insurance)). Toponym resolution resolves candidate place names to a unique identifier, dealing with the problem of multiple instances of the same place name (for example London, UK vs. London, Ontario, Canada).

The Irchel Geoparser was designed to provide a modular, out of the box functional Python library which could easily be incorporated into text processing chains using default settings. As such, it is well suited to teaching applications introducing for example graduate students to the geoparsing process. At the same time, because of its modular design, the entire processing chain is configurable to particular use cases - for example by swapping gazetteers to use those more appropriate to historical documents or retraining language models underlying toponym resolution modules on a specific corpus. Thus, the Irchel Geoparser is also well suited for use in research. Finally, we also aimed to provide a well documented Python library with an open licence allowing not only experimentation but also reproducibility of results.

State of the Field

Although several geoparsers already exist, they have limitations that restrict their practical applicability. Systems such as the Edinburgh Geoparser ([Grover et al., 2010](#)), CLAVIN ([Berico](#)

Technologies, 2012), GeoTxt (Karimzadeh et al., 2019), and the Mordecai family (Halterman, 2017, 2023) have contributed significantly to the field, but they generally require substantial effort for deployment and offer only limited flexibility in the customisation of the geoparsing process. Adapting these tools to a specific task, whether by incorporating new models for languages other than English, using specialised gazetteers or aligning the pipeline to a domain typically involves separate forks of the software, where this is possible, and these are often not well documented.

The Irchel Geoparser addresses these limitations by providing a clearly defined, interchangeable geoparsing pipeline packaged as an installable Python library. Its built-in modules offer strong default performance for English text and all stages of the pipeline can be adapted or exchanged entirely. This makes the Irchel Geoparser suitable both as a ready-to-use tool and as an experimental platform for researchers developing new geoparsing methods. By developing a modular architecture, we make it possible for other researchers to integrate customised modules addressing individual elements of the geoparsing workflow - for example recent work using LLMs to enhance toponym recognition (Yan & Lee, 2024) or resolution (Anuradha et al., 2025; Hu et al., 2024; Yan & Lee, 2024).

Software Design

The Irchel Geoparser treats geoparsing as a pipeline composed of recogniser and resolver modules that adhere to simple, well-defined interfaces. Any Python class implementing these interfaces can be used within the pipeline, allowing users to integrate custom methods while relying on the system for data handling, caching, and tracking. In this structure, documents and geoparsing results are managed within a local SQLite database, enabling persistent storage for long-term workflows, by maintaining a project structure in which raw text, recognised toponyms, resolver outputs and auxiliary metadata are stored together.

The built-in toponym recogniser uses spaCy's named-entity recognition functionality to identify toponyms based on entity labels including Geopolitical Entity (GPE), Location (LOC), or Facility (FAC) (Honnibal et al., 2020). Because it builds directly on spaCy models, users can easily switch languages, use fine-tuned models or incorporate custom pipelines by specifying a different model. The built-in toponym resolver incorporates a transformer-based disambiguation method presented in prior work (Gomes et al., 2024). It retrieves candidate locations from a gazetteer, embeds both the mention context and candidate descriptions using a transformer model, and ranks candidates based on embedding similarity. Users can choose to use the default model that has been fine-tuned using English news articles linked to GeoNames, or they can supply their own transformer checkpoints, which can then be trained for use with any configured gazetteer.

The gazetteer layer itself is also customisable. While a GeoNames-based configuration is provided by default, users can create gazetteer configuration files to build custom gazetteers, combining different data sources as required. This makes integrating gazetteer data created for, for example, historical georeferencing (Grossner et al., 2023) or by national mapping agencies, straightforward. Since modules only depend on the gazetteer interface, switching between different gazetteers requires no additional code changes.

Research Impact Statement

The Irchel Geoparser has been used in both research and teaching contexts. For example, it supported a study on geographic bias in global climate disaster reporting, in which thousands of news articles were analysed and mapped to reveal patterns in the geographic focus of media attention (Kong & Purves, 2025). Hanny et al. (2025) also used the Irchel Geoparser to analyse the locations of disasters described in social media content before carrying out a spatial-temporal analysis of emotions. Anderson & Giordano (forthcoming) integrated the

88 Irchel Geoparser in a forthcoming publication to explore the locations described in holocaust
89 testimonies. It was also used in a number of Master's theses, for instance to investigate how
90 newspaper discourses reflected the return of the wolf to Switzerland in the last two decades
91 (Besse, 2025) and to explore global discourses of disaster resilience (Haus, 2025).

92 Furthermore, the straightforward installation, comprehensive documentation, and a complete
93 demo available as a Docker image (Gomes, 2025) also make the Irchel Geoparser an accessible
94 tool for teaching and for illustrating key concepts in geographic information retrieval and we
95 have used it in multiple workshops and courses.

96 AI Usage Disclosure

97 The Python library described in this paper was developed over a period of about two years.
98 Generative AI was used to assist in software development, for example in brainstorming for
99 potential architectures, to automate and assist in the writing of functions and classes, to
100 create tests and to generate documentation. For brainstorming ChatGPT and Claude were the
101 primary models used, and for coding Claude. Due to the rapid development of these models
102 during the software development period, we do not give specific version numbers here. All
103 code and documentation was reviewed and validated by the authors, who made the core design
104 decisions. Claude was used in the initial drafting of parts of this manuscript. The final version
105 was reviewed and edited by the authors, who take full responsibility for its content.

106 Acknowledgements

107 The development of the Irchel Geoparser was supported by the GIS Hub, a project funded
108 by the Digital Society Initiative (DSI) of the University of Zurich. Additional funding was
109 provided by the Public Data Lab of the Digitalization Initiative of the Zurich Higher Education
110 Institutions (DIZH). We thank Michele Volpi for his involvement in developing the first version
111 of the Irchel Geoparser, and Nicolas Spring for his contributions as an open-source collaborator.

112 References

- 113 Anderson, C. J., & Giordano, A. (forthcoming). Natural Language Processing and the Spatial
114 Humanities: Comparing Geoparsing of Holocaust Testimonies With GPT-4.1, Deep Learning,
115 and Rules-Based NLP. *Forthcoming*. Forthcoming.
- 116 Anuradha, I., Sumanathilaka, D. K., Mitkov, R., & Rayson, P. (2025). Toponym Resolution:
117 Will Prompt Engineering Change Expectations? *Proceedings of the 15th International
118 Conference on Recent Advances in Natural Language Processing*, 95–104. [https://doi.org/
119 10.26615/978-954-452-098-4-011](https://doi.org/10.26615/978-954-452-098-4-011)
- 120 Berico Technologies. (2012). *CLAVIN (Cartographic Location And Vicinity INDEXer)*.
- 121 Besse, M. (2025). *The wolf debate in Switzerland: A spatial and temporal analysis performing
122 geographic information retrieval on newspaper articles* [Master's thesis]. University of
123 Zurich.
- 124 Chesnokova, O., Taylor, J. E., Gregory, I. N., & Purves, R. S. (2019). Hearing the silence:
125 finding the middle ground in the spatial humanities? Extracting and comparing perceived
126 silence and tranquillity in the English Lake District. *International Journal of Geographical
127 Information Science*, 33(12), 2430–2454. <https://doi.org/10.1080/13658816.2018.1552789>
- 128 Gomes, D. (2025). *Irchel Geoparser Demo*. <https://docs.geoparser.app/en/latest/demo.html>
- 129 Gomes, D., Purves, R. S., & Volpi, M. (2024). Fine-Tuning Transformers for Toponym
130 Resolution: A Contextual Embedding Approach to Candidate Ranking. In X. Hu, R.

- 131 S. Purves, L. Moncla, J. Kersten, & K. Stock (Eds.), *Proceedings of the GeoExT 2024:*
 132 *Geographic information extraction from texts workshop* (Vol. 3683, pp. 43–51). CEUR-WS.
 133 <https://ceur-ws.org/Vol-3683/paper6.pdf>
- 134 Grossner, K., Grunewald, S., & Mostern, R. (2023). Bringing places from the distant past to
 135 the present: a report on the World Historical Gazetteer. *International Journal on Digital*
 136 *Libraries*, 24(3), 159–162. <https://doi.org/10.1007/s00799-022-00341-2>
- 137 Grover, C., Tobin, R., Byrne, K., Woollard, M., Reid, J., Dunn, S., & Ball, J. (2010). Use of
 138 the Edinburgh geoparser for georeferencing digitized historical collections. *Philosophical*
 139 *Transactions of the Royal Society A: Mathematical, Physical and Engineering Sciences*,
 140 368(1925), 3875–3889. <https://doi.org/10.1098/rsta.2010.0149>
- 141 Halterman, A. (2017). Mordecai: Full Text Geoparsing and Event Geocoding. *Journal of Open*
 142 *Source Software*, 2(9), 91. <https://doi.org/10.21105/joss.00091>
- 143 Halterman, A. (2023). *Mordecai 3: A Neural Geoparser and Event Geocoder*. arXiv. <https://doi.org/10.48550/arXiv.2303.13675>
- 145 Hanny, D., Schmidt, S., & Resch, B. (2025). Bluesky as a social media data source for
 146 disaster management: investigating spatio-temporal, semantic and emotional patterns
 147 for floods and wildfires. *Journal of Computational Social Science*, 9(1), 16. <https://doi.org/10.1007/s42001-025-00448-x>
- 149 Haus, L. (2025). *Framing Disaster Resilience: A Media Discourse Analysis of British Newspapers*
 150 [Master's thesis]. University of Zurich.
- 151 Honnibal, M., Montani, I., Van Landeghem, S., & Boyd, A. (2020). *spaCy: Industrial-strength*
 152 *Natural Language Processing in Python*. Zenodo. <https://doi.org/10.5281/zenodo.1212303>
- 153 Hu, X., Kersten, J., Klan, F., & Farzana, S. M. (2024). Toponym resolution leveraging
 154 lightweight and open-source large language models and geo-knowledge. *International*
 155 *Journal of Geographical Information Science*, 1–28. <https://doi.org/10.1080/13658816.2024.2405182>
- 157 Karimzadeh, M., Pezanowski, S., MacEachren, A. M., & Wallgrün, J. O. (2019). GeoTxt: A
 158 scalable geoparsing system for unstructured text geolocation. *Transactions in GIS*, 23(1),
 159 118–136. <https://doi.org/10.1111/tgis.12510>
- 160 Kong, I., & Purves, R. S. (2025). Analyzing Geographic Bias of Newspaper Articles Reporting
 161 Global Climate Disasters. *Annals of the American Association of Geographers*, 1–19.
 162 <https://doi.org/10.1080/24694452.2025.2564220>
- 163 Purves, R. S., Clough, P., Jones, C. B., Hall, M. H., & Murdock, V. (2018). Geographic
 164 Information Retrieval: Progress and Challenges in Spatial Search of Text. *Foundations and*
 165 *Trends in Information Retrieval*, 12(2-3), 164–318. <https://doi.org/10.1561/15000000034>
- 166 Teitler, B. E., Lieberman, M. D., Panozzo, D., Sankaranarayanan, J., Samet, H., & Sperling,
 167 J. (2008). NewsStand: a new view on news. *Proceedings of the 16th ACM SIGSPATIAL*
 168 *International Conference on Advances in Geographic Information Systems*. <https://doi.org/10.1145/1463434.1463458>
- 170 Yan, Y., & Lee, J. (2024). GeoReasoner: Reasoning On Geospatially Grounded Context For
 171 Natural Language Understanding. *Proceedings of the 33rd ACM International Conference on*
 172 *Information and Knowledge Management*, 4163–4167. <https://doi.org/10.1145/3627673.3679934>