

AnyPINN: A Modular, Training-Agnostic Python Library for Physics-Informed Neural Networks

Giacomo Guidotto¹, Caterina Millevoi², and Damiano Pasetto¹

¹ Ca' Foscari University of Venice, Venice, IT ² Università di Padova, Padova, IT

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#)
- [Repository](#)
- [Archive](#)

Editor: [↗](#)

Submitted: 05 May 2026

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

AnyPINN is a PyTorch-native library for Physics-Informed Neural Networks (PINNs) as composable `nn.Module` objects. It supports both ODE and PDE problems, forward and inverse alike, with particular depth in the inverse setting where parameters must be recovered from partial observations. The library is organised into four layers (core, problems, lightning, catalog) with a strict one-way dependency direction, so users can engage with exactly the layer that matches their expertise, from running a pre-built example through the bootstrap CLI to defining custom physics and training loops in plain PyTorch. All hyperparameters are plain Python dataclasses: serializable, diffable, and programmatically constructible, which makes systematic ablation studies a matter of iterating over configuration objects rather than editing files. Training is not owned by the library: `Problem.training_loss()` returns a scalar tensor compatible with any optimizer or training infrastructure, including plain PyTorch and PyTorch Lightning.

Statement of Need

Physics-Informed Neural Networks encode differential equation residuals as loss terms, enabling mesh-free, automatic-differentiation-based solutions to both forward and inverse problems (Cuomo et al., 2022; Karniadakis et al., 2021; Lagaris et al., 1998; Raissi et al., 2019).

Despite a growing ecosystem of PINN libraries (F. Chen et al., 2021; Cuomo et al., 2022; Lu et al., 2021), two pain points persist. First, every major library couples physics definition to its own training loop, preventing researchers from using the full PyTorch potential. Second, inverse problems remain second-class: unknown parameters are configuration flags rather than typed objects, and ground-truth tracking during training requires user-written callbacks.

AnyPINN addresses both gaps. The core abstraction, `Problem`, is a standard `nn.Module` whose `training_loss()` returns a scalar tensor, compatible with any optimizer, scheduler, or distributed strategy without library-specific wrappers. For inverse problems specifically, `Parameter` is a first-class type with the same `forward(x)` interface whether it wraps a fixed scalar or a time-varying MLP; switching from fixed to learnable requires changing one configuration line. `ValidationRegistry` binds ground-truth references to parameter names and logs MSE automatically at every training step. `ArgsRegistry` unifies fixed and learnable arguments so that the differential equation callable cannot distinguish between the two, enabling transparent composability.

The target audience ranges from researchers defining novel physics to engineers running parameter sweeps on pre-built models. The bootstrap CLI (`anypinn create`) scaffolds a complete project from a single interactive command, making the library accessible to users who have never written a PINN before. Researchers who need full control can bypass the CLI and the Lightning layer entirely, working directly with `anypinn.core` and plain PyTorch.

State of the Field

The PINN software ecosystem spans more than ten libraries across two major frameworks (Cuomo et al., 2022; Karniadakis et al., 2021). Table 1 compares the libraries most relevant to AnyPINN's scope across several dimensions.

Table 1: Comparison of AnyPINN with existing PINN libraries. “Inverse (first-class)” means the library provides typed parameter objects, unified fixed/learnable interfaces, and built-in ground-truth tracking, not merely the ability to mark a variable as trainable.

Library	Framework	PDE support	Inverse (first-class)?	Scaffold CLI?	Engine-agnostic?
NeuroDiffEq (F. Chen et al., 2021)	PyTorch	2D	No	No	No
IDRLNet	PyTorch	Yes	No	No	No
NVIDIA Modulus (NVIDIA Corporation, 2023)	PyTorch	Yes (3D)	No	No	No
PINA (Coscia et al., 2024)	PyTorch	Yes	No	No	No
DeepXDE (Lu et al., 2021)	TF / JAX / PyTorch	Yes	No	No	No
TensorDiffEq	TensorFlow	Yes	No	No	No
SciANN	TensorFlow	Yes	No	No	No
PyDens	TensorFlow	Limited	No	No	No
Elvet	TensorFlow	Limited	No	No	No
NVIDIA SimNet	TensorFlow	Yes (3D)	No	No	No
AnyPINN	PyTorch	Yes	Yes	Yes	Yes

Several libraries in this comparison are no longer actively maintained: SciANN is explicitly deprecated, TensorDiffEq has seen no activity since 2022, NVIDIA SimNet has been superseded by Modulus (now PhysicsNeMo), and both PyDens and IDRLNet have had no meaningful development since 2023. This further narrows the set of viable options for researchers starting new PINN projects.

Three structural patterns in the existing ecosystem motivate AnyPINN as a separate library rather than a contribution to an existing one.

Framework distribution. More than half of the libraries in the comparison set are TensorFlow-based. PyTorch (Paszke et al., 2019) has become the dominant framework for ML research, and its downstream ecosystem (PyTorch Lightning (Falcon & The PyTorch Lightning team, 2019), Hugging Face, torch.compile, Triton) has grown substantially. DeepXDE has added PyTorch and JAX backends, but was originally designed around TensorFlow idioms. A researcher already working in PyTorch adopting a TensorFlow-native PINN library inherits framework friction unrelated to the science.

Training engine coupling. Every major competitor couples physics definition to its own training loop: NeuroDiffEq exposes a `solve_ivp`-style API; DeepXDE requires calling `Model.train()`; IDRLNet uses a proprietary computational graph; NVIDIA Modulus ships its own distributed Trainer; PINA wraps Lightning with an additional abstraction layer. When the training ecosystem evolves with new optimizers, new precision formats, new distributed strategies, users of these libraries must wait for library-level support or write wrappers around internal abstractions. AnyPINN's Problem is a plain `nn.Module`; the library never owns the training loop.

67 **Inverse problem primitives.** All listed libraries can, in principle, solve inverse problems by
 68 marking variables as trainable. However, none ship `Parameter` as a typed object with a unified
 69 interface, a `ValidationRegistry` for ground-truth tracking, or an `ArgsRegistry` that makes
 70 fixed and learnable arguments indistinguishable to the equation callable. In libraries such as
 71 `DeepXDE` and `NeuroDiffEq`, the distinction between fixed and learnable parameters is visible
 72 at the problem-definition level, and ground-truth comparison requires user-written callbacks.

73 Software Design

74 `AnyPINN` is split into four layers with strict one-way dependency (outer depends on inner,
 75 never the reverse).

76 **`anypinn.core: pure PyTorch.`** The mathematical foundation defines what a PINN problem
 77 *is* without any opinion about training. `Problem` (an `nn.Module`) aggregates `Constraint`
 78 instances, fields, and parameters and exposes `training_loss(batch, log)` returning a scalar
 79 tensor. `Field` is an MLP mapping input coordinates to state variables. `Parameter` inherits
 80 from `Argument` and exposes `forward(x)` whether it wraps a fixed scalar or a learnable MLP,
 81 making it transparent to the equation callable. `ArgsRegistry` is a typed dictionary of
 82 `Argument` instances; `InferredContext` extracts domain bounds from data and injects them
 83 into constraints automatically. Because `Problem` is a standard `nn.Module`, it participates
 84 transparently in the full PyTorch ecosystem: mixed precision, `torch.compile`, distributed data
 85 parallel, and gradient checkpointing require no library-level changes.

86 **`anypinn.problems: ODE and PDE building blocks.`** For ODE inverse problems,
 87 `ResidualsConstraint` enforces $\|\dot{y} - f(t, y)\|^2$ via autograd with support for arbitrary-order
 88 ODEs through chained derivatives. `ICConstraint` enforces $\|y(t_0) - y_0\|^2$ and natively
 89 handles higher-order initial conditions. `DataConstraint` enforces $\|y - y_{\text{obs}}\|^2$ against
 90 observed data. `ODEInverseProblem` composes all three with configurable loss weights and a
 91 selectable loss criterion (MSE, Huber, or L1). For PDE problems, `DirichletBCConstraint`
 92 and `NeumannBCConstraint` enforce boundary conditions on sampled boundary points, and
 93 `PDEResidualConstraint` enforces an arbitrary residual function over interior collocation points
 94 with field-subset scoping for coupled systems. The composable differential operators in
 95 `anypinn.lib.diff` (`grad`, `partial`, `mixed_partial`, `laplacian`, `divergence` and `hessian`)
 96 are available for use inside any constraint definition.

97 **`anypinn.lightning: optional training wrapper.`** A thin PyTorch Lightning (Falcon & The
 98 [PyTorch Lightning team, 2019](#)) layer for users who want batteries-included training. `PINNModule`
 99 wraps any `Problem` as a `LightningModule`; `PINNDataModule` manages data loading and
 100 collocation sampling. Five collocation strategies are available via a configuration string:
 101 "uniform", "random", "latin_hypercube", "log_uniform_1d", and "adaptive" (residual-
 102 weighted, with configurable exploration ratio). Callbacks provide SMMA-based early stopping,
 103 data scaling, and prediction writing. The entire Lightning layer is optional: `anypinn.core` has
 104 no Lightning dependency.

105 **`anypinn.catalog: drop-in problem modules.`** Pre-built ODE functions and data modules for
 106 SIR, SEIR, damped oscillator, and Lotka-Volterra systems. Catalog entries are runnable out of
 107 the box and serve as concrete starting points for the bootstrap CLI. The catalog is designed to
 108 be extended with additional dynamical systems as the library grows.

109 **Configuration as data.** All hyperparameters such as MLP architecture, optimizer settings,
 110 collocation strategy, loss weights, are expressed as keyword-only Python dataclasses. These
 111 objects are plain data: they can be serialized, loaded from configuration files, diffed across
 112 runs, and constructed programmatically in loops. This makes systematic ablation studies i.e.,
 113 sweeping over learning rates, collocation counts, or loss weights, a matter of iterating over
 114 configuration objects in a script rather than editing files by hand.

115 **Input encodings.** `FourierEncoding` and `RandomFourierFeatures` are available as `nn.Module`

116 input encodings, lifting low-frequency MLPs to high-frequency solutions without changes to
117 the training loop.

118 The four-step workflow for a custom inverse problem is:

```
# 1. Define the ODE
def my_ode(x: Tensor, y: Tensor, args: ArgsRegistry) -> Tensor:
    k = args["k"](x) # same interface for fixed or learnable k
    return -k * y

# 2. Configure hyperparameters
@dataclass(frozen=True, kw_only=True)
class MyHP(ODEHyperparameters):
    pde_weight: float = 1.0
    ic_weight: float = 10.0
    data_weight: float = 5.0

# 3. Build the problem
props = ODEProperties(ode=my_ode, args={"k": param}, y0=y0)
problem = ODEInverseProblem(ode_props=props, fields={"u": field},
                             params={"k": param}, hp=hp)

# 4. Train with Lightning or plain PyTorch
optimizer = torch.optim.Adam(problem.parameters(), lr=1e-3)
for batch in dataloader:
    optimizer.zero_grad()
    loss = problem.training_loss(batch, log=my_log_fn)
    loss.backward()
    optimizer.step()
```

119 The bootstrap CLI (anypinn create) generates a complete, immediately runnable four-file
120 project (ode.py, config.py, train.py, pyproject.toml) from a single interactive command,
121 lowering the barrier for users unfamiliar with PINN library APIs.

122 Research Impact

123 AnyPINN ships six working examples spanning epidemiology (SIR with real Italian COVID-
124 19 data, SEIR, reduced-SIR with hospitalization dynamics), mechanics (damped oscillator:
125 recovering damping ratio and natural frequency), and ecology (Lotka-Volterra predator-prey
126 dynamics), plus a minimal exponential decay script (~80 lines, core only, no Lightning). The
127 SIR example reproduces the PINN methodology of Millevoi et al. (2024), serving as a validation
128 benchmark: AnyPINN recovers comparable transmission rate dynamics and state variable
129 trajectories on the same Italian COVID-19 dataset. Each example demonstrates a different
130 dynamical system, scientific domain, and data type, with generated result plots and CSV exports.
131 Neural ODEs (R. T. Q. Chen et al., 2018) offer a complementary approach that replaces the
132 ODE right-hand side with a neural network entirely; AnyPINN instead keeps the equation
133 explicit, which makes the same Problem and Constraint abstractions applicable to both
134 forward solutions and parameter recovery. The PDE constraint layer (DirichletBCConstraint,
135 NeumannBCConstraint, PDEResidualConstraint) and the composable differential operators
136 in anypinn.lib.diff extend the library's applicability beyond ODE systems, providing a
137 foundation for spatiotemporal forward and inverse problems alike.

AI Usage Disclosure

Claude (Anthropic) was used during the development of the AnyPINN codebase to assist with code audit, bug identification, and targeted fixes. All architectural decisions, design choices, and major implementations were authored by the human developer. Claude was not used to generate the core library code or its scientific content.

References

- Chen, F., Sondak, D., Protopapas, P., Mattheakis, M., Liu, S., Agarwal, D., & Di Giovanni, M. (2021). NeuroDiffEq: A Python package for solving differential equations with neural networks. *Journal of Open Source Software*, 6(61), 3008. <https://doi.org/10.21105/joss.03008>
- Chen, R. T. Q., Rubanova, Y., Bettencourt, J., & Duvenaud, D. K. (2018). Neural ordinary differential equations. *Advances in Neural Information Processing Systems*, 31. <https://proceedings.neurips.cc/paper/2018/hash/69386f6bb1dfed68692a24c8686939b9-Abstract.html>
- Coscia, D., Meneghetti, L., Demo, N., Stabile, G., & Rozza, G. (2024). PINA: Physics-Informed Neural networks for Advanced scientific computing. GitHub repository. <https://github.com/mathLab/PINA>
- Cuomo, S., Di Cola, V. S., Giampaolo, F., Rozza, G., Raissi, M., & Piccialli, F. (2022). Scientific machine learning through physics-informed neural networks: Where we are and what's next. *Journal of Scientific Computing*, 92(3), 88. <https://doi.org/10.1007/s10915-022-01939-z>
- Falcon, W., & The PyTorch Lightning team. (2019). *PyTorch Lightning*. Zenodo. <https://doi.org/10.5281/zenodo.3828935>
- Karniadakis, G. E., Kevrekidis, I. G., Lu, L., Perdikaris, P., Wang, S., & Yang, L. (2021). Physics-informed machine learning. *Nature Reviews Physics*, 3(6), 422–440. <https://doi.org/10.1038/s42254-021-00314-5>
- Lagaris, I. E., Likas, A., & Fotiadis, D. I. (1998). Artificial neural networks for solving ordinary and partial differential equations. *IEEE Transactions on Neural Networks*, 9(5), 987–1000. <https://doi.org/10.1109/72.712178>
- Lu, L., Meng, X., Mao, Z., & Karniadakis, G. E. (2021). DeepXDE: A deep learning library for solving differential equations. *SIAM Review*, 63(1), 208–228. <https://doi.org/10.1137/19M1274067>
- Millevoi, C., Pasetto, D., & Ferronato, M. (2024). A physics-informed neural network approach for compartmental epidemiological models. *PLOS Computational Biology*, 20(9), 1–29. <https://doi.org/10.1371/journal.pcbi.1012387>
- NVIDIA Corporation. (2023). *NVIDIA Modulus: A physics-ML platform*. Online documentation. <https://docs.nvidia.com/deeplearning/modulus/index.html>
- Paszke, A., Gross, S., Massa, F., Lerer, A., Bradbury, J., Chanan, G., Killeen, T., Lin, Z., Gimelshein, N., Antiga, L., & others. (2019). PyTorch: An imperative style, high-performance deep learning library. *Advances in Neural Information Processing Systems*, 32. <https://proceedings.neurips.cc/paper/2019/hash/bdbca288fee7f92f2bfa9f7012727740-Abstract.html>
- Raissi, M., Perdikaris, P., & Karniadakis, G. E. (2019). Physics-informed neural networks: A deep learning framework for solving forward and inverse problems involving nonlinear partial differential equations. *Journal of Computational Physics*, 378, 686–707. <https://doi.org/10.1016/j.jcp.2018.10.045>