

FairSense-AgentiX: An Agentic Platform for Multi-Modal Fairness and AI Risk Analysis

Aravind Narayanan^{1*}, Karanpal Sekhon^{1*}, Mahshid Alinoori¹, and Shaina Raza^{1¶}

¹ Vector Institute for Artificial Intelligence, Toronto, Ontario, Canada  ¶ Corresponding author * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: 

Submitted: 08 May 2026

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/)).

Summary

FairSense-AgentiX is an open-source Python platform for automated detection of bias in text and images and for assessing AI deployment scenarios against established risk frameworks. Existing fairness tools such as AI Fairness 360 ([Bellamy et al., 2019](#)), Fairlearn ([Weerts et al., 2023](#)), and LangFair ([Bouchard et al., 2025](#)) address structured datasets or LLM output metrics but cannot audit unstructured content across text and images, explain their reasoning transparently, or connect findings to deployment-level risk frameworks in a single workflow. To address this gap, it implements a multi-node agentic workflow using LangGraph ([LangChain AI, 2024](#))—unlike static classifiers or rule-based checkers—in which a reasoning agent plans its analysis strategy, selects tools dynamically, executes them, and iteratively refines its output through a self-critique loop. The platform supports three analysis modes—text bias detection, image bias detection, and AI risk assessment—and exposes both a Python library and a FastAPI ([Ramírez, 2019](#)) REST/WebSocket server with an accompanying React web interface. FairSense-AgentiX is designed to be LLM-provider-agnostic and tool-swappable, making it applicable in a wide range of research and production contexts. It builds on and substantially extends a prior static pipeline, FairSense-AI ([Raza et al., 2025](#)), by replacing single-pass LLM inference with an orchestrated, self-correcting agentic loop and adding a retrieval-augmented risk assessment module tied to established governance frameworks.

Statement of Need

Three concrete gaps motivate FairSense-AgentiX:

Multi-modal coverage. Real-world content includes job advertisements with embedded images, presentations combining visual and textual claims, and deployment documentation. A text-only auditing tool misses a significant portion of potential bias surface area.

Transparent reasoning. Compliance-oriented use cases—hiring audits, content moderation pipelines, ML model governance—require not just a verdict but an explanation: which exact text spans are problematic, why, and how confident the system is. Black-box classification cannot satisfy this need.

Connection to risk frameworks. Detecting bias is necessary but not sufficient; practitioners also need to understand how a deployment scenario maps to established taxonomies such as the NIST AI Risk Management Framework ([National Institute of Standards and Technology, 2023](#)) and the MIT AI Risk Repository ([Slattery et al., 2026](#)). No prior open-source tool integrates bias detection with retrieval-augmented risk mapping.

FairSense-AgentiX addresses these gaps in a single, cohesive platform that is thoroughly tested

and documented for external adoption.

State of the Field

Fairness auditing of AI systems and content remains an active area of research that lacks mature, end-to-end tooling (Blodgett et al., 2020; Mehrabi et al., 2021). Existing libraries such as AI Fairness 360 (Bellamy et al., 2019) and Fairlearn (Weerts et al., 2023) provide statistical metrics and post-hoc debiasing algorithms for structured datasets but do not address unstructured text or images, and offer no pathway to natural-language explanations or deployment-level risk guidance. More recently, LangFair (Bouchard et al., 2025) introduced LLM-specific fairness metrics for text outputs, but remains restricted to evaluating LLM responses against predefined prompts and provides no image modality, no agentic reasoning, and no connection to deployment-level risk frameworks. Across these tools, existing approaches suffer from one or more of the following limitations: they are restricted to a single modality, they rely on fixed rule sets that cannot adapt to context, they treat analysis as a black-box step with no transparency into reasoning, or they provide no pathway from detected issues to actionable risk management guidance.

Why a new package rather than an extension. The gaps above are architectural rather than additive: AI Fairness 360 and Fairlearn are organised around tabular datasets and statistical metrics, with no abstraction for unstructured content, no orchestrator for multi-step reasoning, and no retrieval layer for external risk taxonomies. Adding these capabilities would require restructuring their core data model rather than extending it. LangFair is similarly scoped to LLM input/output pairs and prompt-level metrics, with no notion of auditing a document or image as a unit. FairSense-AgentiX is therefore best released as a separate package that interoperates with these tools (e.g., exporting structured findings for downstream Fairlearn analysis) rather than as a fork or extension of any one of them.

Software Design

Agentic Orchestration

The core of FairSense-AgentiX is an orchestrator built with LangGraph that implements a Reason–Act–Observe–Reflect (ReAct) loop (Yao et al., 2023). Figure 1 shows the full graph topology.




Figure 1: Orchestrator graph topology showing the seven-node ReAct loop with conditional routing between planning, execution, evaluation, refinement, and finalisation.

The request_plan node uses an LLM to classify the input modality and produce a routing decision with a confidence score and human-readable reasoning. preflight_eval checks whether the planned execution is coherent before committing resources. execute_workflow then dispatches to the selected specialised subgraph, and posthoc_eval scores the returned output on completeness, factual grounding, and span precision. When quality falls below threshold and refinement is enabled, decide_action routes back to apply_refinement, which adjusts the plan and re-enters the loop; otherwise control passes to finalize. This critique-driven refinement loop (Madaan et al., 2023) is the key architectural differentiator: it converts what would otherwise be a single-pass inference into an iterative, self-correcting process.

Specialised Workflow Subgraphs

Three workflow subgraphs handle distinct analysis tasks.

Text bias graph. Accepts raw text, invokes an LLM with structured output to identify bias instances (type, severity, character-level span, explanation), optionally summarises long documents, and produces colour-coded HTML highlighting. Bias categories covered include gender, age, racial, disability, and socioeconomic bias (Blodgett et al., 2020; Mehrabi et al., 2021).

Image bias graph. Runs OCR (Tesseract (Smith, 2007) or PaddleOCR (PaddlePaddle Authors, 2020)) and image captioning (BLIP (Li et al., 2022) or BLIP-2 (Li et al., 2023)) in parallel, merges the resulting text, and feeds it into the same LLM-based analysis pipeline as the text workflow. A Vision-Language Model (VLM) variant uses the full image directly rather than the extracted-text intermediary.

Risk assessment graph. Embeds a natural-language deployment scenario using a sentence-transformer model (Reimers & Gurevych, 2019), performs semantic nearest-neighbour search over a FAISS (Johnson et al., 2021) index of 1,340 risks drawn from the MIT AI Risk Repository v3 (Slattery et al., 2026), retrieves corresponding NIST AI-RMF (National Institute of Standards and Technology, 2023) mitigation recommendations from a second FAISS index per matched risk, and exports results as an HTML table and CSV file.

Tool Registry and Dependency Injection

All tools—OCR, captioning, summarisation, LLM, VLM, embedder, FAISS indexes, formatter, and persistence—are constructed by a central ToolRegistry from configuration and injected into each subgraph at build time. This design decouples graph logic from tool implementation and enables runtime substitution: the entire tool stack can be replaced with lightweight fakes for unit testing without modifying graph code.

LLM providers (OpenAI GPT-4 (OpenAI et al., 2024) and Anthropic Claude (Anthropic, 2024), via LangChain (Chase, 2022)) are selected through environment variables, with no code changes required to switch providers.

Service Layer

FairSense-AgentiX exposes a high-level FairSense Python class with three methods—analyze_text, analyze_image, and assess_risk—that return typed Pydantic result objects (BiasResult, RiskResult). These objects carry execution metadata including run ID, router confidence, refinement count, and quality scores, enabling fine-grained observability. The FastAPI service layer wraps this API and additionally streams agent telemetry events over WebSocket, so clients can observe each reasoning step as it occurs. A React/TypeScript frontend provides a landing page and interactive analysis interface with pre-built demo examples for each mode.

Quality Assurance

The package ships with a comprehensive test suite covering unit tests for every tool, integration tests for each subgraph and the orchestrator, stress tests, and a benchmarking harness that measures p50/p95/p99 latency and memory usage across input sizes. Pre-commit hooks enforce type-checking (mypy) and linting (ruff) before every commit, and GitHub Actions runs the integration test suite on every pull request.

Research Impact Statement

FairSense-AgentiX was developed under the AIXPERT project, a Horizon Europe initiative (Grant Agreement No. 101214389) tasked with building explainable, accountable, and

transparent AI systems. Within this project, FairSense-AgentiX serves as the reference implementation of the agentic fairness-and-risk analysis layer, and is the basis for our ongoing empirical work on critique-driven refinement in fairness auditing.

The platform's research utility derives from three concrete design choices. First, it operationalises the MIT AI Risk Repository (Slattery et al., 2026)—a peer-reviewed taxonomy of over 1,300 documented AI risks—as a queryable, FAISS-backed retrieval system, making the taxonomy directly usable in downstream studies rather than as a static document. Second, it bridges this risk taxonomy to the NIST AI Risk Management Framework (National Institute of Standards and Technology, 2023), providing per-risk mitigation recommendations that can be cited and traced in audit reports. Third, the LangGraph-based orchestrator exposes every intermediate reasoning step (router decisions, tool selections, preflight and post-hoc evaluation scores, refinement-loop deltas), enabling reproducible analysis of agent behaviour—a prerequisite for empirical studies of agentic systems that black-box pipelines cannot support.

The package is distributed on PyPI as `fairsense-agentix` under the MIT License and is actively maintained on GitHub with continuous integration and versioned releases. Developed under the AIXPERT project, FairSense-AgentiX provides the research community with the first open-source, end-to-end auditing harness that combines multi-modal bias detection with retrieval-augmented risk mapping—infrastructure we expect to underpin forthcoming empirical work on agentic fairness auditing.

Limitations

At present, FairSense-AgentiX supports five bias categories (gender, age, racial, disability, socioeconomic), operates over English-language prompts and risk taxonomies, and integrates only with closed-source LLM providers (OpenAI GPT-4 (OpenAI et al., 2024) and Anthropic). Quantitative evaluation against standard fairness benchmarks is not included in this work and is planned as a companion study.

AI Usage Disclosure

No generative AI was used for architectural design or core algorithms. Claude (Anthropic) and GitHub Copilot were used for creating documentation and docstrings, bug fixing, and inline code changes. For this manuscript, these tools were used for drafting, copy-editing, and consistency checking; all factual claims were verified by the authors against the codebase.

Acknowledgements

Resources used in preparing this research were provided, in part, by the Province of Ontario, the Government of Canada through CIFAR, and companies sponsoring the Vector Institute (<http://www.vectorinstitute.ai/#partners>). This research was funded by the European Union's Horizon Europe research and innovation programme under the AIXPERT project (Grant Agreement No. 101214389), which aims to develop an agentic, multi-layered, GenAI-powered framework for creating explainable, accountable, and transparent AI systems.

References

- Anthropic. (2024). *The Claude 3 model family: Opus, sonnet, haiku*. https://www-cdn.anthropic.com/de8ba9b01c9ab7cbabf5c33b80b7bbc618857627/Model_Card_Claude_3.pdf.
- Bellamy, R. K. E., Dey, K., Hind, M., Hoffman, S. C., Houde, S., Kannan, K., Lohia, P., Martino, J., Mehta, S., Mojsilović, A., Nagar, S., Ramamurthy, K. N., Richards, J., Saha,

- 165 D., Sattigeri, P., Singh, M., Varshney, K. R., & Zhang, Y. (2019). AI fairness 360: An
166 extensible toolkit for detecting and mitigating algorithmic bias. *IBM Journal of Research
167 and Development*, 63(4/5), 4:1–4:15. <https://doi.org/10.1147/JRD.2019.2942287>
- 168 Blodgett, S. L., Barocas, S., Daumé III, H., & Wallach, H. (2020). Language (technology) is
169 power: A critical survey of “bias” in NLP. *Proceedings of the 58th Annual Meeting of the
170 Association for Computational Linguistics (ACL)*, 5454–5476. [https://doi.org/10.18653/
171 v1/2020.acl-main.485](https://doi.org/10.18653/v1/2020.acl-main.485)
- 172 Bouchard, D., Chauhan, M. S., Skarbrevik, D., Bajaj, V., & Ahmad, Z. (2025). LangFair: A
173 python package for assessing bias and fairness in large language model use cases. *Journal
174 of Open Source Software*, 10(105), 7570. <https://doi.org/10.21105/joss.07570>
- 175 Chase, H. (2022). *LangChain*. <https://github.com/langchain-ai/langchain>.
- 176 Johnson, J., Douze, M., & Jégou, H. (2021). Billion-scale similarity search with GPUs.
177 *IEEE Transactions on Big Data*, 7(3), 535–547. [https://doi.org/10.1109/TBDATA.2019.
178 2921572](https://doi.org/10.1109/TBDATA.2019.2921572)
- 179 LangChain AI. (2024). *LangGraph: Build resilient language agents as graphs*. [https://github.
180 com/langchain-ai/langgraph](https://github.com/langchain-ai/langgraph).
- 181 Li, J., Li, D., Savarese, S., & Hoi, S. (2023). BLIP-2: Bootstrapping language-image pre-
182 training with frozen image encoders and large language models. *Proceedings of the
183 International Conference on Machine Learning (ICML)*, 19730–19742.
- 184 Li, J., Li, D., Xiong, C., & Hoi, S. (2022). BLIP: Bootstrapping language-image pre-training
185 for unified vision-language understanding and generation. *Proceedings of the International
186 Conference on Machine Learning (ICML)*, 12888–12900.
- 187 Madaan, A., Tandon, N., Gupta, P., Hallinan, S., Gao, L., Wiegrefe, S., Alon, U., Dziri,
188 N., Prabhume, S., Yang, Y., Gupta, S., Majumder, B. P., Hermann, K., Welleck,
189 S., Yazdanbakhsh, A., & Clark, P. (2023). Self-refine: Iterative refinement with self-
190 feedback. *Advances in Neural Information Processing Systems (NeurIPS)*, 36. [https:
191 //openreview.net/forum?id=S37hOerQLB](https://openreview.net/forum?id=S37hOerQLB)
- 192 Mehrabi, N., Morstatter, F., Saxena, N., Lerman, K., & Galstyan, A. (2021). A survey on
193 bias and fairness in machine learning. *ACM Computing Surveys*, 54(6), 1–35. [https:
194 //doi.org/10.1145/3457607](https://doi.org/10.1145/3457607)
- 195 National Institute of Standards and Technology. (2023). *AI Risk Management Framework (AI
196 RMF 1.0)* (NIST AI 100-1). U.S. Department of Commerce. [https://doi.org/10.6028/
197 NIST.AI.100-1](https://doi.org/10.6028/NIST.AI.100-1)
- 198 OpenAI, Achiam, J., Adler, S., Agarwal, S., Ahmad, L., Akkaya, I., Aleman, F. L., Almeida,
199 D., Altenschmidt, J., Altman, S., Anadkat, S., Avila, R., Babuschkin, I., Balaji, S., Balcom,
200 V., Baltescu, P., Bao, H., Bavarian, M., Belgum, J., ... Zoph, B. (2024). *GPT-4 technical
201 report*. <https://arxiv.org/abs/2303.08774>
- 202 PaddlePaddle Authors. (2020). *PaddleOCR: Awesome multilingual OCR toolkits based on
203 PaddlePaddle*. <https://github.com/PaddlePaddle/PaddleOCR>.
- 204 Ramírez, S. (2019). *FastAPI*. <https://fastapi.tiangolo.com>.
- 205 Raza, S., Chettiar, M. S., Yousefabad, M., Khan, T., & Lotif, M. (2025). *FairSense-AI:
206 Responsible AI meets sustainability*. <https://arxiv.org/abs/2503.02865>
- 207 Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence embeddings using siamese
208 BERT-networks. *Proceedings of the 2019 Conference on Empirical Methods in Natural
209 Language Processing (EMNLP)*, 3982–3992. <https://doi.org/10.18653/v1/D19-1410>
- 210 Slattery, P., Saeri, A. K., Grundy, E. A. C., Graham, J., Noetel, M., Uuk, R., Dao, J.,
211 Pour, S., Casper, S., & Thompson, N. (2026). The AI risk repository: A meta-review,

- 212 database, and taxonomy of risks from artificial intelligence. *Patterns*, 101517. <https://doi.org/10.1016/j.patter.2026.101517>
- 213
- 214 Smith, R. (2007). An overview of the Tesseract OCR engine. *Ninth International Conference*
215 *on Document Analysis and Recognition (ICDAR 2007)*, 2, 629–633. <https://doi.org/10.1109/ICDAR.2007.4376991>
- 216
- 217 Weerts, H., Dudík, M., Edgar, R., Jalali, A., Lutz, R., & Madaio, M. (2023). Fairlearn:
218 Assessing and improving fairness of AI systems. *Journal of Machine Learning Research*,
219 24(257), 1–8. <http://jmlr.org/papers/v24/23-0389.html>
- 220 Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct:
221 Synergizing reasoning and acting in language models. *International Conference on Learning*
222 *Representations (ICLR)*. https://openreview.net/forum?id=WE_vluYUL-X

DRAFT