

A Reproducible ML Pipeline for Ovarian Cancer Platinum-Response Prediction with Empirical Quantification of Feature-Selection Leakage

Aaron Michael Anuroop Chintala ¹

¹ School of Systems Biology, George Mason University, Manassas, VA, USA

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: 

Submitted: 22 June 2026

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/)).

Summary

ovarian-cancer-platinum-response is a modular, reproducible Python pipeline for predicting platinum chemotherapy response in ovarian cancer patients using bulk RNA-seq gene expression profiles from the TCGA Ovarian Serous Cystadenocarcinoma cohort ([The Cancer Genome Atlas Research Network, 2011](#)). The software trains and evaluates three supervised classifiers — Logistic Regression, Random Forest, and Support Vector Machine — across three survival thresholds (12, 18, and 24 months), and provides evaluation tools including cross-validated ROC-AUC with bootstrap confidence intervals, SHAP explainability ([Lundberg & Lee, 2017](#)), and Kaplan-Meier survival stratification ([Davidson-Pilon, 2019](#)).

Beyond the core prediction pipeline, the package includes a paired ablation study that empirically quantifies AUC inflation from a widely documented but poorly specified failure mode in genomic machine learning: computing feature selection statistics on the full dataset before the train/test split. The key contribution is the distinction between unsupervised variance-based filtering (mean inflation: +0.01 AUC, non-significant across 30 repeated splits) and supervised F-score-based selection (mean inflation: +0.22 AUC, up to +0.55 for SVM, significant at $p < 0.0001$ in 7 of 9 model/threshold combinations). A label construct-validity investigation further examines whether the OS-threshold proxy labels standard in this literature capture the clinical construct they are assumed to represent.

Statement of Need

Machine learning applied to cancer genomics regularly reports optimistic performance, a problem attributed in part to feature selection leakage ([Teschendorff, 2020](#); [Vabalas et al., 2019](#)). However, published guidance rarely distinguishes between supervised selection methods — which directly use outcome labels, creating a direct information pathway from test fold to training — and unsupervised methods, which use only expression variance and create only an indirect pathway. This distinction matters empirically: our ablation shows the two conditions produce dramatically different inflation magnitudes (mean +0.22 vs. +0.01) on identical data and models. Existing genomic ML tools ([Davis & Meltzer, 2007](#); [Zararsiz et al., 2018](#)) do not include reproducible ablation frameworks for quantifying this effect with paired statistical testing across repeated splits.

A second unmet need is tooling to investigate label construct validity. TCGA-OV lacks a direct platinum-response label; researchers use OS-threshold surrogates without empirically verifying how well they capture the intended clinical construct. The software's `label_investigation` module provides a reproducible framework for this check, applicable to any TCGA cohort with partial clinical response data alongside survival annotations.

Architecture

The software is organised into eight source modules under `src/`: `data_loader` (I/O), preprocessing (survival extraction, platinum-patient identification, censoring-aware label assignment), features (train-only variance selection, z-score scaling), models (classifier factories, CV utilities), evaluation (metrics, results tables), visualization (seven figure types), ablation (leakage conditions, bootstrap CIs, nested CV, paired significance tests), and `label_investigation` (coverage, distribution, consistency, cross-tabulation, exact binomial recall CIs).

All parameters are controlled via `config/config.yaml` (gitignored; a committed template `config/config.template.yaml` is provided). Three CLI entry points with `--help` (`run.py`, `run_ablation_study.py`, `investigate_clinical_labels.py`) enable full pipeline reproducibility from a single command.

Testing

The package includes 52 unit tests across three modules. All tests use synthetic data and pass without the TCGA dataset on disk, enabling CI/CD integration. Three tests are regression tests for bugs that only surfaced on real TCGA data (NaN values in RSEM matrices, mixed-type gene name indices) but were not caught by earlier synthetic-data tests — documenting a real gap between synthetic and real-data coverage.

Results

On the platinum-treated TCGA-OV cohort ($n=214/206/199$ at 12/18/24m), Logistic Regression achieved the highest test-set AUC at the 24-month endpoint (mean 0.70 across 30 repeated correct splits). The ablation study confirmed that unsupervised variance selection introduces negligible leakage while supervised F-score selection inflates AUC by a mean of +0.22, with SVM reaching AUC 1.00 under leaky conditions versus 0.45 under correct conditions at the 12-month endpoint. The label investigation found that the OS-threshold proxy had 0/5 recall (95% CI: [0.000, 0.522]) on the 7 clinically confirmed Progressive Disease patients, consistent with the mechanistic argument that overall survival and first-line treatment response are genuinely different clinical constructs.

Acknowledgements

The author thanks Dr. Christopher Lockhart (George Mason University) for advising the course project from which this software was developed. TCGA-OV data were obtained from the cBioPortal for Cancer Genomics (Cerami et al., 2012; Gao et al., 2013).

References

- Cerami, E., Gao, J., Dogrusoz, U., Gross, B. E., Sumer, S. O., Aksoy, B. A., Jacobsen, A., Byrne, C. J., Heuer, M. L., Larsson, E., & others. (2012). The cBio cancer genomics portal: An open platform for exploring multidimensional cancer genomics data. *Cancer Discovery*, 2(5), 401–404. <https://doi.org/10.1158/2159-8290.CD-12-0095>
- Davidson-Pilon, C. (2019). Lifelines: Survival analysis in Python. *Journal of Open Source Software*, 4(40), 1317. <https://doi.org/10.21105/joss.01317>
- Davis, S., & Meltzer, P. S. (2007). GEOquery: A bridge between the gene expression omnibus (GEO) and BioConductor. *Bioinformatics*, 23(14), 1846–1847. <https://doi.org/10.1093/bioinformatics/btm177>

81 [bioinformatics/btm254](https://doi.org/10.1126/scisignal.2004088)

82 Gao, J., Aksoy, B. A., Dogrusoz, U., Dresdner, G., Gross, B., Sumer, S. O., Sun, Y.,
83 Jacobsen, A., Sinha, R., Larsson, E., & others. (2013). Integrative analysis of complex
84 cancer genomics and clinical profiles using the cBioPortal. *Science Signaling*, 6(269), p1.
85 <https://doi.org/10.1126/scisignal.2004088>

86 Lundberg, S. M., & Lee, S.-I. (2017). A unified approach to interpreting model predictions.
87 *Advances in Neural Information Processing Systems*, 30.

88 Teschendorff, A. E. (2020). A comparison of epigenetic mitotic-like clocks for cancer risk
89 prediction. *Genome Medicine*, 12(1), 56. <https://doi.org/10.1186/s13073-020-00752-3>

90 The Cancer Genome Atlas Research Network. (2011). Integrated genomic analyses of ovarian
91 carcinoma. *Nature*, 474(7353), 609–615. <https://doi.org/10.1038/nature10166>

92 Vabalas, A., Gowen, E., Poliakoff, E., & Casson, A. J. (2019). Machine learning algorithm
93 validation with unequal sample sizes. *PLOS ONE*, 14(7), e0224365. <https://doi.org/10.1371/journal.pone.0224365>

95 Zararsiz, G., Goksuluk, D., Korkmaz, S., Eldem, V., Duru, I. P., Ozcelik, F., Ozturk, A., &
96 Karaagaoglu, A. E. (2018). MLSeq: Machine learning interface for RNA-sequencing data.
97 *Computer Methods and Programs in Biomedicine*, 155, 207–214. <https://doi.org/10.1016/j.cmpb.2017.12.002>
98

DRAFT