

HDH: A Python Library for Hypergraph-Based Distributed Quantum Computing Partitioning

Maria Gragera Garces ¹

¹ University of Edinburgh

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Owen Lockwood](#) 

Reviewers:

- [@MIWdlB](#)
- [@oblivateandsurrender](#)

Submitted: 17 December 2025

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](#)).

Summary

Today's quantum computers are limited by how many qubits a single device can hold. Distributed quantum computing (DQC) works around this by linking multiple smaller devices together so they can jointly run computations too large for any one of them alone, which first requires deciding how to split, or partition, a computation across those devices.

HDH (Hybrid Dependency Hypergraphs) is a Python library that gives researchers a common representation to develop and compare partitioning strategies against. It converts a quantum computation — expressed as a circuit, a measurement-based pattern, a quantum walk, or a quantum cellular automaton — into a hypergraph that captures every way the computation could be split across devices, including hard constraints such as per-device qubit limits that prior abstractions treat as soft penalties rather than requirements. Researchers can run their own partitioning heuristics directly on this representation, or use HDH's built-in capacity-aware baseline, and compare results on a consistent, model-agnostic footing. HDH also imports circuits from popular quantum SDKs (Qiskit, Cirq, PennyLane, Amazon Braket) and exports back to Qiskit and PennyLane, so partitioned results can be turned back into runnable circuits.

Statement of need

HDHs (Hybrid Dependency Hypergraphs) are an abstraction which transforms quantum computation, originating from any quantum computational model (including circuits, measurement-based quantum computing, quantum cellular automata, and quantum walks), to a directed hypergraph that expresses all possible partitions available within the computation. They were originally proposed in ([Gragera Garces et al., 2025](#)) as a unifying approach to quantum distribution, extending the hypergraph abstraction method for partitioning across devices originally proposed in ([Andrés-Martínez & Heunen, 2019](#)). Various partitioning strategies have since been proposed building on that earlier abstraction ([Clark et al., 2023](#); [Escofet et al., 2023](#); [Sundaram & Gupta, 2023](#)), but many are tested on inconsistent hypergraph abstractions, hindering cross-partitioner comparison and improvement.

Having an easy to implement, open-source, and model-agnostic abstraction will enable the fair and consistent cross-comparison of partitioning strategies in future work. Furthermore, HDHs extend this capability beyond the circuit model, addressing a current blind spot in DQC research.

HDH is designed to be used by both distributed quantum architecture researchers and compiler developers building on existing frameworks who require a model-agnostic distribution layer.

State of the field

Existing DQC approaches abstract computations to hypergraphs which are then partitioned, using balanced hypergraph partitioning solvers such as KaHyPar (Schlag et al., 2023).

This framing has two fundamental limitations:

1. It reduces distribution to a balanced partitioning problem that ignores a hard physical constraint: individual QPUs have fixed qubit capacities, and a valid distribution must respect these limits strictly rather than treating them as soft penalties.
2. Existing hypergraph abstractions are model-specific and encode only a subset of possible partition cuts, meaning partitioning strategies are routinely evaluated on inconsistent abstractions, making cross-comparison unreliable and hindering the systematic development of improved heuristics.

While libraries for distributed quantum computing exist (DISQCO (Burt et al., 2026), Qdislib (Tejedor et al., 2025), Optyx (Kupper et al., 2025), DC-MBQC (Xue et al., 2026)), these implement end-to-end distribution pipelines rather than exposing the underlying abstraction as a research tool. No existing library provides a model-agnostic hypergraph abstraction designed specifically to enable the development and fair comparison of partitioning heuristics (the role HDH is built to fill).

Quantum compilation frameworks like Qiskit (Javadi-Abhari et al., 2024), Cirq (Developers, 2025), and PennyLane (Bergholm et al., 2018) provide circuit optimization and device mapping, but they do not offer model-agnostic abstractions for distributed quantum computing. The HDH library is compatible with these SDKs, making it a seamless addition to state-of-the-art quantum software stacks rather than a replacement for them.

Software design

The central design decision in HDH was to separate the abstraction layer from the partitioning layer. Rather than building a monolithic tool that both constructs hypergraphs and partitions them, HDH exposes the HDH as a first-class data structure that any downstream partitioner can consume. This makes the library useful both as a standalone research tool and as a substrate for third-party heuristics.

A hypergraph-based representation was chosen over simpler graph alternatives, as quantum computing models frequently involve operations with more than two inputs or outputs (a Toffoli gate, for instance, acts on three qubits simultaneously), requiring multi-way correlations.

Two further design choices trade added complexity for partitioning flexibility. First, HDHs represent each qubit's state at every timestep as a separate node rather than a single node per logical qubit, and partitioning operates over these timestepped nodes rather than whole qubits. This allows a single qubit's history to be split across devices at whichever point in the computation makes the split cheapest, with its state teleported between them at that boundary, rather than committing the qubit to one device for the circuit's full duration. Second, a multi-qubit gate is represented not as one hyperedge but three: one linking each qubit's pre-gate state to an intermediate state, one spanning all involved qubits at that intermediate point, and one linking each qubit's post-gate state onward. This staging is what allows a partitioner to cut through a multi-qubit gate at either boundary — modelling either a non-local gate executed over the network or a teleportation of one qubit's state immediately before or after the gate — instead of being limited to the coarser choice of which side of a single gate-edge to place a device boundary on.

The library includes a capacity-aware greedy heuristic as a built-in baseline. Existing DQC research typically benchmarks against KaHyPar (Schlag et al., 2023), a general-purpose hypergraph partitioner not designed for quantum hardware constraints. While simpler than

84 KaHyPar, the included heuristic treats per-device qubit capacity as a hard constraint, and
85 returns a feasible assignment whenever one exists rather than trading feasibility off against
86 balance. That makes it a more appropriate DQC baseline and a concrete starting point for
87 researchers developing improved strategies.

88 Finally, HDH is written in Python, the quantum software community's primary language, and is
89 compatible with Qiskit, Cirq and PennyLane from the outset.

90 **From computational models to HDHs**

91 HDHs use the following notation to describe quantum workload dependencies, including
92 predicted elements that represent potential future state transformations based on classical
93 measurement outcomes:

Symbol	Meaning
•	Classical node
●	Quantum node
—	Classical hyperedge
—	Quantum hyperedge
○	Predicted node
--	Predicted hyperedge

Figure 1: HDH symbol legend.

94 Mapping a quantum workload such as a circuit to an HDH involves applying specific
95 correspondences between model elements and hypergraph motifs. This library provides model-
96 specific classes such as the `Circuit` class that enable straightforward conversions to HDHs
97 using mapping tables; any class with a `build_hdh()` method satisfies the library's `Model`
98 protocol and is checked by the same conformance tests:

Motif	Operation
•	Qubit initialisation
•	Single qubit gate
•	Two qubit gate
•	Three qubit gate
•	Measurement

Figure 2: Circuit to HDH mapping table.

99 In the context of DQC, entangling operations in a model can be made non-local (namely
100 non-local gates) and thus partitioned through a quantum network via quantum communication
101 primitives (Wu et al., 2022). Alternatively, qubit states can be individually forwarded through
102 teleportation protocols (Zomorodi-Moghadam et al., 2017). Because HDHs represent both cut
103 types within a single structure, a partitioner is free to combine them within one workload —
104 e.g. keeping a qubit local to a device via non-local gates during one phase of a computation,
105 then teleporting it elsewhere for a later phase — rather than being restricted to whichever
106 single strategy the chosen abstraction happens to support.

107 Unlike prior abstractions, which represent only non-local gates or only teleportation, HDHs
108 capture both (Figure 3).

	Telegate	Teledata	HDH
Gate cut 			
Wire cut 			

Figure 3: Table showing HDH expressivity.

Figure 4 shows the HDH of a six-qubit circuit with a Toffoli gate, a classically conditioned gate and mid-circuit measurements, drawn as a graph. Each gate contributes the hyperedge of its state transformation plus the preceding and following hyperedges that allow pre- and post-teleportation, nodes are possible state transformations rather than qubits or operations, and classical data flow (orange) is part of the hypergraph.




Figure 4: Example circuit and its HDH representation.

Research impact statement

HDH's capacity-aware partitioner is proven optimal on 20 of 24 instances (83%) and averages $1.03\times$ the optimal cut cost (worst case $1.33\times$), measured against exhaustive branch-and-bound search over MQT Bench circuits (Quetschlich et al., 2023) at the tightest feasible capacity (three devices, network overhead 1). It also returns a feasible assignment whenever one exists, which the balance-based partitioners used as DQC baselines do not guarantee.

Carrying teledata and telegate cuts in one structure means the combined formulation automatically matches the best single-mode strategy, whichever that is for a given interaction pattern. Where a qubit's interaction pattern shifts partway through a computation, from one group of qubits to another, exhaustive search gives cut costs of 1, 2 and 3 for two, three and four shifts, against 3, 3 and 6 for the qubit-level formulation used by prior hypergraph

approaches ($1.5\text{--}3\times$ less), matching the teleportation-only formulation throughout. On 30 standard MQT Bench instances, where no such shift occurs, a greedy placer on the combined formulation ties the qubit-level one on every instance while choosing among $8.4\times$ as many possible placements. A simple greedy does not always reach the combined optimum, though: on 3 of those 30 instances, a greedy restricted to teleportation cuts did better.

What has no counterpart elsewhere is that the same partitioner runs unmodified over circuits, MBQC patterns, quantum walks, and quantum cellular automata, exercised across all four by the test suite. Comparing distribution overhead between computational models is not a question a circuit-only library can pose; here it is a matter of swapping the workload builder. A first such study is being developed by the authors.

Every quantitative claim in this section regenerates from benchmarks/.

Early community engagement has been encouraging. The project was presented as a poster at SIGCOMM 2025 (Gragera Garces et al., 2025) (a major networking venue), has received funding through the Unitary Fund microgrant program (dedicated to supporting open source quantum software to benefit humanity) and has already seen external contributors (acknowledged below). Further, we are in discussion with companies in the Distributed Quantum Computing space regarding the library's integration within their stack.

AI usage disclosure

Claude was used during both library development and paper writing.

In the library, Claude generated initial draft code implementations that were subsequently rewritten by the author. It also assisted in producing unit tests, which were validated against expected behaviour across both passing and failing scenarios. These were not always fully rewritten but they were revised and thoroughly tested. Additionally, AI was used to generate inline code comments throughout the library, with the aim of improving readability for contributors and users who may check the source code.

In the paper, Claude was used to assist with wording and polish.

All AI-generated content (both code and text) was reviewed or modified (as per the above descriptions) by the author before inclusion in the present version.

Acknowledgements

We acknowledge contributions from Joseph Tedds, Manuel Alejandro, and Alessandro Cosentino.

We thank Unitary Fund for supporting this project through their quantum microgrant program.

The work of the author is supported by the EPSRC UK Quantum Technologies Programme under grant EP/T001062/1 and VeriQloud.

References

- Andrés-Martínez, P., & Heunen, C. (2019). Automated distribution of quantum circuits via hypergraph partitioning. *Physical Review A*. <https://doi.org/10.1103/PhysRevA.100.032308>
- Bergholm, V., Izaac, J., Schuld, M., Gogolin, C., Ahmed, S., Ajith, V., Alam, M. S., Alonso-Linaje, G., AkashNarayanan, B., Asadi, A., & others. (2018). PennyLane: Automatic differentiation of hybrid quantum-classical computations. *arXiv Preprint arXiv:1811.04968*. <https://doi.org/10.48550/arXiv.1811.04968>

- 166 Burt, F., Chen, K.-C., & Leung, K. K. (2026). A multilevel framework for partitioning quantum
167 circuits. *Quantum*, 10, 1984. <https://doi.org/10.22331/q-2026-01-22-1984>
- 168 Clark, J., Humble, T., & Thapliyal, H. (2023). TDAG: Tree-based Directed Acyclic Graph
169 Partitioning for Quantum Circuits. *Proceedings of the Great Lakes Symposium on VLSI*.
170 <https://doi.org/10.1145/3583781.3590234>
- 171 Developers, C. (2025). *Cirq*. Zenodo. <https://doi.org/10.5281/ZENODO.4062499>
- 172 Escofet, P., Ovide, A., Almudever, C. G., Alarcón, E., & Abadal, S. (2023). Hungarian Qubit
173 Assignment for Optimized Mapping of Quantum Circuits on Multi-Core Architectures. *IEEE*
174 *COMPUTER ARCHITECTURE LETTERS*. <https://doi.org/10.1109/LCA.2023.3318857>
- 175 Gragera Garces, M., Heunen, C., & Marina, M. K. (2025). Distributed Quantum Computing
176 Across Heterogeneous Hardware with Hybrid Dependency Hypergraphs. *Proceedings of the*
177 *ACM SIGCOMM 2025 Posters and Demos*. <https://doi.org/10.1145/3744969.3748417>
- 178 Javadi-Abhari, A., Treinish, M., Krsulich, K., Wood, C. J., Lishman, J., Gacon, J., Martiel, S.,
179 Nation, P. D., Bishop, L. S., Cross, A. W., Johnson, B. R., & Gambetta, J. M. (2024).
180 Quantum computing with Qiskit. *ArXiv e-Prints*. <https://doi.org/10.48550/arXiv.2405.08810>
- 182 Kupper, M., Yeung, R., Poór, B., Toumi, A., Cashman, W., & Felice, G. de. (2025).
183 Optyx: A ZX-based python library for networked quantum architectures. *arXiv Preprint*
184 *arXiv:2512.09648*. <https://doi.org/10.48550/arXiv.2512.09648>
- 185 Quetschlich, N., Burgholzer, L., & Wille, R. (2023). MQT Bench: Benchmarking software and
186 design automation tools for quantum computing. *Quantum*. <https://doi.org/10.22331/q-2023-07-20-1062>
- 188 Schlag, S., Heuer, T., Gottesbüren, L., Akhremtsev, Y., Schulz, C., & Sanders, P. (2023).
189 High-quality hypergraph partitioning. *ACM Journal of Experimental Algorithmics*, 27, 1–39.
190 <https://doi.org/10.1145/3529090>
- 191 Sundaram, R. G., & Gupta, H. (2023). Distributing Quantum Circuits Using Teleportations.
192 *IEEE International Conference on Quantum Software (QSW)*. <https://doi.org/10.1109/QSW59989.2023.00030>
- 194 Tejedor, M., Casas, B., Conejero, J., Cervera-Lierta, A., & Badia, R. M. (2025). Orchestrating
195 quantum-HPC workflows with distributed quantum circuit cutting. *Proceedings of the SC'25*
196 *Workshops of the International Conference for High Performance Computing, Networking,*
197 *Storage and Analysis*, 1898–1906. <https://doi.org/10.1145/3731599.3767547>
- 198 Wu, A., Zhang, H., Li, G., Shabani, A., Xie, Y., & Ding, Y. (2022). AutoComm: A Framework
199 for Enabling Efficient Communication in Distributed Quantum Programs. *55th IEEE/ACM*
200 *International Symposium on Microarchitecture*. <https://doi.org/10.1109/MICRO56248.2022.00074>
- 202 Xue, Y., Yang, R., Liang, Z., & Li, T. (2026). DC-MBQC: A distributed compilation
203 framework for measurement-based quantum computing. *arXiv Preprint arXiv:2601.00214*.
204 <https://doi.org/10.1109/HPCA68181.2026.11408612>
- 205 Zomorodi-Moghadam, M., Houshmand, M., & Houshmand, M. (2017). Optimizing
206 Teleportation Cost in Distributed Quantum Circuits. *International Journal of Theoretical*
207 *Physics*. <https://doi.org/10.1007/s10773-017-3618-x>