

Bristlecone: an F# library for the long-term ecological detective

Andrew C. Martin ¹

¹ Scott Polar Research Institute, University of Cambridge, UK

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: Kalyn Dorheim 

Reviewers:

- [@kMutagene](#)
- [@AndyBunn](#)

Submitted: 25 May 2026

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/))

Summary

Long-term ecology (LTE) seeks to extend our understanding of ecological systems to decades to centuries, beyond the observational period. LTE researchers employ environmental proxies (e.g., microfossils, wood rings) to reconstruct variability in biodiversity and environments over decades to centuries. Disparate dating methods (e.g., radiocarbon dating) and representations of time are thus usually required for interpretation. Although the long view is integral to comprehensively understand ecosystem responses to environmental change, there has historically been limited application of causal inference within LTE (Katherine J. Willis et al., 2007).

Mechanistic modelling approaches may be used to infer the form and strength of ecological processes. In *The Ecological Detective*, Hilborn & Mangel (1997) present a toolkit for confronting ecological models with data. Here, I developed a modelling framework to enable LTE researchers to apply the ecological detective approach to proxy time-series for causal inference. Key capabilities include: a human-readable Domain Specific Language (DSL) for model definition with enforced dimensional correctness; model composition and scaffolding competing hypotheses; and confrontation of models with time-series of disparate dating methods.

Statement of need

Research purpose

Bristlecone's purpose is to enable researchers to build causal knowledge of the role of ecological mechanisms over decades to centuries, through confrontation of models with proxy data. LTE is essential in predicting biodiversity responses to climate change, as we may (indirectly) observe (a) ecological processes that operate over longer timescales (e.g., soil development (McLauchlan et al., 2017)); (b) biodiversity responses to broader environmental envelopes; and (c) ecological stable states and resilience (K. J. Willis et al., 2010). Although LTE research has often been descriptive, causal understandings are essential for establishing predictive capability. Causal understanding may be achieved through combining multiple lines of evidence, including (a) experimentation, (b) identifying spatial-temporal associations in observational data, and (c) mechanistic investigation. Without experimental approaches in LTE, surrogates such as space-for-time substitution can misrepresent the rate and order of processes (Elmendorf et al., 2015); thus, mechanistic investigation is essential.

In contemporary ecology ('neo-ecology'), the 'ecological detective' is an approach for causal inference (Hilborn & Mangel, 1997). The approach requires: (1) examining theory to understand plausible mechanisms; (2) understanding process and observation uncertainties; and (3) confronting plausible alternatives with data. Taking causal mechanisms as processes "*along which a signal can be propagated to produce a response*" (Grace et al. (2025), p5), the ecological detective approach may be applied to infer causal mechanisms using LTE time-series.

Bristlecone is a toolkit that enables LTE researchers to utilise the ecological detective approach through: (1) a human-readable, succinct declarative grammar for the ‘ecological detective’ workflow; and (2) implementations of key components required to use the grammar. The toolkit includes a grammar of time that reflects the dating methods used within LTE and neo-ecology; it thus targets neo-ecologists who would like to integrate neo-ecological and LTE time-series for multi-scale analysis.

Bristlecone occupies a distinct niche: although tools exist for data wrangling and inference, *Bristlecone* is a higher-level grammatical layer focused on correctness in model construction from an ecological perspective. General-purpose inference engines (e.g., Stan (Carpenter et al., 2017; Stan Development Team, 2025), JAGS (Plummer, 2003)) do not support dimensional correctness and hold limited built-in temporal concepts. *Bristlecone* embeds units-of-measure within a DSL: ecological states and flows, error structures, and temporal resolutions are expressed as explicit ecological semantics. The DSL enforces not only dimensional consistency but also ecological model structure, for example: rates must be expressed over time; likelihoods correspond to the correct ecological state; and temporal resolutions match between solvers and models. Stocks / flows, parameters, processes, and environmental forcings must thus be expressed in ecological terms. *Bristlecone* thus formalises the ecological detective workflow using ecological semantics.

State of the field

R is the favoured language of ecologists, increasing in prevalence within high-impact journals from 10.9% in 2013 to 66.9% in 2023 (Gao et al., 2025). Within LTE, key statistical approaches include transfer functions, generalised additive models, and canonical correspondence analyses. Proxy-specific mechanistic approaches have also been applied; the landscape models REVEALS (Sugita, 2007a) and LOVE (Sugita, 2007b) simulate taxon-specific biomass from pollen fluxes. However, mechanistic inference from LTE time-series has been limited (e.g., Jeffers et al. (2011)). In dendrochronology, (semi-)mechanistic approaches are more prevalent, such as inferring temperature/moisture limits to tree growth using VS-Lite (Tolwinski-Ward et al., 2011).

Bristlecone’s two unique contributions are: (1) integrating the ecological detective workflow within a single conceptual framework; and (2) providing ecological clarity through dimensional correctness. *Bristlecone* utilises F#’s type system and expressiveness (Syme, 2020) to enforce dimensionally-consistent ecological models and promote human readability, improving transparency and reproducibility of supplementary material. As the core design relies on F# language features that could not be replicated within R, *Bristlecone* was designed from scratch. However, a compositional approach was sought whereby components may be substituted where beneficial alternatives exist (e.g. optimisers). The F# R type provider (contributors & Martin, 2026) presents an avenue for embedding R libraries and graphics via typed R access.

Bristlecone’s target audience is long-term- and neo-ecologists who wish to explore mechanistic modelling. The library aims to reduce computational knowledge required to utilise the ecological detective approach within LTE research.

Software design

Although F# is uncommon in ecological research, I considered that an F#-based domain-specific language could be intuitive for ecologists with R experience. I sought to include the seven stages of the ‘ecological detective’ workflow: (1) definition of time-series models; (2) specifying model-fitting methods (“estimation engines”); (3) composition of alternative model hypotheses; (4) identifiability testing; (5) model-fitting; (6) model-selection; and (7) diagnostics.

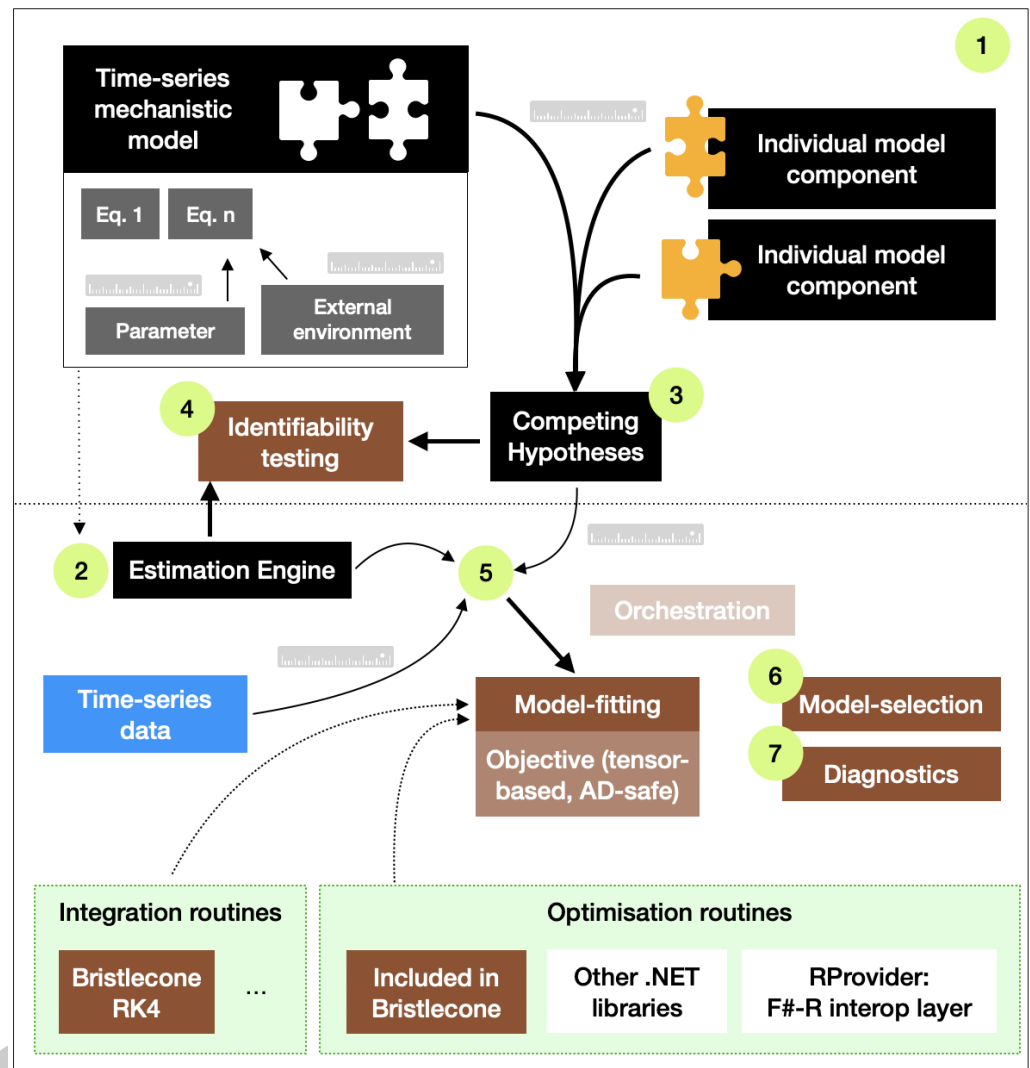


Figure 1 Conceptual framework. Green circles indicate stage of analysis. Box colours: brown = Bristlecone core functions; black = user-defined elements using the Domain Specific Language; blue = user-supplied data; white = external functions.

A core requirement was that ecological models should utilise F# type safety and units of measure such that all models, their composition, and methods are dimensionally consistent. I designed a DSL that ensures dimensional consistency across equations, parameters, external forcings (e.g., climate), and methodological approaches, supporting the scientific process by raising errors during compile-time prototyping.

To demonstrate, fish stocks (see: *The Ecological Detective*, Chapter 10) may be represented in discrete-time by:

```
module Schaefer =

    let B = state<kton> "biomass"
    let I = measure<1> "CPUE" // Catch per unit effort
    let C = environment<kton> "catch"

    let r = parameter "r" Positive 0.01</year> 1.0</year> // Intrinsic growth rate
    let K = parameter "K" Positive 100.<kton> 2000.<kton> // Carrying capacity
    let q = parameter "q" Positive 1e-6<1/kton> 1e-2<1/kton> // Catchability
```

```

let B0 = P K

let dt = Constant 1.<year>
let ``B_est[t+1]`` =
    let n = State B + P r * State B * (Constant 1. - State B / P K)
        * dt - Environment C
    Conditional (n .> Constant 0.<kton>) n (Constant 0.<kton>)

let It = P q * State B

let sigma = parameter "sigma_v" NoConstraints 0.01 1.0
let NLL = ModelLibrary.NegLogLikelihood.LogNormal (Require.measure I) sigma

let model =
    Model.discrete<year>
    |> Model.addDiscreteEquation B ``B_est[t+1]``
    |> Model.addMeasure I It
    |> Model.initialiseHiddenStateWith B B0
    |> Model.estimateParameter r
    |> Model.estimateParameter K
    |> Model.estimateParameter q
    |> Model.useLikelihoodFunction NLL
    |> Model.compile

```

99 Here, 'catch per unit effort' must be dimensionless as required by lognormal observation
100 error, and B_est[t+1] must result in kton. If the model was continuous-time, B would equal
101 kton/year. The entire model pipeline is unit-aware, preventing many model definition errors
102 at compile-time. To compose alternative mechanisms, I designed a pipe-based workflow that
103 turns nested model systems into a set of hypotheses while ensuring dimensional consistency
104 between sub-components:

```

let hypotheses =
    ``base model``
    |> Hypotheses.createFromModel
    |> Hypotheses.apply ``geometric constraint``
    |> Hypotheses.apply ``plant-soil feedback``
    |> Hypotheses.apply ``temperature limitation to growth``
    |> Hypotheses.apply ``N-limitation to growth``
    |> Hypotheses.compile

```

105 I designed Bristlecone's *estimation engine* so that models and data may be combined across
106 time concepts, with engines ensuring temporal consistency at compile-time. Radiocarbon,
107 contemporary calendar, and AD/BC time-series may all be applied or combined with models of
108 various temporal resolutions; the only requirement is that appropriate conversions are configured
109 within the engine. The strength and importance of mechanisms over different timescales (e.g.,
110 observational- versus centennial-time) may thus be investigated by combining ecological models
111 with different (proxy) data sources.

112 Identifiability testing (assessing if the engine and model can infer known parameters) is essential
113 for valid inference. I designed a unit-safe pipeline-based testing suite where synthetic time-
114 series can be generated and applied to test identifiability of model-engine combinations. For
115 model-fitting, I included orchestration for multi-threaded fitting. For model-selection I included
116 Akaike weights, as *Bristlecone* currently focuses on information theoretic approaches. Finally,
117 goodness-of-fit (root mean square error and n-step predictions) and other diagnostics were
118 included. Resultant parameter estimates are unit-typed so unit consistency is carried to further
119 analysis outside *Bristlecone*.

120 Integration with other libraries

121 I designed *Bristlecone* as a high-level workflow; it is thus extensible such that underlying
122 components, such as integration and optimisation routines, may be substituted. I chose to
123 enable gradient-based approaches by integrating automatic differentiation using DiffSharp
124 tensors (Baydin et al., 2015).

125 Three classes of optimisation routines broadly suited to ecological time-series were included,
126 while recognising that more appropriate alternatives for specific problems may be substituted.
127 First, amoeba-based methods have been applied to palaeoecological inference (Jeffers et al.,
128 2011) but are best suited to simpler likelihood surfaces. Second, I included Classical and Fast
129 Simulated Annealing (Lee, 2015; Szu & Hartley, 1987). Fast Simulated Annealing is more
130 performant at identifying global minima; *Bristlecone*'s implementation was previously applied in
131 dendroecology (Martin et al., 2021). Third, I included Monte Carlo methods including Filzbach,
132 a meta-heuristic specific to high-dimensional ecological models (Purves, 2016). As the existing
133 Filzbach library is unmaintained, and F# math libraries lack annealing, I implemented both
134 within *Bristlecone*.

135 Research impact statement

136 *Bristlecone* has been developed openly since 2018, maturing through research requirements. It
137 has been formatted for community use and contribution, including documentation, examples,
138 a benchmark suite, and contributor guidelines. Previous research applied multi-proxy
139 palaeoecological data to infer the role of soil nutrients in plant productivity through the
140 Holocene (Jeffers et al., 2011). Their models and optimisation routine were written
141 from-scratch in verbose C; I reproduced their analysis as a more accessible *Bristlecone* example.

142 *Bristlecone* was integral to Martin (2019) and Martin et al. (2021), where it was applied
143 to determine the role of soil nutrients in controlling Arctic shrub growth (*Bristlecone* v1).
144 Alternative hypotheses of nutrient limitation were confronted with wood ring and nitrogen
145 isotope time-series. That analysis is included as an example, formatted for *Bristlecone* v3.
146 *Bristlecone* is currently being applied to a pan-Arctic synthesis of Arctic palaeoecological
147 records (Martin et al., 2024).

148 AI usage disclosure

149 AI has not been used to write this manuscript.

150 No AI tools were used prior to August 2025. From then, the Microsoft Copilot macOS app was
151 used as an external tool to the development environment to suggest missing tests, refactoring
152 to code snippets, and potential approaches to proposed v3.0/3.1 features. Code suggestions
153 were critically examined and if beneficial implemented manually with human oversight. In
154 some cases, AI code suggestions were pasted into the development environment as a starting
155 point and then manually adapted. No agents have been used.

156 Acknowledgements

157 I thank support from Dr Elizabeth Jeffers and Prof Marc Macias-Fauria.

158 References

- 159 Baydin, A. G., Pearlmutter, B. A., & Siskind, J. M. (2015). *DiffSharp: Automatic differentiation*
160 *library*. <https://arxiv.org/abs/1511.07727>
- 161 Carpenter, B., Gelman, A., Hoffman, M. D., Lee, D., Goodrich, B., Betancourt, M., Brubaker,
162 M., Guo, J., Li, P., & Riddell, A. (2017). Stan: A probabilistic programming language.
163 *Journal of Statistical Software*, 76(1), 1–32. <https://doi.org/10.18637/jss.v076.i01>

- 164 contributors, R., & Martin, A. C. (2026). *RProvider: An f# type provider for r* (Version
165 v3.0.0). Zenodo. <https://doi.org/10.5281/zenodo.21964392>
- 166 Elmendorf, S. C., Henry, G. H. R., Hollister, R. D., Fosaa, A. M., Gould, W. A., Hermanutz,
167 L., Hofgaard, A., Jónsdóttir, I. S., Jorgenson, J. C., Lévesque, E., Magnusson, B., Molau,
168 U., Myers-Smith, I. H., Oberbauer, S. F., Rixen, C., Tweedie, C. E., & Walker, M. D.
169 (2015). Experiment, monitoring, and gradient methods used to infer climate change effects
170 on plant communities yield consistent patterns. *Proceedings of the National Academy of*
171 *Sciences*, 112(2), 448–452. <https://doi.org/10.1073/pnas.1410088112>
- 172 Gao, M., Ye, Y., Zheng, Y., & Lai, J. (2025). A comprehensive analysis of r's application
173 in ecological research from 2008 to 2023. *Journal of Plant Ecology*, 18(1), rtaf010.
174 <https://doi.org/10.1093/jpe/rtaf010>
- 175 Grace, J. B., Huntington-Klein, N., Schweiger, E. W., Martinez, M., Osland, M. J., Feher, L.
176 C., Guntenspergen, G. R., & Thorne, K. (2025). Causal effects versus causal mechanisms:
177 Two traditions with different requirements and contributions towards causal understanding.
178 *Ecology Letters*, 28(4), e70029. <https://doi.org/10.1111/ele.70029>
- 179 Hilborn, R., & Mangel, M. (1997). *The ecological detective: Confronting models with data*.
180 Princeton University Press.
- 181 Jeffers, E. S., Bonsall, M. B., Watson, J. E., & Willis, K. J. (2011). Climate change impacts on
182 ecosystem functioning: Evidence from an *empetrum* heathland. *New Phytologist*,
183 193(1), 150–164. <https://doi.org/10.1111/j.1469-8137.2011.03907.x>
- 184 Lee, C.-Y. (2015). Fast simulated annealing with a multivariate cauchy distribution and
185 the configuration's initial temperature. *Journal of the Korean Physical Society*, 66(10),
186 1457–1466. <https://doi.org/10.3938/jkps.66.1457>
- 187 Martin, A. C. (2019). *Long-term interactions between shrub growth and soil nitrogen availability*
188 *in the western siberian tundra: Application of mechanistic modelling approaches using*
189 *wood-ring data* [PhD thesis]. University of Oxford; University of Oxford.
- 190 Martin, A. C., Bell, L. E., Blake, K., Bradshaw, R. H., Kuoppamaa, M., Pavey, T., Prendin, A.
191 L., Speight, M., Villar, D. A., & Macias-Fauria, M. (2024). The arctic holocene biodiversity
192 database (AHBD). Zenodo. <https://doi.org/10.5281/zenodo.14132367>
- 193 Martin, A. C., Macias-Fauria, M., Bonsall, M. B., Forbes, B. C., Zetterberg, P., & Jeffers, E.
194 S. (2021). Common mechanisms explain nitrogen-dependent growth of arctic shrubs over
195 three decades despite heterogeneous trends and declines in soil nitrogen availability. *New*
196 *Phytologist*, 233(2), 670–686. <https://doi.org/10.1111/nph.17529>
- 197 McLauchlan, K. K., Gerhart, L. M., Battles, J. J., Craine, J. M., Elmore, A. J., Higuera, P.
198 E., Mack, M. C., McNeil, B. E., Nelson, D. M., Pederson, N., & Perakis, S. S. (2017).
199 Centennial-scale reductions in nitrogen availability in temperate forests of the united states.
200 *Scientific Reports*, 7(1), 7856. <https://doi.org/10.1038/s41598-017-08170-z>
- 201 Plummer, M. (2003). JAGS: A program for analysis of Bayesian graphical models using
202 Gibbs sampling. *Proceedings of the 3rd International Workshop on Distributed Statistical*
203 *Computing*, 124, 1–10.
- 204 Purves, D. (2016). Filzbach. In *Filzbach (C code) - GitHub*. [https://github.com/](https://github.com/predictionmachines/Filzbach/tree/master/filzbach)
205 [predictionmachines/Filzbach/tree/master/filzbach](https://github.com/predictionmachines/Filzbach/tree/master/filzbach)
- 206 Stan Development Team. (2025). Stan reference manual. In *Stan Docs*. [https://mc-](https://mc-stan.org/docs/reference-manual/)
207 [stan.org/docs/reference-manual/](https://mc-stan.org/docs/reference-manual/)
- 208 Sugita, S. (2007a). Theory of quantitative reconstruction of vegetation i: Pollen from
209 large sites REVEALS regional vegetation composition: *The Holocene*, 17(2), 229–241.
210 <https://doi.org/10.1177/0959683607075837>

- 211 Sugita, S. (2007b). Theory of quantitative reconstruction of vegetation II: All you need is
212 LOVE: *The Holocene*, 17(2), 243–257. <https://doi.org/10.1177/0959683607075838>
- 213 Syme, D. (2020). The early history of F#. *Proceedings of the ACM on Programming*
214 *Languages*, 4(HOPL), 1–58. <https://doi.org/10.1145/3386325>
- 215 Szu, H., & Hartley, R. (1987). Fast simulated annealing. *Physics Letters A*, 122(3–4), 157–162.
216 [https://doi.org/10.1016/0375-9601\(87\)90796-1](https://doi.org/10.1016/0375-9601(87)90796-1)
- 217 Tolwinski-Ward, S. E., Evans, M. N., Hughes, M. K., & Anchukaitis, K. J. (2011). An efficient
218 forward model of the climate controls on interannual variation in tree-ring width. *Climate*
219 *Dynamics*, 36(11), 2419–2439. <https://doi.org/10.1007/s00382-010-0945-5>
- 220 Willis, Katherine J., Araújo, M. B., Bennett, K. D., Figueroa-Rangel, B., Froyd, C. A.,
221 & Myers, N. (2007). How can a knowledge of the past help to conserve the future?
222 Biodiversity conservation and the relevance of long-term ecological studies. *Philosophical*
223 *Transactions of the Royal Society B: Biological Sciences*, 362(1478), 175–187. <https://doi.org/10.1098/rstb.2006.1977>
- 224
- 225 Willis, K. J., Bailey, R. M., Bhagwat, S. A., & Birks, H. J. B. (2010). Biodiversity baselines,
226 thresholds and resilience: Testing predictions and assumptions using palaeoecological data.
227 *Trends in Ecology & Evolution*, 25(10), 583–591. [https://doi.org/10.1016/j.tree.2010.07.](https://doi.org/10.1016/j.tree.2010.07.006)
228 006