

Olmsted: Rich Interactive Visualization of B Cell Lineages

David H. Rich¹, Elias Harkins¹, Christopher T. Small¹, Kevin Sung¹,
Laura Doecker², Duncan Ralph¹, Amrit Dhar¹, Mackenzie S.
Kopp², Julie Overbaugh², and Frederick A. Matsen IV¹¶

¹ Computational Biology Program, Fred Hutchinson Cancer Center, Seattle, WA, USA ² Human Biology
Division, Fred Hutchinson Cancer Center, Seattle, WA, USA ¶ Corresponding author

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) ↗
- [Repository](#) ↗
- [Archive](#) ↗

Editor: ↗

Submitted: 27 June 2026

Published: unpublished

License

Authors of papers retain copyright
and release the work under a
Creative Commons Attribution 4.0
International License ([CC BY 4.0](https://creativecommons.org/licenses/by/4.0/)).

Summary

Olmsted is an open-source, browser-based application for visually exploring B cell repertoires and clonal family tree data. The immune system learns to recognize new threats by mutating and selecting the cells that produce antibodies, a process that leaves behind a branching evolutionary record much like a family tree. In immunological terms, this affinity maturation of B cell receptor sequences coding for immunoglobulins (antibodies) begins with a diverse pool of randomly generated naive sequences and leads to a collection of evolutionary histories. High-throughput DNA sequencing of the B cell repertoire, combined with computational reconstruction of these evolutionary histories, produces large collections of clonal families and their associated phylogenetic trees. Olmsted enables researchers to scan across collections of clonal families using summary statistics, then interactively explore individual families to visualize phylogenies and amino acid mutations that occurred during affinity maturation.

Statement of Need

Researchers studying B cell responses rely on high-throughput sequencing to characterize antibody repertoires. Computational tools can reconstruct the evolutionary histories of these sequences, producing large collections of clonal families (groups of sequences descended from a common naive ancestor) and phylogenetic trees describing their diversification. However, researchers often lack tools to explore these reconstructions in the detail necessary to choose sequences for functional, structural, or biochemical studies.

Olmsted addresses this gap by combining repertoire-level overview with detailed lineage exploration. Researchers can scan across all clonal families in a scatterplot, then drill down to examine phylogenetic trees with aligned amino acid sequences showing mutations from the naive ancestor. This multi-scale navigation is delivered through a zero-installation web application that keeps data client-side, supporting practical decisions about which sequences to prioritize for downstream experimental studies. Specifically, Olmsted allows users to:

- View all clonal families simultaneously in a configurable scatterplot
- Filter and select families based on biological criteria (mutation frequency, sequence count, V(D)J gene usage)
- Examine phylogenetic trees with aligned sequences showing mutations from the naive ancestor
- Trace the mutational history of individual sequences back to their germline origin
- Visualize paired heavy and light chain sequences together, which is essential for selecting antibodies for expression

41 The submission under review is the Olmsted web application: the interactive, client-side
42 visualization tool in this repository. A companion command-line utility, [olmsted-cli](#), converts
43 common immunoinformatics formats into the Olmsted JSON input format. This utility is
44 a thin, self-contained format shim rather than the scholarly contribution. The substantial
45 engineering and domain modeling of this project live in the web application, which we describe
46 below.

47 State of the Field

48 Several existing tools visualize B cell repertoire and lineage data, but each addresses a different
49 analytical goal and none combines a repertoire-wide overview with interactive, paired-chain
50 lineage exploration. AncesTree ([Foglierini et al., 2020](#)) provides detailed single-tree exploration
51 with amino acid mutation display, but requires Java installation, processes one lineage at a
52 time without a repertoire overview, and handles heavy or light chains separately rather than as
53 pairs. ViCloD ([Jeusset et al., 2023](#)) focuses on large-scale intraclonal diversity analysis across
54 hundreds of thousands of sequences, primarily for characterizing B cell tumors; it analyzes
55 heavy chains only. AIRRscape ([Waltari et al., 2022](#)) enables comparison across multiple
56 repertoires to identify convergent antibody responses at the population level, again for heavy
57 chains only. ImmuneDB ([Rosenfeld et al., 2016](#)) presents collections in paginated list form for
58 database-style querying.

59 Olmsted differs from these tools along several independent axes, each individually uncommon.
60 It integrates repertoire-scale scanning with per-lineage amino acid resolution in one interface:
61 a configurable scatterplot of an entire repertoire is linked to a combined phylogenetic-tree-and-
62 alignment view. In practice this lets a researcher spot a striking family in the overview (an
63 unusually expanded or hypermutated one, say) and drill straight into its mutations without
64 switching tools. It displays paired heavy and light chains side by side, which is essential for
65 antibody discovery, where a functional antibody requires both chains of a pair. It traces full
66 ancestral mutational paths from the naive sequence, and it can overlay model-derived per-site
67 scores on a lineage (see below). Each of these capabilities is rare or absent among existing
68 tools.

69 The most fundamental distinction, however, is structural. Every tool above is either installed
70 and run locally (AncesTree requires a Java installation; ImmuneDB is a database system to be
71 deployed and maintained) or built as the front end of its own analysis pipeline (ViCloD and
72 AIRRscape couple exploration to the specific processing and data they generate). Olmsted
73 is instead a standalone visualization layer, decoupled from any reconstruction method: it
74 consumes a documented, standards-based input format, so lineages built by any upstream
75 pipeline can be explored in it, with no installation and no data leaving the browser. To our
76 knowledge, it is the only installation-free B cell lineage visualization tool that is not tied to a
77 particular upstream analysis pipeline. This separation of concerns lets the visualization improve
78 independently of, and interoperate with, the fast-moving ecosystem of reconstruction methods.

79 Software Design

80 Olmsted is built with React and Redux, with all visualizations implemented as Vega specifications
81 ([Satyanarayan et al., 2016](#)) and rendered through a single wrapper component. The application
82 is deliberately client-side only: it performs no server-side computation, and uploaded data
83 is stored in the browser (IndexedDB, via Dexie) and never transmitted ([Figure 1](#)). This is
84 both a privacy guarantee—sensitive patient-derived sequences never leave the researcher's
85 machine—and an operational simplification, since the public instance at [olmstedviz.org](#) is
86 served as static files with no backend to maintain. The application can also be run locally via
87 Docker.




Figure 1: Dataset management interface.

88 The engineering effort in the application is substantial and domain-specific. It includes: a
89 lazy-loading data model that keeps thousand-family datasets responsive by loading heavy
90 per-tree sequence data only on demand; a field-metadata system that lets dataset-supplied
91 fields dynamically populate axes, color encodings, tooltips, and table columns without code
92 changes; reconciliation of “forest” inputs (multiple disconnected subtrees) into a single rooted
93 tree via consensus synthetic roots; and a Vega-wrapper abstraction that preserves zoom, pan,
94 and brush state across data updates. These decisions, and the guardrails that protect them, are
95 documented in the repository’s DESIGN.md and ARCHITECTURE.md. Correctness of this domain
96 logic is protected by a suite of over 600 automated tests across more than 30 files, run in
97 continuous integration alongside linting and a production build.

98 **Provenance.** The codebase originated in 2018 as a fork of Nextstrain’s Auspice (Hadfield
99 et al., 2018), a viewer for viral-genome phylogeography, and was actively developed through
100 2020 before being shelved. It was revived and substantially rewritten in 2025. Of the roughly
101 80 source modules in the current application, about three-quarters were written after the
102 fork; the remainder are generic framework scaffolding (layout, navigation, URL routing, and
103 store configuration), most of which has itself been rewritten. None of the domain-specific
104 visualizations (the repertoire scatterplot, the combined tree-and-alignment view, the ancestral-
105 lineage tracer, or the V(D)J recombination display) derive from Auspice, whose visualizations
106 target a different domain entirely. Olmsted retains from Auspice a general architectural pattern
107 (a React/Redux single-page application with deep-linkable URL state) rather than reusable
108 visualization code.

109 Data Preparation with olmsted-cli

110 The olmsted command-line tool converts data into Olmsted’s JSON format. Unlike the web
111 application, this data-preparation step requires a one-time install: olmsted-cli is a Python
112 package (Python 3.8 or higher) installed with pip install olmsted-cli (or pipx install
113 olmsted-cli for an isolated environment). It supports two input formats:

- 114 ■ **AIRR format:** The JSON-based standard developed by the Adaptive Immune Receptor
115 Repertoire Community (AIRR Community, 2024; Rubelt et al., 2017; Vander Heiden
116 et al., 2018). The AIRR lineage tree schema is being finalized and we will continue to
117 adapt to changes should they arise.

118 ▪ **PCP (Parent-Child Pair) format:** A CSV-based format with explicit parent-child
119 relationships.

120 Example usage:

```
# Generate a config from your data, then edit it
olmsted build-config -i data.json -o config.yaml
```

```
# Process with the config
olmsted process -c config.yaml -n "My Dataset" -o data-olmsted.json
```

121 The build-config command introspects input data and generates a YAML configuration listing
122 all discoverable fields with inferred types and labels, which the user can edit before processing.
123 These fields will dynamically populate plotting options and hover tooltip information. Once
124 data preparation is done, the resulting file can be shared with collaborators, who then need
125 only the web interface. Plot configurations can also be exported from the web app so they can
126 be recreated later.

127 **Features**

128 **Interactive Visualization**

129 The visualization interface consists of four linked sections:

- 130 1. **Clonal Families Scatterplot (Figure 2):** Each clonal family appears as a point, with
131 configurable axes, colors, and faceting. The filter tool narrows the selection of families
132 displayed. Users can brush-select regions or click individual points.




Figure 2: Example scatterplot.

- 133 2. **Selected Clonal Families Table (Figure 3):** Displays metadata for selected families,
134 including V(D)J gene assignments and a visual representation of the recombination event.
135 The table can be sorted by any column, and families can be starred for easy reference.

The screenshot displays the 'Clonal Family Selection Table' in the Olmsted web application. The table lists various clonal families with columns for Star, Select, Info, Naive sequence, Family ID, Unique Seq Count, V gene, D gene, J gene, Locus, Chain, Paired, and Pair ID. The first five rows are highlighted in yellow, and the sixth row is highlighted in blue. The table shows a selection of clonal families with their respective gene sequences and counts.

Star	Select	Info	Naive sequence	Family ID	Unique Seq Count	V gene	D gene	J gene	Locus	Chain	Paired	Pair ID
★	☐			d1_243428-igl-90165-heavy	74	IGHV3-23*04	IGHD6-13*01	IGHJ4*02	igh	heavy	Yes	pair-d1_24
★	☐			d1_243737-igl-90474-heavy	72	IGHV3-23*01	IGHD5-24*01	IGHJ6*02	igh	heavy	Yes	pair-d1_24
★	☐			d1_71027-igl-71027-heavy	72	IGHV3-15*01	IGHD5-18*02	IGHJ4*02	igh	heavy	Yes	pair-d1_71
★	☐			d1_20468-igl-20468-heavy	48	IGHV3-23*01	IGHD2-2*02	IGHJ6*02	igh	heavy	Yes	pair-d1_20
★	☐			d1_244852-igl-91589-heavy	48	IGHV4-39*01	IGHD3-16*03	IGHJ4*02	igh	heavy	Yes	pair-d1_24
★	☒			d1_22662-igl-22662-heavy	44	IGHV3-30*03	IGHD4-4*01	IGHJ4*02	igh	heavy	Yes	pair-d1_22
☆	☐			d1_151921-igl-151921-heavy	444	IGHV3-23*01	IGHD2-8*01	IGHJ4*02	igh	heavy	Yes	pair-d1_15
☆	☐			d1_174921-igl-21658-heavy	304	IGHV2-5*02	IGHD3-10*01	IGHJ4*02	igh	heavy	Yes	pair-d1_17
☆	☐			d3_7388-igl-7388-heavy	290	IGHV3-49*04	IGHD3-22*01	IGHJ5*02	igh	heavy	Yes	pair-d3_73
☆	☐			d4_208814-igl-208814-heavy	256	IGHV3-48*03	IGHD2-8*01	IGHJ1*01	igh	heavy	Yes	pair-d4_20
☆	☐			d1_247306-igl-94043-heavy	238	IGHV1-3*01	IGHD1-20*01	IGHJ4*02	igh	heavy	Yes	pair-d1_24
☆	☐			d4_203694-igl-203694-heavy	234	IGHV3-48*03	IGHD2-8*01	IGHJ1*01	igh	heavy	Yes	pair-d4_20

Figure 3: Example selected clonal families table.

3. **Clonal Family Tree (Figure 4):** Shows the phylogenetic tree alongside a sequence alignment. Colors indicate amino acid mutations relative to the naive sequence by default, or a heatmap can display contiguous per-site data. The tree and alignment support zooming and panning, as well as focusing on a subtree based on a selected subroot node.




Figure 4: Example tree and alignment clonal family visualization.

4. **Ancestral Sequences (Figure 5):** For a selected leaf, displays the complete mutational path from the naive sequence, traced through the reconstructed ancestral sequences along the inferred phylogenetic tree. For this and the clonal family tree view, mutations can be colored by an arbitrary numeric value supplied for each mutation in the input data and displayed as a heatmap (Figure 6). An example is a per-site, per-amino-acid selection score such as those produced by a deep amino acid selection model (DASM) (Matsen et al., 2026).



Figure 5: Example ancestral sequence visualization.




Figure 6: Clonal family tree and alignment with mutations colored by a per-site metric as a heatmap, rather than by amino acid identity.

Research Impact Statement

Earlier versions of Olmsted supported the interactive exploration of B cell lineages in studies of HIV-1-specific antibody responses by the Overbaugh group (Doepker et al., 2020, 2021; Simonich et al., 2019; Williams et al., 2018), where it was used to navigate and interrogate reconstructed clonal families and to select antibodies for functional characterization. Tracing ancestral paths from a naive sequence to observed antibodies is central to studying affinity maturation trajectories, such as the development of broadly neutralizing antibodies against HIV. By making these trajectories, the underlying phylogenies, and paired heavy/light chains directly explorable in the browser, Olmsted lowers the barrier to prioritizing sequences for expression and biochemical study. The revived tool is positioned to serve the broader B cell phylogenetics community through its use of the AIRR Community's data standards (AIRR Community, 2024;

159 Rubelt et al., 2017; Vander Heiden et al., 2018), enabling a standards-based, interoperable
160 analysis pipeline. The application also supports emerging analyses: per-site, per-amino-acid
161 selection scores from deep amino acid selection models (DASM) (Matsen et al., 2026) can be
162 overlaid on lineages as a heatmap, turning Olmsted into a viewer for model-derived quantities
163 as well as observed mutations.

164 AI Usage Disclosure

165 The 2018–2020 development of Olmsted predates the availability of AI coding assistants and
166 was written conventionally. The 2025 revival was carried out with substantial assistance from
167 agentic AI coding tools (Claude Code). These tools were used to modernize the application's
168 dependencies and JavaScript frameworks, replace the legacy server-side data pipeline with
169 client-side processing and browser-based (IndexedDB) storage, and add the test suite and
170 continuous integration described above. AI assistance was also used in preparing repository
171 documentation and in drafting portions of this manuscript. All AI-assisted changes to code
172 and text were reviewed, edited, and tested by the authors, who take full responsibility for the
173 content.

174 Acknowledgements

175 We thank Trevor Bedford, James Hadfield, and other authors of Nextstrain, on which Olmsted's
176 source code is based. We are grateful to the AIRR data representation working group for their
177 cooperative development of the AIRR lineage schema. This work was supported by National
178 Institutes of Health grants R01 AI146028, R01 GM113246, R01 AI120961, R01 AI138709, U19
179 AI117891, and U19 AI128914, and by the Howard Hughes Medical Institute.

180 References

- 181 AIRR Community. (2024). *AIRR clones and trees schema*. <https://docs.airr-community.org/>.
- 182 Doepker, L. E., Danon, S., Harkins, E., Ralph, D. K., Yaffe, Z., Garrett, M. E., Dhar, A., Wagner,
183 C., Stumpf, M. M., Arenz, D., Williams, J. A., Jaoko, W., Mandaliya, K., Lee, K. K., Matsen,
184 F. A., & Overbaugh, J. M. (2021). Development of antibody-dependent cell cytotoxicity
185 function in HIV-1 antibodies. *Elife*, 10, e63444. <https://doi.org/10.7554/eLife.63444>
- 186 Doepker, L. E., Simonich, C. A., Ralph, D., Shipley, M. M., Garrett, M., Gobillot, T.,
187 Vigdorovich, V., Sather, D. N., Nduati, R., Matsen, F. A., & Overbaugh, J. M. (2020).
188 Diversity and function of maternal HIV-1-specific antibodies at the time of vertical
189 transmission. *J. Virol.*, 94(8), e01594–19. <https://doi.org/10.1128/JVI.01594-19>
- 190 Foglierini, M., Pappas, L., Lanzavecchia, A., Corti, D., & Perez, L. (2020). Ancestree: An
191 interactive immunoglobulin lineage tree visualizer. *PLoS Comput. Biol.*, 16(7), e1007731.
192 <https://doi.org/10.1371/journal.pcbi.1007731>
- 193 Hadfield, J., Megill, C., Bell, S. M., Huddleston, J., Potter, B., Callender, C., Sagulenko, P.,
194 Bedford, T., & Neher, R. A. (2018). Nextstrain: Real-time tracking of pathogen evolution.
195 *Bioinformatics*, 34(23), 4121–4123. <https://doi.org/10.1093/bioinformatics/bty407>
- 196 Jousset, L., Abdollahi, N., Verny, T., Armand, M., De Septenville, A. L., Davi, F., & Bernardes,
197 J. S. (2023). ViCloD, an interactive web tool for visualizing B cell repertoires and analyzing
198 intracolon diversities: Application to human B-cell tumors. *NAR Genom. Bioinform.*, 5(2),
199 lqad064. <https://doi.org/10.1093/nargab/lqad064>
- 200 Matsen, F. A., Dumm, W., Sung, K., Johnson, M. M., Rich, D. H., Starr, T. N., Song, Y. S.,
201 Fukuyama, J., & Haddox, H. K. (2026). Separating selection from mutation in antibody
202 language models. *Elife*. <https://doi.org/10.7554/eLife.109644>

- 203 Rosenfeld, A. M., Meng, W., Luning Prak, E. T., & Hershberg, U. (2016). ImmuneDB: A system
204 for the analysis and exploration of high-throughput adaptive immune receptor sequencing
205 data. *Bioinformatics*, 33(2), 292–293. <https://doi.org/10.1093/bioinformatics/btw593>
- 206 Rubelt, F., Busse, C. E., Bukhari, S. A. C., Bürckert, J.-P., Mariotti-Ferrandiz, E., Cowell, L. G.,
207 Watson, C. T., Marthandan, N., Faison, W. J., Hershberg, U., Laserson, U., Corrie, B. D.,
208 Davis, M. M., Peters, B., Lefranc, M.-P., Scott, J. K., Breden, F., AIRR Community, Luning
209 Prak, E. T., & Kleinstein, S. H. (2017). Adaptive immune receptor repertoire community
210 recommendations for sharing immune-repertoire sequencing data. *Nature Immunology*,
211 18(12), 1274–1278. <https://doi.org/10.1038/ni.3873>
- 212 Satyanarayan, A., Russell, R., Hoffswell, J., & Heer, J. (2016). Reactive vega: A streaming
213 dataflow architecture for declarative interactive visualization. *IEEE Transactions on*
214 *Visualization and Computer Graphics*. <https://doi.org/10.1109/TVCG.2015.2467091>
- 215 Simonich, C. A., Doecker, L., Ralph, D., Williams, J. A., Dhar, A., Yaffe, Z., Gentles, L., Small,
216 C. T., Oliver, B., Vigdorovich, V., Mangala-Prasad, V., Nduati, R., Sather, D. N., Lee,
217 K. K., Matsen, F. A., & Overbaugh, J. (2019). Kappa chain maturation helps drive rapid
218 development of an infant HIV-1 broadly neutralizing antibody lineage. *Nat. Commun.*,
219 10(1), 2190. <https://doi.org/10.1038/s41467-019-09481-7>
- 220 Vander Heiden, J. A., Marquez, S., Marthandan, N., Bukhari, S. A. C., Busse, C. E., Corrie,
221 B., Hershberg, U., Kleinstein, S. H., Matsen IV, F. A., Ralph, D. K., Rosenfeld, A. M.,
222 Schramm, C. A., AIRR Community, Christley, S., & Laserson, U. (2018). AIRR community
223 standardized representations for annotated immune repertoires. *Frontiers in Immunology*,
224 9, 2206. <https://doi.org/10.3389/fimmu.2018.02206>
- 225 Waltari, E., Nafees, S., McCutcheon, K. M., Wong, J., & Pak, J. E. (2022). AIRRscope: An
226 interactive tool for exploring B-cell receptor repertoires and antibody responses. *PLoS*
227 *Comput. Biol.*, 18(9), e1010052. <https://doi.org/10.1371/journal.pcbi.1010052>
- 228 Williams, K. L., Wang, B., Arenz, D., Williams, J. A., Dingens, A. S., Cortez, V., Simonich, C.
229 A., Rainwater, S., Lehman, D. A., Lee, K. K., & Overbaugh, J. (2018). Superinfection
230 drives HIV neutralizing antibody responses from several B cell lineages that contribute to a
231 polyclonal repertoire. *Cell Rep.*, 23(3), 682–691. <https://doi.org/10.1016/j.celrep.2018.03.082>
232