

PPanGGOLiN V2: technical enhancement and extended functionalities for prokaryotic pangenome analysis

Jérôme Arnoux ^{1*}, Jean Mainguy ^{1*}, Adelme Bazin ², Guillaume Gautreau ³, Téo Lemane ¹, David Vallenet ¹, and Alexandra Calteau ^{1¶}

¹ LABGeM, Genomique Métabolique, CEA, Genoscope, Institut François Jacob, Université d'Évry, Université Paris-Saclay, CNRS, France.  ² Laboratory of Biology and Modeling of the Cell, Ecole Normale Supérieure de Lyon, CNRS, Université Claude Bernard Lyon, Université de Lyon, France  ³ Université Paris-Saclay, INRAE, MalAGE, Jouy-en-Josas, France  ¶ Corresponding author * These authors contributed equally.

DOI: [10.xxxxxx/draft](https://doi.org/10.xxxxxx/draft)

Software

- [Review](#) 
- [Repository](#) 
- [Archive](#) 

Editor: [Frederick Boehm](#) 

Reviewers:

- [@hkaspersen](#)
- [@ramadatta](#)
- [@mjsull](#)

Submitted: 24 June 2026

Published: unpublished

License

Authors of papers retain copyright and release the work under a Creative Commons Attribution 4.0 International License ([CC BY 4.0](#)).

Summary

The exponential growth of genomic data, particularly for microbes, has made pangenomic approaches a gold standard for large-scale comparative genomics. By capturing the full genomic diversity of a species rather than relying on a single reference, pangenomics has transformed microbial genomics, revealing the adaptive potential of bacteria and the evolutionary dynamics underlying functional diversity. Among available tools, PPanGGOLiN distinguishes itself through its graph-based model coupled with statistical gene partitioning.

Here we present PPanGGOLiN v2, which introduces substantial improvements across three dimensions: new analytical features that expand what users can investigate, a comprehensive software architecture redesign that improves maintainability and extensibility, and performance improvements that address the computational demands of ever-growing genomic datasets.

Statement of need

In the last decade, the number of sequenced microbial genomes has exploded from a few thousand to several million. While this wealth of data offers immense potential, traditional genome-centric approaches reach their limits when exploring and interpreting variation at this scale ([Land et al., 2015](#); [Parks et al., 2018](#)). A pangenome condenses all genomic information from a set of genomes into a single structure, capturing gene presence/absence and genetic context, such as mobile genetic elements or functional annotations, at the species or clade level ([Tettelin et al., 2005](#)).

PPanGGOLiN ([Gautreau et al., 2020](#)) addresses these needs through a graph-based model that clusters genes into families based on sequence similarity and encodes neighborhood information, compressing genomic diversity into a single structure. Gene families are then partitioned using a Bernoulli Mixture Model coupled with a Markov Random Field accounting for both family occurrence and genomic neighborhood. The software suite further integrates PanRGP ([Bazin et al., 2020](#)), for identification of Regions of Genomic Plasticity (RGPs) and their insertion spots, and PanModule ([Bazin et al., 2021](#)), for finding conserved modules.

PPanGGOLiN v2 builds on this foundation by adding new analytical features: genomic context search, RGP clustering, pangenome projection, and metadata association ([Figure 1](#)) together with a full architectural redesign and a modernized development workflow, making the tool

40 more capable, maintainable, and open to community contributions.

41 State of the field

42 Pangenomics has transformed prokaryotic genome analysis by moving beyond a reference-
43 centric approach toward representations that capture the entire genomic diversity of microbial
44 populations. Two pangenome graph models are widely used. Sequence-level graphs, like Bifrost
45 (Holley & Melsted, 2020) or vg toolkit (Garrison et al., 2018), encode nucleotide-level variation,
46 providing small variants resolution, but at high computational and interpretive complexity.
47 Gene-level graphs, such as Panaroo (Tonkin-Hill et al., 2020) or pangene (Li et al., 2024),
48 instead cluster genes into families to form graph nodes, with edges encoding genomic contiguity.
49 This representation neglects fine-scale variation but remains highly informative about gene
50 presence/absence and mobile elements at a fraction of the cost.

51 Most gene-level tools rely on a predefined occurrence threshold to split gene families into
52 different partitions, leading to inconsistent predictions depending on genome quality and diversity.
53 PPanGGOLiN instead uses a statistical method relying on both gene family occurrence patterns
54 and graph topology to partition families into: *persistent*, *shell*, and *cloud* genomes. Among
55 dozens of gene-level pangenomic tools, only mOTUp (Buck et al., 2022) and micropan
56 (Snipen & Liland, 2015) also use statistical partitioning.

57 Additionally, PPanGGOLiN builds on its graph-based representation and fine-grained partitioning
58 to identify RGP, spots, and conserved modules; analyses unique to PPanGGOLiN.

59 New features

60 Projection: pangenome annotation of external genomes

61 The projection feature annotates external genomes using a previously computed pangenome,
62 without recomputing the full structure. Input genes are assigned to existing gene families
63 and partitions based on sequence similarity; unmatched genes are considered genome-specific
64 and assigned to the cloud partition. From these assignments, PPanGGOLiN further predicts
65 RGP, spots, and modules present in the projected genome. This is useful for comparing newly
66 sequenced genomes against an established reference pangenome.

67 In Version 2, PPanGGOLiN introduces a Proksee-compatible output format (Grant et al., 2023),
68 to visualize the circular genomes annotated with pangenome information, from the pangenome
69 itself or via projection, allowing further interactive analyses within Proksee.

70 Genomic context extraction

71 A genomic context refers to a set of genes conserved within similar genomic regions across
72 multiple genomes. The pangenome graph already encodes neighborhood relationships, making
73 it well suited for extracting such context. The complete method is described in detail by
74 Arnoux et al. (2026). This feature provides a species-level view of conserved genomic
75 neighborhoods, letting users identify functionally related genes and, for example, accurately
76 reconstruct metabolic pathways.

77 RGP clustering

78 A new feature allows comparative analysis of RGP across genomes in the pangenome.
79 Clustering is based on gene content similarity, where two RGP are considered related if
80 their genes belong to the same families, quantified using a Gene Repertoire Relatedness (GRR)
81 score.

82 Clustering is performed through graph modeling: each RGP is a node, and an edge is added
83 between two nodes if their GRR (min or max) exceeds a threshold (0.8 by default). Each
84 connected component of the resulting graph corresponds to a cluster of RGPs sharing a
85 common gene repertoire. Results are exportable as a tabulated file or in formats compatible
86 with graph visualization tools, such as Gephi(Bastian et al., 2009) and Cytoscape (Shannon
87 et al., 2003), facilitating the exploration of the diversity and dissemination of mobile genetic
88 elements across genomes.

89 Metadata

90 Another key addition is the ability to associate metadata with any pangenome element, including
91 genes, contigs, genomes, gene families, edges, RGPs, spots, and modules. Metadata is provided
92 as a tabulated file requiring at minimum the target element's identifier, with no restriction on
93 the type or number of additional fields, making the format highly flexible. Each entry is linked
94 to a named source, allowing an element to carry multiple annotations from different sources.


95 For performance and portability, metadata is stored directly in the pangenome file. While
96 not used directly in pangenome computation, it is propagated to all outputs, facilitating
97 downstream interpretation and exploration.

98 Combined with RGP clustering, and pangenome projection, metadata association bridges the
99 gap between pangenome structure and biological interpretation, letting users connect genomic
100 variation to functional, ecological, or experimental data.

101 Software design

102 PPanGGOLiN is built around a central data model in which the HDF5 pangenome file serves
103 as the reference data object for the entire system. Analyses can be performed step by step or
104 through workflow commands executing the pipeline in a single run, providing flexibility and
105 convenience (Figure 1). Once executed, the HDF5 pangenome file supports multiple uses:
106 export to other formats for downstream analyses, input for additional PPanGGOLiN commands,
107 and programmatic access through the Python API. This central-object design also enhances
108 reproducibility: parameters are resolved consistently (command line, then configuration file,
109 then defaults), validated across analysis steps, and stored within the pangenome file, enabling
110 analyses to be reliably reproduced and compared.

(i) Pangenome construction



(ii) Pangenome exploitation




Figure 1: Overview of the PPanGGOLiN v2 workflow. Rounded boxes represent software commands. Commands marked with * are new in v2. Commands fall into two sections. (i) *Pangenome construction*: from genomic data (annotated or not), to progressively build the pangenome graph and partition gene families into *persistent*, *shell*, and *cloud*. Two independent analyses then follow: *rgp* and *spot* identify RGP and their insertion spots, while *module* identifies conserved genomic modules. Results are stored in a single HDF5 file (*pangenome.h5*), the central data object. External metadata can be associated with any pangenome element via *metadata **, using a tabulated file. (ii) *Pangenome exploitation*: the HDF5 file serves as input for four categories of downstream commands. *Pangenome analysis*: *msa* computes multiple sequence alignments for gene families; *rarefaction* computes the pangenome's rarefaction curve; *rgp_cluster ** clusters RGP based on shared gene content. *Pangenome-guided analysis*: *align* maps an external genome or protein set to pangenome families; *context ** extracts conserved genomic neighborhoods around genes of interest; *projection ** annotates a new genome using an existing pangenome. *Pangenome output*: *draw*, *fasta*, *write_pangenome*, and *write_genomes* export results in various formats for downstream analyses and visualization.

Technical enhancements

PPanGGOLiN was originally developed as a research project. Despite robust and user-friendly workflows, its codebase required a complete redesign to meet long-term software engineering standards.

Several targeted improvements address performance and reproducibility. Prodigal (Hyatt et al., 2010) was replaced by Pyrodigal (Larralde, 2022), reducing I/O overhead during annotation. The HDF5 pangenome file structure was redesigned to drastically reduce file size (Figure 2) and improve API access. Pangenome file reading was optimized, reducing loading time from minutes to seconds for large datasets. Memory consumption during sequence writing was reduced by reading sequences directly from HDF5 files rather than loading them into memory. A configuration file system now facilitates PPanGGOLiN integration into computational platforms (see Research impact).

Finally, documentation was fully revised and is now hosted on ReadTheDocs, lowering the barrier for new users. PPanGGOLiN v2 is distributed via Bioconda (Grüning et al., 2018) and PyPI, ensuring broad accessibility across computing environments.




Figure 2: Pangenome HDF5 file size (GB) generated by PPanGGOLiN v1.2.74 vs v2.3.0 across increasing genome counts. PPanGGOLiN v2 consistently produces smaller files, supporting the impact of the v2 file structure redesign.

Research impact statement

PPanGGOLiN is widely adopted in microbial genomics, with over 300 citations on Google Scholar (August 2026). Its results are integrated into MicroScope (Vallenet et al., 2020), a microbial genome annotation and analysis platform, and a dedicated Galaxy (The Galaxy Community, 2024) module has been developed to broaden accessibility. The tool is a core dependency for PANORAMA (Arnoux et al., 2026), enabling cross-species biological system predictions and pangenome comparisons, and supports PanGBank (Mainguy et al., 2026), a pangenome database. A companion project further extends this ecosystem by integrating PPanGGOLiN pangenomes — including RGP, spots, and modules — alongside functional metadata into a

Neo4j graph database, enabling advanced graph-based queries across pangenome elements (Arnoux et al., 2024). PPanGGOLiN is actively maintained and evolves through community contributions, bug reports, and feature suggestions. In 2023, it received the Open Science Award for Open Source Research Software from the French Ministry of Higher Education and Research.

AI usage disclosure

AI tools were used in a limited capacity during later development. GitHub Copilot (mainly Claude Sonnet 4.6) assisted with code review suggestions, test writing, and documentation. Claude Sonnet 4.6 (Anthropic) assisted in designing and refining the workflow overview figure (Figure 1) and in manuscript proofreading. Core software design, implementation, and architectural decisions were developed independently without AI assistance. All AI-generated outputs were systematically reviewed and validated by authors.

Acknowledgments

This research was supported by the CFR PhD program of the French Alternative Energies and Atomic Energy Commission (CEA) for JA and the BlueRemediomics project for JM, funded by the European Union under the Horizon Europe Program, Grant Agreement No. 101082304.

We thank the PPanGGOLiN user community for their contributions, feedback, and bug reports, which have continuously improved the tool and shaped this new version.

References

- Arnoux, J., Bonifati, A., Calteau, A., Dumbrava, S., & Gautreau, G. (2024). Integrating complex pangenome graphs. *2024 IEEE 40th International Conference on Data Engineering Workshops (ICDEW)*, 350–354. <https://doi.org/10.1109/ICDEW61823.2024.00052>
- Arnoux, J., Mainguy, J., Bry, L., Grado, Q. F. de, Hoblos, Y., Vallenet, D., & Calteau, A. (2026). Panorama: A robust pangenome-based method for predicting and comparing biological systems across species. *PLOS Computational Biology*, 22(7), e1013856. <https://doi.org/10.1371/journal.pcbi.1013856>
- Bastian, M., Heymann, S., & Jacomy, M. (2009). Gephi: An open source software for exploring and manipulating networks. *Proceedings of the International AAAI Conference on Web and Social Media*, 3(1), 361–362. <https://doi.org/10.1609/icwsm.v3i1.13937>
- Bazin, A., Gautreau, G., Médigue, C., Vallenet, D., & Calteau, A. (2020). panRGP: A pangenome-based method to predict genomic islands and explore their diversity. *Bioinformatics*, 36, i651–i658. <https://doi.org/10.1093/bioinformatics/btaa792>
- Bazin, A., Médigue, C., Vallenet, D., & Calteau, A. (2021). *panModule: Detecting conserved modules in the variable regions of a pangenome graph*. bioRxiv. <https://doi.org/10.1101/2021.12.06.471380>
- Buck, M., Mehrshad, M., & Bertilsson, S. (2022). mOTUpan: A robust bayesian approach to leverage metagenome-assembled genomes for core-genome estimation. *NAR Genomics and Bioinformatics*, 4(3), lqac060. <https://doi.org/10.1093/nargab/lqac060>
- Garrison, E., Sirén, J., Novak, A. M., Hickey, G., Eizenga, J. M., Dawson, E. T., Jones, W., Garg, S., Markello, C., Lin, M. F., Paten, B., & Durbin, R. (2018). Variation graph toolkit improves read mapping by representing genetic variation in the reference. *Nature Biotechnology*, 36(9), 875–879. <https://doi.org/10.1038/nbt.4227>

- 177 Gautreau, G., Bazin, A., Gachet, M., Planel, R., Burlot, L., Dubois, M., Perrin, A., Médigue, C.,
178 Calteau, A., Cruveiller, S., Matias, C., Ambroise, C., Rocha, E. P. C., & Vallenet, D. (2020).
179 PPanGGOLiN: Depicting microbial diversity via a partitioned pangenome graph. *PLOS*
180 *Computational Biology*, 16(3), e1007732. <https://doi.org/10.1371/journal.pcbi.1007732>
- 181 Grant, J. R., Enns, E., Marinier, E., Mandal, A., Herman, E. K., Chen, C., Graham, M.,
182 Van Domselaar, G., & Stothard, P. (2023). Proksee: In-depth characterization and
183 visualization of bacterial genomes. *Nucleic Acids Research*, 51, W484–W492. <https://doi.org/10.1093/nar/gkad326>
- 184
- 185 Grüning, B., Dale, R., Sjödin, A., Chapman, B. A., Rowe, J., Tomkins-Tinch, C. H., Valieris,
186 R., & Köster, J. (2018). Bioconda: Sustainable and comprehensive software distribution
187 for the life sciences. *Nature Methods*, 15(7), 475–476. <https://doi.org/10.1038/s41592-018-0046-7>
- 188
- 189 Holley, G., & Melsted, P. (2020). Bifrost: Highly parallel construction and indexing of colored
190 and compacted de bruijn graphs. *Genome Biology*, 21(1), 249. <https://doi.org/10.1186/s13059-020-02135-8>
- 191
- 192 Hyatt, D., Chen, G.-L., LoCascio, P. F., Land, M. L., Larimer, F. W., & Hauser, L. J. (2010).
193 Prodigal: Prokaryotic gene recognition and translation initiation site identification. *BMC*
194 *Bioinformatics*, 11(1), 119. <https://doi.org/10.1186/1471-2105-11-119>
- 195 Land, M., Hauser, L., Jun, S.-R., Nookaew, I., Leuze, M. R., Ahn, T.-H., Karpinets, T.,
196 Lund, O., Kora, G., Wassenaar, T., Poudel, S., & Ussery, D. W. (2015). Insights from 20
197 years of bacterial genome sequencing. *Functional & Integrative Genomics*, 15(2), 141–161.
198 <https://doi.org/10.1007/s10142-015-0433-4>
- 199 Larralde, M. (2022). Pyrodigal: Python bindings and interface to prodigal, an efficient
200 method for gene prediction in prokaryotes. *Journal of Open Source Software*, 7(72), 4296.
201 <https://doi.org/10.21105/joss.04296>
- 202 Li, H., Marin, M., & Farhat, M. R. (2024). Exploring gene content with pangene graphs.
203 *Bioinformatics*, 40(7), btae456. <https://doi.org/10.1093/bioinformatics/btae456>
- 204 Mainguy, J., Lemane, T., Bazin, A., Arnoux, J., Gautreau, G., Médigue, C., Calteau, A.,
205 & Vallenet, D. (2026). PanGBank: A large-scale resource of precomputed microbial
206 pangenomes built with PPanGGOLiN. *bioRxiv*. <https://doi.org/10.64898/2026.08.05.742796>
- 207
- 208 Parks, D. H., Chuvochina, M., Waite, D. W., Rinke, C., Skarszewski, A., Chaumeil, P.-
209 A., & Hugenholtz, P. (2018). A standardized bacterial taxonomy based on genome
210 phylogeny substantially revises the tree of life. *Nature Biotechnology*, 36(10), 996–1004.
211 <https://doi.org/10.1038/nbt.4229>
- 212 Shannon, P., Markiel, A., Ozier, O., Baliga, N. S., Wang, J. T., Ramage, D., Amin, N.,
213 Schwikowski, B., & Ideker, T. (2003). Cytoscape: A software environment for integrated
214 models of biomolecular interaction networks. *Genome Research*, 13(11), 2498–2504.
215 <https://doi.org/10.1101/gr.1239303>
- 216 Snipen, L., & Liland, K. H. (2015). Micropan: An r-package for microbial pan-genomics. *BMC*
217 *Bioinformatics*, 16(1), 79. <https://doi.org/10.1186/s12859-015-0517-0>
- 218 Tettelin, H., Massignani, V., Cieslewicz, M. J., Donati, C., Medini, D., Ward, N. L., Angiuoli,
219 S. V., Crabtree, J., Jones, A. L., Durkin, A. S., Deboy, R. T., Davidsen, T. M., Mora,
220 M., Scarselli, M., Margarit y Ros, I., Peterson, J. D., Hauser, C. R., Sundaram, J. P.,
221 Nelson, W. C., ... Fraser, C. M. (2005). Genome analysis of multiple pathogenic isolates of
222 streptococcus agalactiae: Implications for the microbial "pan-genome". *Proceedings of the*
223 *National Academy of Sciences of the United States of America*, 102(39), 13950–13955.
224 <https://doi.org/10.1073/pnas.0506758102>

- 225 The Galaxy Community. (2024). The galaxy platform for accessible, reproducible, and
226 collaborative data analyses: 2024 update. *Nucleic Acids Research*, 52, W83–W94. <https://doi.org/10.1093/nar/gkae410>
227
- 228 Tonkin-Hill, G., MacAlasdair, N., Ruis, C., Weimann, A., Horesh, G., Lees, J. A., Gladstone,
229 R. A., Lo, S., Beaudoin, C., Floto, R. A., Frost, S. D. W., Corander, J., Bentley, S. D., &
230 Parkhill, J. (2020). Producing polished prokaryotic pangenomes with the panaroo pipeline.
231 *Genome Biology*, 21(1), 180. <https://doi.org/10.1186/s13059-020-02090-4>
- 232 Vallenet, D., Calteau, A., Dubois, M., Amours, P., Bazin, A., Beuvin, M., Burlot, L., Bussell,
233 X., Fouteau, S., Gautreau, G., Lajus, A., Langlois, J., Planel, R., Roche, D., Rollin, J.,
234 Rouy, Z., Sabatet, V., & Médigue, C. (2020). MicroScope: An integrated platform for
235 the annotation and exploration of microbial gene functions through genomic, pangenomic
236 and metabolic comparative analysis. *Nucleic Acids Research*, 48, D579–D589. <https://doi.org/10.1093/nar/gkz926>
237

DRAFT